A non-Russellian treatment of the referential-attributive distinction

John-Michael Kuczynski
University of California, Santa Barbara

Kripke made a good case that "…the phi…" is not semantically ambiguous between referential and attributive meanings. Russell says that "…the phi…" is always to be analyzed attributively. Many semanticists, agreeing with Kripke that "…the phi…" is not ambiguous, have tried to give a Russellian analysis of the referential-attributive distinction: the gross deviations between what is communicated by "…the phi..", on the one hand, and what Russell's theory says it literally means, on the other, are chalked up to implicature. This paper shows that, when the phenomenon of implicature is scrutinized, there is overwhelming reason to doubt that a Russellian analysis can succeed. A positive, non-Russellian analysis is proposed: it is shown that, if definite descriptions are treated as referring expressions, it is easy to deal with the referential-attributive distinction. When "…the phi…" is functioning attributively, the definite description is seen as referring to some object described in an understood, antecedent existence claim.

1. Introduction

Let us refer to sentences containing definite descriptions as "D-sentences." Russell's Theory of Descriptions concerns the semantics — that is, the literal meaning — of such sentences. The theory of descriptions is essentially this. Any sentence of the form

(i) "…the phi…"

means

(ii) *exactly one object has phi and…that object…*

or, more formally,

Pragmatics & Cognition 12:2 (2004), **253–294**.
ISSN 0929–0907 / E-ISSN 1569–9943

(ii_F)[1] *for some x, given any y, y has phi iff y=x and…x…*

So, for example:

(iii) the king of France is bald

means

(iv) *exactly one person is a king of France, and that person is bald.*

or, more formally,

(iv_F) *for some x, given any y, y=x iff y is a king of France, and x is bald.*

Keith Donnellan noted that D-sentences are typically used to convey or communicate propositions of two very different kinds (Donnellan 1966: 237). (This is not to say that such sentences are *ambiguous* in terms of their literal meanings, only that such sentences can be *used to communicate* two very different kinds of propositions.) Consider the sentence

(1) The person who murdered Smith is surely insane.

This sentence could be used to convey some existence claim, namely:

(2) *Some one person murdered Smith, and that person (whoever it is) is surely insane.*

On the other hand, (1) could be used to convey a proposition that is (de re) about a *particular* person: the speaker might have some particular person in mind to whom he wishes to attribute the property of being insane. In such a case, (1) will convey a proposition of the form:

(3) *So and so is insane.*

When a D-sentence is being used to convey an existential generalization, let us say that it is operating *attributively.* And when a D-sentence is being used to convey a statement that is (de re) about a particular object, let us say that it is operating *referentially* (Donnellan 1966: 237).[2] Also, let us say that, if a sentence S has a proposition P for its literal meaning, then S *semantically encodes* P.

As I will be using the term 'convey', a sentence-token may *convey* a proposition other than, or additional to, its literal meaning. If you and I are robbing a bank, and I say "the cops are coming", I am *conveying* or *communicating* that we should hurry up. But the literal meaning of my words is not *we should hurry up*. I will use the term 'communicate' as a synonym for 'convey', thus defined. So what is "conveyed" by a sentence token need not coincide with what is *semanti-*

cally encoded in that sentence-token — i.e., with that token's literal meaning. In saying that D-sentences are typically used to *convey* both referential and attributive propositions, I have said nothing about what their *literal meanings* are.

But of course it is worth knowing what the *literal meanings* of D-sentences are. Accordingly, my purpose here is to answer the question: Granted that D-sentences sometimes *convey* referential propositions and sometimes convey attributive propositions, what is the correct *semantic* analysis of such sentences? Of course, one answer to this question is to say that, at the level of literal or semantic meaning, definite descriptions are sometimes functioning referentially and sometimes are operating attributively: "the person who murdered Smith is insane" would be semantically ambiguous between (2) and (3).[3]

Kripke (1977: 178) came up with a powerful argument against this view. If "the" were really *semantically* ambiguous between referential and attributive meanings, then we would expect there to be languages that had *different* words corresponding to these two different meanings, just as there are languages that have different words corresponding to the two meanings of "bank". Kripke says he would be "very surprised" if there were such a language (Kripke 1977: 179). Kripke is right: as a matter of empirical fact, there are no such languages. There is no natural language L such that the L-translation of "the guy who murdered Smith is insane" (attributive use) would involve one article while the L-translation of "the guy who murdered Smith is insane" (referential use) would involve a different article.[4] Given this, it is extremely hard to believe that the referential-attributive distinction has to do with the *semantics* of D-sentences. Accordingly, I will operate on the assumption that "the" is not ambiguous.

Here is the plan of this essay. In Section 2, I will try to show that some *apparent* counter-examples to Russell's theory are *actual* counter-examples to it, notwithstanding some valiant efforts made by neo-Russellians. In Section 3, I will show that Russell's theory does not naturally apply to definite descriptions occurring in extended discourses: given any extended discourse containing definite descriptions, e.g. Henry Kissinger's book *the White House Years*, Russell's theory has the absurd consequence that it is one single existence claim. In Section 4, I will set forth a positive analysis of the referential/attributive distinction that assumes the falsity of Russell's view and that assumes that definite descriptions are just what they appear to be, namely referring expressions.

2. An argument against the Theory of Descriptions

Outline

Everyone agrees that there are many *apparent* counter-examples to Russell's theory. But Russellians say that when pragmatics — more specifically, the phenomenon of *implicature* — is taken into account, these apparent counter-examples prove to be *merely* apparent. In this section, I will show that, on the contrary, consideration of pragmatics actually gives *additional* force to these apparent counter-examples and makes it very hard to believe that they are anything other than *actual* counter-examples. The argument I will present will take for granted some Gricean views on pragmatics. These views are (in their broader outlines) almost universally held. But, strictly speaking, my argument will have a conditional form: *if* Grice's views are at least approximately right, *then* such and such follows about Russell's theory.

2.1 Two classes of apparent counter-examples to Russell's theory

The Russellian is, of course, right that what a sentence literally means is not always what it communicates. But that fact cannot be used to exonerate every semantic theory. To take a trivial example, suppose someone argued that "grass is green" really meant *snow is white.* That person's semantic theory would conflict with the linguistic data, and in such a case the conflict would clearly have a semantic basis. I argue that, in the case of Russell's theory, the conflict has just as much of a semantic basis as it does in the just mentioned theory.

There are many classes of apparent exceptions to Russell's theory. I will focus on two of these classes, and on the Russellian (or neo-Russellian) attempts to neutralize these apparent exceptions.

Here is an instance of the first of these two classes of apparent exceptions. Imagine the following. We are at a party. We both see a red-haired man talking to a blond woman. I say to you:

(a) the red-haired man is a lawyer.

The proposition that I have communicated to you is *object-dependent* with respect to the red-haired man in question. If Bob is the red-haired man in question, then in order for the proposition I've *communicated* to you to be correct, *Bob* must be a lawyer. If Fred is the red-haired man in question, then in order for the proposition I've communicated to you to be correct, Fred must be a

lawyer. Basically, the identity of the proposition depends on the identity of the red-haired man: the proposition in question is thus *de re* and concerns that man (Donnellan 1966: 237).

Now according to Russell's theory, the literal meaning of (a) is given by

(b) there is exactly one red-haired man and that man is a lawyer.

Yet, (b) is not semantically *de re* about anyone; it *describes* someone, but it doesn't *refer* to anyone.

Sentences like (a) are by no means unusual. In fact, D-sentences are very often used to communicate object-dependent propositions. So here we have a prima facie conflict between what D-sentences actually do, on the one hand, and what Russell's theory says they do, on the other.

Russellians try to neutralize this conflict by saying the following:

> It is true that no token of (a) semantically encodes an object-dependent proposition; and it is also true that many tokens of (a) communicate an object-dependent proposition. But that is no threat to Russell's theory. Semantics is not pragmatics. Russell's theory gives the semantics of D-sentences, not their pragmatics: it tells us what such sentences literally mean, not what they are used to communicate.
>
> Why does (e.g.) (a) communicate a proposition that is de re about a certain red-haired man, when it does not semantically encode any such proposition? Given the context, it is de rigueur to interpret (a) as bearing an object-dependent proposition — just as a waiter who hears "I'd like a vodka tonic" knows that, given the context, it is de rigueur of him to interpret that sentence as bearing a command.

Let us refer to this doctrine as "contextualism$_o$"[5] — note the subscript, which stands for "object-dependent". The basic idea behind contextualism$_o$ is this: Let P be the proposition that, according to Russell's theory, (a) semantically encodes; and let P' be the proposition that, under the circumstances described, (a) communicates: after contextual factors are taken into account, we'd expect (a) to communicate P' *given* that it semantically encodes P. Basically, (a) *implicates* (to use Grice's term) an object-dependent proposition but semantically encodes some existence claim.[6]

Here is an instance of the other of the two classes of apparent exceptions we will be concerned with. Consider the sentence

(c) the T.V. is broken.

According to Russell's theory, this is synonymous with:

(d) there is exactly one T.V. and it is broken.

But (c) has absolutely never been used to communicate the proposition *there is one T.V. and it is broken.* And this, of course, is a problem for someone who says that (c) is synonymous with "there is exactly one T.V. and it is broken". Some Russellians (though not all[7]) have responded to this by saying the following:

> Quantifiers are generally context-sensitive. If you say to your friend "there is no beer", you obviously don't mean there is no beer in all of existence; you mean there is no beer in the refrigerator or there is no beer in the fraternity house. If you say "everyone came to my lecture" you don't mean every single individual — all six billion of them — came to your lecture; you mean every one in the department came to your lecture. Basically, quantifiers are sensitive to context and respond to context by becoming abridged in scope.
>
> Given this, it is no mystery why "the T.V. is broken" never communicates the proposition there is exactly one T.V. and it is broken. According to Russell's theory "the T.V. is broken" is synonymous with some sentence beginning with a "there is" (or, in any case, with some equivalent expression e.g. "for some"). So, in effect, a "there is" (or a "for some") is implicit in "the T.V is broken", supposing that Russell is right. Now this "there is" — like any other "there is" in English — is context-sensitive: it responds to context by becoming appropriately — and perhaps severely — restricted in scope. So if you say "the T.V. is broken", the "there is" implicit in that sentence ranges only over some highly restricted domain — e.g. items in the house or items that are salient at the moment.[8]

Let us refer to this doctrine as "contextualism" (no subscript). Contextualism may be summed up as follows: quantifiers — in particular, "there is" (and its equivalents: e.g. "for some") — are context-sensitive. Given this, it is no mystery why e.g. "the T.V. is broken" is never used to communicate (inter alia) the proposition that there is exactly one T.V. in all of existence; the "there is" implicit in "the T.V. is broken" responds to context by becoming appropriately abridged in scope — just like the "there is" in "there is no beer" — so that what is meant by that token is not the absurd proposition that there is only one T.V. in all of existence and that it is broken, but only the perfectly reasonable proposition that there is one T.V. *in a certain context* (e.g. in the immediate vicinity) and that it is broken.

2.2 Can pragmatics help Russellians?

I will now show that contextualism and contextualism$_o$ are inconsistent with some exceedingly plausible and almost universally held principles of pragmatics.

In literally hundreds of publications, Russellians have taken it for granted that contextualism and contextualism$_o$ are satisfactory ways of dealing with the apparent counter-examples just mentioned. But what has not been done on a single occasion (at least not in any well known publication) has been to see if these doctrines actually square with the accepted principles of pragmatics. For example, consider Stephen Neale's book *Descriptions*, which is the most recent defense of the Theory of Descriptions, and probably the most complete one ever written. There, Neale acknowledges that "…the phi…" may express a proposition that is object-dependent with respect to the bearer of phi. But he sees that this has no bearing on whether Russell's theory is correct. His argument is simply: Russell's theory is a semantic theory; what is communicated by a sentence is not necessarily its literal meaning — implicature may be involved. Therefore the apparent counter-examples just cited are not actual counter-examples.

Of course, Neale is right that what is communicated is often an implicature, and not literal meaning. But that doesn't mean that Russell's theory is off the hook. For *some* deviations between semantic theory, on the one hand, and what sentences communicate, on the other, *do* have a semantic, as opposed to a pragmatic, basis. We gave a trivial example of this earlier. If someone said that "grass is green" means *snow is white*, that person's semantic theory would be inconsistent with the linguistic data, and the conflict would obviously have a semantic basis. Here is a less trivial example. Pre-Frege semanticists said that "someone" is a referring expression. This conflicts with the fact that "someone is tall; someone is short; therefore someone is both tall and short" does not express a valid argument. Of course, someone who wanted to hold on to the idea that "someone" is a referring expression could say:

> I grant that the sentence "someone is tall; someone is short; therefore someone is both tall and short" doesn't communicate a valid argument. But that doesn't mean that "someone" isn't a referring term; for what is communicated by a sentence is often an implicature, and not its literal meaning.

Obviously such an argument would have essentially no plausibility.

So when one plays the pragmatics card to exonerate some semantic theory, one has to be willing to see if the accepted truths about pragmatics

and implicature — e.g., that conversational implicature is *cancelable*, that conventional implicature doesn't obscure the truth-conditional content of the relevant sentence (in a moment these claims will be clarified) — actually bear out one's claim. Consider the semantic theory mentioned a moment ago: (*) "snow is white" means *grass is green*. It is a linguistic datum that (ceteris paribus) "snow is white" communicates the proposition *snow is white*, not *grass is green*. The proponent of (*) must say that implicature is responsible for this deviation between literal and implied meaning. But conventional implicature isn't at work; for conventional implicature doesn't mask the truth-conditions semantically associated with a sentence; in other words, a conventional implicature will not make people be wrong about what the truth conditions of a sentence are. The sentence 'shysters make a lot of money' differs in its conventional implicatures from "lawyers make a lot of money". But it is patently obvious that those two sentences have exactly the same truth-conditions. And this point generalizes without limit: conventional implicature simply does not hide truth-conditions.

What no Russellian does — and this includes Neale — is to see if the principles of pragmatics actually bear out the thesis that "….the phi…" is an existence claim. Obviously context interacts with semantics. But it interacts according to certain laws — laws that it is the business of pragmatics to describe. (Grice himself has given us some of these laws.) I submit that when we consider these laws, and consider what kinds of interactions between semantics and context *would* occur if Russell's theory were correct, we find that what *would* thereupon be communicated is quite dramatically different from what *is* communicated. Again, no Russellian, and definitely not Neale, bothers to attend to these interactions. They say, rightly: Russell's theory is semantics, pragmatics is not semantics, and implicature is not literal meaning. But they don't take the further step of seeing if the laws of pragmatics make communicated meaning deviate from literal meaning *in the right way*, i.e. in a way that is consistent with the theory in question (see Neale 1990: Chapter 3). The purpose of the next few pages is to show that when we consider the laws of pragmatics, and pit them against Russell's theory, the result is an *unacceptable* deviation between Russell's theory and what is communicated by "…the phi…".

2.3 The concept of implicature

Semantics studies what sentences literally mean. Pragmatics studies what is communicated by sentences but is *not* part of their literal meanings. So it

studies what is *suggested* or *implicated* by sentences. What is suggested or implicated by a sentence-token is typically referred to as an *implicature.*

It is generally agreed that, ultimately, there are two kinds of implicature: conversational and conventional.[9] I will not attempt to define these terms; I will illustrate their meanings through examples.

2.3.1 *Conversational and conventional implicature*

Consider the following exchange:

A: Where can I get a bite to eat?
B: There's a McDonald's down the road

B has not *explicitly* told A that he (A) can get a bite to eat at the nearby McDonalds; B has only *implicated* this. B's words *literally* mean only *there's a McDonald's down the road;* they *don't* have as their literal meaning *you can get a bite to eat down the road.* Nonetheless, given the context, and given certain understandings that govern human conversation, B can reasonably interpret A's words as an answer to his question.[10]

Now suppose I say:

(1) shysters make a lot of money.

Obviously, I will be communicating a certain lack of respect for lawyers. Now, if say

(2) lawyers make a lot of money,

I will *not* be communicating any such lack of respect.

Nonetheless, there doesn't seem to be any *semantic* difference between (1) and (2); they seem to be perfectly synonymous; they necessarily have the same truth conditions and — what is probably more on point[11] — the propositions they express consist of the same objects and concepts arranged in the same way.[12]

The disrespectful message communicated by (but not semantically encoded in) (1) is an example of a *conventional* implicature. What is the difference between conventional and conversational implicatures? The sentence "there is a gas station around the corner" *can* be used to communicate the proposition *you can get gas around the corner.* But it will do so *only in certain conversational contexts*. But the sentence 'shysters make a lot of money" bears a message about lawyers *regardless* of the conversational context. The word "shyster" is

conventionally used to denigrate lawyers; so sentences containing that word denigrate lawyers regardless of the conversational context.

2.3.2 *Some facts about conversational implicature*

To help us evaluate contextualism and contextualism$_0$, we should make a few observations about both kinds of implicature — in particular, about how such implicatures can be distinguished from the literal meanings of the sentence-tokens that bear them. Let us start with conversational implicature.

Conversational implicature can always be *cancelled* (Grice 1975: 159).

A: Where can I get gas?
B: There's a gas station right around the corner. But it's closed.

B has *cancelled* the implicature that his first sentence would otherwise have borne.

Given that conversational implicatures are cancelable, it follows that people are able to distinguish literal meaning from conversationally implicated meaning. If B were to say "there's a McDonald's down the road, but you can't get a bite to eat there", that statement would not register with any competent, cognitively normal speaker as a contradiction. Why does it not so register? Because, at some level, any cognitively normal English speaker in A's position knows that the words "there's a McDonald's down the road" do not, strictly speaking, mean *you can get a bite to eat at the McDonald's.* To see this, suppose that the words "there's a McDonald's down the road" *did* register as having for their literal meaning the proposition *you can get a bite to eat at the McDonald's down the road;* in other words, suppose people *took* the words "there's a McDonald's down the road" as meaning *you can a bite to eat at the McDonald's down the road.* In that case, if B said

(a) there's a McDonald's down the road, but you can't get a bite to eat there,

that would register as being every bit as self-contradictory as the words

(b) you can get a bite to eat at the McDonalds down the road, but you cannot get a bite to eat at the McDonalds down the road.

But the simple truth is that (b) does, whereas (a) does not, register as being self-contradictory.

The fact that conversational implicatures are so neatly cancelable means that, at some level, people have a solid grip on what is conversationally implicated as opposed to literally stated. This has some important consequences. First, conversational implicatures do not cause consistent sentences to register

as self-contradictory. Also, conversational implicatures do not cause self-contradictory sentences to register as consistent: if a self-contradictory statement *implicates* a true statement, it will always be as clear as day that, first of all, the statement in question is *technically* self-contradictory, and therefore, second, that the consistent message that was communicated was only implicated. Suppose you say "that plane is a mile long". Obviously you are implicating a proposition that may well be correct, viz. *that plane is very long* or *the plane is too long to fit into the hangar.* Now suppose you say "that plane is a mile long, but that plane is not a mile long". It will be as clear as day that, technically, this sentence is self-contradictory. Of course, the sentence "that plane is a million miles long, but it is not *really* a million miles long", may (like practically any string of verbiage) be used to implicate a true, and therefore consistent, statement (e.g. *that plane is large*, *but it will still fit into the airplane hangar*). But no one would ever have any temptation to suppose that this consistent statement was literally meant by the sentence in question.

2.3.3 *Some facts about conventional implicature*

Let us now turn to *conventional* implicature. Much of what we said about conversational implicature applies to conventional implicature.

Suppose I am in the unfortunate position of having to inform you that your mother has died. Of course, there are many English locutions that can transmit this information. Suppose I say:

(c) your mom croaked.

Obviously, I will be communicating a certain lack of respect for your mother and, therefore, for you. On the other hand, if I say

(d) your mother passed on,

I will *not* be communicating any such lack of respect.

Nonetheless, there doesn't seem to be any *semantic* difference between (c) and (d); they seem to be perfectly synonymous; they necessarily have the same truth conditions and — what is probably more on point — the propositions they express consist of the same objects and concepts arranged in the same way.

Conventional implicatures are not cancelable, at least not in any neat way (see Mey 1993: 103–106). Suppose I said "your mother croaked" and then tried to cancel the implicated insult by adding "your mother was a dear woman": the effect would not be a neat cancellation of the offensive implicature. It wouldn't

really be clear what was going on. But even though conventional implicatures are not neatly cancelable, there is generally no problem separating what is conventionally implicated in a sentence from what is semantically encoded in it. It is clear enough what the literal meaning of "your mother croaked" is and it is also clear what is implicated by it. What is literally meant is *your mother has ceased to live*; what is implicated is some kind of a cavalier attitude towards this fact.

Given this, some important principles follow. For one, conventional implicature will not make a self-contradictory statement sound consistent. "Your mother croaked, but she did not die" registers as quite self-contradictory.

Conventional implicature *can* make a consistent sentence (i.e., one that semantically encodes a consistent proposition) sound paradoxical. But it will be obvious that the paradox is pragmatic in nature — that it is a "Moore"-paradox and not a semantic paradox. Suppose I say "your mother croaked, but her death is a matter of the gravest importance". That sentence is, at some level, paradoxical. For, in saying "your mother croaked", I am implicating that I have a frivolous attitude towards her death, whereas in saying "I have the gravest respect for her", I am stating that my attitude towards her death is not frivolous. But no one would ever dream of denying that "your mother croaked, but her death is a matter of the gravest importance" could be perfectly true; in fact, almost everyone who heard it would acknowledge that if the first conjunct was true, then so was the second. However, it would be clear that anyone who uttered that sentence would be making a "Moore paradoxical" statement (like "snow is white, but I don't believe that snow is white").

There is another way to put this. Given a sentence that is consistent in terms of its literal meaning, and that consequently has consistent truth-conditions, conventional implicature cannot make that sentence sound as though it has *in*consistent truth-conditions. Everyone agrees that "your mother croaked, but her death is a matter of grave importance" is *at some level* paradoxical; but everyone can see right away that its *truth*-conditions are consistent (are not like those of "x is a triangle and x is also a circle"). It is patently obvious that what is paradoxical about that sentence has nothing to do with its truth-conditions, but with some kind of truth-conditionally inert innuendo.

Also, conventional implicatures do not counter-act or undercut the literal meanings of the sentences that carry them. Suppose someone in a position of authority asks me a yes-no question, and I respond by saying "yup". My response is informal and thus implicates a certain lack of deference towards my interlocutor. If I'd responded by saying "yes", or "affirmative", my response

would have been more formal and thus would have implicated a greater degree of respect. Surely "yup" and "yes" are synonymous and differ only in conventional implicature. Notice that the disrespect implicated in my saying "yup" in no way obscures the semantics of my statement. In fact, the cavalier quality associated with saying "yup", instead of "yes", would actually make it doubly clear that I was tendering an affirmative answer to your question.

2.4 An important point about contextualism

Before we can evaluate either contextualism or contextualism$_o$, we need to make explicit one point about contextualism. When we say that quantifiers are context-sensitive, what we mean is that *all by themselves* — i.e., without explicit restricting information having to be added: without extra *verbiage* having to be added — quantifiers respond to contexts of utterance by becoming appropriately abridged in scope. If you open the refrigerator and you say to your roommate "there's no beer", it will be clear from context alone that you mean there's no beer in the refrigerator. You won't have to add some restricting clause for it to be clear that you want the existential quantifier to range only over objects in the refrigerator. You won't have to say, for example, "there's no beer *in the refrigerator*". In any case, in so far as you did have to do this — in so far as you did have to say "there is no beer *in the refrigerator*" in order to keep the "there is" from being understood to range over everything in existence (as opposed to just the objects in the refrigerator) — to that extent the "there is" would obviously not be sensitive to *context*. If you had to *add* some rider like "in the refrigerator", that would mean that context was *failing* to effect the right restrictions on the quantifier; it would mean, in other words, that the quantifier was *not* responding to context by becoming abridged in scope — was *not* "context-sensitive." To sum up, in so far as explicit restrictions must be put on a quantifier to restrict its range, to that extent the quantifier in question isn't *context*-sensitive.

Consequently, when the contextualist is explaining why my utterance of "the red-haired man is a lawyer" is not taken to mean *there is exactly one red-haired man x in all of existence and x is a lawyer*, he cannot say that this is connected with the presence of verbiage in the Russellian paraphrase of this sentence; i.e., he cannot say that this is because, in the Russellian paraphrase of that sentence, the quantifiers are subject to any *explicit* restrictions — are *explicitly* abridged in scope by verbiage. He is going to say that it is entirely due to the *context-sensitivity* of the quantifiers in that paraphrase.

Once again consider my tokening of "the red-haired man is a lawyer". The proposition communicated by this token is obviously not: *there is exactly one red-haired man in all of existence, and that individual is a lawyer.* The contextualist deals with this by saying that the "there is" in the Russellian paraphrase of "the red-haired man is a lawyer" is *context-sensitive*, and that, as a result of this context-sensitivity, its scope is restricted. So what ends up being communicated is (say) *there is exactly one red-haired man* in this room *and that individual is a lawyer* (not *there is exactly one red-haired man* in all of existence…). Now, to say that the "there is" becomes thus abridged in scope as a result of its *context-sensitivity* is to say precisely that no verbiage must be *added* to it — that no *explicit* restrictions must be placed on it — in order to make it be appropriately abridged in scope. So the contextualist cannot say that the Russellian paraphrase of "the red-haired man is a lawyer" is (e.g.) "there is exactly one red-haired man *in this room* and that individual is a lawyer"; for if he said that this, or any other sentence containing *explicit* restrictions on the "there is", were the right paraphrase of "the red-haired man is lawyer", he would be conceding that context *by itself* was not enough to make the "there is" appropriately abridged in scope — he would be conceding that extra *verbiage* had to be added to restrict the "there is" and, by implication, that *context* was not able to effect that restriction. So one would be abandoning contextualism if one said that the "there is" that (according to Russell's theory) is implicit in "the red-haired man is a lawyer" were restricted by the presence of extra *verbiage* in the Russellian paraphrase of that sentence.

So the contextualist paraphrase of a given utterance of "the red-haired man is a lawyer" is not going to be (e.g.) "there is exactly one red-haired man x next to that sofa, under that chandelier, near that door and x is a lawyer". It is simply going to be: "there is exactly one red-haired man x, and x is a lawyer".

In general, given any D-sentence, the contextualist paraphrase of a given token of that sentence is not going to be different, in terms of verbiage, from the orthodox Russellian paraphrase.[13]

Before we close the present argument, a cautionary note is in order. When discussing contextualism, we must take special care to distinguish sentences from propositions. The contextualist admits — indeed, he insists — that the *proposition* encoded in a D-sentence like "the red-haired man is a lawyer" will vary with context. But he attributes this variation to the context-sensitivity of the quantifiers in the paraphrase of that D-sentence: he will attribute this to the fact that these quantifiers — all by themselves, without extra verbiage having to be added — respond to context by becoming appropriately abridged in scope.[14]

The *paraphrase* of a D-sentence doesn't change with context, although (according to contextualism) the *proposition* encoded in it does.

2.5 Why contextualism cannot save the Theory of Descriptions

Having done all of this preparatory work, we can now evaluate contextualism and contextualism$_0$. Let us start with the former.

The contextualist paraphrase of:

"the T.V. is broken"

is simply

"there is exactly one T.V. and it is broken".

In other words, the contextualist says that these two sentences are perfectly synonymous. But consider the sentence:

(i) "the T.V. is broken. But there are many television sets, so I shall buy another".

This certainly doesn't register as even minimally self-contradictory. Now if contextualism is right then (i) is synonymous with:

(ii) *"there is exactly one T.V. and it is broken. But there are many television sets, so I shall buy another."

Indeed (ii) *registers* as a full on contradiction: and it obviously *is* a contradiction.

Here we have a major problem for the contextualist. What is communicated by a D-sentence in one environment does *not* register as a contradiction, while what is communicated by the Russellian paraphrase of that D-sentence in an exactly similar environment *does* register as a full on contradiction. This is a problem for the contextualist because, as we have seen, he is committed to the view that (e.g.) "the T.V. is broken" is a mere notational variant of "there is exactly one T.V. and it is broken"; and it boggles the intellect how mere notational variants could communicate such radically different propositions: one a self-contradiction, the other a perfectly consistent proposition.

Now there is still hope for the contextualist. For he can, at this point, appeal to the pragmatics/semantics distinction; he can say that what is *semantically encoded* in "the T.V. is broken" is identical with what is semantically encoded in "there is exactly one T.V. and it is broken", but that (in the context above) one of those sentences carries an implicature that the other does not carry.

But given what we've noted about implicature, this counter-move doesn't work. There are two kinds of implicature: conventional and conversational. Could the implicature in question be conversational?

No. (ii) registers as self-contradictory; (i) does not. This difference cannot be due to conversational implicature: for conversational implicature doesn't make a consistent sentence register as self-contradictory.

So, with regard to the fact that "the T.V. is broken" doesn't *communicate* that there is exactly one T.V. in existence, this fact cannot be chalked up to *conversational* implicature.

Could the difference between (i) and (ii) be due to *conventional* implicature? No. Remember that conventional implicature will not make a self-contradictory (or otherwise inconsistent) sentence register as consistent. So the reason why (i) registers as consistent, while its Russellian paraphrase does not, has nothing to do with conventional implicature.

Given that it has nothing do with either conventional or conversational implicature, it is obvious that it has nothing to do with any kind of implicature; for those are the only two kinds.[15] Since the difference between (i) and (ii) is not due to implicature, it is hard to avoid the conclusion that it is due to *semantics.* The only difference between (i) and (ii) is that one of them contains a D-sentence where the other contains the Russellian paraphrase of that sentence. So the semantic divergence between (i) and (ii) must be due to a semantic divergence between the D-sentence and its Russellian paraphrase. An exactly similar argument could be given to show that *any* D-sentence differs semantically from its Russellian paraphrase, and that Russell's theory is therefore false.

Let us sum up. Consider the statement

(A) I have two points to make. First, there is exactly one T.V. and it is broken. Second, there are many T.V. sets in existence, so I shall buy another.

(A) is obviously self-contradictory. Now consider the statement:

(B) I have two points to make. First, the T.V. is broken. Second, there are many T.V. sets in existence so I shall buy another.

(B) is *not* self-contradictory. We've seen that if contextualism is right, then (B) is simply a notational variant of (A) — just as "that stick is a yard long" is a mere notational variant of "that stick is three feet long". Remember that if S is a consistent statement, neither conversational nor conventional implicature — and therefore *no* kind of implicature — can make S sound inconsistent. So no kind of implicature could explain why (B) makes a consistent statement,

while (A) does not. Thus, it is *not* implicature that accounts for this difference. So semantics must be responsible: those statements are *semantically* different, contrary to what contextualism says.

Kent Bach is a Russellian but is *not* a contextualist. He holds that the literal meaning of "the T.V. is broken" is simply *there is exactly one T.V. and it is broken*: the *there is* is *un*restricted. Bach's (1987: 103–104) argument is this. Consider the following sentence:

(*) "the man who walked across Antarctica backwards is extraordinarily determined".

According to Russell's theory, in its non-contextualist form, (*) is synonymous with

(**) "there is exactly one x such that x is a man who walked backwards across Antarctica, and x is extraordinarily determined".

Surely, Bach claims, the "there is" that is explicit in (**) is not subject to any context-based restriction. Therefore the "there is" that, according to Russell's theory, is implicit in (*) is not subject to any such restriction. So, if Russell's theory is right, then the proposition semantically encoded in (*) is simply:

(***) There is exactly one x such that x is a man who walked backwards across Antarctica, and x is extraordinarily determined.

By parity of reasoning, if there were only *one* T.V. in the entire world, then the "there is" implicit in "the T.V. is broken" would not be subject to any context-based restriction. So, under that circumstance,

(i) "the T.V. is broken"

would mean

(ii) *There is exactly one T.V. and it is broken.*

Now we don't want to say that the semantics of (i) depends on how many T.V.'s there are in the world; we want (i) to have a *consistent* semantics — one that doesn't change depending on how many T.V.'s there are. Given that (i) would mean (ii) in a world where there was only one T.V.; and given also that the semantics of (i) wouldn't change if more T.V.'s were added to that world; it seems to follow that the literal meaning of (i) in *our* world — where there are several T.V. sets — is simply *there is exactly one T.V. and it is broken.* This point obviously generalizes to all other D-sentences.

As for the fact that tokens of (i) never *communicate* (ii), Bach deals with this (in typical neo-Russellian fashion) by saying that Russell's theory concerns the literal meaning of D-sentences, and that sentences often don't communicate what they literally mean.

Let us evaluate this argument. Bach is probably right to say that it would be wrong to make the semantics of "…the phi…" depend on how many phi's there were. (Presumably it would be wrong to make "the T.V. is broken" undergo a shift in meaning if, all of a sudden, every T.V. but one was annihilated.) So he has made a good case that *if* Russell's theory, in any of its various forms, is correct, it is not the *contextualist* form that is correct: it must be the "orthodox" form. But given what we've already seen, it is clear that Bach's (quite orthodox) version of Russellianism is not viable. Again, Bach's view is:

(i) "the T.V. is broken"

means

(ii) *There is exactly one T.V. and it is broken*

and is thus synonymous with

(iii) "there is exactly one T.V. and it is broken".

But we have set forth spent pages of arguments to the effect that "the T.V. is broken" is not synonymous with "there is exactly one T.V. and it is not broken". So we have already, in effect, refuted Bach's view.

2.6 Why contextualism$_o$ cannot save the Theory of Descriptions

Given what we've said about contextualism, it will not be hard to make a case against contextualism$_o$; in fact, we will hardly have to add anything to what we've already said.

Once again, the contextualist says that no token of

(i) "the T.V. is broken"

is *de re* about any T.V.; but he will admit — what really cannot be denied — that such a token may well communicate a proposition that is *de re* about a T.V. And he will say that this communicated proposition is an implicature.

Now a contextualist$_o$ may or may not be a contextualist. But it doesn't really matter. As we've seen, if he *is* a contextualist, then he will say that the paraphrase of (i) is

(ii) "there is exactly one T.V. and it is broken".

And if he is *not* a contextualist, then he will say that the paraphrase is

(ii) "there is exactly one T.V. and it is broken" .

So, either way, the contextualist$_o$ will maintain (or should maintain) that the right paraphrase of (i) is (ii) — he will maintain that (i) and (ii) are perfectly synonymous. But they are not synonymous, as we've already seen. So the contextualist$_o$ — whether or not he is also a contextualist — is simply in error as to the actual form of (i).

In this connection we can rebut an argument that is often given *for* Russellianism. If you say "the blue haired man is drinking", and the man in question is green-haired, your listener will be reluctant to give his full assent to what you are saying if he knows the man's actual hair color. (Importantly, he will also be reluctant to dismiss it as simply wrong, whereas he *could* dismiss its Russellian paraphrase as simply wrong.) This is often taken to favor Russell's account. Stephen Neale writes:

> The referentialist can say nothing useful here, but the Russellian can provide a theoretical explanation of the aforementioned tension: What has been left out by the referentialist is the Gricean distinction between the proposition expressed and the proposition(s) meant…According to Grice, "what, in such a case, the speaker has said may be false, but what he meant by be true" (Neale 1990: 142). The proposition expressed by an utterance of "the F is G" is still descriptive, but the speaker and hearer are willing to entertain the idea that some particular individual b is uniquely F in order to communicate an object-dependent proposition about b…Applied to Donnellan's example [which is, in the relevant respects, like our example involving the green-haired man], the proposition [semantically encoded in] my utterance of 'smith's murderer is insane" is false; but the proposition I intend to communicate is true…Thus the Russellian-Gricean has, if only a rudimentary way, an account of the conflicting pretheoretic intuitions we typically have when presented with cases involving misdescription. The possibility of misdescription does not advance the case for a semantically referential interpretation in the least; indeed, the unitary Russellian analysis has the edge here (p. 92).

It is true that one would be reluctant to give one's full assent to your utterance of "the man over there with the blue hair is drunk". But that is not the decisive fact. The decisive fact is that one would also be very reluctant to characterize as simply *false* what you are saying, *whereas one could unexceptionably so characterize the Russellian paraphrase of that statement.*

To see that this so, it will help if we draw a consequence of a point earlier made about implicature. We observed that, for any sentences S1, S2, and S3, if S1 and S2 are synonymous, then "S1 and S3" will register as having inconsistent truth-conditions if and only if "S2 and S3" registers as having inconsistent truth-conditions. Implicature won't change this. Basically, implicature will never make a consistent sentence register as having inconsistent truth-conditions and will never make an inconsistent sentence register as having consistent truth-conditions.

A corollary of this is that if S1 and S2 are synonymous, then it will be appropriate to respond to S1 with a "that is false" if, and only if, it is appropriate to respond in that way to S2. So implicature cannot make it appropriate to respond to a true sentence with a "that is false" or to a false sentence with a "that is true". The truth-conditions of "shysters make a lot of money" and "lawyers make a lot of money" are obviously the same; the two sentences differ only in implicature. Consequently, "shysters make a lot of money but it is not the case that lawyers make a lot of money" is obviously self-contradictory. I would be guilty of blatant incoherence if I said "that is false" in response to the one but not the other.

A: Lawyers make a lot of money.
B: That is false.
A: But surely shysters make a lot of money.
B: *That is true.

Of course, a given person might not wish to countenance use of a vulgar term like "shyster"; so his reaction to a sentence like "shysters make a lot of money" may be very different from his reaction to "lawyers make a lot of money". But it is clear to any speaker of English that I could not *technically* be correct to say "that is false" in response to the one unless I was also correct to say that in response to the other.

What we just said about "that is true" and "that is false" obviously applies to statements like "technically that is false", "that is probable", and so forth. One would be guilty of blatant incoherence if one said "that is probable" in response to "shysters make a lot of money" but "that is not probable" in response to "lawyers make a lot of money".

With these points in place, suppose there is a man standing near us whose hair *seems* to be blue because of the light, but is *actually* green. (In fact, there is no blue haired man anywhere near us.) I know that this gentleman's hair is

really green, but you do not know this. I also happen to know that the man in question is a lawyer. Given this, consider the following exchange

> You: There is exactly one blue-haired man over there, and that man is a lawyer.
> Me: Technically, what you said is false. There is no blue-haired man over there. *A fortiori* it is not the case that there is a blue-haired man over there who is a lawyer.

There is no denying that what I said is *technically* beyond reproach; you said (inter alia) that there was a blue-haired man over there; I rightly denied this. There are no semantic controversies here. Now consider the following exchange:

> You: The blue-haired man is a lawyer.
> Me: *Technically, what you said is false. There is no blue-haired man over there. *A fortiori* it is not the case that there is exactly one blue-haired man over there *and* that that man is a lawyer.

There is no denying that in *this* dialogue my response rings less true. There is no denying that my response, in this dialogue, is markedly abnormal. For here it is incumbent on me to give a much more qualified response: "well, the man you have in mind is indeed, a lawyer; but, in fact, his hair is green". Saying simply "that is false; there is no blue-haired man over there" registers as deeply abnormal: there is no denying *this*.

Of course, the dialogues are exactly the same except that you use a D-sentence in the second dialogue where, in the first, you use its Russellian paraphrase. So obviously the fact that my statement registers as technically unexceptionable in the first dialogue, and as less than unexceptionable in the second, has to do with this one difference. Now the Russellian will concede all of this: but he will *deny* that it has anything to do with there being a semantic difference between "the blue-haired man over there is a lawyer" and "there is exactly one blue-haired man over there, and that man is a lawyer". He *will* say that this difference is due to some kind of *implicature*. But this move is not feasible. Recall what we said about implicature. If two sentences S1 and S2 differ *only* in implicature, then "technically, that is false" (or "that is true") will be an appropriate response to the one if and only if it is an appropriate response to the other.

Neale is right in one matter: one's reluctance to give one's full assent to your utterance has to do with the presence of descriptive information in the definite description. But this fact ultimately doesn't favor a non-referential analysis of definite description; for a phenomenon exactly parallel occurs in connection

with demonstratives, and there are overwhelming independent reasons to believe demonstratives are referential (Kaplan 1989: 512–516). When I say to you "you are tall" my statement is *de re* about you. Now suppose (once again) we are at a party. We see what looks to be a woman in the corner. I happen to know it is a man dressed as a woman. But you do not know this. You say to me: "she was extremely rude to me earlier". (As it happens, the person in question was rude to you earlier; and I know this.) Now I will unquestionably feel a certain reluctance to assent to what you are saying. (I will also feel a reluctance simply to reject what you are saying.) I will probably not say (unless I wish to deceive you) "yes, she was rude to you earlier". But I will also be hesitant to say simply "no, you are wrong". Given that demonstratives are (probably) directly referential, we want to avoid explaining this uneasiness by becoming descriptivists about them; by saying that, e.g., utterances of "you are tall", "she is tall" are not *de re* about anyone. Given the distinction between a demonstrative's "character" and its "content" we can easily account for this uneasiness without becoming descriptivists about them. Demonstratives do have a kind of descriptive content ("character"); but this content does not figure into what is asserted when a demonstrative is used. When you say "she is drunk" (when "she" is being used demonstratively, not anaphorically), you are not asserting that some female is the object of your attention; but you are (roughly speaking) presupposing it. The uneasiness that I feel (in the hypothetical dialogue above) in assenting to your utterance of "she is drunk" — when the person in question is a male — would thus seem to correspond to my rejection of what you have presupposed, not of what you have asserted. An exactly analogous position can be taken with regard to "wrong" definite descriptions, e.g., "the man with the blue hair is drunk" (when the man in question has green hair): my uneasiness at assenting to this statement — even when I know the man in question is drunk — can be (and, given what we've said, should be) seen as corresponding to my rejection of a false presupposition of yours, and not to my rejection of what you have actually asserted.

Actually, what we just said about demonstratives not only defuses Neale's point; it provides an additional positive justification for a referential analysis of definite descriptions. Again, suppose I say, while pointing to a man dressed as a woman,

(i) She is drunk.

And suppose you know that the person in question is really a man. Your reaction to that will be one of neither wishing to give unqualified assent or of unqualified dissent: it will be just like your reaction to an utterance of

(ii) The blue-haired man over there is drunk

when you know that the person in question has green hair (but is drunk). In any case, your reaction to (ii) will be a lot more like your reaction to (i) than it would be to a sentence that is uncontroversially wrong, e.g.,

(iii) There is exactly one blue-haired man over there, and that man is drunk.

Under the circumstances described, it would be quite inappropriate to respond to (i) by saying simply "you're wrong: there is no woman over there". (A much more qualified response is necessary, e.g.: "well, the person you're referring to is drunk; but that person is actually a man".) And, under the circumstances described, it would be quite inappropriate to respond to (ii) by saying simply: "you're wrong: there is no man with blue hair over there". (Again, a much more qualified response is demanded: "well, the person you're referring to is drunk but…") But, under the circumstances, it would be appropriate — at least relatively appropriate — to respond to (iii) by saying simply "you're wrong: there is no such man over there". Technically, such a response would probably be unexceptionable. So a sentence containing a "wrong" definite description seems to behave much more like one containing a "wrong" demonstrative than it does like a false existence-claim; and this suggests that definite descriptions are means of securing reference rather than means of asserting existence.

3. Another argument against the Theory of Descriptions

In this section, I will set forth my second argument against Russell's theory. That argument is basically this: Russell's theory has the absurd consequence that discourses of any length — e.g. Gibbon's history of Rome — become *single*, *indissoluble* existence claims: so that unless one has read (e.g.) Gibbon's history in its entirety, one has not taken in a single proposition, but only some ill-formed fragment thereof.

In some cases, what a definite description refers to is determined by some antecedent sentence. The Theory of Descriptions tends to assign the wrong meanings to sentences containing definite descriptions of this type. Consider the following passage (for easy reference later on, I will number the sentences):

> (i) Two men were walking down the street, a tall one and a short one. (ii) The short man wanted to play basketball. (iii) The tall man wanted to play golf. (iv) The short man decided to play golf with the tall one.

Let us refer to this passage as S, and let us suppose that S is the beginning of an 800 page story. What is the meaning of (say) sentence (ii)? What is the correct semantic analysis of that sentence? Interpreted literally, the Theory of Descriptions says that it is this:

> (ii') There is exactly one short man and he wanted to play basketball.[16]

But that is patently not what (ii) is saying. A Russellian might respond by going contextualist: "(ii) is saying that *within a certain context* there is exactly one short man who wanted to play basketball". But that won't work. As we saw earlier, the contextualist is in complete agreement with the orthodox Russellian on the question of what *verbiage* paraphrases any given D-sentence. The orthodox Russellian says that (ii') gives the meaning of (ii), and the contextualist Russellian agrees. The difference is that, according to the contextualist, the "there is" in (ii') is context sensitive. So according to the contextualist S is synonymous with:

> ($i'_{contextualist}$) Two men were walking down the street, a tall one and a short one. ($ii'_{contextualist}$) There was exactly one short man and he wanted to play basketball. ($iii'_{contextualist}$) There was exactly one tall man and he wanted to play golf. ($iv'_{contextualist}$) There was exactly one short man and exactly one tall man, and the former decided to play golf with the latter.

Let S2 be *this* passage. Plainly S2 is not synonymous with S. S2 does, whereas S does not, *convey the message* that there is exactly one tall man and one short man. In any case, S2 has a *strong tendency* to convey such a message, whereas S has no such tendency. And this is evidence that they are not synonymous.

An objection might be made to this:

> For the sake of argument, suppose that Russell is right, and that S is synonymous with S2. Given how absurd it would be to assert that there was exactly one tall man in existence, wouldn't someone reading S figure out that it was supposed to convey that there was within a certain context exactly one tall man (as opposed to all of existence)? Under that circumstance, mightn't implicature explain why people read S in the way they do?

What kind of implicature would it be? Conventional? No — conventional implicature doesn't hide literal meaning: conventional implicature doesn't make two synonymous sentences even *appear* to diverge in terms of their

truth-conditions. But S and S2 very much appear to diverge in this respect. Nor could conversational implicature be responsible. Whether a conversational implicature attaches to a sentence, or string of sentences, is contingent on the circumstances in which that sentence is tokened. Even if we imagine S and S2 being tokened in exactly similar circumstances, they would still, almost invariably, diverge in respect of what they communicated. So conversational implicature simply couldn't be what is responsible for that divergence.

Admittedly, the Russellian does have a compelling response to all of this: We should think of S, along with the story it introduces, as one big existence claim. So S's logical form is:

> For some x, some y, x is a tall man and y is a short man and y is not identical with x, and x wanted to play basketball and y wanted to play golf, and y decided to play golf with x…

This proposal does seem to capture the meaning of S. But there is a problem with it. S is just the beginning of an 800 page discourse; and (we may well suppose) there are tokens of "the tall man" and "the short man" in page 800 that are co-referential with those in S. So if the proposal under consideration is correct, then this 800 page discourse becomes one big existence claim. In general, if that proposal is correct, then given any extended discourse that contains co-referential definite descriptions, that discourse is just one giant existence claim.[17] A corollary of this is that if you read the first 400 pages of (say) *The White House Years* and then stop, you haven't taken in a single well-formed proposition; you've only taken in an ill-formed fragment of some proposition. But that is plainly false.

Given that we don't want to treat the 800 page story introduced by S as one big existence claim, it follows that we don't want to treat "the short man" in (ii) as *bound* by the existential quantifier (implicit) in (i). For if we regarded the definite description in (ii) as thus bound, then we'd have to do the same with (iii) and (iv) — we'd have to do the same with any occurrence of "the short man" on page 800 that refers to the short man introduced in (i). So we'd turn the story in question into an 800 page existence claim. But, for reasons we've seen, we don't want to do that. This point is crucial: it means that, even if an individual is *introduced* through an existence claim, subsequent mention of that individual needn't be bound by the existential quantifier in that introductory claim.

The idea that S introduces some giant existence claim has another counter-intuitive consequence. Consider the following passage:

> (i) A big dog was chasing a little dog. (ii) The big dog bit the little dog. (iii) The little dog howled in pain.

Let N be some discourse that consists *only* of that passage; and let N' be some 800 page discourse that *begins* — but (of course) does not end — with that passage. If the proposal in question were correct, N and N' would have *no well formed part in common* in virtue of their both containing (i)–(iii). For, when it occurs in N', (i)–(iii) corresponds to some ill-formed fragment of an existence claim. (*For some x, y [x is a big dog and y is a little dog and x is chasing y and x bit y and y howled in pain* — this is ill-formed: there is no right bracket.) By contrast, as it occurs in N, that passage (that same verbiage) *does* correspond to a well-formed proposition. So these exact same words would semantically encode different propositions in N and N'. For in N these words *would* have encoded a complete propositions; whereas in N' these same words would not *by themselves* encode a complete proposition. So it follows (vacuously) that these words would encode *different* propositions according as they occur in N or N'. But that is hard to believe; and, concomitantly, it is hard to believe that these words encode *no* proposition (in isolation) when they were to occur in N'.

4. A positive analysis of the referential-attributive distinction

Now, at last, I will set forth a positive analysis of the referential/attributive distinction. Once again let's consider S (remember that S is the beginning of an 800 page discourse):

> (i) Two men were walking down the street, a tall one and a short one. (ii) The short man wanted to play basketball. (iii) The tall man wanted to play golf. (iv) The short man decided to play golf with the tall one.

We've seen reason to reject any kind of a Russellian treatment of this passage. We want to view (iv) as semantically encoding a proposition of the form *a decided to play golf with b*.

One point is crucial, for I think that, in it, lies what might be a correct analysis of the referential/attributive distinction. One cannot *understand* (ii) unless one has taken in (i). Why is this? Because (i) is what assigns a referent to the definite description in (ii): so one doesn't know what the definite description in (ii) refers to unless one has taken in (i). So *by itself* (ii) doesn't get across a complete thought; for one cannot know who the definite description in (ii) refers to unless one has taken in (i): one cannot know *to whom* the property of

wanting to play golf is being attributed. So it is not (ii) by itself that gets across a complete thought, but only (i) and (ii) in conjunction with each other. Now the thought jointly conveyed by (i) and (ii) is a kind of existential generalization (or at least is close to one in form). That proposition is *a tall man and a short man were walking down the street and the tall man wanted to play golf* — and that, it would seem, is an existence claim.

For reasons we've seen, (i) and (ii) convey different propositions; they do not convey different parts of a single proposition (of a single existence claim). At the same time, there is no way to understand (ii) except in terms of (i). Further, taken jointly, (i) and (ii) express an existence claim. So *in understanding* (ii) one is taking in an *existence claim.* At the same time, the proposition *semantically encoded* in (ii) is not an existence claim; for that proposition is of the form *so and so* wanted to play golf. So here we have a case where, in understanding a sentence that semantically encodes a proposition that is *not* an existential generalization, one is nonetheless taking in an existential generalization — a case where, given a sentence conveying a proposition of the form *so and so has phi*, that sentence is *conveying* an existence claim.

Let us sum up this leg of the argument. For reasons we've seen, (ii) semantically encodes a proposition of the form *so and so has phi.* At the same time, there is no way for one to understand (ii) unless one has taken in (i). (i) and (ii), jointly, get across an existence claim. So in understanding (ii) one is taking in an existence claim. So, in effect, (ii) is *conveying* an existence claim — a proposition whose form differs from its own — even though it semantically encodes a proposition whose form is mirrored by its own (cf. Strawson 1990: 117).

With this point in place, we are, I think, in a position to understand how it is that a D-sentence can convey an existential generalization. Consider the following sentence:

(A) The next person to win the Boston Marathon will be in good shape.

Typically, this sentence would be used to convey some attributive — some existential — proposition, namely:

(B) *Some one person will win the Boston Marathon and that person will be in good shape.*

For reasons we've seen, we don't want the proposition *semantically* encoded in (A) to be some existence claim. We want (B)'s surface structure to mirror its deep structure; and this will not be the case if (B) is the deep structure of (A). So we want (A)'s deep structure to be some proposition of the form *so and*

so will be in good shape. But we also want to account for the fact that (A) gets across an existence claim.

To this end, here is what we might say. We must regard any typical token of (A) (any token of (A) that gets across (B)) as parasitic on some *understood* existence claim, some claim like:

(C) *Some one person will win the Boston Marathon.*

So the relation of (A) to (C) is precisely analogous to the relation of — referring back to passage S above — sentence (ii) to sentence (i) *except* that (C) is merely *understood* whereas (i) is explicitly stated. So (C) may be thought of as *introducing* a certain person — just as (i) introduces a certain person. And then (A) *refers* to that person — just as, in (ii), "the tall man" refers to someone introduced by (i). So the proposition semantically encoded in (A) is of the form *so and so will be in good shape*; the definite description in (A) is, in effect, functioning referentially. But it is (C) that assigns a referent to this definite description. (So, in this case, the existence claim has a function in some ways analogous to that of pointing;[18] its purpose is to *pick out* the right individual. If I point to someone and say "that man is a true philosopher", I am picking out a certain individual. It seems that (C) is serving the purpose of *picking* out the individual concerning whom (A) is making an assertion.)

So there is no way to understand (A) except in terms of (C). For one needs to take in (C) to know what the definite description in (A) is referring to: one cannot understand (A) except as being parasitic, so to speak, on (C). Operating jointly, (C) and (A) express an existence claim. So, in understanding (A), one is taking in an existence claim: in effect, (A) *conveys* an existence claim. At the same time, (A) semantically encodes a proposition that is *not* an existence claim — semantically encodes a proposition of the form *so and so will be in good shape.* So (A) *conveys* an existence claim while semantically encoding a referential proposition (or at least a proposition whose structure is that of a referential proposition).

Another example might clarify what I am proposing. Consider the sentence (Kripke's example):

(*) The man who murdered Smith is insane.

This sentence could quite naturally be used to convey either a referential proposition or an existence claim. For reasons discussed, we want the proposition semantically encoded in (*) to have the form of a referential proposition (to have the form *so and so is insane*) *regardless* of whether (*) is being used to

convey a referential proposition or an existential generalization. But we must reconcile this with the fact that it is conveying an existential generalization. To this end, here is what we must say. First of all (*), under all circumstances, semantically encodes a proposition whose basic form is *so and so is insane*. But in some cases — when it is operating attributively — it is parasitic on some *understood* existence claim, namely:

(**) *Some one person murdered Smith.*

When this happens, the definite description in (*) is being assigned its referent by this understood existence claim. Under this circumstance, one cannot understand (*) except in terms of (**). So under this circumstance, to understand (*) one must take in (**). The proposition *jointly* stated by (*) and (**) is an existence claim. So in understanding (*) one is taking in an existence claim (or, at any rate, some proposition very much like an existence claim). At the same time, the proposition semantically encoded in (*) has the form of a referential proposition.

When does a D-sentence convey an existential generalization and when does it convey a referential proposition? Consider the sentence:

(#) The president of the U.S.A. in 2345 will have a lot to deal with.

Why would this sentence almost invariably convey an existence claim? The answer seems to be roughly this. Anyone hearing this sentence would know that the *speaker* didn't have any idea to whom the definite description was referring. As auditors, we tend to be charitable; we try to read plausible meanings into the utterances of others. *A fortiori* we usually try to assign *some* meaning to other people's utterances: it is only in extreme cases that we assume that others are just mouthing words. Oftentimes this is accomplished by imputing *unstated* meanings to the speaker. In keeping with this, anyone hearing a token of (#) will try to assign it a plausible meaning. Given that no one knows who the U.S. president in 2345 will be — and given that pretty much everyone knows that nobody knows this — anyone hearing (#) will not be assigning *any* meaning (or at least no relevant or plausible meaning) to this utterance if he assumes that it is operating *referentially*, i.e., if he assumes it is being used to get across a proposition of the form *so and so will have a lot to deal with*. So anyone hearing (#) will impute some *unstated* meaning to the speaker such that, given this unstated meaning, (#) makes sense; and a candidate for that unstated meaning would surely be:

(##) Some one person will be president of the U.S.A. in 2345.

In fact, arguably, in order for (#) to convey *any* proposition, it must be understood as being parasitic on some claim like (##).

Again consider the sentence:

(*) The man who murdered Smith is insane.

If there is someone who is widely believed to have murdered Smith, and if the speaker can best be interpreted as having this particular person in mind, then (*) will be functioning referentially. If there is no such person — or, in any case, if the best way to make sense of the speaker is *not* to see him as referring to that person — then people will take his utterance of (*) to be parasitic on an understood existence claim (namely, *some one man murdered Smith*); and, therefore, this utterance will be operating attributively.

No doubt the following counter-objection will be made:

> *Are you suggesting that anaphoric incomplete descriptions are always used referentially? This would seem to be plainly not so. Suppose I say: "Here is a purely hypothetical scenario [the sentences are numbered for easy reference]: (i) A man and a woman went to Vienna one Summer. (ii) The woman really wanted to go to Oslo instead. (iii) But the man insisted on Vienna. (iv) The woman felt as though her husband was trying to estrange her from her parents, who lived in Oslo". Let us refer to this passage as Z. In Z the definite descriptions plainly can't be referential, since the whole anecdote is purely hypothetical — so here there is no man or woman to speak of, to refer to. And yet you would have us referring to them: you would say that (i) "introduces" two people; and that the definite descriptions in (ii)–(iv) refer to these people.*

I grant that the definite descriptions in Z are not referring to anybody. But that has no real bearing against my position. Proper names are paradigm cases of referring expressions. But "Hamlet" doesn't *refer* to anybody, and neither does "Luke Skywalker" or "Malvolio". Nevertheless, these are all names — they are *empty* names. The fact that there can be empty names in a work of fiction doesn't mean that names are not to be analyzed as referring expressions. Just as there can be empty names in a work of fiction, so there can be empty definite descriptions; and by reasoning analogous to that just given, this in no way indicates that definite descriptions are not to be analyzed referentially.

Also, there is a compelling reason to maintain that Z's logical form is *not* given by some large existence claim, and that Z's logical form *does* consist of discrete (though semantically interconnected) propositions. Z could be used in different ways. It could be used to tell a *true* story, a *fictional* story, or to run a hypothetical scenario of some kind. It could be used to tell a *whole* story

(or, what may or may not be different, to run a *whole* hypothetical); or just to *initiate* some larger story (of any one of the three kinds just mentioned); or it could form a middle part of some story. We've seen reasons to hold that, if Z (or some homonymous series of sentences) were used to state a *factual* story, then its logical form would *not* be given by some big existence claim. Now we probably don't want to say that the syntactical structure of the logical form of Z changes depending on whether Z is being used to state fact or fiction (or hypothesis). That would be very odd. Just to make this point clear: consider the sentence

(i) The longest street in Riverdale is Main Street.

Sentence (i) could be used in fact or fiction (or, what may or may not be different from fiction, in a hypothetical scenario). We don't want to say that, if (i) is used to initiate a work of fiction it has the logical form "there is some street x in Riverdale such that for any street y in Riverdale, x is longer than y just in case x is distinct from y, and x is Main Street", while if (i) is used to state fact it has the form *a has phi.* That would be peculiar. For if it were the case that the *syntactical* analysis of a sentence depended on whether it occurred in fiction or non-fiction, then if we uncovered some ancient parchment, we could presumably tell whether it stated myth or fact by doing a syntactical analysis of it.

Of course, it is, it practically goes without saying, highly reasonable to suppose that the *semantics* of (i) depends on whether it is being used to state fact, fiction, or hypothesis. But it is radically less plausible to suppose that the *syntax* of the underlying proposition is thus dependent. (That would be to make all D-sentences radically ambiguous.) Given what we saw in connection with passage S, there is some reason to suppose that if Z were used to state fact, its logical form would be given by a series of *distinct* propositions — that, in effect, the definite descriptions would not be bound by an existential quantifier (implicit) in (i). So, given this, it is not unreasonable to hold that the logical form of (ii), for example, is probably some proposition of the form *a wanted to go to Oslo instead.* So given a (token of) a D-sentence that is operating attributively, like

(#) The president of the U.S.A. in 2345 will have a lot to deal with,

there is no objection to representing the proposition thereby conveyed thus:

(#") (a) Some one person will be president of the U.S.A. in 2345. And (b) the person who will be the president of the U.S.A. in 2345 will have a lot to deal with

with the qualifications, first, that (a) and (b) are distinct (though semantically connected) propositions, and, second, that it is (b) alone that gives the proposition semantically encoded in (#).

There is one last objection to consider:

> *You say that the semantics of "the guy who murdered Smith is insane" is always some proposition of the form* alpha has psi*, but that, when that sentence is used attributively, it communicates a proposition of the form* exactly one object has phi, and that object has psi *(*exactly one person murdered Smith, and that person is insane*). So the semantics of "the phi has psi" — whether or not that sentence is used attributively — corresponds only to the second conjunct in* exactly one object alpha has phi, and alpha has psi. *And the semantics of "the guy who murdered Smith is insane" — whether used attributively or not — corresponds only to the second conjunct in* exactly one person murdered Smith, and that person is insane.
>
> *But we know from Frege and Russell that the logical form of (e.g.)* exactly one person murdered Smith, and that person is insane is: (*) *for some x, given any y=x iff y murdered Smith, and x is insane. (In other words, the logical form of the proposition encoded in "exactly one person murdered Smith, and that person is insane" is (*).) Your theory requires that* exactly one person murdered Smith, and that person is insane *be neatly separable into two conjuncts: one being* exactly one person murdered Smith, *the other being* that person is insane. *But that proposition is really an existence claim, and thus isn't separable into two conjuncts (the main connective is the existential quantifier, not a conjunction). There is nothing well-formed in (*) corresponding to the second conjunct of* exactly one person murdered Smith, and that person is insane. *So your theory fails to associate any well-formed semantic entity, and thus any proposition, with "the guy who murdered Smith is insane." But obviously that sentence* does *have a proposition for its meaning.*

This is a significant criticism, and there is a lot to say in response to this. I can only touch on a couple of key points.

First of all, when it is said that such and such is "the logical form" of some sentence, that absolutely does not mean that such and such is the *semantics* of that sentence. To give the logical form of a sentence is to give a perspicuous representation of its inferential structure. So "the logical form" of "that is a sphere" is given by any statement that gives a perspicuous representation of the inferential structure of the underlying proposition. So the logical form of that sentence is given by (i) *that is an object O such that, for some distance D, and some point P, any point on O's surface is D from P* and also by (ii) *that is a closed three-dimensional figure of uniform curvature*. In fact, there are as many different representations of the logical form of "that is a sphere" as there are

mathematical analyses of the concept of a sphere: infinitely many. But it is fairly clear that neither (i) nor (ii) gives the *semantics* of "that is a sphere". To deny this would be ludicrous: it would make a mathematical prodigy of any person who understood "that is a sphere". Also, it would make "that is a sphere" be ambiguous as many times over as there are mathematical analyses of the concept "sphere": so it would make that sentence be infinitely ambiguous. But it isn't ambiguous at all.

So to say that (*) gives "the logical form" of "exactly one person murdered Smith, and that person is insane" is to say only that (*) is a perspicuous representation of the truth-conditions of that sentence, and is *not* to say that it gives the semantics of that sentence. This is the first step in our response to the objector.

The next point is one similar to one Kripke made (cited earlier). Many pronouns and pronominal expressions can occur *either* referentially (as in "he [pointing at John] is a true friend") or anaphorically (as in "if a person gives you the shirt off his back, he is a true friend"). Some philosophers and linguists have maintained that such pronouns are semantically *ambiguous* between referential and anaphoric meanings. But if this is really the case, then we'd expect there to be languages that had *one* set of pronouns corresponding to the referential uses of "he", "she", and so forth, and a *different* set of pronouns corresponding to the anaphoric uses of "he", "she" and so forth. But, echoing Kripke, I would be "very surprised" if this were the case. My guess is that in *any* language, the word corresponding to the referential "he" is the same as that corresponding to the anaphoric "he". And intuitively it is hard to believe that "he" has totally different meanings in "he [pointing to John] is a true friend" and "if a person gives you the shirt off his back, he is a true friend". So we would do well to see both those occurrences of "he" as having the same meaning; consequently we would do well to see both those occurrences of "he is a true friend" as having the same semantic profile. Now the sentence "he [pointing at John] is a true friend" obviously has a complete proposition for its semantics; it is a well-formed, self-contained semantic unit. If, as we've just been arguing, "he [pointing at John] is a true friend" has the same semantic structure as the italicized part of "if a person gives you the shirt off his back, *he is a true friend*", then the italicized part is a well-formed semantic unit.

The Frege-Russell view is that the "he" in "if a person gives you the shirt off his back, he is a true friend" is a bound-variable ("for any man x, if x gives you the shirt of his back, x is a true friend"). Now obviously the "he" in "he [pointing at John] is a true-friend" is not a bound variable. So the Frege-Russell view

makes "he" be ambiguous and is thus to be regarded with suspicion. The obvious way to disambiguate "he" is to see "he" as having a referential role in both "he [pointing at John] is a true-friend" *and* in "if a person gives you the shirt off his back, he is a true friend". The only difference is that in the latter case, the "he" refers to somebody merely described, whereas in the other case, it refers to somebody pointed to.[19]

Exactly similar reasoning supports the view that "that person is insane" has the *same* semantic profile in both "exactly one person murdered Smith, and (a) *that person is insane*" and also in (b) "that person [referring to some particular person] is insane". I highly doubt that there is some language L such that L-translation of (a) is different from the L-translation of (b). It is obvious that in "that person [referring to some particular person] is insane", the expression "that person" is not a bound-variable. But according to the Frege-Russell view, "that person" when it occurs in "exactly one person murdered Smith, and that person is insane" is essentially a bound-variable: "for some x, given any y, y=x iff y murdered Smith, and x is insane". So the Frege-Russell view ambiguates "that person" and is therefore suspect. Again, the obvious way to disambiguate it is to see it as operating referentially in both cases, with the qualification that in the one sentence it is referring to somebody merely described in a previous clause ("anaphoric" use),[20] whereas in the other case, it is referring to somebody picked out in some more robust, and possibly ostensive, way ("referential" use). Most linguists, it should be said, analyze anaphora along these lines (e.g., MacCawley 1993: 415–430; Lyons 1977: 657–677). Once it is granted that both uses are, after a fashion, referential, the way is clear for our analysis of the referential-attributive distinction.

4.1 How our analysis relates to Strawson's

Peter Strawson (1990: 213–234) was the first person to contest Russell's analysis of definite descriptions. I'd like to end with a few words on my debt to him, and also on why Strawson's analysis prevails against some neo-Russellian attacks on it.

Russell holds that definite descriptions correspond to *assertions* of existence. Strawson holds that they correspond to *presuppositions* of existence: if you say "the king of France is bald", you are not *asserting*, but *presupposing* that there is a king of France.

What does it mean to say that, in using a definite description, one is *presupposing* existence? An analogy might help here. If I say, pointing to a certain

elephant, "that elephant looks tired", I am not asserting, but taking it for granted that there exists a contextually salient elephant. If there is no such elephant, then my words are abortive — they fail to say anything that can be true or false. There is no denying this: indexical expressions (like "that elephant") simply must be analysed referentially (Kaplan 1989: 512–516). According to Strawson, definite descriptions are similar to indexical expressions; if an utterance of "…the phi…" is to make a true or false statement, then there must be a phi — the existence of a phi is presupposed. (If '…the phi…' *asserted* that there is a phi, then that sentence would be false, not abortive, in the absence of a phi.)

Strawson's analysis corresponds much more closely to our linguistic intuitions than does Russell's. If someone said "there is exactly one king of France, and that person is bald", it would be completely appropriate to respond by saying: "what you've said is false, for there is no king of France". But if someone said "the king of France is bald", it would not be quite appropriate to say either "that is true" or "that is false"; neither response feels quite right. So "the king of France is bald" evidently doesn't meet the conditions needed to make a true or false statement; and the obvious explanation for that failure is to say that "the king of France" is a referring term that doesn't refer to anything.

Russsellians have no choice but to admit that "the king of France is bald" registers as abortive, as opposed to true or false. But they say that this is because what it *implicates* is a statement that presupposes the existence of a king of France: what is *semantically encoded* in that sentence is an existence-claim. Russell is right about the semantics of that sentence; and Strawson is naively making semantic claims on the basis of data that can be explained in terms of what is implicated, as opposed to literally meant (Blackburn 1984: 308–309).

Given some extremely basic facts about implicature, we have seen that this Russellian move has no validity. Conversational implicature can always be cancelled. But the apparent *presupposition* of existence embodied in "the King of France is bald" is not cancelable; it is inseparable from that sentence. Suppose you said "the king of France is bald, but I am not taking it for granted that there is a king of France". That would be nonsense — and it would bear no resemblance to what goes on when you say "there is a gas station around the corner, but you can't get gas there, since it's closed". The same applies to any other attempted cancellation of that presupposition. (You may verify this by trying different ways of canceling the presupposition; the result is never anything that resembles "there's a gas station around the corner, but you can't get gas there, since it's closed" or any other clear case of cancellation.) Clearly the

presupposition is not cancelable. So the kind of implicature Russellians have in mind cannot be conversational.

So if any kind of implicature can save Russell's theory from Strawson's attack, it must be conventional implicature. But conventional implicature is not up to this task. What conventional implicature adds is always truth-conditionally inert — Frege called it "coloring" or "tone" — and couldn't possibly account for the data that Strawson adduces. Adding mere "coloring" or "tone" to a sentence S will not make it inappropriate to say "that is false" in response to S unless such a reaction was already appropriate. It is appropriate to say "that is false" in response to "shysters make more money than investment bankers" only in so far as it is appropriate to say "that is false" in response to "lawyers make more money than investment bankers". So it could not be *conventional* implicature that makes it inappropriate to say "that is false" in response to "the king of France is bald", given that it *is* appropriate to say "that is false" in response to "there is exactly one king of France, and that person is bald".

So neither conventional nor conversational implicature account for the data Strawson has in mind. So *no* kind of implicature is up to the task. Thus semantics must be what is responsible. Strawson was right all along.

My analysis obviously owes a great deal to Strawson's analysis. On my analysis, when "…the phi…." is operating attributively, that is because "the phi" refers to someone described in an understood, and thus unasserted, existence claim: so existence is *presupposed* — just as it is on Strawson's account. And, of course, when "…the phi…" is operating referentially, existence is presupposed, for one can refer only to an object whose existence is presupposed.

But there is one matter that Strawson is not clear about. Existence can be presupposed in two very different ways. Suppose you and I see a red elephant next to a tree. In that case, we have *de re* awareness of the red elephant. Under that circumstance "…the red elephant next the tree…" does indeed presuppose the existence of an elephant, but this has nothing to do with some understood existence claim. The definite description doesn't refer to an object *described* in any way; *a fortiori* it doesn't refer to an object described in some implicit existence claim. The object referred to is given to us through sense perception, not through an existence claim, or any other kind of proposition.

On the other hand, in an attributive use of a sentence of the form "…the guy who murdered Smith…", the expression "the guy who murdered Smith" refers to somebody who is given to us through an existence claim — not through sense perception or through any other modality that sustains *de re* awareness of any kind.

Of course, any case of reference involves a presupposition of existence: you can only refer to something whose existence you take for granted. But the presupposition of existence can be associated *either* with a *proposition* of some kind (an existence claim) *or* with something of an entirely different nature (e.g., sense perception). Strawson doesn't register this important fact. Once this fact about presupposition is made clear, the way is clear to saving Strawson's analysis from various attacks to which it would otherwise be vulnerable.

We've already seen one case of this. In an attributive use of "…the guy who murdered Smith…", the definite description *seems* not to be operating referentially; so Strawson's referentialist analysis is threatened. But if we keep in mind that something merely described can be referred to, then we can say that the definite description in an attributive use of "…the guy who murdered Smith…" refers to somebody described in an implicit existence claim. And Strawson's analysis is no longer threatened.

In this way we can neutralize another attack on Strawson and, more generally, on the view that definite descriptions are referring expressions. When you say (*) "the first president of the United States might have been a woman", you are not saying that George Washington might have been a woman. At first, this fact seems to bear *against* a referentialist account of definite descriptions.[21] But when it is remembered that an object merely *described* can be referred to, this fact ceases to be much of a threat to Strawson's position.

Before we can see how exactly this works, we must note a fact about the term "might". "Such and such might be the case" can be paraphrased as: "it is possible that: such and such is the case". So (*) can be paraphrased as "it is possible that: the first president of the United States was a woman". Of course, "it is possible that" is a modal operator; and like any operator, it can take "wide-scope" or "narrow-scope" when it is associated with a molecular proposition.

In the case of (*) the understood statement would be: *there was a person who was a unique president of the United States before anyone else.* And the modal operator — the "it is possible that" — is given wide-scope. (We will see in a moment why it is given wide-scope.) The result is that, when you say "the first president of the United States might have been a woman", what is communicated is: *it is possible that*: *there was a person who was uniquely president of the United States before anyone else, and that person was a woman.* So the definite description in (*) refers to the person *described* in the understood existence claim, and *not* to the actual first president of the United States. And that is why (*) does not typically convey the proposition that *George Washington* might have been a woman. So all the facts fall into place when we realize that

presuppositions can be associated with *statements* of a certain kind: we get to hold onto the well-supported idea that definite descriptions are referring expressions, and *also* to the fact that (*) does not say that George Washington might have been a woman.

The question arises: why does the modal operator automatically get wide-scope in this case? The reason is simple: when we are confronted with a sentence that can be taken to convey *either* something reasonable *or* something unreasonable, we take it to mean something reasonable. If the "it is possible that" were given narrow scope, the claim communicated by (*) would not be reasonable. (The claim would then be: *someone was uniquely a president of the United States before anyone else, and it is possible that: that person was a woman.* That claim is not a particularly reasonable one.) If the "it is possible that" is given wide-scope, then (*) communicates a reasonable proposition (a true one, in fact); if that operator is given narrow scope, then (*) communicates a far less reasonable proposition. Thus we give it wide-scope.

5. Conclusion

D-sentences are used to communicate both referential and attributive propositions. But D-sentences are not semantically ambiguous between referential and attributive meanings. So "the T.V. is broken" *either* has some proposition of the form *a has phi* for its meaning *or* it has some existence claim for its meaning: it doesn't have both. Russell's theory says it has an existence claim for its meaning. But Russell's theory is wildly discrepant with what "the T.V. is broken" is actually used to say. And this discrepancy almost certainly has a *semantic* basis. Attempts to put the blame on implicature simply aren't consistent with some basic facts — for example, that conversational implicatures are cancelable and that conventional implicature reinforces, and does not obscure, literal meaning. If we assume that definite descriptions are referring expressions, we can naturally account for the attributive uses of D-sentences. When "…the phi…" operates attributively, that is because there is an understood, but unstated, antecedent existence claim; and "the phi" refers to the entity described in that claim.

Notes

1. The F-subscript stands for "formal".

2. The page references relating to Donnellan correspond to Martinich 1990.

3. This is Howard Wettstein's view. See Wettstein (1981).

4. I owe this point to Tim Stowell, professor of Linguistics, UCLA. I have verified it with Susanna Cumming, Professor of Linguistics UCSB. Professor John Du Bois (UCSB Linguistics Department) said that he knows of no language that has one word for the referential "the" and another for the attributive "the". Professor Stowell did point out that some languages don't have definite articles. But he said all those that do have such articles have *one* word corresponding to both the attributive "the" and the referential "the". (I am treating different inflections of e.g. German 'der' as one word — counting 'der', 'den', 'dem' as different versions of one word, and not as altogether different words. This is standard practice.)

5. In this paper I will make heavy use of the term "contextualism" (and also of "contextualism$_o$"). I must make it clear that I will *not* be using this term in the standard way. Ordinarily, the term "contextualism", at least in the philosophy of language, refers to the position that it is not sentences *per se* that have truth-conditions but *utterances* of sentences (see, for example, Recanati 2003). According to that doctrine, what has meaning is not the *type* "snow is white", but particular tokens (utterances or inscriptions) of that sentence. I will be using the term "contextualism" to refer to a completely different doctrine, namely that, when contextual factors are taken into account, D-sentences behave in the way that Russell's theory predicts that they behave, despite first appearances to the contrary. I will use the term "contextualism$_o$" to denote a similar doctrine.

6. See Grice (1989); Kripke (1977); Neale (1990: Chapter 3).

7. Kent Bach does not take this line. We will deal with Bach's view at the end of Section 2.

8. See MacCawley (1993: 212–213); Neale (1990: Chapter 3).

9. Grice (1975) is the *fons et origo* of the basic ideas of modern pragmatics and, in particular, of the concepts of conversational and conventional implicature.

10. See Mey (1993: 99–103); Grice (1975); Neale (1990; Chapter 3).

11. It seems that statements that necessarily have the same truth-conditions needn't be synonymous. There are many different — actually, *infinitely* many non-trivially different — analyses of a statement like "alpha is a sphere". To give but two: it can be analyzed as either "alpha is such that, for some distance E, for some point P1, given any two points P2 and P3 such that P2 and P3 are on alpha's surface, P2 and P3 are E from P1". Or it can be analyzed as: "alpha is a closed three-dimension figure of uniform curvature". These two statements, though necessarily equivalent — indeed, though they can be known *a priori* to be necessarily equivalent, unlike "that is water" and "that is H_2O" — are presumably not *synonymous*; for they consist of different concepts and are structurally very different.

12. They differ in 'tone" or "coloring", but not in literal meaning. See Dummett (1973; Chapter 1).

13. By 'orthodox Russellianism' I mean Russellianism as it was prior to being modified or supplemented by neo-Russellian doctrines like contextualism.

14. Just to be clear on this point: He will say that the "there is", and any other quantifiers in the sentence in question, will range over an appropriately restricted class of objects — that they will range over e.g. the objects that are right next to a certain sofa — so that what is said is not that there is in all of existence one red-haired man… but only that there is exactly one red-haired man next to some sofa. And he will say that this restriction is due to context, not the presence of extra verbiage.

15. The claim that conversational and conventional implicature are the only two kinds may seem a bit bold; but it is really a tautology. Trivially, whether an implicature attaches to a sentence either is, or is not, contingent on the conversational context. If it is, then that implicature is, by definition, 'conversational'. If it is not, then it is, by definition, 'conventional'.

16. Actually, it isn't clear whether, according to Russell's theory, the verb here should be "is" or "was". This in itself poses a major problem for Russell's theory. Every Russellian paraphrase of a D-sentence begins with a "there is" (or some synonym). For the theory to have any plausibility at all, the tense of the verb "to be" in the "there is" must change according to context. According to Russell's theory, the literal meaning of (i) "the first creature to live on dry land died millions of years ago" is given by (ii) "there is exactly one creature x such that x uniquely lived on dry land before any other creature y, and x died million years of ago". But obviously the "is" at the beginning of (ii) should be a "was": otherwise we are simply given the wrong meaning. In general, the Russellian must say that the tense of the "there is" must vary with circumstances. But this move creates other, probably insuperable, problems for the Russellian. Suppose someone says "the founder of the World Bank wrote over forty books between 1970 and 1980". Is the correct paraphrase "there *is* some x such that x uniquely founded the World Bank…" or is it "there *was* some x such that…". If the proposal just considered — the tense on the "there is" varies with context — is correct, then *we don't know how to interpret that sentence until we know the discursive context.* But that is absurd. Obviously you understood that sentence perfectly well without having any idea whether in my view, or in actuality, the founder of the World Bank is still in existence. Also, if that proposal were correct, then one couldn't legitimately utter a D-sentence unless one had knowledge of the relevant historical facts. I myself do not know whether the founder of the World Bank is still alive. If we countenance the 'tense-contextualism' just described — and we must, if Russell's theory is to have any hope — then I cannot legitimately say "the founder of the World Bank wrote over forty books between 1970 and 1980" unless I know whether that individual is still in existence. But I *can* legitimately token that sentence without having that historical knowledge.

17. Or, in any case, that discourse comprises an existence claim that is at least as large as the distance is great between the outermost co-referential definite descriptions in that discourse.

18. See Kaplan (1990: 323): *"If pointing can be taken as a form of describing, why not take describing as a form of pointing?"* (Kaplan's italics.) So (C) picks out an individual — points to him, as it were — by describing him; and the definite description in (A) refers to that individual.

19. This analysis of anaphora is, of course, controversial; and I obviously have only given a modicum of support for it here. A great deal of support for it is found in Evans' (1980) landmark paper. (It must be said that Evans' view is not quite coincident with the one we have been describing.) Evans' insights have been developed and substantiated by Kent Bach (1987: Chapter 11). Linguists tend to be sympathetic to this analysis of anaphora. See, for example, MacCawley (1993: 415–430); and also Lyons' (1977) discussion of anaphora.

20. Evans (1977: 224) writes: "Pronouns are often used as referring expressions, and it is not particularly surprising that some of them should have their reference fixed by a description recoverable from the antecedent, quantifier-containing, clause".

21. In fact, this is one of Gareth Evans' arguments in favor of Russell's analysis: Evans (1982: Chapter 2, Section 4).

References

Bach, K. 1987. *Thought and Reference.* Oxford: Clarendon Press.

Blackburn, S. 1984. *Spreading the Word: Groundings in the Philosophy of Language.* Oxford: Clarendon Press.

Donnellan, K. 1966. "Reference and definite descriptions". In Martinich (ed), 235–247.

Dummett, M. 1981. *Frege: Philosophy of Language*, 2nd ed. Cambridge, MA: Harvard University Press.

Evans, G. 1980. "Pronouns". *Linguistic Inquiry* 11: 337–362.

Evans, G. 1982. *The Varieties of Reference.* Oxford: Clarendon Press.

Garfield, J. L. and Kitely, M. (eds). 1993. *Meaning and Truth.* New York: Paragon House.

Grice, P. 1975. "Logic and conversation". In Martinich (ed), 149–160.

Grice, P. 1989. "Presupposition and implicature". In P. Grice, *Studies in the Ways of Words.* Cambridge, MA: Harvard University Press, 269–282.

Kaplan, D. 1989. "Demonstratives". In J. Almog, H. Wettstein, and J. Perry (eds), *Themes From Kaplan.* Oxford: Oxford University Press, 481–563.

Kaplan, D. "Dthat" 1990. In Martinich (ed), 316–329.

Kripke, S. 1977. "Speaker's reference and semantic reference". In Garfield and Kitely (eds), 162–188.

Lyons, J. 1977. *Semantics.* Cambridge: Cambridge University Press.

MacCawley, J. D. 1993. *Everything that Linguists Have Always Wanted to Know about Logic*, 2nd ed. Chicago: The University of Chicago Press.

Martinich, A. P. (ed). 1990. *The Philosophy of Language*, 2nd ed. Oxford: Oxford University Press.

Mey, J. 1993. *Pragmatics: An Introduction.* Oxford: Blackwell.

Neale, S. 1990. *Descriptions.* Cambridge, MA: The MIT Press.
Recanati, F. 2003. *Literal Meaning*. Cambridge: Cambridge University Press.
Russell, B. 1993. "On denoting". In Garfield and Kitely (eds), 87–99.
Salmon, N. 1986. *Frege's Puzzle*. Cambridge, MA: The MIT Press.
Strawson, P. 1990. "On Referring". In Martinich (ed), 219–234.
Wettstein, H. 1981. "Demonstrative reference and definite descriptions". *Philosophical Studies* 40: 241–257.

Author's address

John-Michael Kuczynski
2665 South Bayshore Drive, Suite 1101
Coconut Grove Florida
33133 Florida
U.S.A.

Email: jsbach@jps.net

About the author

John-Michael Kuczynski is a PhD candidate in philosophy at the University of California, Santa Barbara. He has published a book advocating an anti-realist view of sense-perception (*Elements of Virtualism*, 2003), and has several articles forthcoming, or already published, in the *Journal for Theoretical and Philosophical Psychology, Cognition and Communication, Kriterion*, *Teorema*, and other journals.