The Moral Structure of Legal Obligation

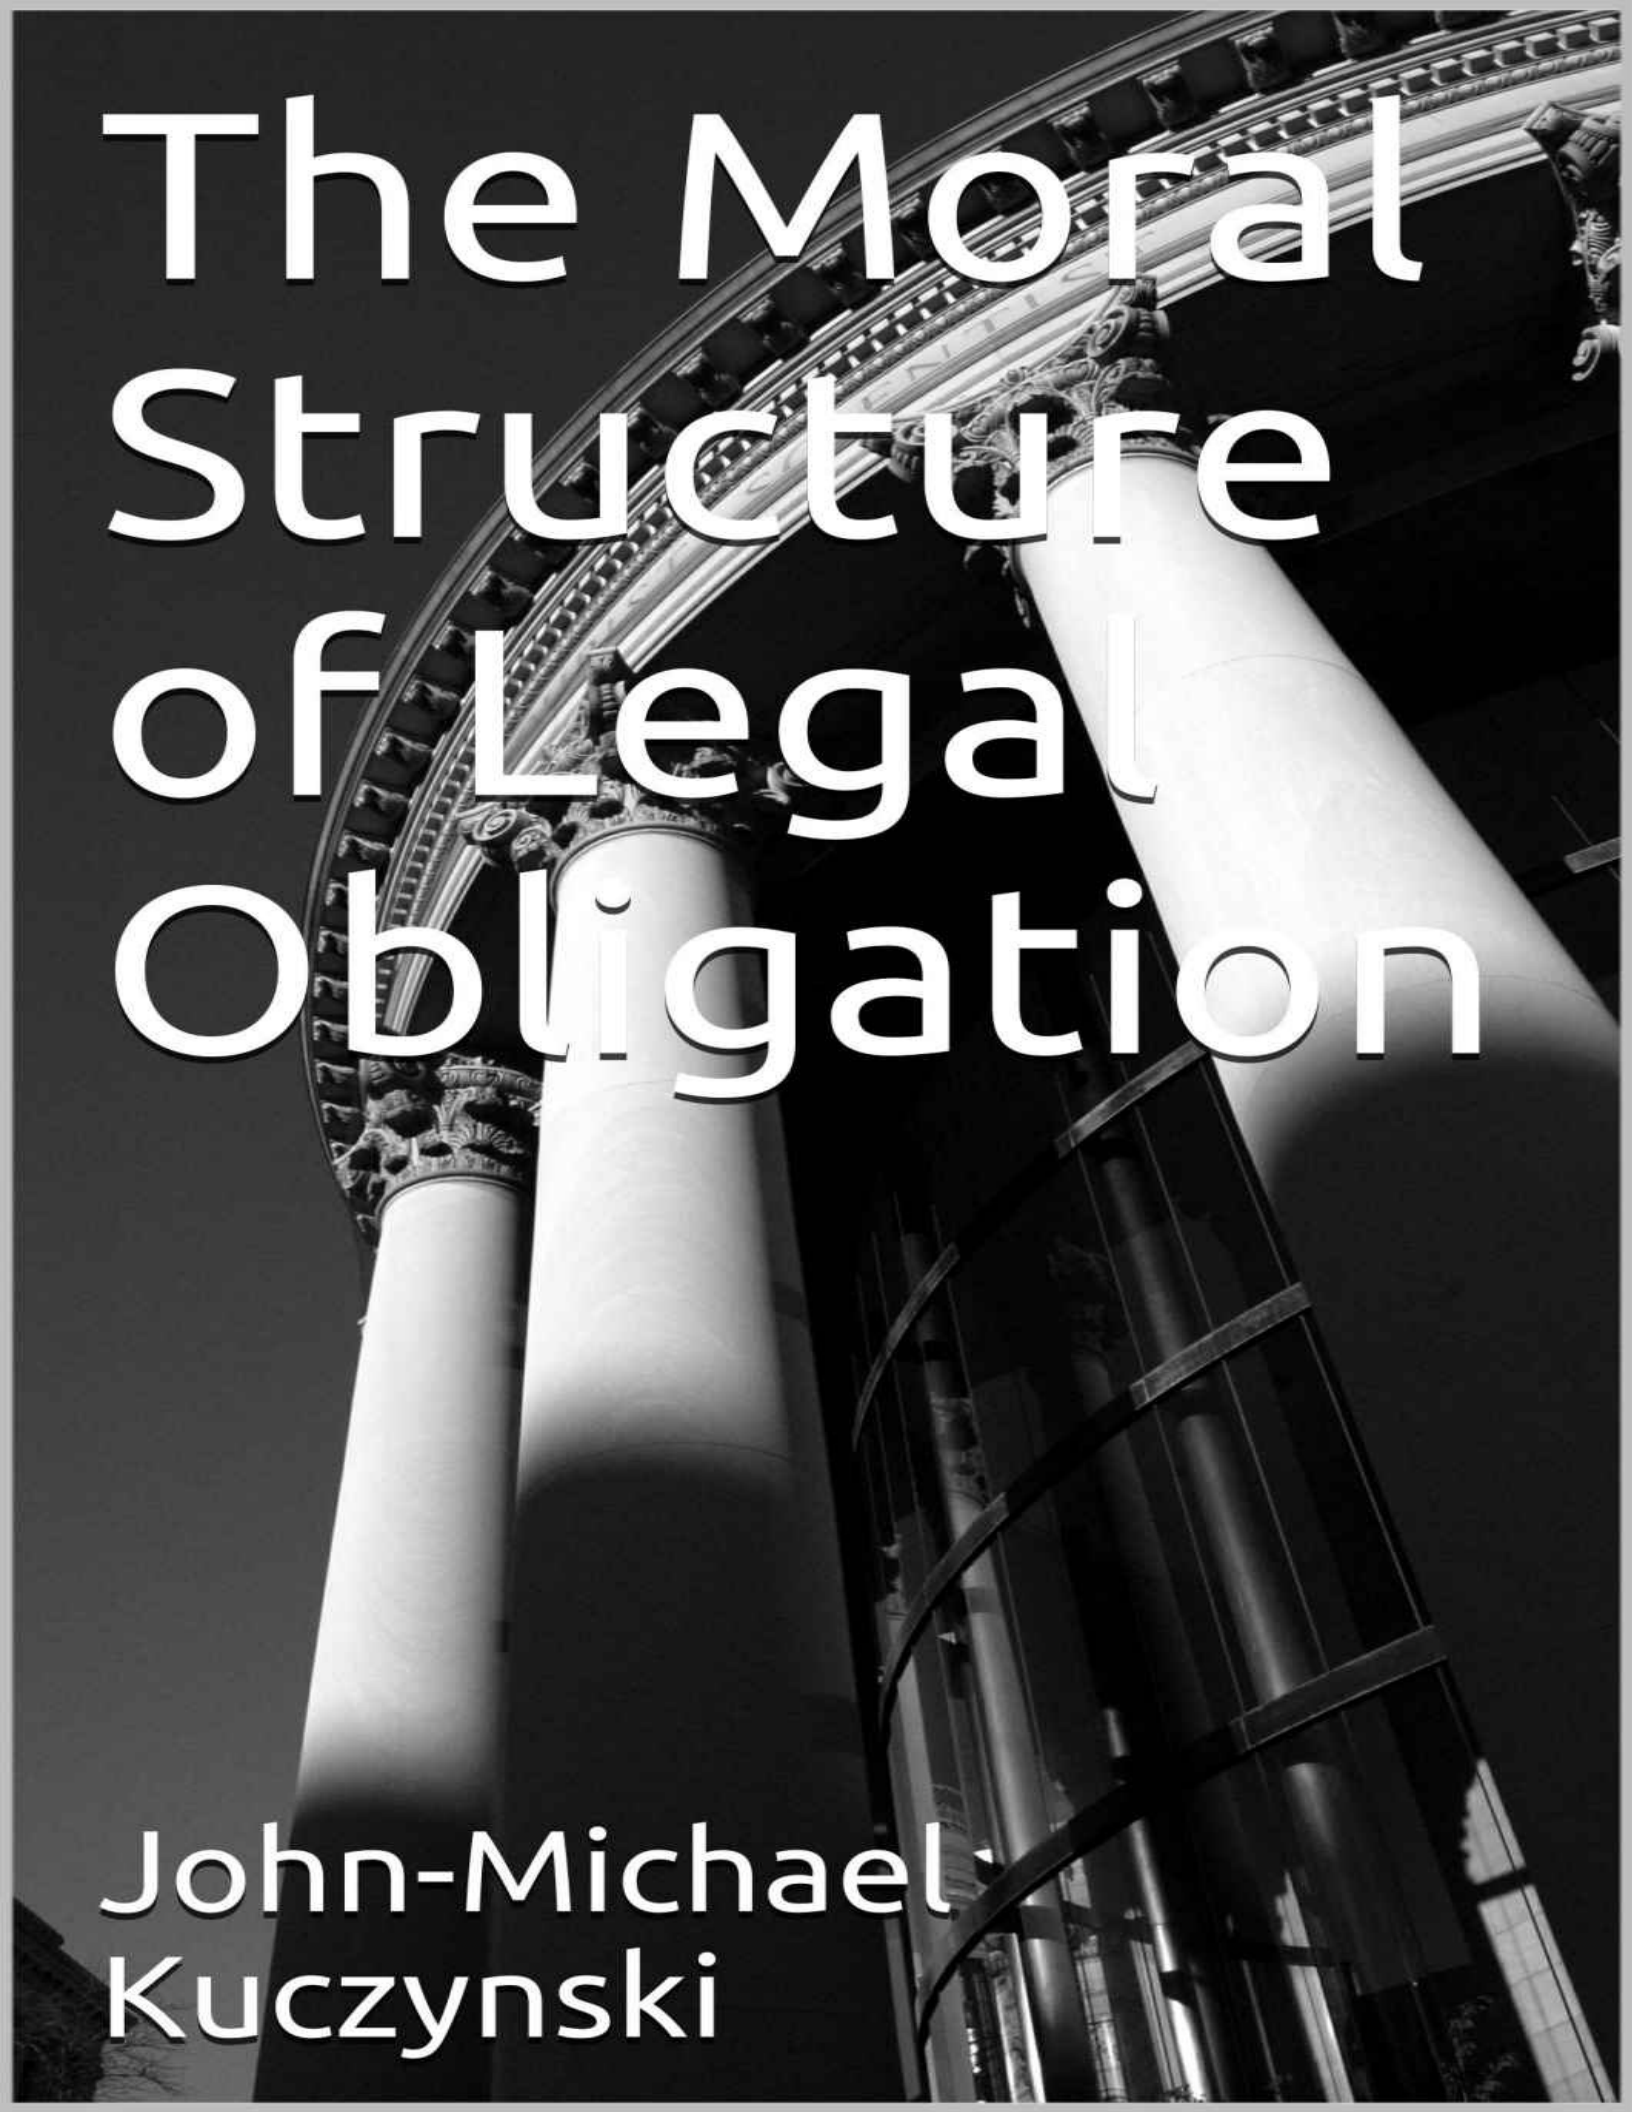

John-Michael Kuczynski

The Moral Structure of Legal Obligation

by

John-Michael Kuczynski, PhD

Fellow, Bard College

TABLE OF CONTENTS

Preface

What are laws, and do they necessarily have any basis in morality? The present work argues that laws are governmental assurances of protections of rights and that, consequently, the concepts of law and legal obligation must be understood in moral terms. There are, of course, many immoral laws. But once certain basic truths are taken into account – in particular, that moral principles have a "dimension of weight", to use an expression of Ronald Dworkin's, and also that principled relations are not always expressed by perfect statistical concomitances – the existence of iniquitous laws poses no significant threat to a moralistic analysis of law. Special attention is paid to the debate between Ronald Dworkin and H.L.A. Hart. Dworkin's over-all position is argued to be correct, but issue is taken with his argument for that position. Hart's analysis is found to be vitiated by an impoverished conception of morality and also of the nature of government.

Our analysis of law enables us to answer three questions that, at this juncture of history, are of special importance: Are there international laws? If not, could such laws exist? And if they could exist, would their existence *necessarily* be desirable? The answers to these questions are, respectively: "no", "yes", and "no."

Our analysis of law enables us to hold onto the presumption that so-called legal interpretation is a principled endeavor, and that some legal interpretations are truer to existing laws than others. At the same time, it accommodates the obvious fact that the sense in which a physicist interprets meter-readings, or in which a physician interprets a patient's symptoms, is different from the sense in which judges interpret the law. So our analysis of law enables us to avoid the extreme views that have thus far dominated

debates concerning the nature of legal interpretation. On the one hand, it becomes possible to avoid the cynical view (held by the so-called "legal realists") that legal interpretation is *mere* legislation and that no legal interpretation is more correct than any other. On the other hand, it becomes possible to avoid Blackstone's view (rightly descried by Austin as a "childish fiction") that judges merely discover, and do not create, the law.

Chapter 1: Summary of the present work

If Smith and I were stranded on a desert island, and we thus both fell outside the scope of any legal system, I would have a moral, but not a legal, obligation to refrain from punching him. Therefore, my *actual* legal obligation not to punch fellow citizen Smith (in other words, the obligation that I have to refrain from punching him, given that he and I both fall within the jurisdiction of some one legal system) is distinct from my moral obligation not to do so. But – so I argue in the present work -- my legal obligation not to punch Smith is *a* moral obligation: it is a moral obligation to comply with a system of rules on which people depend for protections of their most basic rights: a system which builds highways and hospitals, prevents people from killing and stealing from one another, and holds physicians and other professionals accountable to certain ethical standards.

One may have a legal obligation to commit an immoral act. But, in such a case, one's legal obligation is not morally hollow. One's legal obligation to commit an immoral act is identical with one's moral obligation to comply with a system of the kind just described. Discharging that obligation *involves* committing an immoral act, but that obligation is not itself morally negative or morally hollow.

This last point demands elucidation. It is one thing to have an obligation to do X, and it is quite another to have an obligation the fulfillment of which *involves* doing X. A surgeon has an obligation to heal his patient. Fulfilling that obligation *involves* causing his patient to suffer. But a surgeon doesn't have an obligation to cause his patient to suffer. (In fact, he has an obligation to minimize his patient's suffering. A surgeon would be violating his professional obligations if he chose not to use anesthetic when there was no medical reason not do so.) Similarly, fulfilling one's legal obligations may *involve* committing an immoral or morally neutral act. But it doesn't follow that legal obligations are ever immoral or amoral. If I have a legal obligation to do X, that means that I have a moral obligation the fulfillment of which *involves* my doing X. It does *not* mean that I have a moral obligation to do X. That is why a legal obligation to commit an immoral act may be morally wholesome. Hence the following argument is invalid -- "given that one has legal obligations to commit immoral acts, it follows that legal obligations are not moral obligations" -- and once this is seen, one of the more obvious obstacles to accepting a moralistic analysis of law has been removed.

Let us discuss a positive reason to accept such an analysis. A government is a success precisely to the extent that it serves the interests of those whom it governs. The same is therefore true of laws, given that laws are among the instruments of government. It follows that the concept of law is to be understood in moral terms, even though there are immoral laws.

Governments often behave in an unspeakably evil manner towards their own constituents; and oftentimes this misconduct involves, or even coincides with, the issuing of laws. On this basis, many have concluded that the concept of law is *not* to be understood in moral terms. But this line of thought involves a failure to make three important distinctions:

(i) The distinction between individual and institutional frames of evaluation;
(ii) the distinction between functional and purely factual frames of evaluation; and
(iii) the distinction between suspensions of protections of rights, on the one hand, and positive violations of rights, on the other.

The *personal* success of a head of government may not involve his serving the interests of his constituents, and it may even involve his thwarting those interests. But a head of government is a success *as* a head of government of state precisely to the extent that he serves the interests of his constituents. Given this, it follows that the objective of government – of *any* government - is to protect the interests of the governed. Stalin failed *as* a head of government precisely because he did not serve the interests of his constituents. It follows that the objective of the Soviet *government* was to serve the interests of its constituents – even though that may not have been the objective of Stalin or of any of the other senior officials composing that government. If a dictator issues a law that harms his constituents, he fails *as* a law-maker, even though issuing that law may help him achieve *personal* success of some kind – even though, for example, it may help him hold onto power. In general, one fails *as* law-maker if one issues a law that hurts those who are subject to it. This shows that the purpose of a legal system *per se* is to serve the interests of those who are subject to it – even though that may not be the objective of all, or even any, of the individuals who created that legal system. Distinction (i) thus removes one of the major reasons to reject a moralistic analysis of law.

Now let us discuss (ii). Even if everyone alive suddenly succumbed to liver-disease, and each person's liver stopped removing toxins from his blood-stream, there would still be a significant sense in which livers were to

be understood as removers of toxins. If one accepts a strictly fact-based frame of evaluation, one is blind to this obvious fact; and one is *not* blind to it if one accepts a frame of evaluation that allows for functional categories. An analogue of this line of thought holds of law and government. Even if every government were thoroughly corrupt and exploitative, there would still be a significant sense in which governments were to be understood as protectors of rights, and in which the laws issued by governments were to be understood as attempts to provide protections of rights. Like hearts and livers, governments and the laws they issue are to be understood in functional, and not strictly statistical-empirical, terms.

The significance of distinction (iii) is to be understood in terms of the following argument. If a certain legal system doesn't prohibit anyone from violating Smith's rights, then Smith falls outside the jurisdiction of that system. There is legal obligation only where there is moral protection, and the moral core of one's legal obligations lies in this fact.

But there is an obvious objection to this argument. People have legal obligations under governments that harm them. Given this, it seems straightforwardly false to say that legal obligations under a government categorically presuppose receipt of moral protections from that same government. Here is where distinction (iii) becomes relevant. Suppose that Officer Smith beats me without cause. Even while I am being beaten, the government continues to protect my rights: anyone caught attempting to burgle my house or kidnap my children will be arrested; and it may even be other police officers who, acting in the name of duty, stop Smith from further beating me. Thus, Smith's misconduct is a case where a positive violation of my rights is superimposed on the continued existence of governmental protections of those same rights. The same thing is true of *any* case where one has legal obligations under a government that harms one. Thus, even

though one may have legal obligations towards a government that harms one, it doesn't follow that legal obligation doesn't presuppose moral protection. So far as one thinks otherwise, one is failing to make distinction (iii).

Like the present author, Ronald Dworkin argues that the concept of law is to be understood in moral terms. His argument is that judicial decisions are often made on moral, as opposed to narrowly statutory, grounds. This argument is less than probative. If the judge rules that the money is Smith's, then the bank-manager is acting illegally if he refuses to give Smith the key to the safety-deposit box. If the judge rules that the money is *not* Smith's, then the bank-manager is acting illegally if he *does* give Smith the key. The judge may or may not make his decision on the basis of a moral principle. In either case, the law is what the judge says it is. This is the antithesis of the idea that there is any *necessary* relationship between law and morality, even though it is obviously compatible with the idea that the two may sometimes overlap.

We've already seen why law is to be understood in terms of morality, notwithstanding that Dworkin's argument for this conclusion is less than cogent. But if our criticism of Dworkin is correct, it seems to follow the legal interpretation (so-called) is mere legislation, and thus isn't interpretation at all. But legal interpretation is *not* mere legislation, and this is perfectly compatible with the fact that the law is what the judge says it is.

A brief detour through the history of science may help to indicate the broader outlines of my argument for this. The layperson's concept of temperature has no application outside of what are, from a physicist's viewpoint, extremely narrow horizons. When physicists identified heat with molecular motion, they were not identifying the content of the layperson's concept of temperature. Rather, they were replacing the lay-person's concept with one that is consistent with it, within its narrow sphere of applicability,

but is more precisely defined and also has a much wider sphere of application. Such principled extensions of existing concepts are referred to as "delineations" or "precisifications." Any judicial ruling is *ipso facto* an extension, not an identification, of existing law. But judicial rulings are *principled* extensions of the law, i.e. they are delineations of it, and are not cases of law being created out of whole cloth. Just as some delineations of the concept of temperature ("a body's temperature is the sum of the velocities of its constituent molecules divided by the number of those particles") are truer to pre-theoretic data than others ("heat is dephlogistinated matter"), so some interpretations of law are truer to existing law than others.

Blackstone said that judges identify, and do not create, law. Austin described Blackstone's position as a "childish fiction." Given the argument just outlined – and, in fact, independently of that argument -- we must agree with Austin. But given that same argument, we must *dis*agree with the nihilistic view, held by the so-called "legal realists", that judicial interpretation is mere legislation, and isn't answerable to objective standards.

Chapter 2 Basic concepts

Do laws necessarily have some kind of moral basis? Yes – or so the present work will argue. But from Augustine to the present, every attempt to identify a necessary connection between law and morality has been flawed, giving credibility to the position that there is no such connection.

If we are to see the true nature of law and legal obligation, we must first note a fact about morality. Good and bad are not binary concepts. They have "a dimension of weight", to use an expression of Ronald Dworkin's. It would

be *better* to steal penicillin from the pharmacy than it would be to let your child succumb to his life-threatening infection. So given only those two choices, the right thing to do is to steal. But it doesn't follow that stealing from Walgreen's is good in an absolute sense. It is not. Theft and deception are not good, even though under many circumstances they are *less bad* than the alternative.

If we say that, under the circumstances just described, stealing from Walgreen's is good in an absolute sense, rather than merely being *less bad* than the alternative, we are led to say that there is *no* moral value to obeying the law. The law comes to be morally hollow or, at best, to have a highly circumstantial connection with morality.

But a binary conception of morality can have the exact opposite consequence. If we think in terms of absolute right and absolute wrong, we cannot say that a law is *at all* moral unless obedience to it is morally *de rigueur*. So if we concede (as we surely must) that stealing from Walgreen's is not a *wholly* good thing, then we end up embracing the absurd and evil view that saving your child's life is less important than your not committing some misdemeanor. Laws are falsely elevated to the status of moral verities.

So if we think of morality in binary terms, then either laws end up being morally arbitrary or they become unsurpassable excellences to be obeyed at all costs. But if we think of morality in terms of *weight,* then we can see laws for what they are: things that have *some* basis in morality, but that must sometimes be disobeyed if some weightier moral imperative is to be upheld. The idea that laws can have *no* moral value derives from the tendency, just discussed, to think that whatever is *comparatively* good or bad is *absolutely* so.

Another confusion concerning morality underlies a failure to see the moral basis of law. There is a difference between having *moral value*, on the

one hand, and being deserving of *moral credit* or *commendation*, on the other. A rabbit doesn't deserve moral *credit.* It is not in the same category as somebody who works for the Red Cross or Doctor's Without Borders. Being an amoral creature, a rabbit is ineligible for any kind of moral commendation. But a rabbit's life has *moral value.* That is why it would be wrong to torture or needlessly kill a rabbit. Thus, what is good does not always deserve commendation. To have moral value and to be commendable are not the same thing.

In fact, something that has moral value may not only be unworthy of commendation, but may even deserve con*dem*nation. A bit of fiction will remove the appearance of paradox from this last point, and will thus help unveil the moral underpinnings of law. Smith is a talented pianist. He is at a phase of his pianistic development where it is especially important that he practice. The only piano to which Smith has access is in his tiny apartment, which he shares with several other people. Given his work-schedule, the only time Smith can practice is late at night. There is thus no way that Smith can practice without preventing his roommates from sleeping. Moreover, each of his roommates has a peculiar medical condition whereby the sound of piano-playing causes permanent physical and psychological damage. For various reasons, it is not an option for Smith to find other roommates.

Smith doesn't want to harm anyone. But like everyone else, he wants to flourish as a human being. In his case, this involves playing the piano. So Smith *does* practice late at night. In doing so, he experiences ecstasies and sharpenings of awareness that fulfill his inner nature as nothing else could.

In practicing, Smith is flourishing. Human flourishing is surely a good thing[1], and the same is therefore true of Smith's practicing. At the same time, as Smith well knows, his practicing is ruining the lives of his roommates. So in practicing, Smith is doing something condemnable. Thus

Smith's practicing has moral value *and* is condemnable.

Suppose that some law were enacted that permitted Smith to practice under these circumstances. That law would protect a moral good – Smith's flourishing – while also protecting behavior that is positively condemnable.

The moral basis of law is easily overlooked if we identify the good with the commendable and the bad with the condemnable. There are laws that protect moral goods and, in the same stroke, practically guarantee the occurrence of the most despicable and condemnable behaviors. There are laws that, while not licensing positively evil conduct, fail to encourage or protect commendable action. On the basis of these facts, many have inferred that laws can protect things that have no moral *value* and, consequently, that laws can have no connection to morality. But such an inference embodies two muddles. First, it fails to distinguish the concept of having moral value from the concept of being deserving of moral credit. Second, it assumes that anything having moral value cannot be condemnable.

The failure to make these distinctions is yet another expression of the binary moral thinking discussed earlier. If we think in terms of absolute right and wrong, then Smith's practicing is either totally lacking in moral value or it becomes an unimpeachable excellence; and any law that allows Smith to practice is either completely divorced from morality or it becomes a moral verity unto itself. Both of those of views clash with our moral intuitions. But if we think in terms of weight, we can see such a law for what it is: something that is probably wrong and should be abolished, but that also protects *a* moral good (Smith's flourishing).

There is, I believe another confusion – one that has nothing to do with morality *per se* – that, in my judgment, has retarded efforts to see the moral basis of law. Outside of micro-physics[2], principled relations are not typically expressed by strict and invariable concomitances. There might be

F's that are not G's. It could even be that there is only one F that is a G, and a billion F's that are not. But that *by itself* doesn't mean that there isn't a principled relationship between a thing's being an F and its being a G.

There are two reasons for this. First, principled relations are often given by *ceteris paribus* propositions – propositions of the form *all other things being equal.*..If, all other things being equal, F's are G's, then there is a principled relationship between a thing's being an F and its being a G, even though in circumstances where all other things are not equal, an F may not be a G.[3] So *given only* that there are morally empty laws, it doesn't follow that there isn't a principled relationship between a thing's being a law and its protecting a moral good. Given only that a match may not ignite after being struck, it doesn't follow that there isn't a principled relationship between a thing's having the property of being a struck match and its having the property of being on fire.[4]

There is another reason – one that, I think, has a more direct bearing on the concept of law -- why principled relationships don't always correspond to perfect statistical regularities. A brief detour may help. Many planes are defective and therefore don't fly. But there is certainly a principled relation between a thing's being a plane and its being something that flies. When we are dealing with concepts having a *teleological* or *functional* component, principled relationships can fail rather dramatically to correspond to perfect statistical correlations. Because something is a plane in virtue of, at least in part, what it is *supposed* to do, and not merely in virtue of what it *actually* does – because, in other words, the concept *plane* is not a strictly descriptive category, and has a functional component -- there are cases where planes don't fly, *even though* there is an important relationship between the property of being a plane and the property of being something that flies.

The concept of *law*, I believe, has a teleological component. If this is

right, then something is a law in virtue of, at least in part, what it is *supposed* to do and not merely in virtue of what it actually does. I will argue that, in fact, there is a significant sense in which laws are *supposed* to protect moral goods. An assassin who kills only the wicked is not, on that account, better as an assassin than one who kills the good. But a law that does good is, on that account, better as a law than one that fails to do so. It thus seems that it is the purpose of laws to do good (notwithstanding that this is often *not* the objective of the person, or persons, who *make* the law). Once this is granted, it is not hard to show that there is a moral basis to law, even though there are laws that fail to do any good and even though there are laws that are created by people who have no intention of doing good.[5]

Granting this, I will argue that, as a matter of conceptual necessity, a legal system *as a whole* is not only supposed to do good, but must *succeed* in doing so. There are legal systems (e.g. that of Stalin) that seem to do nothing but harm. But such a view, I will argue, involves our comparing such legal systems *not* to the *absence* of a legal system, but rather to better legal systems (e.g. that of the present-day Swedish government). When we wish to assess whether legal systems are categorically good things, our benchmark must be the absence of a legal system, and not the presence of a superb legal system.

It might be appropriate to give an illustration of the principle at work here. Even a bad doctor can do much good – he knows how to set a broken leg, to perform CPR, and so on. It is only compared to a good doctor that he can be regarded as doing harm. (When a doctor "does harm", it is almost always by *failing* to do good – by failing to do what a good doctor would have done in the same circumstances.[6]) When we wish to assess whether doctors are, as a rule, good things, we must compare the presence of doctors with the *absence* of doctors; we mustn't compare bad doctors with good doctors. Similarly, when we wish to assess whether legal systems are

categorically good things, our benchmark must be the absence of law altogether, and not the presence of some superlative system of law. We will find that, as a matter of analytic necessity, even the worst legal systems do a fair amount of good in absolute terms. (They also do a great deal of evil. But when we take into account the non-binary nature of morality – the fact that goods and bads *outweigh*, but do not *cancel*, one another – that fact is perfectly consistent with the fact that, in absolute terms, legal systems always do a non-trivial amount of good.)

I will also argue that, leaving aside those laws that result from patent deviations of the normal law-making mechanisms, any actual law *does* in fact protect some moral good, even though, in many cases, that good is dwarfed by some evil with which the law in question is incompatible.

Ronald Dworkin is one of the great exponents of the view that law and morality have some kind of non-trivial relationship to each other. The first part of the present work is spent showing why Dworkin's defense of this position is unsatisfactory. The rest is an attempt to show that, even though Dworkin's arguments are less than probative, the position he is defending is a reasonable one. We will find that, even though Dworkin's arguments are not always decisive, some of the points involved in those arguments are indispensable to a proper understanding of the deep but delicate relationship between the concepts of law and morality. In particular, we will find that Dworkin's discussion of the difference between rules and principles gives us the key to seeing the moral structure of legal obligation.

Chapter 3 Can there be morally empty laws and legal systems?

It is an empirical fact that there are immoral laws and legal systems.

But it does not follow that the relationship between law and morality is entirely a matter of happenstance. There would be sheer anarchy, and hence a *de facto* absence of law, in any society permitting unrestricted murder, theft, and arson.[7] So any legal system worthy of the name must forbid such iniquities and must to that extent coincide with the dictates of morality.

Further, even though many laws are unjust, there is a sense in which they are always *supposed* to be just.[8] A military campaign that kills only the wicked is not, on that account, more successful *as a military campaign* than one that kills the innocent. By contrast, a legal system that promotes the ends of justice is, on that very account, more successful *as a legal* system than one that fails to promote those ends. So there would seem to be a sense in which any legal system necessarily *aspires* to be just, even though many fall sadly short.

Nonetheless, some hold that any relations between law and morality *are* entirely contingent. Of course, there are good laws and legal systems, and this fact may even have deep roots in human nature. It may be rooted in a desire that people generally have, to some degree or other, to be virtuous. It may also be rooted in facts about human frailty: since any given person is vulnerable to practically any other person, it is *ceteris paribus* in one's interest to live under the dominion of a fair legal system.[9] Still, according to the view in question, there is nothing self-contradictory or absurd in the thought of an entirely wicked or amoral legal system. To the extent that we are prone to see laws in terms of morality, we are guilty of some kind of blunder.[10]

§ Ronald Dworkin is one of those who thinks that there are important, necessary relations between the concepts of law and morality. In connection with this, Dworkin argues against a position that he refers to as "positivism".

He says that positivism consists of three theses.

(i) Laws are either applied or not applied, and can no more be interpreted than the rules of chess. When a judge "interprets" the law, he is creating new law; and when a judge exercises his so-called "judicial discretion", he is either applying existing statutes or he is creating new law on the spot.[11]

(ii) The law is nothing other than a set of rules. Consider the various statutes that fill our law books. Those are the law, and nothing else is. Any principle that hasn't been converted into a statute is no part of the law. For example, consider the reasonable ethical principle that one shouldn't be allowed to profit from his own wrongdoing. If there isn't a statute explicitly forbidding one to profit from wrongdoing, then a judge who bases a decision on that principle is simply making up new law.

(iii) For something to be a law, it has to have a certain "pedigree", i.e. it has to issue from a certain kind of authority. So a moral principle is no part of the law unless the relevant authority wants it to be.[12]

Dworkin argues that each of (i)-(iii) is false. Here two separate issues arise. First, are Dworkin's criticisms of (i)-(iii) cogent? Second, supposing that they are cogent, what relevance does that have to the question of whether laws must have an element of morality? Dworkin's implicit answer is: "Given that (i)-(iii) are wrong, it follows that law necessarily embodies some kind of morality."

In the present work, I will argue that Dworkin's criticisms of (i)-(iii) are

not cogent and that (i)-(iii) are in fact correct. I will also argue that, even if (i)-(iii) were completely wrong, it would in no way follow that there were a significant necessary connection between law and morality. The question of whether morality and the law are internally related is independent of the question of whether (i)-(iii) are correct.

But I will also argue that, despite everything just said, *there is* a deep, internal connection between the concepts of law and morality. Dworkin has misidentified the nature of this connection, and his attempts to prove the existence of such a connection are less than probative. But despite these facts, and despite the fact that there are wicked legal systems, it is impossible on purely conceptual grounds that any legal system should altogether fail to incorporate *some* kind of morality. In connection with this, we will find an objective basis for our pre-theoretic intuition that there is *some* sense in which one is always supposed to obey the law.

§ It may be appropriate to say a word on the basic structure of the present work. In chapter 4, we will argue that Dworkin's arguments against (i)-(iii) are fallacious; and, further, that (i)-(iii) are correct; but, finally, that (contrary to what Dworkin believes) the truth of (i)-(iii) in no way warrants the rejection of a moralistic conception of law.

In chapter 5, we will identify what I believe to be the misunderstandings concerning morality that have thwarted attempts to identify the relationship between law and morality. We will focus on the fact, already discussed, that moral principles have a dimension of weight.

In chapter 6, we will discuss the defense of positivism found in Hart's classic paper "the Separation of Law and Morals." We will find that two of Hart's apparently more compelling arguments are not valid, since they incorporate an untenable, binary conception of morality and also a

misunderstanding of the nature of government.

In chapter 7, we will develop the points concerning government outlined in the previous chapter. We will argue that, given a plausible conception of the nature of government, it follows quite directly that, even though there are morally unsatisfactory laws and legal systems, laws must be seen as attempts to protect moral goods and also that, as a matter of conceptual necessity, a legal system cannot completely fail to do a non-trivial amount of good.

In chapter 8, we will argue that, in light of some general semantic principles, we must regard laws as being what we will call *narrow-scope* assurances. Thanks to this seemingly technical and minor point, our analysis, unlike other moralistic analyses of law, will not have to be revisionist to any degree. In other words, our analysis will never force us to regard as a non-law anything that we pre-theoretically regard as a law, or to regard as a law anything that we pre-theoretically regard as a non-law.

In chapter 9, I will discuss the plausible point, recently made by an eminent philosopher of law, that the laws of a brutal dictator can be as lacking in moral content as the demands of a gunman. I will discuss my reasons for holding that, even though a brutal dictator may be in the same psychological category as a gunman, it would only be in rather singular circumstances that the dictator's laws were in the same moral category as a gunman's threats. In this section, we will discuss some important insights regarding causality, due to Phillip Pettit and developed by Frank Jackson (in collaboration with Pettit), that we will find to have a manifold relevance to the nature of law.

In chapter 10, we will delineate the moral structure of legal obligation. We will see that legal obligation is *a* kind of moral obligation; and we will already have seen (in previous chapters) that all laws have a moral component. But, in this chapter, we will find that the relationship between

these two facts, though deep and necessary, is anything but straightforward. In connection with this, we will find reason to reject wholesale some *prima facie* reasonable attempts to find a moral basis for legal obligation.

In chapter 11, the topic of international law will be discussed. We will make a case for each of the following four claims. First, there is currently no international law. Second, there *could* be international law, at least in principle. Third, that the conditions under which international could exist would be bear extremely little resemblance to current international conditions. Fourth, and most importantly, the existence of international law wouldn't necessarily be desirable from an ethical standpoint.

Finally, in chapter 12, we will discuss the concepts of legal interpretation and of international law. It will be argued that, *pace* Blackstone, the sense in which laws are "interpreted" is not comparable to the sense in which a coded transmission or an EKG-printout is interpreted. Paradigm cases of interpretation involves the *identification* of pre-existing realities. In correctly interpreting the data generated by an EKG, a cardiologist is identifying the pre-existing fact that patient Smith has a valve defect. In correctly interpreting an intercepted message, a code-breaker is identifying the pre-existing fact that certain people have transmitted a certain message ("we are going to attack at dawn"). In each of these cases, interpretation consists in *identifying* an existing reality.

What judges do is not comparable to this. Judges do not identify, but rather extend, existing law. That is why, no matter how a judge rules, his ruling is *ipso facto* what the law henceforth requires. (Of course, a later ruling may overturn the first. But that only confirms the point that judicial rulings define, and therefore do not identify, the law.)

But even though judicial interpretations are extensions, and not identifications, of existing law, such interpretations are not *mere* legislation,

and are answerable to objective standards. Consequently, some legal interpretations are truer to existing law than others. The Supreme Court is not *making up* law (at least not necessarily). But it isn't interpreting law, at least not in any paradigmatic sense of the word "interpret."

Chapter 4 Dworkin on the separation of law and morality

Before we consider what Dworkin says, let us ask: what do (i)-(iii) have in common? Why do they form *one* doctrine, as opposed to three separate ones?

Each of (i)-(iii) is a different way of saying that, from a strictly moral standpoint, what counts as law is quite arbitrary. Consider the various regulations involved in military life or in the administration of some bureaucracy. These regulations are obviously not moral laws. It isn't a law of morality that you must use a number-two pencil when filling out a certain form. (ii) basically says that "the law" consists of nothing other than such regulations. Of course, (ii) is *consistent* with the existence of just laws. But it is also consistent with the existence of laws that are immoral or lacking in any kind of morality, whether positive or negative.

Let us turn our attention to (iii). The laws of morality obviously don't have anything to do with earthly authority. Their basis lies in Pure Reason or in the Form of the Good.[13] (iii) says that laws (in the legal sense) *must* issue from some earthly authority. So (iii) says, or suggests, that morality is one thing, law another.

Here a qualification is in order. (iii) doesn't say that for something to be a law, it is necessary *and sufficient* that it issue from a certain kind of earthly

authority. It is left open whether it must satisfy additional conditions. So depending on what, if any, additional conditions there are, (iii) is compatible with the view that laws must be moral. For example, it could be said that a law is something that (a) meets certain standards of morality and (b) issues from certain kinds of earthly authorities. (This is, in fact, what I will argue to be the case.) But given that moral principles don't have earthly origins, (iii) certainly *suggests* that law and morality are distinct.

(i) is closely connected with the other two. Rules are things that people – congressman, presidents, dictators, and so on – make up. (i) says that, when a judge goes beyond these man-made things, he is going beyond the law itself. This straightforwardly implies that rules exhaust the law, i.e. it entails the truth of (ii). It also implies that, if morality should ever become part of the law, it is only by the grace of some rule or ruling.

We may illustrate this with a bit of fiction. A certain judge must decide a case. None of the rules are on point. Since the rules are silent, the judge bases his decision on some ethical principle. If (i) is right, then by going beyond the rules, the judge was going beyond the law. This means that his decision was based on an ethical principle that was not *already* a part of the law: that principle became a part of the law only in so far as it was incorporated into the judge's ruling. So (i) implies that morals are a part of the law only when rules and rulings allow them to be. If it is granted – as it surely must be – that rules and rulings may be completely amoral, it would certainly *seem* to follow that (i) entails that the law needn't embody any morality. (Later I will argue that, in fact, even if (i) is true, it does *not* follow that laws needn't embody any morality. Hence the word "seems" in the second to last sentence.)

In conclusion, each of (i)-(iii) is a way of saying that laws are things that people make up, and no more have to embody any morality than do the army-

regulations concerning the kind of boot-polish one may use.

§ But if each of (i)-(iii) should turn out to be false, it wouldn't follow that laws *do* have some kind of necessary connection with the moral law. Suppose that laws don't have to have a certain pedigree. It doesn't follow that they must be moral. That would follow only if it could be shown that their *not* having to issue from an earthly authority entailed that they *did* have to emanate from the Form of the Good, or some such. But, as we will see, there is a third option.

Suppose that (i) is false, i.e. that when judges "interpret" statutes, they are *not* necessarily making up new law. It doesn't follow, as Hart[14] pointed out, that judicial interpretation provides a contact point between law and morality. Suppose that the principle governing the judge's interpretations is this:

(*) I must interpret the statutes in the way that is the most convenient from the viewpoint of the nefarious mega-corporation that is secretly running this country. To the extent that a given statute is ambiguous or vague, I must delineate it in the way that best serves the interests of that corporation.

In that case, presumably, the falsity of (i) wouldn't demand that there be some intersection between the law and morality.[15]

Finally, suppose that (ii) is wrong. That by itself doesn't mean that there is any component of *moral* principle to the law. In effect, we just saw this. Suppose that (*) is the principle underlying judicial interpretation, but that (*) were not itself one of the statutes. In that case, the law would include something other than various statutes, and (ii) would be false. But the falsity

of (ii) plainly wouldn't entail that there was an element of *moral* principle to the law. The principle would be purely pragmatic.

Given what we've just said, it follows that even if *all* of (i)-(iii) were false, it wouldn't necessarily follow that the law embodied any kind of morality. So we must remember that, even if Dworkin's criticisms of (i)-(iii) go through, that by itself doesn't entail that there is an element of morality to the law.

Dworkin's Criticisms of (i)-(iii)

Let us now consider Dworkin's criticisms of (i)-(iii). First of all, we must distinguish laws from what we call "the law."[16] When we talk about *particular* laws, we are not referring to diffuse ethical principles, but to specific statutes – we are referring to penal code 76.5 subsection B, or some such. But unless it could be shown that such particular laws exhausted *the* law, it would not follow that the latter was nothing but a set of rules.[17]

In Dworkin's view, it is a matter of empirical fact that the law includes principles that are not explicitly mentioned anywhere in the rule books. In defense of this view, he cites two landmark cases.

In *Riggs v. Palmer* a New York court had to adjudicate the following problem. A young man stood to inherit a great deal of money from his still alive and healthy grandfather. The young man decided not to wait, and killed his grandfather.[18]

Obviously there was a law forbidding murder, and the young man was found guilty of breaking *that* law. But, at the time, there was no statute saying that one forfeits one's inheritance if one kills the testator. Nonetheless the courts decided that the young man *had* forfeited his inheritance. The principle at work was that one should not be allowed to profit from wrong doing or illegal action.

Dworkin also discusses the case of *Henningsen v. Bloomfield Motors.* Henningsen bought a car. He signed a contract absolving the car-maker from any criminal liability in the event that he (Henningsen) should be injured as a result of some defect in the car that the automaker could have foreseen. The car did have such a defect, and that did lead to Henningsen's sustaining severe injuries. Henningsen sued Bloomfield motors. The courts decided that the contract was void, and found for Henningsen. The principle at work was that a valid legal contract must meet certain minimum standards of decency and fairness. The contract in question did not meet these standards: hence the court's decision.[19]

On the face of it, these two cases provide us with clear evidence that moral principles *are* part of the law, and also that the law is not confined to the rules that are written down in the books.

§ But, of course, given only what we've said, the positivist by no means has to concede defeat. There are two paths he can take. First, he can say: "the judges in those cases *were* going beyond the law; they were doing the *right* thing, not the legal thing." As we will later see, there is much to be said for this position.

The positivist could also say that the ethical principles guiding the judges were really *de facto* rules. There is a *rule* in our legal system requiring judges to nullify absurdly immoral contracts; there is a *rule* requiring judges not let a murderer benefit from his crime.

Here Dworkin gives us a penetrating (and, I believe, entirely correct) discussion of the distinction between rules and principles. Principles have a "dimension of weight." Rules do not. A rule either applies or it doesn't. Principles can apply to varying degrees.

Suppose there is a law saying that three people must witness the signing

of a will if that will is to be valid. This rule either applies or it doesn't. If only one person witnesses the signing of the will, then that rule voids that will. If three or more people witness the signing, then that rule does not void the will (though some other rule may).[20]

A consequence is that a coherent legal system cannot have two conflicting *rules.* It cannot have one rule saying that the signing of a will must be witnessed by three people if that will is to be valid, and another rule saying that only two people must be present. One of those rules would have to go.

Of course, every rule has a certain degree of vagueness.[21] What if, in the future, there are cyborgs? (A cyborg is a half-robot, half-person.) Do *they* count as possible witnesses from the standpoint of the law just discussed? What about people with autism? People who are extremely drunk? People who are slightly drunk, but still more or less functional?

But for a rule to be vague or indeterminate is not for it to have a dimension of weight. If cyborgs are people, then they *are* viable witnesses. If not, not. The question is *whether* the rule applies, not how much.[22]

Principles are not binary and have a dimension of weight. It is not a *rule* that people cannot profit from their own wrong-doing. If this were a rule, then all sly deal-making, all wheeling and dealing, would be against the law. There would be no commercial activity at all, and our economy would grind to a halt.

But this doesn't mean that connivance and manipulation are to be given free reign. Judges *give weight* to the principle that business should be conducted with decency. But it is to be balanced against other competing principles. Among them are:

the right of people to engage in financial dealings with one another should

not be unduly restricted*;*

and

people are entitled to pursue prosperity and personal fulfillment.[23]

Our legal system comprises the principle that people should be allowed to pursue prosperity. It also comprises the principle that business must be conducted with a certain decency and that, consequently, the pursuit of prosperity should be subject to certain restraints. But our legal system is not incoherent on account of comprising both those principles. This is because, unlike rules, principles are not binary, and apply to varying degrees. Consequently, the existence within a single legal system of two opposed principles does not constitute a *contradiction*, though it does create certain *tensions* that judges must manage.

Since principles have a dimension of weight, they are not rules. It is therefore not in the cards to say that ethical principles, such as those motivating the judges in *Riggs* and *Henningsen*, are actually rules. So (ii) is wrong.[24]

For these very reasons, it follows that (iii) is wrong. Moral principles are not created by governments or by any other earthly authority. We just saw that they are a part of the law. Therefore (iii) is wrong.

§ Let us put forth the reasoning behind (i) before we consider what Dworkin says about it. Rules are not interpreted: they are either applied or not applied.[25] You don't *interpret* a mathematical rule (e.g. a function from numbers to other numbers); you either apply it or you don't. For the same

reason, you don't *interpret* a legal rule: you either apply it or you don't.[26] This isn't to say that no *intelligence* is needed to know how to apply the law: just as intelligence may be needed to find an integral, so intelligence may be needed to know how to apply the law. But the concept of *interpretation* doesn't apply in either case. That concept applies only where there is no definite fact of the matter – only where a certain indeterminacy prevails. So legal interpretation is necessary precisely where there is an indeterminacy or gap in the law, and it therefore consists in creating law where there was previously an absence of law. Put another way, judicial "discretion" is necessary precisely when there is a gap in the law, and the exercise of such discretion therefore consists in the inventing of new law.

Dworkin's criticism of (i) is a masterpiece of philosophical analysis.[27] First of all, the term "discretion" has different meanings. Sometimes when we say that a judgment involves "discretion", we mean only that it is not mechanical and embodies sensitivity to the relevant principles.

On this disambiguation of the term "discretion", the fact that judges sometimes exercise discretion in no way supports (i). Given only that a problem cannot be solved mechanically, and that a certain amount of "discretion" is involved, it doesn't follow that there is no right answer. To say otherwise would be to say that there are no truths outside of first order sentential logic.[28] The fact that judges have to exercise discretion means only that some understanding of the relevant legal precepts is involved. It doesn't imply that the law is *absent* and that the judge must therefore fill the void with a new law of his own making.

To reinforce this point of Dworkin's, I would briefly like to go beyond what Dworkin explicitly says. The term "indeterminate" has two meanings: one ontological, the other epistemological. If we say that Smith's whereabouts are "indeterminate", we mean that we don't *know* where Smith

is; we don't mean that there is no objective fact of the matter. If we say that it is indeterminate whether a person who is 5 ft 10 inches is tall, we mean that there is no objective fact of the matter: the concept of tallness leaves the matter open.[29] It is true that, in some cases, the law is "indeterminate". But that doesn't necessarily mean that there is no *fact* as to what the law demands. It could mean simply that, as of yet, it is *unknown* what the law demands. Therefore, the existence of "indeterminacy" in the law doesn't entail the existence of gaps in the law, unless the indeterminacy in question is ontological. By itself, the fact that such indeterminacy calls for judicial discretion does not show that when a judge exercises discretion, he is making up new law. It would have to be shown that the indeterminacy in question was ontological, not epistemological. Otherwise the discretion involved would be the kind that characterizes *any* search for objective truth (outside of first order sentential logic, and other domains where all issues can be settled through mechanical decision-procedures[30]). So the positivist's argument seems to involve an equivocation on terms like "indeterminacy" and "discretion"; it shifts between the epistemological and ontological interpretations of those words.

Let us now return to what Dworkin himself explicitly says about the concept of discretion. There is another meaning of the term "discretion". Sometimes when we say that a decision is at so and so's discretion, we mean that so and so's decision on the matter cannot be overturned. If an umpire decides that a pitch was a strike, then that decision cannot be overturned.

On this disambiguation of the term "discretion", the fact that judges exercise discretion in no way supports (i). Whether the pitch was a strike or a ball is at the umpire's discretion. But that doesn't mean that there was no fact of the matter. It means that the umpire is the one who decides what the fact of the matter is. Given only that it is at the judge's discretion how to apply the

law, it doesn't follow that there is no fact as to what the law is. It follows only that the judge is the one who decides what that fact is.

There is yet another disambiguation of the term "discretion". Suppose the law requires that somebody convicted of such and such be given a sentence of anywhere from five to twenty years in jail. In this case, it is at the judge's discretion whether to hand down a sentence of five years, or seven, or twelve. But the existence of this sort of judicial discretion provides no support for (i). If the judge imposes a sentence of twelve years, he is following, not creating, the law.

There is one last disambiguation of the term "discretion", and this one is by far the most important. Consider *Henningsen v. Bloomfield motors.* The judge made a decision. But the decision was obviously not predetermined by the statutes. In fact, if we look at the statutes, we find not only that they did not determine that decision, but that they nearly enough ruled it out. Going by what the rule-books say, the contract was legal. But the judge decided on grounds of principle that the contract should be thrown out. Evidently, it was in the judge's power to override the statutes: it was at his discretion. This seems to mean that the law was the judge's to invent. Of course, it will be said that the judge "interpreted" the law, and did not invent it. But such a statement appears to distort the facts.

This, presumably, is the sort of thing that defenders of (i) have in mind. The judge was called on to "interpret" the law. But really he usurped it and replaced it with a new law. So legal interpretation is legal creation; when a judge exercises his "discretion", he is making up new law.

Dworkin's criticism of this argument is cogent. The very point the positivist is trying to establish is that no ethical principle has a place in the law, unless it is incorporated into some statute. Therefore the positivist cannot argue that, since the judge based his decision on principles not already

embodied in the statutes, he was therefore going beyond the law. That would be question-begging, as it is the very point that the positivist wishes to establish.

Critiquing Dworkin's Critique

Our background knowledge about American culture and mores makes it hard to be objective about cases like the ones Dworkin discusses. As Americans, we idolize the renegade detective who is out for justice and is willing to break the rules. We spurn the rule-bound police-chief who forces him to do things "by the book", making it practically impossible to catch the bad guy.[31] Such background knowledge tends to color our analysis of the cases that Dworkin discusses. But the following imaginary scenario gives us a way of filtering out those biases.

You are a famous judge. Because of your well-known legal prowess, a foreign government asks you to settle a case in their court-system. You don't know anything about the customs of this other country. In particular, you don't know if they are extremely liberal or aggressively traditional. You are told at the outset that your job is to apply their legal system as strictly as possible. Like anyone else who breaks their law, you will be severely punished if it is felt that you are replacing their legal system with your own conception of justice. The case you must decide is an exact analogue of *Henningsen v. Bloomfield motors.* How would you decide?

It seems to me that the path of caution would be to rule *against* the plaintiff -- to uphold the contract and dismiss the suit. Your reasoning would be: "I have no idea how this culture would react to principled *violations* of the law; so I'd better not break it – I'd better be a stickler for it. I'm going to go by the book."[32]

Some of us may have the intuition that the judge in *Henningsen* was just applying the law, even though he was overriding the statutes. Dworkin's argument relies on the validity of this intuition. But as we've just seen, if we keep the law the same, and vary the nature of the cultural background, our intuitions become very different. This suggests that the intuitions to which Dworkin is appealing are not hewed to the narrowly legal realities, but to extra-legal cultural realities.

§ But there is a much graver problem with Dworkin's analysis. The traditional picture is this. In some cases, the rules are silent. Such cases go to a judge. The judge makes a decision. His decision represents an *interpretation* of the law. The judge is not making up new law. This is the picture that Dworkin is defending. I would like to argue against it.

Let us start with a truism: what is legal is what is permitted by the law, and what is illegal is what is forbidden by the law. Given this, suppose that the judge in *Riggs* had decided that the young man was entitled to the money. In that case, anyone carrying out the judge's decision would have been acting legally. It would have been legal for the bank manager to open the safe-deposit for the young-man; and it would have been *illegal* for embittered police officers and indignant relatives of the deceased to prevent the young man from claiming his fortune.

Now suppose the judge had decided that Riggs was *not* entitled to the money. In that case, it would have been *illegal* for the bank-manager to open the safe-deposit box for the young man; and it would have been *legal* for the police to prevent the young man from attempting to open the safe-deposit box.

If the judge really had been interpreting the law, then *one* possible decision on his part would have resulted in an avalanche of illegal activity –

of people allowing the young man to have access to money which was not his, or of people depriving the young man of money which *was* his. But this is not how it works. It doesn't matter what the judge decided: anyone carrying out the judge's decision was *ipso facto* acting legally, and anyone defying it was *ipso facto* acting illegally. If the judge decides that Smith is entitled to the money, then it is illegal for law enforcement personnel, and by extension the general public, not to abide by that decision. If the judge decides that Smith is *not* entitled to the money, then it is illegal for law enforcement personnel, and by extension the general public, not to abide by *that* decision. So what can legally be done by law enforcement personnel, and by extension the general public, is determined by the judge's ruling.

Suppose a bailiff or sheriff's deputy says the following to a judge who has just decided some case where the rules are silent:

"I appreciate the sincere effort you have made to interpret the law. Unfortunately, you misinterpreted it. You see, contrary to what you said, the law requires such and such. Therefore, I am not going to abide by your ruling, and will regard as a law-breaker anyone who tries to stop me."

Such a display would be ludicrous. But far from being ludicrous, such behavior would often be *de rigueur* if judges interpreted, and did not create, the law. If judges are really interpreting the law, then they are attempting to *identify* what the law already demands and, therewith, what legal obligations people already have. An erroneous decision would *mis*identify those legal obligations: it would *mis*identify the legal obligations of law enforcement personnel (including the bailiff) and, by extension, of every civilian. In that

case, unless the judges in this country happen to be superhumanly incapable of error, it would often be illegal for a bailiff *not* to defy the judge's ruling. But no matter what the judge decides, the legal obligations of bailiffs and law enforcement personnel, and consequently of the general public, line up with that decision. Legality, and therefore the law, necessarily mould themselves around the judge's decision. This entails that the judge's decision creates, and does not identify, the law. The positivist is victorious.

How is an anti-positivist to deal with this? There are a couple of paths he can take. For the sake of completeness, I am going to consider *all* the possible options, not just those that an anti-positivist would seriously consider taking.

Henceforth, the words of an imaginary anti-positivist interlocutor will be put in bold-face. This will eliminate the tedious locutions associated with indirect discourse (e.g. "it might be said that…", "in response one might say…"). Let us now consider what an anti-positivist might say:

Suppose the judge makes the wrong decision. I grant that the acts of those carrying out his decision are not illegal. But that is because the judge's decision, though wrong, *made* those various acts legal. So relative to the law prior to the judge's decision, they were illegal. But relative to the law after the judge's decision, they were legal.

But that is exactly what a positivist would say. The positivist says: the judge's ruling *creates* law. The anti-positivist says: the judge's ruling *exposes* existing law. According to the passage just considered, the judge *creates* law.

Remember Euthyphro's question: Is an act pious because the gods like it,

or do they like it because it is pious? Does the judge choose a certain decision because it is legal, or is that decision legal because he chooses it? The passage just considered says: it's legal because he chooses it. That *is* legal positivism. So the anti-positivist couldn't (and wouldn't) take that line of thought.[33] Let us move on.

Suppose the judge makes the wrong decision. So suppose that, in actuality, the law entitles the young man to the money, but that the judge wrongly rules otherwise. In that case, the police officers, bank-personnel, and so on, who were carrying out the judge's decision *were* in fact acting against the law. The police officers preventing the young man from claiming the money were acting against the law.

But they were also acting *in accordance* with the law. A given act can be legal in one respect and illegal in another. Relative to one set of laws (those relating to inheritance), the police were acting illegally. But relative to a different set of laws (those relating to disobeying judges), the police were acting legally. So, technically, the police who stood between the young man and the safety deposit box *were* acting illegally. But they were also acting legally.

There is no contradiction in this. A comparison may help illuminate the principle involved. An act can be moral in some respects and immoral in others. It is a moral law, presumably, that one must keep one's promises.[34] Suppose I make a promise to a Mafia don to kill a totally innocent person named Smith. If I abstain from killing Smith, I am breaking my promise, and thus acting immorally. But I am also acting morally; for it is a moral law that one mustn't kill innocent people. Similarly, the police officers are acting legally relative to one set

of laws, and illegally relative to another.

Actually, the comparison can be taken further. Ultimately, the right thing to do is to abstain from killing Smith. But that doesn't mean that I am totally absolved from having to keep my promise. Moral laws don't *cancel* each other; they simply *outweigh* each other. It is *still* bad of me to break my promise. But it would be much worse to kill Smith. So *on balance* it is moral to not kill Smith – even though, in not killing him, I am acting both morally and immorally.

For a similar reason, it is *on balance* legal for the police to carry out the judge's instructions. Like moral principles, laws have weight. And if one's acts are *on balance* to be legal, one must act in accordance with the law that has the most weight.

This is not a position that an anti-positivist could (or would) take. The concept of legality is binary. It makes no sense to say that "on balance" an act is legal. It is either legal or it isn't. If a legal system characterizes a given act as both legal and illegal, then that system is simply incoherent. The idea that legality comes in degrees, or that some laws outweigh others, is a non-starter.

Of course, some laws are more *important* than others, and this is reflected in the fact that different crimes are given different punishments. The law prohibiting murder is very important; that is why the penalty for breaking it is so severe. The law prohibiting littering is less important; that is why the penalty for breaking it is less severe. But both littering and murder are 100% illegal. So the conception of legality put forth in the passage above is not viable.[35]

To fix our ideas, suppose that the young man is entitled to the

money, but that the judge wrongly rules otherwise. It doesn't follow that the police officers carrying out the judge's orders are acting illegally. You see, every law is hedged with implicit exception-clauses. For example, as a citizen, you are entitled to vote. But that privilege is suspended if you are waving a gun around when you show up at the polling station. In *Riggs*, the young man was entitled to the money. But that doesn't mean that he can claim it *no matter what.* If he drives his car through the front door of the bank, he cannot claim it (at least not that minute). If he refuses to sign the relevant paper-work, he cannot claim it. Similarly, if a judge rules that he cannot claim it, then he cannot claim it. So everyone who was carrying out the judge's orders was acting legally: their actions fell under one of the exception-clauses hedging the law entitling the young man to the money. What's the problem here?

The position described amounts to this: laws hold – except when judges decide they don't. You're entitled to the money – unless the judge decides otherwise. It is legal to do X so long as the judge says it is; otherwise, it isn't. But that is hardcore positivism: the law is what the judge says. Let us move on.

Suppose that the judge makes the *wrong* decision: he decides to let Smith go free when, in fact, the law requires that Smith be jailed. You are right, of course, that the narrowly *legal* realities are created, not revealed, by the judge's decision. But, contrary to what you say, it doesn't follow that the judge created *the law.* A later judge may overturn the first judge's decision, thereby bringing the legal realities into

alignment with the law. The law is sometimes down, but never out. An inept judge may give legal reality temporary dominance over the law. But, sooner or later, the law bounces back. The falsity of the first judge's decision is revealed by the transient nature of the legal realities it created.

This position doesn't help the anti-positivist at all. The first judge's decision determines the legal realities, at least for the time being. A second judge's decision may create new legal realities. But this only confirms the positivist's point: legal reality is created, not revealed, by judicial decisions.

There is no denying that, at any given time, legal reality – what is legal and what is not -- is determined by the rules and the rulings. These rules and rulings may be jettisoned later, but only by a fresh rule or ruling. So, at any given time, the law is a nullity to the extent that it diverges from legal reality and, therefore, from the rules and rulings.

You don't say that someone is ugly if he has a sour expression for one second, or that somebody is stupid if he makes one thoughtless comment. Beauty and intelligence are four-dimensional properties. They are properties not of instantaneous states of affairs, but of ongoing processes. The same is true of the law. You don't say that the law is defined by every spasm of judicial misconduct. Legal reality is only an expression of the law, and is not identical with it. Just as a person's comments may not always express his true intellectual capacity, so the legal realities needn't always express the true nature of the law. In the instant, legality may diverge from the law. But in the arena of the four-dimensional, they tend to coalesce. The legal realities created by erroneous decisions tend to meet a fitting end on the slaughter-bench of legal history.[36]

There are two ways to interpret this. We can say that long-term patterns of judicial activity tend to be better *expressions* of the law than isolated judicial incidents (just as, where individuals are concerned, long term behavior patterns are a better *indicator* of intelligence than are isolated behaviors). Or we can we say that the law *is* a long-term pattern of judicial activity.

Let us consider each possibility. Suppose that the law *is* a long-term pattern of activity. Long term patterns of judicial activity are constituted by the decisions that judges and juries make, along with the legal realities they define. If the law *is* a four-dimensional object, then its nature is fixed by momentary legal realities. They alone are what is operative in the here and now. The law ends up being epiphenomenal on these momentary legal realities.

Further, if we think of the law in this way, then whether a judicial decision is right or wrong depends on how consistent it is with other events: even the most atrocious decision can be historically vindicated, provided that events down the line are correspondingly atrocious. This is a form of legal nihilism: anything goes, provided that history plays along. This is the opposite of anti-positivism.

Let us now consider the other interpretation: long-term patterns of judicial activity *express* the law but do not constitute it. Just as the long-term behavior patterns of a person are more likely to be consistent with her character than any one act, so long-term patterns of judicial activity are likely to be *consistent* with the real nature of the law. But just as a medical condition (e.g. Parkinson's or chronic fatigue syndrome) may prevent an intelligent person from *ever* expressing her intelligence in its fullness, so

contingences (e.g. evil judges, incompetent attorneys) may minimize the extent to which the true character of the law is expressed in long term patterns of judicial activity.

If that is right, then legality may diverge practically without limit from the law. It may be illegal for a slave to leave the plantation: but *the law* permits it – nay, *the law* forbids the plantation-owner from having slaves to begin with. The problem is that legal reality fails to express the true nature of the law. It may be illegal for women to vote. But it is contrary to *the law* to prohibit women from voting. Again, legal reality fails to express the true nature of the law.

Suppose this is right. In that case, what do we mean by *the law*? It is always an empirical question what kind of law there is in a certain society. If X is *illegal* in society S, that fact is, at the very least, extremely strong evidence that X is contrary to the law in S. If X is legal in S, that is, at the very least, extremely strong evidence that X is *not* contrary to the law in S. The anti-positivist must say that the legal realities are *merely* evidence of the true character of law – and just as clues found at a crime-scene can be grossly misleading as to the identity of the killer, so facts about legality can be grossly misleading as to the true nature of the law in S.

This move is less than entirely satisfying. I understand perfectly how the killer might plant evidence that incriminates someone else. Killing is one thing; depositing hair or clothing at the scene of the crime is something else. The reality is one thing; the information through which we learn about it is something else.[37] This ontological fact corresponds to the epistemological fact that evidence is typically defeasible. We can imagine finding new evidence (e.g. a video-recording of Jones committing the crime) that trumps the evidence we previously found incriminating Smith. But facts about legality are not comparably defeasible. If insider trading is illegal on

Monday, and legal on Tuesday, it is very hard to believe that the change was a purely evidential one. Surely *the law* changed.

§ Let us distill the points just made. When we describe an act as being "against the law", we mean that it fails to comply with some standard or falls short of some *norm.* But in some cases, "the law" has a non-normative meaning; it denotes the legal *process*. In the statement: "the law operates slowly but surely", the expression "the law" denotes something which *happens*; it denotes judicial *activity*, not a norm in terms of which such activity is to be evaluated. The same is true of statements like "I'm going to put my faith in the law" and "the law is easily manipulated." Nobody puts their faith in a set of norms: they put their faith in the judicial processes that are supposed to realize those norms. A set of norms cannot be manipulated: what can be manipulated are the human processes that are meant to realize those norms.

The debate between the positivist and the anti-positivist must be understood in terms of the *normative* meaning of the expression "the law." If we take "the law" to denote judicial *activity*, then every judicial decision, no matter how absurd, is automatically a part of "the law." But this is specifically what the anti-positivist rejects; the anti-positivist says that a judicial decision *can* fall sadly short of what the law actually requires. So when the anti-positivist talks about "the law", he must have the *normative* meaning of that expression in mind.

The judges in *Riggs* and *Henningsen* did indeed base their decisions on principles. But this only shows that principles have a role in judicial *activity*; it doesn't show that they are part of the law in the *normative* sense. The positivist will happily admit that principles are involved in judicial *activity.* What the positivist denies is that principles can be a part of "the law" *in the*

normative sense without being incorporated into rules. Let us move on.

In a baseball game, everybody abides by the umpire's ruling. But it doesn't follow that the umpire himself *makes* a given pitch be a strike or a ball. The reality is independent of what the umpire decides. The ball either was, or was not, inside the strike-zone.

At the same time, *somebody* has to have the final word – otherwise every baseball game would disintegrate into anarchy. So even though the umpire's ruling determines whether *for all intents and purposes* the pitch was in the strike zone or not, there is still an independent reality. Further, even though the umpire can be in error in any given case, he cannot *generally* be in error. If his decisions are patently false on a regular basis, he will be dismissed.

Something exactly analogous holds with respect to judicial decisions. Obviously judges make mistakes. But the buck has to stop somewhere. Otherwise there would be anarchy, as every civil dispute would lead to mayhem. So the judge's decision defines what the law is *for all intents and purposes* – just as *for all intents and purposes* an umpire's call defines whether a pitch was a strike or a ball. This allows a judge to be wrong on any given occasion. But the judge cannot *generally* be wrong. If a judge's decisions are patently false on a regular basis, then he will be dismissed, for the same reasons as a habitually incompetent umpire. So the over-all pattern of a judge's decisions is necessarily consistent with the real nature of the law, even though any given decision may deviate from it.

Legality* is what the law is *for all intents and purposes*. The umpire's decision defines the status of the pitch *for all intents and purposes*, even though a given call isn't necessarily true to the actual nature of the pitch. Similarly, the judge's decision defines the status of the law *for all intents

***and purposes*, even though any given decision may not be true to the actual nature of the law. Thus it is no more absurd to distinguish legality from the law than it is to distinguish an umpire's decision regarding a pitch from the pitch itself.**

There is a very clear distinction between the actual trajectory of the ball and what the umpire *says* the trajectory was. Anyone (who is suitably placed) can see the pitch with his own eyes. So anyone can compare the umpire's call with the corresponding reality.

A related point is that the reality doesn't mould itself to the umpire's calls. Even if the umpire is wrong all the time, that has no effect on where the ball actually goes, and (leaving aside easily removed accidental facts) it has little or no effect on where anybody thinks the ball goes.

In the case of the law, the reality *does* mould itself to the decisions that judges make. Suppose that a series of judges were to rule that those of middle-eastern descent could not own land. Whatever the law was before these rulings were made, the law would certainly come to fit these rulings: it would *become* illegal for middle-easterners to own land, even if it wasn't before.

Our own history illustrates this point. When judges ruled that segregation was Constitutional, segregation was permitted by the law. When they ruled that it was unconstitutional, it was forbidden by the law. It is an historical datum that the law moulds itself to judicial rulings.

This is inconsistent with the idea that the law is a reality that transcends the rulings. For Berkeley, a physical object's *esse* lies in its *percipi.* This is just a way of saying that there are *no* physical objects.[38] To take an anti-realist view of X just *is* to say that the nature of X is a function of our

representations of X.[39] Enough judicial misrepresentations of the law create a new law of which they are a correct representation. For purely metaphysical reasons, this makes it hard to believe that the law has any reality outside of judicial rulings.

This makes it clear why the objector's comparison is spurious. A series of bad calls on the part of an umpire doesn't affect where the ball goes, or even (leaving aside facts easily idealized away from) where people think the ball goes. But a supposedly deviant judicial ruling becomes a correct one, provided that enough similar rulings follow it. If some crackpot judge rules that people of middle-eastern descent cannot own land, that ruling becomes the law, provided that enough similar rulings follow it.

Suppose that today (Feb. 13, 2006), a judge were to make the barbaric ruling that those of middle-eastern descent couldn't own land. It would be patently obvious to everyone (except a few bigots) that the ruling was contrary to the law (even if there were some narrowly statutory justification for it). This fact would be *expressed*, but not *constituted*, by that ruling's being overturned by later judicial rulings.

Here we must recall what we said earlier. Our intuitions about a given judicial decision often embody background knowledge about the culture. To ensure that our intuitions track the purely legal realities, and not extra-legal ethical and cultural realities, we must consider the same judicial decision while varying the cultural background.

Given this, imagine a country – call it Acirema --where the same legal rules as ours are operative, but where the background mores are very

different. There it is taken for granted that a certain amount of xenophobia is a good – indeed, a necessary – thing. In Acirema, a crackpot judge rules that those of middle-eastern descent cannot own land. The judge's ruling is upheld by later rulings. Consequently, the first ruling is solidified into law. Are we to say that, in Acirema, the judges all happen to be wrong about what the law is? Even if the initial judge was wrong, the law eventually moulds itself to his decision. But, for reasons we've discussed, this isn't consistent with the idea that there is any reality (at least not any legal reality) relative to which that ruling was either true or false.

§ These points can be illustrated with a chapter from our own history. In 1900, it was illegal for women to vote. This is an empirical fact. If a woman tried to vote, she would have run into an impenetrable wall of statute and law enforcement. But if there are any principles at all in our legal system, we may be sure that all of them have always demanded that women have the right to vote. Nothing could be more undemocratic, and thus more contrary to the principles (if any there be) inherent in our legal system, than denying women the right to have a say in who the nation's leaders will be. So if there was ever a case where legal realities clashed with principles inherent in the law, this was it.

The positivist says that in 1900 it was against the law for women to vote. Period. Fortunately, those laws were later overturned. But that isn't because the law finally managed to express itself; it isn't because some principle buried deep within the law finally rose to the surface. It is because the law itself changed. The change consisted in some judge's making a new ruling or in some legislature making a new rule.

According to the anti-positivist, our legal system isn't just a collection of rules: it also includes principles. As we just saw, if our legal system includes

any principles, they have always demanded that women have the right to vote. So the anti-positivist must say that denying suffrage to women was contrary to the principles inherent in our legal system, and was thus in some sense contrary to the law. This would seem to mean that, all along, the law *did* allow women to vote. But this seems strange: it is an empirical fact that in 1900 it was illegal, and therefore contrary to the law, for women to vote. So the anti-positivist's position is given by the self-contradictory and therefore false proposition that, even though it was illegal for women to vote in 1900, the *law* nonetheless allowed them to.

What does the anti-positivist say here? There are a couple of options. First, he could say that, deep down, it was in fact legal in 1900 for women to vote. But judges, police officers, and so on, didn't know it: the relevant facts about the law were too far below the surface.

But such a view would be far too revisionist to be taken seriously. Also, the term "deep down" seems not to apply to the concept of legality. As we've noted, something is either legal or it isn't. The binary nature of legality doesn't allow any place for the concept of a *deep down*. There is no *deep down* where everything is black or white. If we wish to say that *deep down* such and such is true of the law – e.g. that *deep down* U.S. law in 1900 did permit women to vote, even though they could not *legally* do so – then we must implausibly distinguish legality from the law.

Let us move on (I remind the reader that bold-faced passages are those of an imaginary anti-positivist interlocutor).

The law did indeed change. But that does nothing in the way of discrediting anti-positivism. You have tendentiously chosen a case where, under the banner of legal principle, judicial activity really did usurp

existing law. But there are other cases where principle-based decisions *are* revelations, not usurpations, of existing law.

By this reasoning, the great legal decisions that define U.S. jurisprudential history were cases of judges making up laws. It isn't clear how, in any respect relevant to our inquiry, the changing of laws relating to suffrage are different from the changing of laws relating to slavery or abortion. If the objector is right, then it seems to follow that practically every significant decision made by the Supreme Court was a case of judicial usurpation of the law.

This is not a sophistical use of slippery-slope reasoning. A slippery slope argument purports to show that, given certain apparently reasonable premises, cases we *know* to differ in a certain respect from a paradigm do *not* in fact so differ. A person with no hair at all is a paradigm of baldness. A slippery slope argument purports to show that a person with a full head of hair is also bald.[40] Nothing comparable to this is going on in the present case. I am not trying to prove that some judicial decision that differs from the paradigm does *not* so differ. Rather, I am discussing the paradigm itself. If there are *any* principles at all in our legal system, one of them is that citizens must be allowed to choose who their leaders are and must therefore be allowed to vote. If there was ever a paradigm case of a legal *principle* at work, it was the abolition (in the U.S.) of the laws prohibiting women from voting. When we look at that paradigm, what we find is that it involved a massive change in legal reality. So if the decision to give women the right to vote is a case of a principle inherent in our system of law finally managing to make itself heard – and that decision *is* such case, if anything is – then we

must distinguish legal reality from the law. *The law* didn't forbid women from voting: but it was still *illegal* for them to do so. But, as we've seen, it doesn't appear possible to make this distinction.

You've been talking a lot about empirical facts ("it is an empirical fact that women couldn't vote in 1900"). Well, here's an empirical fact for you. The Constitution is the supreme law of the land. If you say otherwise, then you need to take a class on American government. The Constitution is not just a bunch of statutes: it is loaded with principles. You've mentioned a couple yourself. I gather you'll say that the Constitution is *not* the law. But then who's being absurdly revisionist now?

I do not wish to advocate any kind of anti-realism or nihilism about the concept of Constitutionality. There are facts about what the principles in the Constitution require, and (I believe) those principles never did countenance atrocities like slavery.[41] But it is also a matter of empirical fact that slavery was legal. Given this fact, if we wish to maintain that the principles in the Constitution are a part of "the law", then we must either distinguish legality from what is permitted by law – a move we have already seen to be a non-starter – or we must distinguish different strata of legal reality: the surface legal realities permitted slavery, but *deep down* the law did not permit slavery.

We've already seen the problem with this. There are no strata of legality. Something is either legal or it isn't. If a legal system comprises contradictory

rules, then it is *ipso facto* incoherent. U.S. law was evil, but not incoherent, in virtue of permitting slavery. So we cannot say that, while permitting slavery, U.S. law *also* prohibited it (on some deeper level).

Here there is an obvious possible counter-move for the anti-positivist:

Contrary to what you just said, U.S. law *was* incoherent in virtue of permitting slavery. In due course, certain political and judicial events removed that incoherence. Indeed, much legal progress consists in the removal of such incoherencies. To take another example, it has always been unconstitutional, and therefore illegal, to prohibit women from voting. But for a long time, the law did indeed prohibit it. Therefore, it was legal, and also illegal, for a woman to vote. Fortunately, that incoherence was eventually removed by some legislators and judges who were able to see into the depths of our legal system. In general, the Supreme Court is in the business of removing contradictions from our legal system.[42]

The problem is that there isn't any significant sense in which women could legally vote in 1868. Even if there is such a sense, it is not remotely comparable to the sense in which they could *not* legally do so. Obviously two very different concepts are at work here. We would simply render the term "legal" ambiguous if we were to use it to mark both of them.

It seems reasonable to suppose that the spirit, if not the letter, of the Constitution has never countenanced the denial of suffrage to women: the *principles* inherent in the Constitution are, I believe, inconsistent with something so obviously undemocratic. But it is also an empirical fact that

women did not have the right to vote in 1868. So if we are to say that the aforementioned principles have always been a part of American law, then we must distinguish the law from legality: it was *illegal* for women to vote, but it was against *the law* to deny them that right. But we have already seen that legality cannot be distinguished from the law. Hence, just as the positivist argues, those principles were *not* a part of the law – until (come 1920[43]) they were incorporated into rules and rulings.[44]

§ This gives us some insight into the nature of the obligation that judges have towards the principles in the Constitution. Judges *do* have an obligation to uphold those principles. But, as paradoxical as it may sound, the obligation in question is not of a strictly *legal* nature, as I would now like to argue.

It makes no sense to say that the Constitution imposes a "legal obligation" on a judge to do X, but that the judge can legally *not* do X. Such a position would eviscerate the notions of legality and of legal obligation. Given this, suppose we say that the principles inherent in the Constitution imposed a *legal* obligation on judges (in the ante-bellum South) not to recognize slavery. In that case, any judge who upheld any law permitting slavery was acting illegally. By implication, the same is true of any member of law-enforcement and, by extension, any civilian who complied with any law permitting slavery.

But there is no significant sense in which this was the case. Any member of law enforcement, and by implication any civilian, who defied those laws was acting illegally. (Somebody who tried to help slaves escape was acting illegally, just as in our time somebody stealing a car is acting illegally.) Anyone who abided by them was acting legally. (Somebody who *did not* assist a slave in his effort to escape was acting legally.) These are historical data. So if we say that the Constitution imposed a *legal* obligation on judges

in the ante-bellum South to uphold laws permitting slavery, then we are falsely saying that there *was* some significant sense in which slavery was illegal. Thus, while U.S. judges probably do have *an* obligation to uphold the principles in the Constitution, that obligation is not of a *legal* nature. We must find some other way to understand the relationship between the law, on the one hand, and the principles inherent in the Constitution, on the other.

It is a matter of plain historical fact that the presence of principles in the Constitution *has* given people freedoms that they otherwise would not have had. So the principles just discussed *do* have an important relationship to legal reality. But this doesn't mean that those principles are *already* in the law. The relationship between the principle and the law is not one of identity, but of causality: the presence of principles in the Constitution *causes* judges to make rulings that they might otherwise not make.[45] As we saw a moment ago, if we say that those principles are *already* in the law – regardless of what judges and legislatures do – then we put ourselves in the absurd position of having to say that legal reality may not always coincide with the law.

Is my analysis overly revisionist?

I just argued that the U.S. Constitution is not part of U.S. law. As the objector said, this is likely to register with many as blatantly false and as thus warranting the rejection of my analysis.

After putting forth my own positive analysis of law, we will find a way to deal with all of the data thus far considered. Right now, I can only make a few anticipatory remarks.

First of all, we must at all costs hold onto the truth that what is legal is what is permitted by law and that what is illegal is what is forbidden by law.

As we've seen, Fuller's analysis is inconsistent with this and must therefore be rejected. At the same time, I do admit that, in so far as my analysis requires that I deny the Constitution a place in the structure of U.S. law, that redounds to the discredit of my analysis (even though, I have argued, it does not, strictly speaking, warrant its rejection). Let me now outline how my analysis of law will satisfy these desiderata.

I will argue that laws are governmental assurances of protections of rights. (Neither this contention, nor any of my arguments for it, presuppose the truth of any of the anti-Dworkin or anti-Fuller points I have made so far; indeed, my arguments for it are, ultimately, Dworkin- and Fuller-friendly. So no questions will have been begged, at least not against those two scholars.) Further, I will argue that judicial "interpretations" of the verbiage found in the Constitution consists not in analyzing assurances that already exist in a completed form in the Constitution, but rather consist in *completing* the assurances that are initiated, but not completed, by that verbiage. So those rulings complete Constitutional assurances.

Here it may help to focus on a few platitudes. A so-called assurance that isn't understood isn't an assurance at all. If you don't speak English, and I say "I assure you that I will protect you from the local Mafia chieftain", I haven't *actually* assured you of anything, though I have attempted to do so. So far as a (so-called) assurance is vague or unclear, it is no assurance, but merely an attempted assurance. Consequently, so far as the verbiage in the Constitution requires interpretation – in other words, so far as it is unclear *what* those words are meant to assure, or who they are meant to give those assurances to – *no* assurance has been provided by them. (There has been only an attempt to give an assurance.) Since, for reasons that we will see, laws *are* assurances (on the part of the government) of protections, it follows that, in so far as the verbiage in the Constitution requires judicial

clarification, that verbiage fails to provide assurances and thus fails to express laws. From this, it follows that judicial "interpretations" of the Constitution actually *complete* the incomplete assurances found in the Constitution. Given that laws are governmental assurances, it follows that it is the verbiage in the Constitution *plus* the needed judicial completions (or "interpretations", as they are usually called) that constitute laws. From this, finally, it follows that when a judge's ruling embodies a delineation or interpretation of some obscure part of the Constitution, that ruling does not create any kind of *contradiction* (at least not any *legal* contradiction) with the contents of the Constitution.

Let me sum up the obscure and dense argument-outline just provided. So far as they are in need of interpretation, the contents of the Constitution *ipso facto* fail to provide assurances. So far as they fail to provide assurances, they *ipso facto* fail to be law. The supplementation of judicial rulings is needed for those contents to *become* law. Consequently, *pace* Fuller, those rulings do not *contradict* those contents.

Fuller's picture is this: we have the Constitution, on the one hand, and we have judicial rulings, on the other. Both are obviously a part of U.S. law. When a judge rules a certain way, his ruling is law. (On this I am in agreement with Fuller.) But judges may misinterpret the Constitution. When a judge does so, the law created by the judge's misinterpretation is inconsistent with the law found in the Constitution. Thus U.S. law becomes incoherent.

My picture is this. First, because Fuller's picture implicitly requires that we distinguish legal reality from the law, one of Fuller's premises must be wrong and must therefore be rejected. Given this, there are two possible sub-pictures (within my picture), both of which model the relevant data equally well. One is to reject the idea that the U.S. Constitution is a part of U.S. law –

it is to be treated not as law *per se*, but as some kind of paradigm or exemplar. Its undeniable relevance to U.S. law is not one of identity or overlap, but is rather one of *influence* or *causality*. This picture does require us to swallow a slightly bitter pill – that of denying that the U.S. Constitution is part of the law. But there is independent corroboration for that position, as we have seen; and later, in Chapter 12, we will find additional corroboration for it.

The other sub-picture is this. We hold onto the idea that the U.S. Constitution is a part of the law. We argue (on grounds that we make sure not to beg questions against Dworkin or Fuller) that laws are governmental assurances of rights. On this basis, we argue that, so far as the Constitution requires interpretation -- i.e. so far as it is not unclear what assurances it provides (or to whom they are being provided) -- judicial interpretations of the Constitution *complete* the incomplete proto-laws/proto-assurances found in the Constitution. If that is right, then such interpretations cannot possibly *contradict* the contents of the Constitution. In that way, we avoid the position, taken by Fuller, that misinterpretations of the Constitution categorically create *contradictions* within U.S. law. As we've said, we *want* to avoid that position; for we have seen it to require an acceptance of the absurdity that what is legal may be different from what is permitted by law, and that what is illegal may be distinct from what is prohibited by law.

I do not deny that there *can* be contradictions within a legal system. Surely there can be. In fact, given a legal system of any complexity, it is not improbable that it will comprise a contradiction. (There is nothing contradictory in the supposition that a police officer might have orders to arrest people who drive over 65 mph and also *not* to arrest such people. What is contradictory is not the supposition that such orders, or the corresponding laws, exist, but that they can both be followed.) Nor do I deny that judicial

interpretations (so-called) of the Constitution may *in fact* contradict it. What I deny is this: when the contents of the Constitution are genuinely obscure – when the letter of the Constitution is unclear – those contents do not amount to law and, consequently, judicial "interpretations" of those contents cannot *contradict* the law found in the Constitution. Such interpretations are needed to turn those contents into law, and thus cannot be at variance with them.

I am going to end this outline with a statement even more obscure and paradoxical than those made in the last few pages. Despite everything just said, I believe that *there can* be misinterpretations even of the obscure parts of the Constitution. But, in this context, the word "interpretation" does not have the meaning that it ordinarily bears. When we make it clear what its actual meaning (in this context) is, we will find, I propose, that all of the propositions just put forth hang together.

But I must reiterate my caveat that the last few pages are meant only to outline an argument that cannot be given with any fullness, or even stated with any clarity, until we've had time to develop some positive (not to say positivist) theses about law and government. [46]

Are (i)-(iii) correct?

In light of what we have found, let us reconsider (i)-(iii). (ii) turns out to be right, at least in spirit. What is legal, and is therefore consistent with law, is what the rules permit. When the statutes are silent, a judge's ruling fills the gap. Either way, legality is defined entirely by rules and rulings.

This establishes the truth of (iii). Statutes and judicial rulings are, by definition, things that have a certain pedigree; they come from the legislature or the bench. Since legality is defined entirely by rules and rulings, legality is a matter of pedigree. So, therefore, is the law, unless we can find a way to

divorce legality from *the* law.

On the relevant disambiguation, (i) also turns out to be correct. Suppose a ruling does not expose what is already in the statutes. In that case, by supposition, the ruling goes beyond the statutes. It is a matter of empirical fact that, no matter what the judge decides, his decision will define legality, at least for the time being. If the judge decides that Smith must go to jail, then law-enforcement, and by extension everyone else, is legally obligated to abide by that. If the judge decides that Smith is to be set free, then law-enforcement, and by extension everyone else, is legally obligated to abide by *that.* Perhaps the judge's decision is later reversed by another judge, and the principle motivating the second decision is contrary to that motivating the first. But it is an empirical fact that, for the time being, the line between the legal and the illegal is drawn by the judge's word, so far as it is not drawn by statute. If some principle comes to the rescue, that is only thanks to the ruling of some later judge. Where the judge isn't simply carrying out a statute, what he decides is *constitutive* of the law, and therefore cannot be an interpretation of it. So if by "legal interpretation" we mean the act of speaking for the law where the statutes are silent, then so-called acts of legal interpretation are indeed acts of laying down the law, and not of creating it.

If by "legal interpretation", we mean the act of delineating what is already present in the statutes, then legal interpretation is not an act of laying down the law, and (i) is therefore false. But this is no problem for the positivist, since it is consistent with the positivist thesis that the law is nothing but a set of rules (and rulings). On the other hand, if by "legal interpretation", we mean something other than delineating the content of the statutes, then so-called legal interpretations create, and do not reveal, legal reality. In that case, (i) turns out to be correct. So on no disambiguation of the terms "discretion" and "interpretation" is there any significant problem for

(i)-(iii).

In conclusion, (i)-(iii) turn out to be correct. Dworkin's arguments to the contrary seem to involve an equivocation on the term "the law". Sometimes "the law" refers to legal *culture* – to an ongoing pattern of judicial activity, with all of its motivating principles. Sometimes "the law" refers to what is legal at a given time. If we take "the law" in the first way, then the presence of principles in the law is completely compatible with the positivist view. If we take "the law" in the second way, then just as the positivist alleges, principles have no place in the law except where rules and rulings give them such a place. Either way, the positivist is right.

The irrelevance of (i)-(iii) to the moral status of law

But for the sake of argument, let's suppose that (i)-(iii) are false. As we've already seen, that doesn't necessarily reconnect morality with the law. Let us extend the points we made earlier on this topic.

Let's suppose that (i) is false, i.e. that the law includes principles, not just rules. In that case, there could be genuine legal interpretation (as opposed to legal creation that is misleadingly described as interpretation). Such interpretation would consist in making explicit what the principles inherent in the law demand. But it does not follow that the law incorporates any kind of morality, for the principles in question needn't be moral ones. They might be along the lines of: "do whatever would most please the board of directors of the evil mega-corporation that is running this country." The falsity of (i) is thus compatible with the absence of any morality from law. The same consideration shows the same to be true of (ii).

We don't need to look to science fiction to find relevant examples. No doubt there were statutes in Stalin's Russia and Hitler's Germany that needed

to be supplemented by judicial consideration of principle. I doubt that the principles involved were always ethical ones; they were probably more along the lines of *tow the party line*.[47] As Hart wrote:

[I]ntelligent [principle-based] decisions which we oppose to formal or mechanical decisions [those that are mere applications of existing rules] are not necessarily identical with decisions defensible on moral grounds. We may say on many a decision: "Yes, that is right; that is as it ought to be", and we may mean only that some accepted purpose or policy has been thereby advanced; we may not mean to endorse the moral propriety of the policy or the decision. So the contrast between the mechanical [narrowly statutory] decision and the intelligent [principle-based] one can be reproduced inside a system dedicated to the pursuit of the most evil aims. It does not exist as a contrast to be found only in legal systems which, like our own, widely recognize principles of justice and moral claims of individuals.[48]

What about (iii)? Let's suppose that not all law issues from legislators, judges, evil boards of directors, and so on. Does it follow that some law issues from morality herself?

What is the argument here? The basic idea seems to be this:

Suppose some law doesn't issue from some kind of earthly authority – it doesn't come from a king or a legislature, or any other such thing. Then where *does* it come from? By supposition, it must have some *non*-earthly basis. What kind of *non*-earthly entity could do the trick? Not the laws of arithmetic. Not Riemann's geometry. The only relevant non-earthly

entity is the form of the Good, or some such.

But there is a third alternative. The law doesn't issue from the form of the Good *or* from an earthly authority. It issues from an earthly *non*-authority. Laws tend to reflect the outlook of the people they govern, even in the most autocratic of societies. At any given time, the Pharaoh can kill any given person: his subjects will acquiesce to that. But Pharaoh wouldn't last long if he required his subjects to spurn their gods. Consequently, the laws issued by Pharaoh will embody a certain respect for those religious views. So here the laws *do* have some kind of basis in something that is not the form of the Good *or* the dictate of an authority figure: the law has a basis in the outlook of those who have *no* authority.

The anti-positivist might say that, by virtue of having the power to depose the Pharaoh, and thus to enforce a certain consistency between their religious outlook and the Pharaoh's laws, the peasants *are* authority figures (collectively, if not severally). In that case, it may be true that if the law doesn't issue from an earthly "authority", then it must come from something non-earthly (presumably, the form of the Good, or whatnot). But this position extends the notion of authority to the point of emptiness. Since an "authority" figure is now *defined* as one that can affect the law, the position in question amounts to the innocuous tautology that, if the law doesn't come from an earthly authority, then it does *not* come from an earthly authority. If the anti-positivist is to avoid trivializing his position, he cannot define "authority" in the way proposed. But in that case, as we saw, his thesis is simply false.

In conclusion, there doesn't appear to be any significant connection between the (putative, but not actual) falsity of (i)-(iii) and the truth of some kind of moralistic conception of the law.

Have we begged the question against the anti-positivist?

If the arguments thus far presented are to deserve any credence, we must show that they are not destroyed by the following important objection (due to Thomas Holden):

You take it for granted that, in times past, there were *laws* in the U.S. permitting slavery and prohibiting women from voting. Admittedly, that is probably what a historian would say. But in this context, it is quite question-begging. After all, as an anti-positivist, I would say that those so-called laws were not laws at all, on the grounds that they lacked the relevant moral qualifications.

I can see why one might take those pseudo-laws to be actual laws. After all, they had some of the same formal or genetic properties. Like real laws, they were enacted by congress and then duly deposited in the so-called *law*-books, then to be cited and debated by so-called *law*yers, judges, and professors. But it would be folly to say, on these grounds alone, that there were *actual* laws prohibiting women from voting.

A comparison (borrowed from Gareth Evans[49]) might be illuminating. Whales have many of the phenomenal and functional properties of fish. Given this, it is understandable why one might think that whales are fish. But they are not. The pseudo-laws prohibiting women from voting were functionally much like real laws. But given only that information, it would be as ludicrous to say that those so-called laws were *real* laws as it would be to say that whales must be fish since they look and act like fish.

Your insistence there were actual *laws* prohibiting women from

voting nearly enough *is* a kind of functionalism about the law. After all, your grounds for that insistence can be nothing other than the fact that those so-called laws had more or less the same social role as real laws. But anti-positivism just *is* the position that this sort of functionalism is not the right approach. An anti-positivist would say that your conception of the law is guilty of the same sort of blunder as the attempt to analyze qualia along functionalist or behaviorist lines. The essence of qualia is phenomenology, not functional role. Similarly, says the anti-positivist, the essence of law is morality, not socio-functional role.

So your acquiescence to the supposed historical datum that there are laws prohibiting women from voting is entirely question-begging. After all, that acceptance is by motivated by an acquiescence to a viewpoint according to which laws are individuated by their social-functional role, as opposed to their moral qualifications: and that is exactly the viewpoint that the anti-positivist rejects.[50]

That, incidentally, is why Dworkin goes to such trouble to reject the idea that laws must have a certain *pedigree*. He wishes to *reject* the idea that the question whether x is a law is to be answered in terms of what sort of social institution x originates from. So, in rejecting this idea, Dworkin is *rejecting* a functional conception of law. And, in taking it for granted that there were *laws* prohibiting women from voting – on the grounds that the prohibitions in question had a certain social role (and, in particular, a certain pedigree) – you are merely assuming, and not arguing for, the sort of functionalism rejected by Dworkin.

There are two points to make here: a purely historical point and a philosophical point. The historical point is that - leaving aside Augustine, Aquinas, and a few other ancients -- no anti-positivist has gone so far as to

say that there were no *laws* prohibiting women from voting in the U.S. We will soon discuss the reasons that, in all likelihood, have motivated resistance towards such wholesale revisionism.

Lon Fuller[51] did deny that many of the so-called "laws" enacted by the Nazis were actual laws. But there were independent grounds for that view. The "laws" issued by the Nazis *did* lack many of the genetic and functional properties of things that are uncontroversially laws. We will discuss this more fully in a moment.

§ Let us now give a more philosophical response to the objector's point. In order to avoid begging the question against the objector, let us introduce a new term. A law* (note the asterix) is the kind of thing that a *historian* would regard as a law; it is the sort of thing that people *suppose* to be laws. So it is a historical datum, and thus uncontroversial, that there were *laws** prohibiting women from voting. What is here in question is whether those laws* were laws.

Let us momentarily talk about the concept of law (no asterix). Presumably the concept of law is a "natural kind" concept, at least in an (in this context) appropriately extended sense of that term, and thus corresponds to a unified and fundamental explanatory category.[52] When philosophers study the concept picked out by the term "law", they want to study a concept of the sort just described: a concept that has independent explanatory value. They want to study a concept that has proven its mettle in fields to which the concept of law is relevant – fields like history, economics, and social psychology.

Given these points, let me now state the position I wish to defend. If we say that, because of their moral deficiencies, the laws* prohibiting women from voting were not laws, then we end up throwing the philosophical

concept of law out of alignment with its counterparts in other disciplines. We end up studying an artifact of our own analyses, and not something that has any independent explanatory value. Further, supposing that there were no laws forbidding women from voting, all of the questions addressed within the philosophy of law are recreated within the philosophy of law*. This suggests that the so-called analysis proposed by the objector isn't an analysis at all, but only a terminological change masquerading as such.

§ Let C_1 be the concept corresponding to the term "law*". So the laws* prohibiting women from voting fall under C_1, as do the morally wholesome laws that prohibit insider-trading. For the sake of argument, suppose that philosophers decide that the concept corresponding to the term "law" is not C_1, but is rather some more restrictive concept C_2, to be defined thus: the extension of C_2 includes all and only the *morally wholesome* things that are in the extension of C_1. So the law* against insider trading presumably falls in the extension of C_2, but the laws* that prohibited women from voting do not. C_2 is thus the concept that corresponds to the term "law" from the viewpoint of the objector's extreme form of anti-positivist.

It is readily seen that, under these circumstances, the term "law" would be quite useless, and would have to be replaced by the term "law*." Consider the laws* that prohibit air-traffic controllers from striking. Are they laws, in the objector's sense, or are they just laws*? Can we use the word "law" to refer to them, or must we stick with the word "law*"? We can ask these very questions *mutatis mutandis* in connection with the laws* prohibiting illegal aliens from having drivers' licenses, the laws* requiring police to Mirandize suspects, and the laws* prohibiting the formation of certain kinds of

monopolies.

If it means what the objector suggests it does, the word “law” is useless as a device of reference *until* all of these delicate questions are answered. But once those questions are answered, that word would be stripped of much of its discursive usefulness. After all, it is these very questions that historians, sociologists, and even economists aspire to answer.[53] So the term “law” would be (relatively) useless until those issued had been settled, and would thus be useless in the way of helping to settle them.[54] In connection with this, it is readily seen that C_2 does not pick out an explanatorily significant class of objects, since analogues of the points just made about the term “law” apply to that concept.[55]

To avoid these problems, we’d have to let the neutral term “law*” take the place the of the term “law.” But then all of the problems studied in the philosophy of law would simply be displaced onto the philosophy of law*. In particular, we would want to know what property was had by all and only those things that are laws*. It would not be an option to say that there were no laws* forbidding women from voting. All of the problems that the objector hoped to solve would therefore re-surface in connection with this other term. This would appear to confirm our earlier suggestion that the objector is not analyzing concepts, but merely re-labeling them.

§ Of course, sometimes it *is* necessary to revise or even replace lay-concepts, and to make corresponding changes in nomenclature. There is no one scientific expression corresponding to the lay concept of temperature. [56] (Or so it has been argued -- and so we will assume for expository purposes.) Consequently, the lay-term “temperature” has been replaced by a number of different expressions.[57] That is because such a replacement was necessary in order to prevent explanatory breakdowns and to facilitate the

expression of known truths.[58]

But for the reasons discussed a moment ago, the objector's proposed conceptual revision, along with the terminological changes thereby demanded, would make it practically impossible to express known truths. In connection with this, it seems that they would *create*, not prevent, explanatory breakdowns. After all, it is hard to see how those expressive breakdowns could be without consequence for our ability to operate with the concepts in question. Also, if our nomenclature is so unwieldy, that is presumably because it embodies some sort of confusion – because it doesn't track the concept-defining fault-lines which it is the business of explanation to expose.

So there is no evidence that anything like the conditions that demanded a rational reconstruction of the concept of temperature obtain in connection with the concept associated with the term "law."

§ Of course, it cannot be said *a priori* that the concept of law is entirely unamenable to the sort of revisions to which other concepts have been profitably subjected. Some have said that sociological concepts must be replaced with psychological concepts.[59] Some have gone so far as to say that, ultimately, the only explanatorily legitimate concepts are those belonging to physics.[60] Perhaps, then, the concept of law *does* need to be replaced with concepts of some more fundamental kind.[61] Perhaps I was wrong to say that no rational reconstruction of that concept was needed.

But if so, that is not likely to help the anti-positivist. Suppose it turned out that the concept picked out by the term "law" had to be replaced by, or re-articulated in terms of, concepts belonging to physics. Surely those re-articulations wouldn't correspond to the sort of *moral* re-articulation proposed by the objector. Rather, those re-articulations would involve

morally neutral concepts like *mass*, *spin*, and *charge*. The same thing *mutatis mutandis* would be true if we replaced the concept picked out by the term "law" with concepts belonging to biology, cognitive psychology, or even disciplines that *do* make at least some allowance for the existence of moral categories (e.g. psycho-analytic psychology[62]). A cognitive scientist would re-articulate the concept of law in terms of concepts like *algorithm*, *syntax*, and the like, and not concepts like *morality* and *fairness*. So if the concept of law is to be subjected to scientific or logical revision, there is no evidence that the revision would be the one proposed by the anti-positivist, and considerable evidence that it would *not* be that one.[63] [64] [65]

In conclusion, we have not begged the question against anti-positivism, or even the objector's extreme version of that doctrine, by taking for granted the truth of the historian's judgment that there really were laws preventing women from voting. There is independent evidence for the explanatory usefulness, even indispensability, of that concept and also for the explanatory sterility, even deleteriousness, of the objector's proposed replacement of that concept. There is thus no reason to believe that the objector's morally purified concept of law is anything other than an artifact of his own analysis.

§ Fuller[66] argued that laws must have a basis in morality. In connection with this, he denied that some of the so-called laws enacted by the Nazis really were laws. But he had *independent* grounds for this. For example, the "laws" he discusses were often never even revealed to the public. From a non-moralistic, purely descriptive viewpoint, this does distinguish those so-called laws from the things we *know* to be laws. Obviously those Nazi resolutions had a very different social role from, say, the law prohibiting insider trading or even the U.S. suffrage laws of the 19th century. After all,

the latter were publicly made, publicly known, and publicly discussed.

Fuller tried to show that laws must have a moral basis by showing that, whenever a so-called law fails to have any such basis, that fact correlates with its lacking some property that laws are *generally* expected to have, even by those who deny any moral component to the law. Fuller never denied that a so-called law was a real law *solely* on the grounds that it was immoral. In connection with this, Fuller never proposed a revision of historical data; he didn't deny that there were *laws* permitting slavery or prohibiting women from voting. He always focused on cases where the existence of an actual law was already in doubt, even to those who didn't share his moralistic conception of law.[67] So far as I know, no anti-positivist (apart from the ancients: Augustine, Aquinas…) advocates any revision of the historical data on the scale of that proposed by the objector.

Chapter 5 The moral background

We have considered some arguments that attempt to show that law must incorporate some kind of morality. We have seen that those arguments are fallacious. But my objective in this work is to show that, nonetheless, a moralistic conception of law is a defensible one.

The idea that laws and legal systems can be morally empty is rooted in a number of background misunderstandings. There is no way to give a compelling argument for a moralistic conception of law until those background misunderstandings are cleared away. This chapter will be concerned solely with eliminating those misunderstandings and, therewith,

the conceptual debris that hides the moral structure of law. No positive arguments concerning the relationship between law and morality will be given in this chapter.

Causal versus conceptual relations between law and morality

I must first make it clear what exactly I will be trying to show. I do not wish to argue that there is a *causally* necessary relationship between the law and morality. I wish to argue that there is purely *conceptual* relationship between them – one that is entirely internal to the concepts of law and morality, and doesn't require the mediation of any sort of contingent fact. It is *analytic* that laws incorporate a certain morality.

This brings us to a point that Hart makes. We are inclined to see as *conceptually* necessary features of the law what are in fact by-products of contingent, though perhaps deep, facts about human biology.[68] A comparison may help. There will inevitably be a certain justice to the practices of a Merrill Lynch or a Goldman-Sachs. The management probably wouldn't last long if it were totally iniquitous in its practices. But this isn't because there is some *internal* connection between the concepts of justice and money-brokerage. The connection is contingent. People are vulnerable to one another. The managing directors are people: they aren't bullet-proof; they have to sleep. It is therefore not an option for them to be totally unscrupulous in their relations with the stock-holders. But that *would* be an option if they were bullet-proof, had superhuman strength, and so on. So the connection between stock-brokerage and morality, though deeply rooted in facts about human nature, is nonetheless contingent.

For similar reasons, governments may not have the option of being limitlessly unscrupulous in their dealings with their subjects. Therefore the

laws issued by governments tend to reflect a certain morality. But perhaps this wouldn't have to be the case if the members of government were bullet-proof, didn't have to sleep, and so on. So given the biological nature of the human organism, it may be *causally* necessary that the law have points of connection with morality. But it doesn't follow that it is *conceptually* necessary that the law incorporate some kind of morality.

My purpose is to show that, even if we set aside all of the causal or otherwise contingent factors just discussed, we will find that all legal systems must incorporate a certain morality.

Incorporating, as opposed to merely being consistent with, a moral principle

We must now make a further distinction, one not made by Hart or Dworkin, that raises the bar even higher for anyone trying to show an internal connection between law and morality. A legal system can be *consistent* with some moral precept without *embodying* it.

Let's suppose that some omnipotent despot wants to make the legal system in his country as immoral or amoral as is consistent with his desire to stay in power. His legal system will not allow people to own weapons of mass destruction or to kill one another without restriction. Otherwise the despot would lose power in an instant (since his subjects could use those weapons to depose him). So that system will be consistent with at least some moral principles (e.g. by prohibiting possession of fire-arms, it will prevent people from slaughtering one another). But it would be false to say that the system *embodied* those principles. In this case, the connection between the despot's legal system and the moral principles in question seems to be a matter of macro-engineering – a mere by-product of the game-theoretic constraints under which the despot is operating. There doesn't seem to be any

sense in which that legal system embodies a recognition of those moral principles. For reasons of combinatorics, it just happens to be on the right side of them.

This does not mean that, in order for a legal system to embody morality, its creators must have the intention to promote some kind of morality. A legal system can be inherently moral even if all the people behind it are scoundrels who are acting entirely out of self-interest. Here is an illustration. X, Y, and Z are tired of fighting with one another, and wish to regulate their mutual relations by means of some legal code. Each of these parties acts entirely in his own interest. But none of the three has more power than either of the others. Further, each *fears* that, in due course, he will become the weak one. Consequently, each of the three wishes that the legal system in question be fair and not allow the strong to trample on the weak. So even though each of its creators is an utter cad, the legal system they jointly create is inherently moral: it isn't merely *consistent* with morality; its very purpose is to impose a certain morality. [69]

In light of the points made in the penultimate paragraph, suppose it turns out to be conceptually necessary that any legal system be *consistent* with certain ethical principles. It does not follow that legal systems must *incorporate* such principles, at least not in the sense intended by those like myself and Dworkin. Someone who is trying to show that there is necessarily an element of morality to the law surely does not mean that, owing to the necessities described by Boolean algebra, any legal system must be *consistent* with the law. What she means is that any legal system must incorporate a certain *respect* for ethical principles.

I wish to show that it is conceptually necessary that laws satisfy certain minimum moral requirements and that laws not only be *consistent with*, but *incorporate* an acknowledgement of, certain moral truths.

“Outweighed” does not mean “cancelled”

As Dworkin said, principles have “a dimension of weight”. A consequence is that principles can be outweighed, but never nullified. The principle that citizens should be allowed to pursue prosperity can be outweighed, but not cancelled, by the principle that business dealings should meet certain minimum standards of decency.

Dworkin was discussing legal principles, but his insight applies to moral principles as well. The principle that confidences should be honored can be outweighed, but never completely put out of commission. The same is true of all other moral principles.

I will argue that legal obligation is moral obligation. In cases where it seems otherwise, one’s moral obligation to obey the law is *outweighed* by a contrary moral obligation.

A bit of fiction may illuminate the points just made. Suppose you are Smith’s psychoanalyst. As Smith’s therapist, you have a duty to keep everything that he says in confidence. This duty is of a moral, and not merely a legal or professional, nature. It would be immoral of you to amuse people at cocktail parties with tales of Smith’s pathologies. Further, although Smith is a disturbed and even wretched person, he has always been decent to you. Moreover, he has been sincere and unremitting in his willingness to listen to your unflattering insights into his own character, and in his attempts to make the changes to his personality demanded by those insights. You have thus developed a deep therapeutic and personal rapport with Smith – even a certain respect for the man -- and your obligation to keep his confidence has its roots in the inexpressible understandings that constitute the most pure and valid of intimacies.

One day Smith tells you that he is going to kill all fifty people at his workplace.[70] You have two conflicting ethical duties. You can keep Smith's confidence, or you can tell the authorities about what Smith is planning to do. Clearly, the right thing to do is to alert the authorities. The lives of those fifty people are more important than your rapport with Smith.

But it doesn't follow that it suddenly becomes perfectly moral to violate Smith's trust in you. It does not. A violation of trust is always a bad thing. In this case, it is an especially bad thing, given the deep understandings that you and Smith share. But the alternative is worse. Betraying Smith's confidence was "the right thing to do" because it was the *least wrong* thing to do, not because it was a positive good.

In this case, we might say that it was "moral" of you to go to the authorities. But this is a loose way of speaking. In this context, "moral" really means: "comparatively moral" or "less immoral than the alternative." Something that is less immoral than the alternatives may be the "right", or perhaps the "just", thing to do. But nothing that is immoral ever becomes moral by being less immoral than the available alternatives.

It seems to me that terms like "right" and "just" often denote the concept of being *comparatively* moral. The psychoanalyst did the "right" thing, the "just" thing, because it was moral *on balance* to break Smith's confidence.

Just as nothing is moral in virtue of being *less* immoral than the alternatives, so nothing is immoral in virtue of being less moral than the alternatives. Suppose that, instead of going to the authorities, you decided to honor Smith's confidence in you. In this you are doing something moral: it is always moral *not* to betray a confidence, especially one having such deep roots. But honoring your commitment to Smith would have been far *less* moral than saving the lives of those fifty people. It was comparatively, but not absolutely, immoral.

Just as “right” and “just” sometimes denote the concept of being *comparatively* moral, so “wrong” and “unjust” sometimes denote the concept of being *comparatively* immoral. You were “wrong” to keep Smith’s confidence; i.e. doing so was far less moral than the alternative.

§ Let us map these points onto the topic of law. Suppose that there is a law L preventing therapists from divulging the intentions of their patients, even in extreme situations like the one just discussed. As we’ve just seen, a therapist would always have *some* moral obligation to obey L, even though he would sometimes have a stronger obligation to break that law. If we think of moral obligations in absolute terms, we end up seeing L as having, at most, a highly circumstantial moral value. In connection with this, we see one’s legal obligation to obey L as having, at most, a similarly circumstantial moral basis.

But if we keep in mind the fact that obligations have a dimension of weight, we can see L, and also our legal obligation to obey it, as *categorically* having moral value. In cases where it would be “wrong” to obey L, that moral value is not altogether absent, but is merely *outweighed* by a greater moral value. In due course, I will argue that there is always *a* moral obligation to obey the law and that, when it seems otherwise, that is because that moral obligation has been eclipsed – not extirpated – by a contrary moral obligation.

§ The points just made are closely related to my larger objectives in this work. I will argue that *as a rule* laws protect non-trivial moral goods. (As we will discuss, “as a rule” doesn’t always mean “100% of the time.”) There is thus a *principled* relationship between a thing’s being a law and its protecting a non-trivial moral good.

There are, we will see, laws that protect no moral good at all. But they are in the nature of flukes – degenerate or pathological cases of laws. For this reason, we will argue, their existence does not demand that we deny the existence of a principled relationship between law and morality.

Leaving aside degenerate cases, whenever there is a law that *seems* to do no good, what is really going on is that the good done by that law is outweighed, not cancelled, by a good with which that law is incompatible. In connection with this, one's moral obligation to obey that law is outweighed, but not extirpated, by a contrary moral obligation.

I will argue that, given *any* law at all, there is always *a* moral obligation to obey it. So there is *a* moral obligation to obey even the degenerate laws that do no moral good at all. (It is not yet possible to state my reasons for this, even in outline form.) This is because, I will argue, legal obligation *is* a kind of moral obligation.

At the same time, we will see that my *legal* obligation not to (for example) kill Fred does not coincide with my *moral* obligation not to do so. My legal obligation not to kill Fred is *a* moral obligation; but it is not *identical with* my moral obligation not to kill Fred. We are dealing with two different moral obligations that are discharged in the same way.

In relation to this, we will see that, leaving aside degenerate cases of law, one always has two, quite distinct moral obligations to obey the law. Since any non-degenerate law *does* protect a moral good, one *ipso facto* has *a* (defeasible) moral obligation to obey it – *in addition to* the moral obligation that coincides with one's legal obligation. In the degenerate case of a law that does no non-trivial moral good at all, one has exactly *one* moral obligation to obey the law; and that moral obligation coincides with one's legal obligation. So, as a rule (though not 100% of the time), there are *two* moral foundations to one's legal obligation to obey the law.

In the frequent cases where it is appropriate to break the law, that is because the moral obligation – or, more often, the obligations -- to obey it have been outweighed, but not cancelled, by some contrary moral requirement.

I hope that, given only these obscure and threadbare remarks, it is already starting to become clear how a non-binary conception of morality may well be useful in helping us see the moral composition of law.

"Immoral" versus "condemnable"

Let us now clear up another confusion that obscures moral structure of law. In some situations, it is not logically possible to avoid doing something that is morally less than perfect. We just described one: it is not logically possible for the therapist to keep his patient's confidence and to alert the authorities. So, as a matter of logic, it is necessary to do something that is morally sub-standard – that is *immoral.* But this doesn't mean that, in virtue of doing that thing, the psychiatrist is *himself* acting immorally. It would be wrong to say that, under those circumstances, the psychiatrist had acted immorally – even though the psychiatrist had done something that was, in absolute terms, immoral.

Like practically every key term in ethical discourse, the word "immoral" is highly ambiguous. Sometimes it denotes the property having negative moral value, and sometimes it denotes the very different property of deserving moral *condemnation.* The word "moral" is characterized by a corresponding ambiguity. Given only that someone commits an immoral act, it doesn't follow that she is deserving of moral condemnation: given the circumstances, it might have been logically impossible to avoid doing something immoral. When we say that a *person* has acted "immorally", we

don’t usually mean that she has done something immoral: we mean that she has done something morally *condemnable.* The psychiatrist who alerts the authorities is doing the best thing possible under the circumstances. Therefore he is not deserving of moral condemnation, even though the thing he did was, for the reasons we discussed earlier, morally sub-standard.[71] By contrast, somebody who has plenty of money of his own, but still steals money from his roommate, *has* acted “immorally”, because he has done something deserving of moral *condemnation*; unlike the psychiatrist, he did *not* do the best thing available under the circumstances.

As I hope to show, the moral basis of law is easier to see if we keep in mind the difference between having negative moral value and being deserving of condemnation. It is also helpful to keep in mind the correlative difference between having positive moral value and being deserving of commendation. There are many laws that permit condemnable behavior. But it doesn’t follow, and as a rule it is not the case, that such laws protect nothing of moral value.

Chapter 6 Hart on the separation of law and morality

In a previous chapter, we considered, Dworkin’s brilliant attempt to find a categorical, as opposed to contingent, connection between law and morality. We found that Dworkin’s attempt was less than probative.

There are two ways to respond to this. The first is to say that a moralistic conception of law is wrong and that any attempt to argue otherwise is therefore doomed to failure. The second is to say that a moralistic conception of the law is correct, but that Dworkin’s arguments for such a

position must be supplemented.

I take the second position. H.L.A. Hart takes the first. Hart's position is that laws, and legal systems, needn't embody any morality. The essence of Hart's viewpoint is found in his classic paper "The Relationship between Law and Morality."[72] That paper is no doubt one of the most significant attempts to undermine a moralistic conception of law. For that reason alone, a careful examination of it is in order.

But there is another reason to examine Hart's paper. Although, in my view, Hart's non-moralistic conception of the law is not correct, Hart cogently argues against a number of plausible attempts to identify the moral basis of law. Hart thus helps to clarify what the actual moral basis of law might be.

§ The basic structure of Hart's paper is this. A number of arguments have been given for thinking that law and morality must overlap. These arguments are all failures.

Of course, given only that certain arguments *for* a conclusion fail, it doesn't follow that the conclusion is false. But Hart's implication is that, indeed, law and morality don't necessary overlap (though, for various reasons, they may be likely to). The hidden premise seems to be that, if it were at all possible to prove that laws must have a moral component, it would be by way of the arguments that Hart attempts to shoot down (and, in some cases, *does* succeed in shooting down). By way of anticipation, that hidden premise is false. Let us now go through Hart's article.

§ The first argument that Hart considers is simple. Laws give rights. Because there is a law prohibiting murder, I have a legal right not to be murdered. Because there is a law prohibiting theft, I have a legal right not to be stolen

from. So the connection between law and morality is obvious: laws are protections of moral rights.

In response to this, Hart makes a compelling point. Legal rights are not always *moral* rights. Romans had the legal right to own and beat slaves. Surely these rights are not moral rights.

§ Hart now turns to another argument concerning the relationship between morality and law. No law will be entirely free from vagueness or ambiguity. Consider a law forbidding the use of motor vehicles in the park. Do jet-powered roller skates count? What about toy cars? Suppose that one day Jones uses jet-powered roller skates in the park. Judge Smith must decide whether Jones is in violation of the aforementioned law. Since the law itself is indeterminate, Judge Smith must find some way to "precisify" it (to use a term made popular in recent work on vagueness). But he must find a *principled* way of doing so: he would be abrogating his duties as a judge if he let an *arbitrary* precisification of the law lead to somebody's being fined or jailed. Hence law and principle necessarily intersect.

Here is a similar argument. All language is vague (outside of number theory, formal logic, and a few other empirically sterile domains). Thus, all laws are necessarily vague – or, at least, are necessarily *stated* vaguely.[73] This requires that the application of law cannot generally be mechanical, and must be principle-driven. That, in turn, requires that law and principle convene. Thus moral principle *is* necessarily a part of the law.

Hart's criticism of this is compelling. Not all principles are *moral* principles. The principle at work in the judge's precisification might not be a *moral* principle. Instead, it might be: *whenever the rules are unclear, I must precisify them in the way that best serves the fuehrer.* So, while the vagueness of legal language may indeed require that law convene with *some* kind of

principle, the principle needn't be of a moral nature.

§ Let us discuss the next argument that Hart considers. Before the Nazi atrocities occurred, a German jurist by the name of Gustav Radbruch held that law is one thing, morality another. The Nazis came along and enacted unspeakably wicked "laws." Radbruch felt that, in some way, the Nazis were helped in this by a kind of positivist mentality towards the law. If it had generally been accepted that genuine laws could not be utterly wicked, then the Nazis' directives would have been seen for what they were -- the manipulations of lying thugs – and would never have been dignified as "laws".

Hart points out that this is not so much an argument as it is an exhortation to change the way we use words like "law" (and "Recht", and so on). It does seem to be an empirical fact that the Nazis enacted evil laws.

There is another point to make in response to Radbruch (one that Hart is not so explicit about). Suppose we make it a matter of definition that "laws" are good. Despots are still going to *describe* the directives they issue as "laws". In fact, they would welcome the positive insinuation carried by that term. Thus, thanks to the proposed change in nomenclature, there would be a presumption that the "laws" issued by a Hitler or Stalin were good. But if we *don't* reserve the term "law" for morally good directives, then nothing[74] of moral consequence follows from the fact that a despot issues some "law": the Hitlers and Stalins of the world are deprived of that verbal reinforcement.

§ Hart considers two other arguments to the effect that law must incorporate morality. First, the application of laws must be more or less uniform. This is practically a truism, as the following story illustrates. Judge Smith gives person x a fifty year sentence for breaking law L; then Smith

gives y a $10 fine for breaking L; and, finally, Smith awards a medal of honor to z for breaking L. All other judges behave in a similarly erratic manner with regard to L.

Under this circumstance, there are various judges making various decisions that are in no way regulated and that, consequently, cannot be regarded as applications of any law or rule.[75] So, under this circumstance, L simply doesn't exist. This shows that, where there isn't a certain uniformity in the application of a law, there is no law to begin with. So where there is law, then a principle of *fairness* – like treatment for like offences – must be observed.[76]

Hart's criticism of this is powerful. Uniformity in the application of law doesn't necessarily embody any morality. Suppose the law says that everybody must wear a cotton uniform. Citizen Smith is violently allergic to cotton. But Judge Jones, ever a stickler for the law, forces Smith to wear the cotton uniform. In applying the law with such unremitting uniformity, Judge Jones is doing something highly immoral. So the uniformity demanded by law doesn't necessarily embody any morality.[77]

§ Hart now considers a more significant argument. Given only that there are individual laws that lack moral value, it doesn't follow that the same must be true of legal systems as a whole. Further, it seems that, as a whole, any legal system *must* incorporate a certain morality. It is hard to see how there could be *any* law in a context where there were no laws forbidding murder, theft, etc. If those iniquities were unrestricted, the same would be true of everything else. It is hard to imagine a society where murder and theft were unpunished but where parking laws were enforced. The mechanisms that enable the latter to be enforced – e.g. the health of parking enforcement personnel – would no longer be intact. So if there are to be *any* laws, then

certain basic moral rights – e.g. the right to life – must be protected. Thus, the roots of any legal system must lie in prohibitions which coincide with moral prohibitions. So even though individual laws may be evil, legal systems as a whole must have a certain connection to morality. [78]

Hart's response is compelling. We must not confuse causal or natural[79] necessity with conceptual necessity. If human beings had impenetrable shells, and were thus physically invulnerable to one another, legal systems *wouldn't* have to forbid mutual assault.[80] (A legal system needn't forbid what cannot be done.) Given *any* prohibition that seems to be intrinsic to the law, creative use of science-fiction can show us a legal system that does not incorporate that prohibition. This suggests that the aforementioned connection between morality and legality is contingent, not analytic.

We can use a point that we made earlier to reinforce this powerful argument of Hart's. Even if there is an *analytic* connection between the existence of a legal system and the prohibition of certain wrongs, that legal system doesn't necessarily *incorporate* that principle. Speaking in terms of combinatorics, it would be hard to imagine a social arrangement that was completely opposed to every moral interest of *everyone* involved. But it doesn't follow that any such an arrangement will embody a *respect* for morality or, therefore, that it *incorporates* any sort of morality.

Critiquing Hart's critique

Let's suppose that all of Hart's criticisms are decisive. As we've seen, that doesn't entail that there is no moral component to the law, but only that various attempts to prove as much have failed. So Hart hasn't given us any positive reason to believe that laws needn't have a moral component.

Of course, if Hart's criticisms were decisive, that would certainly *confirm* his non-moralistic conception of law. But although some of Hart's criticisms are decisive, some of them are, I believe, invalid. When we expose the non-sequiturs involved, the initial appeal of Hart's positivism quickly fades.

Hart plausibly says that legal rights are not always moral rights. That statement is ambiguous. Something that is *a* good or right thing may not be the best thing under the circumstances – it may even be the *worst*. We've seen many examples of this. If the psychotherapist were to *not* go to the authorities, he would be doing something very wrong. It could truly be said that the psychotherapist does not have *the right* to withhold knowledge of the patient's murderous plan from the police. That is true in the sense that, under the circumstances, the alternative is far *more* right, but not in the sense that keeping the patient's confidence is without moral merit. When we talk about what people have "the right" to do, we often have in mind what would be the *most* moral – or the least immoral – possible course of action. But the *most* moral (or least immoral) path is seldom the only one that has *any* moral value.

By virtue of owning slaves, Roman citizen Plautus enjoys many goods of which he would otherwise be deprived. Some are narrowly hedonistic, but some are valuable even by the most non-hedonistic of standards. Because he owns slaves, Plautus has the leisure to reflect on philosophical problems and to write deep treatises. *By themselves*, such things are good. By itself, it is *right* that Plautus has a chance to develop his gifts as a thinker, and Plautus has *a right* to develop his abilities. Everybody surely has a moral right to flourish. The evil Roman laws relating to slavery protect Plautus' right to flourish. Those laws were evil *not* because they didn't protect *anybody's* rights, but the rights they obliterated outweighed those they protected.

Given only that some law does a great injustice, it doesn't follow that it protects *no* moral right. It could be that the right it protects is outweighed by an incompatible right. This is the case with the laws just discussed in connection with Plautus. Hart's argument falsely presupposes that the concept of a right is an absolute, as opposed to a comparative, notion.

§ But given only what we've just said, Hart has no reason to concede defeat, since he has a powerful counter-response at his disposal. Perhaps the laws we just spoke of in connection with Plautus do some good (even though, on balance, they do a massive amount of bad). But it doesn't follow that *all* laws do good. Couldn't we imagine a *hypothetical* law that didn't do anyone any good at all?

Suppose that, while in a trance, the King (or whoever or whatever the relevant law-making entity happens to be) proclaims that no one is allowed to wear green shirts indoors. Though made in a state of near-unconsciousness, the King's proclamation precipitates the governmental apparatus that ordinarily converts the King's proclamations into law. So it becomes a law, at least for a while, that no one is allowed to wear green shirts indoors. Let L be this law.

L doesn't serve anyone's interests. It doesn't even gratify some petty whim on the part of the King. So Hart does seem to be right that there are, or at least could be, laws that have *no* positive moral content at all – even after we allow for the idea that any gratification is a good, albeit one that may be grossly outweighed by some incompatible good.[81]

A bare outline of some contentions that will play an important part in this work

Much of the remainder of this work is going to be spent responding to the point just made. So during the next few pages, I would like to outline the viewpoint I will defend in regards to that point. Please bear in mind that the next few pages are not my argument proper, and are merely meant to indicate the general conceits driving my analysis.

First of all, I grant that L *doesn't* have any moral value. I am not going to take the revisionist approach of denying that L is a law. So I concede Hart's point that there are *some* laws (hypothetical, if not actual) that protect no moral good at all. But there are a few reasons why this does not warrant acceptance of Hart's conception of law.

As a matter of conceptual necessity, there are laws only where there is government. To understand law, we must therefore understand government. As we will see, a government is, almost by definition, something that protects and serves the interests of those under its dominion. This suggests that legal system exists to help the government discharge its protective functions. It therefore suggests that something is a law in virtue of its protecting, or being supposed to protect, a moral good.

Obviously there are evil governments. But an evil government, I will argue, is one that *abuses* its role as protector and provider. It is not one that altogether fails to protect and provide.[82] Because people depend on government for fundamental moral protections, people are very much at its mercy. Governmental evil exploits, and thus presupposes, this sort of dependence. So the existence of governmental evil presupposes that, at bottom, governments are in the business of providing protections and services.

What we just said about governments is true of legal systems. Just as there are evil governments, so there are evil legal systems. In fact, few things do as much harm as a corrupt legal system. But what makes it possible for a

legal system to be so capable of wrong-doing is precisely that it has such a vital role in the protection of people's rights. A slight defect in an organ that you depend on for your well-being – your heart, brain, lungs, or liver – will have consequences far more catastrophic than a massive defect in a finger-nail or thigh-muscle. It is precisely *because* legal systems provide such important moral protections that, when they become pathologized, they do so much harm.

It is true that a dictator's most despicable machinations may involve the legal system. But this is precisely *because* people depend on it so much for important moral protections, and are thus so much at its mercy. It is thus because the legal system has such deep roots in morality that it may become an especially useful vehicle for an unscrupulous official's machinations.

These remarks expose what I believe to be a systemic defect in positivist thought. If a hypothesis is counter-exampled by pathological or degenerate cases, it doesn't necessarily follow that it is wrong. In fact, if the *only* counter-examples to a hypothesis are pathological cases, that may actually redound to its credit. There is no denying that *as a rule* livers are indispensable to our well-being, and that *as a rule* they do good. Of course, a disease-ridden liver needn't do its owner much good, and may even lead to a speedy death. But, given only that fact, it would be absurd to reject the thesis that, as a rule, livers perform important metabolic functions. The fact that *diseased* livers do so much harm almost (though not quite[83]) confirms the fact that, as a rule, livers do great metabolic good. After all, a diseased *fingernail* does virtually no harm – and that is plainly related to the fact that a non-diseased fingernail does so little good (as least as compared with a liver or heart).

Given only that many legal systems do harm, it doesn't follow that *as a rule* legal systems don't do great good. It is a possibility that the legal

systems that do great harm are pathological cases, and that we must therefore question their fitness as possible counter-instances to moralistic conceptions of law. What we just said about legal systems applies, so I will argue, to individual laws.

Even if we set aside the points just made, we will find that, in absolute terms, *every* legal system does a great deal of good. In *comparative* terms, some legal systems do a great deal of harm: the Soviet legal system did far less good (per capita) than the Swedish system. So in relative terms the Soviet system did little good. But in absolute terms, the Soviet system did much good (or so I will argue). The view that governments can be *fundamentally* lacking in moral value involves a tendency to regard comparative truths, like the one just discussed, in absolute terms.

We tend to think of principled relations as being expressed by constant conjunctions. There is a principled, as opposed to haphazard, connection between fire and heat because there is heat wherever there is fire. A moment ago we saw one reason why a principled relation needn't be expressed by a constant conjunction. Now I would like to discuss another (related) reason for this.

There are many planes that don't fly. But it would be absurd to infer from this that there isn't a principled, internal relation between the property of being a plane and the property of being something that flies. The concept of a plane has a functional or teleological component. Where such concepts are concerned, principled relations between properties are not necessarily, or even usually, expressed as constant conjunctions of properties.

If I am not mistaken, this fact brings to light another reason why the existence of morally hollow laws like L doesn't necessarily warrant an acceptance of Hart's positivism. Sometimes laws don't do any good. Sometimes they do great harm. Sometimes they do great harm because that is

precisely what their creators *intended.* Despite all this, it is clear that laws are *supposed* to do good. There are two different ways to deal with these facts.

Here is the first way. We can say that, to the extent that we expect laws to do good, our expectation doesn't correspond to anything essential to the concept of law. Laws are supposed to do good only in the sense that *we* suppose them to do good – perhaps because of some blunder on our part.

Here is the other way. We can say that a law that doesn't do good is like a plane that doesn't fly. Obviously politicians create laws that do harm. Sometimes these politicians even *intend* for them to do harm. But such cynical manipulations *exploit* the public's correct supposition that a law is something that is supposed to do good. To the extent that the public is guilty of a blunder, it concerns the psychology of the law-maker, and not the concept of law.

I believe that the second viewpoint is the right one, and will spend a great deal of time arguing for it.[84] Hart holds the first viewpoint. In Hart's view, when a legal system incorporates any kind of morality, that doesn't reflect anything essential to the concept of law. If the line of thought just presented is correct, Hart's viewpoint is quite misguided. The *function* of a legal system is to protect. It is no more a matter of happenstance that, in actuality, legal systems generally do good than it is a matter of happenstance that, in actuality, most hearts keep people alive and that, in actuality, most planes fly.

The points just made about principles can, I think, be understood in terms of Dworkin's distinction between rules and principles. It is because principles have a "dimension of weight", as Dworkin puts it, that the existence of an F that is not a G doesn't mean that there isn't a principled relationship between a thing's being an F and its being a G. Since principled relationships have weight, they are expressed by tendencies, not by perfect

concomitances. A thing's *tendency* to have G, given that it has F, can be over-ridden, even though that tendency is extremely strong and even though, consequently, the connection between the properties in question is deep. As we've seen, the positivist infers from the existence of wicked laws that there is no principled relationship between law and morality. But this assumes falsely that principles, like rules, are binary. So we find a reason distinct from that cited by Dworkin why the positivist involves a failure to register the logical distinction between rules and principles.[85]

§ Let us end this chapter by making a point about L – the fictitious law that prohibits the wearing of green shirts indoors. On the face of it, L seems to be in the nature of a fluke: a product, not of the standard process, but of a pathological variant thereof. For reasons already discussed, and to be elaborated later on, this makes it questionable whether L is really a counter-instance to the thesis that *as a rule* laws protect moral goods.

At the same time, there is no denying that L *is* a genuine counter-instance to the thesis that *something must protect a moral good to be a law.* As we will see, the existence of laws like L makes it considerably more difficult than it would otherwise be to delineate the moral structure of legal obligation. By the same token, if we *dismiss* laws like L, and regard them as not *really* being laws, then we will end up with a revisionist and tragically simplistic conception of law. For this reason, even though the existence of morally empty laws like L is a stumbling block for any moralistic analysis of law, *especially* my own, I will not even entertain the notion that L is not a genuine law.

Chapter 7 The concept of government as a basis for a moralistic conception of law

We've seen some reason to believe that not every individual law does good. Nonetheless, I will now attempt to show, first, that all *legal systems* do good (even though many do unspeakable harm) and, second, that there is a sense in which all individual laws *incorporate* a certain morality.

Ironically, the best way to begin the process of establishing these two claims is to consider a viewpoint that is as much opposed to them as possible. I am referring to the so-called *command-theory* of law. Some historical background is needed.

Some thinkers, including Thomas Aquinas, have held that genuine laws cannot be immoral.[86] According to this view-point, the so-called "laws" that prevailed in the Jim Crow south, or in Nazi Germany, were either not really laws at all or they were in fact moral.

The problem with such a view is that is (or seems to be) plainly false. Bentham described it as "nonsense on stilts"[87], and Austin described it as "stark nonsense."[88] The existence of immoral laws would seem to be a simple fact of life.[89] Earlier, in chapter 4, we gave a less question-begging argument as to why the view in question cannot be accepted.

Wishing to distance himself as much as possible from such a view, John Austin proposed a very different theory. Laws are commands. In any society, some entity has the power. Sometimes that entity is a single person; sometimes it is an organization. Laws are simply orders issued by that entity.

In Austin's view, the difference between the orders of a gunman and the laws that govern a nation is only one of degree. The authority of the gunman is transient and limited in scope. It applies only to the few unfortunates in the convenience store, and lasts at most a few hours. That

alone is why his orders do not qualify as laws. The authority of a Stalin applies to millions of people, and lasts for decades. That alone is why Stalin's demands do qualify as laws.[90]

§ There are some well-documented problems with the command-theory of law. First, that theory seems to presuppose that societies have a very simple top-down structure. On top, there is a big man who is answerable to nobody -- and then there is everybody else. But this picture applies to few societies, if any. Power is always distributed over a number of people. So the notion of a command has to be stretched if the command-theory is to be applied to any actual society. It would have to be stretched beyond recognition to apply to democratic systems where, from a strictly legal standpoint, nobody commands anybody.[91]

There is another problem with the command-theory. Some laws serve the purpose of giving people certain powers. There are laws saying that certain kinds of relationships count as marriages. These laws don't have the form "do this" or "don't do that." They have the form: "*if* you and your partner wish to have certain privileges, then do such and such."[92] A similar point applies to the laws governing the making of contracts, the writing of wills, and the formation of businesses. Where laws of this kind are concerned, there is no threat and there is no command. So the command-theory appears not to apply.[93]

There is a related point. Certain kinds of activities presuppose the existence of law. There couldn't be marriages or business-partnerships where there were no laws. This isn't because lawlessness would *threaten* the existence of such things. It isn't that, but for the presence of law, such things could exist only for a short amount of time. A business-partnership or marriage is *itself* a legal arrangement. There don't have to be laws for people

to fall in love, or to have profitable arrangements with one another. But a marriage or business-partnership is, by definition, the conferral of a certain *legal* status on what would otherwise be mere camaraderie.[94]

The very existence of money and, more generally, of property seems to presuppose the existence of law. It isn't that, if there were no laws, people wouldn't be able to hold onto their money or land. It is that for something to be money or property *is* for it to have a certain legal status. Technically, a piece of money *is* a legal certificate: it is a legal entitlement to a certain share of that nation's wealth.

A related point (of inestimable importance, as we will see) is that laws give us rights and powers, whereas the orders of a gunman do not. Thanks to the law, I have the *right* not to be attacked by those who wish me ill. A gunman's threats do not bestow any rights on his hostages.

To sum up, the command-theory seems to presuppose that laws are only there to crack the whip when people step out of line. But laws sometimes have a more constructive function, as we've just seen.

There is a third problem with Austin's position. It seems to be inconsistent with the attitude we have towards the law. As Hart puts it, the command-theory doesn't distinguish between being *obliged* to do something, on the one hand, and being *obligated* to do it, on the other. If a gunman asks me for my money, I am obliged, but not *obligated*, to give it to him. By contrast, if *the law* requires me to pay taxes, then I am (supposedly) *obligated* to do so.[95] With this distinction in mind, Hart writes: "Law is surely not the gunman situation writ large, and legal order is surely not to be thus simply identified with compulsion."[96]

§ As we've seen, one of the stumbling blocks for the command-theory is the existence of laws which have a *constructive* function, e.g. laws which

allow people to form business partnerships or which give special benefits to those who have a certain number of children. What do these laws have in common with the laws that prohibit theft and murder?

Both give (or protect) *rights*. (They protect moral rights by giving legal rights.) The law forbidding murder protects an existing moral right that people have not to be murdered. The law forbidding theft protects a moral right not to be stolen from. The law forbidding speeding protects a moral right to travel freely and safely. These points suggest that all laws protect *moral* rights. This in turn suggests that:

® A law is a protection of a moral right.

There are many reasons why ® is wrong. But, I believe, ® is right in spirit; and it points us in the right direction. Before showing why ® is wrong, I propose to show why it is not *wholly* wrong. This will make it possible to produce a more adequate analysis of the concept of law.

§ We must begin by making a distinction. It is one thing to have a legal obligation to act a certain way. It is quite another to be causally required by a legal system to act in a certain way.

Let us illustrate this with a story. Martians colonize the Earth. There are laws regulating how Martians interact with one another. (These parallel the laws that, in actuality, regulate human interaction. For example, Martians are not allowed to kill one another, except in self-defense). But the Martian laws do not give human beings any protections at all. Any Martian can do anything to any person, and have the full protection of the Martian legal system. If a human resists a Martian in any way, he will instantly be vaporized by a Martian execution squad. Further, Martian law doesn't

regulate relations among people: so far as Martian law is concerned, any person can do anything to any other person. Martian law gives no protections to people.

In this scenario, even if the Martian and human communities are geographically coincident, we are not dealing with a *legally* unified community. The Martian legal system has dire consequences for human beings, but it doesn't impose *legal* obligations on them. For example, Martian law requires certain Martians (members of an analogue of our animal control) to kill any human who is anywhere near Martian eateries. Thus, Martian law being what it is, a *causal* condition for a human being's remaining alive is that she not be anywhere near a Martian eatery. But Martian law no more imposes a *legal obligation* on humans to stay away from such places than human law imposes a *legal obligation* on rats to stay away from delicatessens. To use Hart's terminology, Martian law *obliges*, but does not *obligate*, people to behave in certain ways.

An extension of the previous story will clarify my meaning. One day, Martian Gloxo decides to beat up human being Smith. Smith fights back. The Martian then says: "you have a legal *obligation* not to resist in any way; therefore you are breaking the law." It is obvious that the human doesn't have a *moral* obligation to let the Martian beat him. But does Smith have even a narrowly *legal* obligation to permit this?

No. Martian law being what it is, Smith will suffer dire consequences for defending himself against Gloxo. But it would be wrong to say that Smith had a *legal obligation* to acquiesce to Gloxo. Martian laws would seem not to define any legal obligations for Smith, even though they *oblige* him to act in certain ways.

By giving *no* protections to humans, Martian law fails to *recognize* humans. Consequently, humans have no obligations under Martian law.

If a legal system does not *recognize* you, then you have no obligations under it (though it may *causally oblige* you to act in certain ways). If a legal system does not offer you any protections at all, then it does not recognize you. Therefore, if a legal system does not offer you any protections at all, you have no obligations under it.

§ Now suppose that Martian law changes. The Martian legal system is still biased in the worst possible way against human beings. Human beings cannot own land (on a commercially significant scale). Nor can they operate motor vehicles, vote for officials, and so on. But now human beings *do* have certain protections from the Martians. Martians *will* be punished for attempting to take a human's belongings or for attempting to cause physical harm to a human. Further, Martian law prevents humans from attacking or otherwise violating other humans.

Under these changed circumstances, does Martian law impose legal obligations on humans? Suppose that human being Brown sees human being Jones attack a Martian police officer. Does Brown have a *legal* obligation to report Jones to the Martian authorities? The answer seems to be: yes. There is plainly a sense in which Brown is *bound by law* to report Jones' behavior to the authorities.

This is not to say that Brown has a *moral* obligation to report Jones to the authorities, or that the law imposing the aforementioned obligation on Brown is a good or fair one. But it does seem clear that Brown is under some kind of legal obligation to report Smith.

What has changed? In the first scenario, the Martian legal system didn't give *any* protections to humans. This is no longer the case in the second scenario. As previously noted, if a legal system doesn't protect your rights to some degree or other, then it does not impose obligations on you. This does

not mean that, unless a legal system protects your rights, you have no *moral* obligation to uphold any legal obligations it may impose on you. It means that it doesn't impose any legal obligations on you to begin with.

§ To deepen our understanding of the nature of legal obligation, it will help if we momentarily consider the years following the Civil War, but preceding the reforms brought on by the Civil Rights Movement. During those years, the law was grossly unfair to African-Americans. At the same time, African Americans did have *some* legal protections. It was not legal for any American citizen – whether of African or European ancestry - to kill an African American or take his property.

Because African Americans depended on this legal system for these vital protections, it imposed obligations on them. An African American who witnessed a crime had a legal (not to say a moral) obligation to report it to the authorities. African Americans had a legal (not to say a moral) obligation to pay taxes.

This brings us to the root-problem with the command-theory. For there to be a legal obligation, even of the most morally hollow or perverse kind, a directive has to be a part of a system which protects at least *some* of one's rights. The gunman's orders are not a part of a system that protects the hostages' rights. That is one reason (not the only one, as we will see) why they don't qualify as laws.

My position can be embodied in what might initially seem to be nothing more than glib rhetoric. We are familiar with the slogan "no taxation without representation." I would generalize this by saying: "no legal obligation without representation." Suppose that government G provides no services at all for you. It doesn't let you travel on the roads it has built; it doesn't let you take advantage of its health-care and plumbing facilities; it provides you with

no legal protections at all (anyone can beat you or steal from you with impunity); you are totally shut out.

Under these circumstances, you surely don't have any sort of obligation, ethical or legal, to pay taxes to G. For the same reasons, you don't have *any* legal obligations to G: you are not *subject* to its laws, even though those laws may impose severe causal constraints on how you may live your life. I do not mean that the legal obligations that you *do* have to G are morally empty, but rather that you don't have any such obligations to begin with.

G's failure to grant you any protections amounts to a failure on G's part to *recognize you* at all. Since G does not even recognize your existence, you have no legal obligations towards it. This is not an ethical, but a narrowly logical, point. X cannot have legal obligations to Y if Y doesn't recognize X's existence. There is no legal recognition without an extension of certain protections. (This, I think, is a truism. Legal recognition *is* the conferral of protections.) So it is a truism, or at least follows from truisms, that you have no legal obligations to a government unless that government gives you certain basic protections – unless you depend on it for protection of your rights.

I would like to consider a possible objection to one of the points just made (I will put it in the mouth of an imaginary interlocutor):

You say that legal recognition of X *is* the conferral of legal protection on X. But that isn't true. Suppose that legal system L has a law saying that everyone must be evil towards Smith, and that under no circumstances is one to give any kind of aid or comfort to Smith. In that case, L *recognizes* Smith, but it grants him no protections.

When we say that some legal system "recognizes" X, we don't mean that it *acknowledges* X's existence. The term "recognition", as it occurs in the expression "legal recognition", is a technical term. Talmudic and Koranic law forbid the consumption of pork. So those systems of law *recognize* the existence of pigs, in the non-technical sense of "recognize." But pigs have no legal standing under those systems of law. The nation of Saudi Arabia doesn't officially recognize the state of Israel. But there are Saudi laws that refer to Israel. When we say that legal system L doesn't "recognize" X's existence, what we mean is precisely that X doesn't have any standing under L. That means that X doesn't have any rights or protections under L. So it is a truism that legal recognition, in the relevant sense, consists in the conferral of protections.

§ Now let us consider the sort of law that *most* resembles a gunman's orders – the kind of laws that John Austin had in mind. Stalin issues a general order to the following effect. If one knows of a case where somebody X spoke ill of Stalin's leadership, then one has an obligation to denounce X to the authorities. It doesn't matter whether X is a close relative or a best friend. It is guaranteed that, once denounced, X will be killed or at least sent to a labor camp. Failure to denounce X is tantamount to agreeing with X's blasphemy, and is itself punishable by execution or forced labor.

This is about as evil a law as we will ever find. But it is still *a* law. Why? Not merely because it was an *order* issued by somebody who had the power to enforce obedience, but because it was issued by an entity (the Soviet government) that, while being evil, guarantees innumerable important rights.

Soviet citizens depended on the government for many protections: protection from rape, theft, murder, and the like. Of course, sometimes the

Soviet Government was itself guilty of such atrocities. But that doesn't change the fact that Soviet citizens depended on the government for those protections, and usually received them. That is the reason (or part of the reason) why the evil laws issued by that government defined legal obligations for Soviet citizens. A gunman's hostages don't receive protections from him. That is the reason (or part of the reason) why those demands don't correspond to legal obligations.

Before we move on, we must make one point explicit. "Law" is a relational term. American laws don't apply to everybody; they don't apply to people in Saudi Arabia or Denmark (except, possibly, servicemen and diplomats who are stationed there). American laws are laws *for* certain people, but not others. There are no such things as laws *simpliciter*, only laws *for* such and such people. So the term to be defined is not "law", but "law for X."

§ The principle that falls out of our discussion seems to be this:

(*) A directive D issued by X defines a legal obligation for Y if Y depends for the protection of his rights on X.

The directives issued by the United States government define legal obligations for me because I depend for the protection of my moral rights on that government.

But as it stands, (*) is incorrect. Suppose that you live in a rough neighborhood, and your older brother Bob protects you from the local hoodlums. But for Bob, you would either be dead or have to suffer constant torment. In exchange for this protection, Bob makes certain reasonable

demands of you. You cannot play your stereo late at night, and you must do certain chores. Even under these circumstances, it doesn't seem that the demands that Bob makes of you have the force of law. Depending on what the exact nature of a given demand is, it might be dishonorable or outright immoral of you not to comply with it. There might even be a severe penalty for not complying with it, e.g. Bob might withdraw his protection. But by no stretch would such a demand constitute a law.

It is, I think, pretty clear what the problem is. Laws are issued by the government. Bob is not the government; and in making demands of you, Bob is not acting on behalf of the government. That is why his demands do not have the force of law.

Thus (*) is false. For something to be a law for somebody, it is *necessary* but not *sufficient* that it be a directive issued by something that protects at least some of that person's moral rights. What are the remaining conditions?

In this section, I will go through a number of different answers to the question just posed. These answers will be referred to as (L1), (L2), and so on. Each of these will be false, except (I will argue) for (L5). So (L5) is my official answer to the question just asked.

§ We are now in a position to hazard a first guess (albeit one we will come to reject) as to what laws are:

(L1) L is a law for X, i.e. it defines a legal obligation for X, provided that two conditions are met. First, L is a protection of a moral right had by somebody or other (not necessarily X himself). Second, L is issued by the *government* on which X depends for protection of his rights.

The italicized word distinguishes (L1) from the erroneous (*). But let us say a little more about the raison d'être for (L1). The racist laws of the Jim Crow South defined legal (not moral) obligations for African Americans because they satisfied two conditions. First, those laws protected *somebody's* freedoms – they protected the powers of non-African American southerners. Second, evil though those laws were, African Americans depended on the government that issued them for vital moral protections.

§ But there are a number of problems with (L1). Some of these require only that we rephrase (L1). Others warrant the complete rejection of that thesis. Let us start by discussing one of the former problems.

First of all, what is meant by the term "government"? Given certain defensible (though, in my view, erroneous) conceptions of what a government is, (L1) turns out to be false or viciously circular.

Governments provide laws. At least arguably, this is an essential fact about government. So it seems reasonable to suppose that a government in a certain area is, by definition, something which imposes laws in that area. If this is right (I myself doubt whether it is), then (L1) is viciously circular, as it defines "law" in terms of itself.

(L1) is, arguably, guilty of another form of vicious circularity. We tend to think that a government in a certain area is not just the *de facto*, but the *official*, authority in that area. The Mafia was surely not the *government* in Sicily, because even though it may have been the *de facto* authority, it never had *official* status.[97] If this is right (I myself do not believe that it is), then an "official" authority is the one that is recognized by law. In that case, (L1) is viciously circular, as it defines "law" in terms of itself.

There is yet another reason to think that, given a correct analysis of the concept of government, (L1) must be rejected. What does the U.S.

government do for its citizens? It prevents them from killing one another; it builds roads and hospitals; it provides an army that protects them from foreign invaders. The U.S. government thus seems to be in the business of protecting certain basic moral rights – the right to life, safety, and so on. The same is true of every other government worthy of the name. Given this, it seems reasonable to say that a government in a certain area is, by definition, something which issues protections of rights to people in that area. If that is right (as I will argue it is), then (L1) amounts to:

L is a law for X, i.e. it defines a legal obligation for X, provided that two conditions are met. First, L is a protection of a moral right had by somebody or other (not necessarily X himself). Second, L is issued by something on which X depends for protection of his rights.

Thus defined, the demands that brother Bob makes of you count as laws, since you depend on Bob for protection of your rights. But then (L1) is false, since those demands are obviously not laws.

§ The problems just discussed are not insuperable. Ironically, the best way to deal with them is to show the *falsity* of an analysis of the term "government" that, if correct, *would* free (L1) of the problems just discussed.

The analysis I have in mind is one that I will refer to as Weber*'s thesis. Note the asterix. Weber*'s thesis is similar to, but *not* identical with, a thesis actually put forth by Max Weber.[98] Weber*'s thesis is this: The government in a particular area is simply the entity which has a *monopoly on the use of force* [99] *in that area*.[100] If this is right, then a government doesn't have to protect rights or (therefore) lay down laws. Thus, Weber*'s

thesis would seem to give us a way of ridding (L1) of circularity.

But, I would argue, Weber*'s thesis is false. More exactly, though it may be *extensionally* correct, it is *intensionally* incorrect. It may be true that, in actuality, each government is the thing which has a monopoly on the use of force in its geographical area. But nothing is a government *in virtue* of its having a monopoly on the use of force, as the following story illustrates.

Let S be some nation-state. On a mountain-top near S, there lives an invincible ogre. Every now and then, the ogre goes into S and does many bad things. The ogre isn't interested in administrating human affairs. So after each rampage, he returns to his mountain-top.

Surely the ogre is not the *government* of S. At the same time, the ogre has all the power. If it wanted to, it could obliterate all of S's forces with a single sweep of its paw. For whatever reason, it doesn't do so. Perhaps it wants S's infra-structure – including its law enforcement – to remain intact, so that S's economy will continue to flourish and produce many goods for his taking. Perhaps he is just too lazy to deal with the tedium of state-craft.

It might be suggested that, if the ogre were a more continuous presence in S, then it would qualify as the *de facto* government of S. But that doesn't seem to be true. Suppose the ogre grows bored of his mountain-top, and decides to stay in S continuously. He sleeps and eats there. Every day he commits another horror. The one thing the ogre does *not* do is participate in the management of S. He doesn't prevent citizens of S from killing one another or from vandalizing historical monuments. He doesn't fix problems with the sewage system or help the police solve crimes. Even though he has total power, and uses it liberally, he does nothing that is even indirectly constructive.

It seems to me that the ogre is *not* the government of S, or even a part of that government. A government is not simply something which wields power.

It has a more constructive function. It provides services; it distributes protections and freedoms.

Of course, there are bad governments (e.g. those of Stalin and Pol Pot). But their relationships to their constituents are very different from the relationship between the ogre and the citizens of S. The average Soviet Citizen had some assurance that he could walk down the street without being killed or raped. He had some assurance that he would not starve, and that he would have shelter, heating, and plumbing. To the imperfect but non-trivial extent to which he had such assurances, it was thanks to the Soviet government. The ogre doesn't provide such assurances to the citizens of S.

When we look at governments, we tend to compare them with other governments, and not with the *absence* of government. But here we are concerned not with the difference between good and bad government, but with the difference between the presence and absence of government. There is an enormous difference between an atrocious government, like of that of Stalin or Mao, and *no* government. The difference lies entirely in the extent to which certain basic rights of people (life, shelter, food) are protected. When we compare Stalin's government with other governments, it is easy to fall prey to the illusion that Stalin's government is just something which has a monopoly on power, and does nothing for the welfare of the Soviet citizens. If our standard of comparison is the Swedish or British government, that is a correct characterization of Stalin's government. But in absolute terms, i.e. if our standard of comparison is *no* government, it is a false characterization.

§ The government is a guarantor of rights. But not just *any* guarantor of rights is the government. Brother Bob, who protects you from the neighborhood hoodlums, is not the government. Why not? Because, I propose, he is not the *main* guarantor of rights in the area. The government in an area is the *main* guarantor of rights in that area. So, in (L1), we must

replace the term “government” with the term “main guarantor of rights in the area”:

(L2) L is a law for X, i.e. it defines a legal obligation for X, provided that two conditions are met. First, L is a protection of a moral right had by somebody or other (not necessarily X himself). Second, L is issued by the main guarantor of rights in the area.

But (L2) is false.[101] The same is therefore true of (L1), since (L2) is simply an explication of (L1).

Suppose that government G is dominated by evil wealthy land owners. The government provides many special services for these landowners. But they are brokered through secret backroom meetings, with a wink and a nod. In consequence of these under-the-table agreements, the wealthy land owners are given special police protections, tax-exemptions, and so on.

These protections are not *laws*. There is no *law* requiring land owners to be given these special considerations. But if (L2) is correct, then these protections – given, as they are, by the main rights-broker in the area – *do* count as laws. So (L2) is wrong.

A possible response would be to say that G’s protections of the landowners were *de facto* laws. But this threatens to trivialize the notion of the law. If we regard as law every backroom deal among potentates, then we obliterate the very real distinction between the rule of law, on the one hand, and the rule of influence and wealth, on the other. There is a clear sense in which the rule of law is absent in a country where the day belongs not to

codified statutes, but to secret deals made among big shots. Thus, under-the-table protections given to V.I.P.'s don't count as laws. So (L2) is too strong: it counts as laws things that are not laws.

§ (L2) is also too weak: it *fails* to count as laws things that clearly are laws. In almost any place where some atrocity is accepted, there is a law on the books prohibiting it. Wife-beating is technically illegal in practically every country. So is the exploitation of child-labor. But in many cases, those laws are not enforced. Sometimes they *cannot* be enforced, given cultural and economic realities.

We don't want to say that, in virtue of being without consequence, such laws don't *exist.* That would, once again, obliterate the distinction between the rule of law and the rule of custom. If what is done is *ipso facto* law, and what is not done is *ipso facto* not law, then it becomes trivially true that every nation is subject to the rule of law. But such a proposition is never trivial. Indeed, it is often false.

There is another way to look at this. Suppose we say that laws must be obeyed or enforced to exist. In that case, we are suggesting that laws exist only in so far as they are followed. If that is right, then a law would seem to collapse into the behaviors that are supposedly acts of *following* that law. But then there isn't anything that is being followed. If there is no opposition between the law and the act of following it, then the law drops out. There are just behavior patterns: the law ends up being a kind of Humean regularity. [102]

But that isn't what laws (in the legal sense) are. A law isn't a macro-characteristic of patterns of behavior.[103] It is something in terms of which behavior is to be evaluated.[104] So we must grant that there are, or at least can be, laws that nobody follows or enforces and, therefore, that (L2) is false.

§ Why do unenforced child-labor laws still count as *laws*? Because those laws are *official* governmental assurances that certain rights will be protected. Such a law is an official promise that the rights of children will be upheld.

Why are under-the-table deals between dictators and V.I.P.'s *not* laws? Because they are not official. Even though the government is issuing a very real protection, the protection is not *official.* This suggests that we should replace (L2) with:

(L3) L is a law for X, i.e. it defines a legal obligation for X, provided that two conditions are met. First, L is an *official* protection of a moral right had by somebody or other (not necessarily X himself). Second, L is issued by the main guarantor of rights in the area.

But (L3) is not satisfactory. In effect, we have already seen why. An "official" protection would seem to be one that is sanctioned by law. In any case, that is the most obvious definition of "official". So (L3) seems to collapse into:

(L4) L is a law for X, i.e. it defines a legal obligation for X, provided that two conditions are met. First, L is a protection *that is sanctioned by law* of a moral right had by somebody or other (not necessarily X himself). Second, L is issued by the main guarantor of rights in the area.

Of course, (L4) is not acceptable. The term "law" is the one that (L4) is

meant to analyze. Since that term occurs in the *definiens*, (L4) is viciously circular. The same is therefore true of (L3), since (L4) is merely an explication of (L3).

§ But the circularity can be removed. Once we do so, we will at last have what might be a creditable hypothesis as to what laws are.

Let Smith be the dictator of a country whose politics are secretly dominated by backroom deals with wealthy sugar plantation owners. As we've seen, if Smith makes a secret deal with plantation owner Jones, whereby the government will give Jones various special protections, that does *not* constitute a law. But suppose that, instead of giving these protections to Jones through a secret deal, Smith goes on national television and says:

I hereby promise to give such and such protections to plantation owner Jones. Jones has my word that the instruments of government will ensure that Jones enjoys the aforementioned privileges and, consequently, that anyone who tries to strip Jones of those protections will be fined, jailed, or executed. I hereby demand that all military, law-enforcement, and administrative personnel abide by this.

Let us take the story a little further. After hearing this announcement, the country *does not* revolt in horror at this brazen admission of governmental corruption. Perhaps people are simply too afraid. Consequently, the instruments of government – the police, armed forces, and so on – do provide the very protection that Smith publicly and explicitly promised to Jones.

In *this* case, it seems to me that Smith's assurance of protection *does* qualify as a law. Suppose that bureaucrat Brown resists Smith's directive. In

that case, Brown could, I think, be accused not simply of failing to kow-tow to the *de facto* powers that be, but of breaking the law.

An unenforced statute forbidding wife-beating *is* a law, because the creation of that statute amounts to an explicit and public assurance on the part of the government that certain rights will be protected. It is an assurance that the government fails to live up to: but it is still an assurance.

There is law where there is a public and explicit assurance of some kind of protection of a right. Hence:

(L5) L defines a legal obligation for X, i.e. is a law for X, exactly if the following condition is met. L is an explicit and public[105] assurance of a protection of a right, and the entity providing that assurance is the main guarantor of rights in the area.

In (L5), the concepts occurring in the *definiens* -- *explicit*, *public*, *moral good*, *assurance*, *protection*, and so on -- can all be understood independently of the concept of law. So their occurrence in the *definiens* constitutes no circularity.

§ Suppose that the politburo convenes and enacts a "private law" L: a law that no one is to know of, except those in the politburo. To fix our ideas, suppose that L prohibits people from owning parrots.

Under those circumstances, does the man in the street have a *legal* obligation not to own parrots? I am not asking whether he has a *moral* obligation not to own them, but only a narrowly legal one. Does the man in

the street have a *legal* obligation to report parrot-owners to the authorities? The answer to these questions is plainly “no”. A “legal” obligation whose existence is unknowable is not a legal obligation at all.[106]

Of course, there have been regimes that enacted “private laws” – laws that nobody (apart from the relevant political insiders) could know about – and then imprisoned or executed people on the basis of those so-called laws. But those were not laws. Given only that a dictator describes something as a “law”, it doesn’t follow that it is a law. These so-called laws were mere resolutions. Describing them as “laws” was a deliberate deception.[107]

Suppose that you are the one weak person in a society of ten powerful people. (They are hundred foot giants. You are a six inch dwarf.) Every now and then, these Powerful Ones retire into a room. You are not allowed in. All you hear is an indistinct grumbling. Eventually the Powerful Ones come out, and say:

We just came up with a new law. But we’re not going to tell you what it is -- unless you break it, in which case you will be severely punished.

The “law” they came up with is that you are not allowed to wear tennis shoes indoors. Let us set aside the question of what kind of *moral* obligations you might be under. In virtue of the fact that the Powerful Ones just passed this “law”, are you under any *legal* obligation not to wear tennis shoes indoors?

It seems clear that there is *no* such obligation. Of course, there may be causal consequences to your wearing tennis shoes indoors: the others may be beat or kill you. But here we are led back to the distinction between *causal* and *legal* requirements, between being *obliged* and being *obligated*. The sense in which you must not wear tennis shoes indoors is entirely different

from the sense in which I must not drive 66 mph in a 65 mph zone. The second “must” defines an obligation (albeit a morally rather empty one); the first “must” defines a mere causal condition for your avoiding some harm. You are no more *obligated* by the secret resolution of the others to refrain from wearing tennis shoes indoors than you are *obligated* not to drink from a glass that, unbeknownst to you, contains a rare and deadly virus. In both cases, we are dealing with causal, not legal, requirements. We may conclude that a secret resolution defines no legal obligation, and thus corresponds to no law.[108]

§ It is true, of course, that most people (here or anywhere) don’t know most of the laws. In fact, it is probably beyond the intellectual capacity of most people, perhaps everyone, to grasp *all* of the niceties of our legal system. But that has nothing to do with their *unavailability*. Rather, it has to do with the fact that, once available, they are not easy to comprehend.

The laws relating to the tax-code are perfectly *available.* Right now, I could walk over to a library and examine them. I would not be able to understand many of them. But that has nothing to do with their availability. Von Neumann’s proof[109] that quantum physics cannot be made deterministic through the introduction of hidden variables is available to everyone, even though only a few experts understand it. But the resolutions of the Star Chamber, or of Stalin’s politburo, are unavailable in the sense that outsiders aren’t even given the chance to try to understand them. That is why those resolutions do not define legal obligations or, therefore, laws.[110]

§ (L5) says that a law must be a public and explicit promise of a protection, provided that promise is made by the entity which is the primary guarantor of rights in the area. It doesn’t say that the promise must be kept. So (L5) is

consistent with the existence of unenforced laws.

Let us sum up. A government in an area is the main guarantor of rights in that area. L defines a legal obligation for X if L is a public and explicit assurance of the protection of a right (not necessarily one of X's rights) made by the government in that area.[111]

An alternative to (L5): reviving (L2)

The purpose of this section is to show that, notwithstanding what we said a moment ago,

(L2) A law is a governmental assurance of a protection of a right.

may in fact provide a correct analysis of law.

Here it may be objected that I am providing two different analyses of the concept of law, and that, consequently, at least one of those analyses must be wrong. I must, it will be said, pick *one* analysis and stick with it. My response is this. Any analysis of the concept of law must be consistent with a certain hard core of intuitive data (e.g. that there are laws prohibiting murder). But there are *different* ways of modeling that data, and *ceteris paribus* there can be no principled way of choosing the one model over the other.

Also, we will find some reason to believe that (L2) and (L5) are actually equivalent so that, leaving aside the point just made, I am not guilty of inconsistency in holding that *both* of them are viable analyses of law.

Let us briefly review what we said about (L2). Suppose a government official, e.g. a generalissimo or a senator, provides under-the-table assurances of protections of rights to a wealthy land owner. Such an assurance is obviously not a law but, at the same time, it is provided by the government.

On this basis, we chose to replace (L2) with (L3). But we found reason to believe that the occurrence of the word “official” (L3) amounted to vicious circularity, since “official” presumably means “sanctioned by law”, or some such. This led us, in due course, to (L5).

Let me now say why, notwithstanding the argument just given, (L2) may be correct. Somebody who is an official isn’t always acting *as* an official. If, in a fit of rage, police officer Smith punches a friend of his, he isn’t (necessarily) acting *as* a police officer. Smith could just be acting as a belligerent drinking partner who happens to be a police officer by profession. Of course, he could commit such an act *as* a police officer. Suppose that Smith punches his friend while on duty, and then subsequently arrests his friend and files a report falsely stating that he had punched him in order to keep him from committing some crime. In that case, in punching his friend, Smith is acting *as* a police officer; and his act is one of *police*-misconduct, as opposed to garden-variety misconduct.

Suppose that Generalissimo Jones gives plantation-owner Brown an under-the- table assurance of some kind. It is not unreasonable to say that, under such a circumstance, Jones is not acting *as* a government official, but rather as somebody who happens to be a government official.

Here is another illustration of the principle I am trying to expose. Suppose that Aaronson is the Secretary of the Treasury of country X. As such, Aaronson has access to the country’s funds – he has keys to all the relevant bank vaults. One day, Aaronson decides to walk into one of those vaults and help himself to a million dollars. In doing this, Aaronson is not acting *as* the X’s Treasury-secretary, but rather *as* a private citizen whose profession happens to be that of the Secretary of the Treasury. Because of his position, Aaronson has powers – I am not referring to official governmental powers, but to garden-variety causal powers – that other citizens lack. In particular, he

can walk into a bank vault in X's central bank any time that he likes and put wads of cash in his pockets. These are not among his official powers as Treasury-secretary. They are mere *causal* powers that Aaronson has in consequence of his being Treasury-secretary. Any government official will have *causal* powers that are not among his official governmental powers. (Think of the damage that soldiers or police officers *can* do that they are not permitted to do.)

It seems to me that, when Generalissimo Jones gives plantation-owner Brown an under-the-table assurance, Jones is not acting *as* a government official, but *as* a private citizen who also happens to be a government official. As a consequence of being the Generalissimo, Jones has *causal* powers that are not among his official powers; and when he gives special protection to Brown, he is exercising one of these unofficial powers. He is thus not acting as an agent of the government. Consequently, no *governmental* assurance has been given: and that, it might plausibly be suggested, is why no *law* has been issued.

If this line of thought is correct, then (L2) *is* right. Under-the-table deals are *not* governmental assurances of protections of rights. They are assurances made by people who happen to be government officials but who are not acting *as* government officials. Consequently such an under-the-table assurance is not a *governmental* assurance and is therefore no exception to (L2). So there is no reason to reject (L2) or, therefore, to replace it with (L5).

From one viewpoint (L2) is preferable to (L5). The latter, I feel, has an *ad hoc* and arbitrary quality; for it leaves us with the question "why must a governmental assurance be *explicit* to qualify as a law?" We've seen many reasons to believe *that* such an assurance must be explicit; but we haven't seen any reason to see *why* this must be so. When we think of laws, concepts like *right, justice,* and *protection* come to mind. But the concept of

explicitness does not (so a feeble but persistent intuition tells me[112]). Of course, this by itself doesn't prove anything. When we think of the concept of "real number", we don't necessarily think *sets of ratios*. Concepts have depths, and conceptual analyses are often surprising. Still, there is some expectation that a conceptual analysis – while not, of course, merely registering what our pre-theoretic intuitions tell us – must have a certain fit with those intuitions. It seems to me that, so far as it involves the concept of explicitness, (L5) lacks such a fit. To my mind, the term "explicit" gives (L5) an epicyclical quality. (L2), on the other hand, has just the kind of seamless simplicity that we expect from a good conceptual analysis: "a law is a governmental assurance of protection of a right." At the same time, (L2) *explains* why laws cannot be cryptic – why they cannot be veiled in secrecy and must be explicit. If I do not *convey* to X that I am going to protect him, then I have not assured X of any thing. An assurance is a communication. If the existence of some so-called "law" isn't publicly known, then that law isn't a communication or, therefore, an assurance to the public. The concept of explicitness (in the sense of being non-cryptic and public, not in the sense of being stated with the kind of literalness in terms of which logicians and semanticists think) is built into the concept of an assurance. So (L2) is not only consistent with the fact that laws must be explicit, but actually *explains* it. By contrast, (L5) makes it seem as though that fact is a kind of accident. For these reasons, it seems to me, (L2) is *ceteris paribus* preferable to (L5).

To sum up, the wording of (L5) suggests that it is just an extra fact, so to speak, about laws that they must be explicit, whereas (L2) – at least when duly explicated – provides an intuitive basis for this fact. So, in conclusion, it looks as though (L2) is right.

But now we must remember some points made earlier. We found that, in virtue of containing the word "official", (L3) is viciously circular. The

question arises whether our new defense of (L2) involves a similar circularity. After all, we have found that, so far as it is defensible, (L2) amounts to this:

(L2*) A law is an assurance given by an entity that not only has government powers, but is also exercising in a manner consistent with its *official* duties.

So it looks as though (L2*) is no less circular than (L3), and that we must therefore reject it.

I believe that, despite what we said earlier, (L2) is *not* circular and that it does give a correct analysis of the concept of law; for, I will argue, we can explicate the concept of acting *as* a government official without circularly invoking the concept of law.

Suppose that Rutherford is on soccer team T. There are some things that Rutherford does *as* a member of T (e.g. passing the ball to the center-forward in the context of a match), and there are other things that he does *not* do as such a member (e.g. brushing his teeth before he goes to sleep). We might say that, when Smith brushes his teeth, his doing so is not an "official" act of team T, whereas when Smith scores a goal, his doing so *is* such an act. In this context, then, the concept of something's being "official" doesn't need to be explicated in terms of the concept of *law*. When Smith scores a goal, his act is "official" not because some law says so, but because of the way in which that act is embedded in a certain collective enterprise. When Smith brushes his teeth, his act is not an official act, because his act cannot be understood in such terms.

When Generalissimo Jones has Congress (or his nation's equivalent thereof) convert his will into law, Jones' act is embedded in a certain way in a

certain collective enterprise. When Jones provides an under-the-table agreement, his act is not thus embedded. In the first case, we have an act that has a systemic basis – that engages in the machinery of government in a manner that, at some fundamental level, is in keeping with the operating principles of that entity. In the second case, we have a free-wheel; we have an act that lacks any sort of systemic basis.

If this is right, then the difference between an official and unofficial act – between acting *as* a government official and acting as somebody who merely happens to be such on official – doesn't have to be drawn in terms of the concept of law, and can be drawn in terms of concepts of a generally functional or teleological character. Given this, it follows that (L2) ultimately *is* correct.

Nonetheless, one could argue that, when the Generalissimo acts officially, the "systemic basis" had by his act *does* implicate the concept of law. This is certainly a position that deserves consideration, and the reasons I've given to think otherwise do not constitute a definitive reason for rejecting it. One who thinks that (L2) *is* guilty of circularly exploiting the concept of law has the option of accepting (L5) in its stead -- notwithstanding that I've given my reasons for preferring (L2).

But, in the present context, what is important is that one accept *either* (L2) or (L5). It is a matter of comparative indifference to my larger project which of those two is accepted. A complete defense of the view that (L2) is free of circularity must be given elsewhere, since it would involve a thorough examination of the richly articulated region of conceptual space common to the concepts of function, causality, and collective intentionality.

§ There is one last point to make here. (L2) and (L5) are not necessarily incompatible. My own view is that they are in fact equivalent. (Here we are

going to employ a principle that we will later find to have manifold relevance to the concept of law; and I will illustrate with an example that I will later re-use.) Suppose that you don't speak English, and that I say to you "I assure you that I am going to protect you from the local Mafia chieftain." I have *attempted* to give you an assurance. But I haven't *actually* done so. No assurance has been provided. An assurance is necessarily understood and *a fortiori* is necessarily taken in. So the concept of a cryptic assurance is of dubious coherence; assurances necessarily have a certain degree of explicitness. (This point will be developed in chapters 11 and 12, and will be of central importance to our views on international law and legal interpretation.) Thus, if a government makes an assurance that isn't known to the public, then it certainly hasn't assured the public of anything. In that case, from the public's point of view, *there is no assurance* and thus no law.

If this line of thought is right, then the concepts of publicity and of explicitness[113] are built into the concept of an assurance, a consequence being that (L5) is simply a pleonastic expansion of (L2) – much as "x is a three-sided triangle" is a pleonastic expansion of "x is a triangle." So an adequate explication of (L2) accommodates the expectation (explicitly identified by Fuller) that laws must be public.

At the same time, there is a kind of pragmatic non-equivalence between (L2) and (L5). The wording of (L5) suggests (though it doesn't strictly entail) that a law's being public and explicit is something *additional* to its being an assurance of a certain kind. By contrast, (L2) doesn't carry this misleading implicature. So even though they are (I believe) strictly equivalent, (L2) and (L5) are not pragmatically equivalent; and the fact that (L5) is, whereas (L2) is not, characterized by a certain ungainliness is to be traced to the fact that (L5) does, whereas (L2) does not, bear that misleading implicature.

In conclusion, (L2) and (L5) are probably equivalent. But because of

their pragmatic differences, (L2) has a better fit with our intuitions than (L5). So we may conclude with some, though not complete, confidence that:

(L2) A law is a governmental assurance of a protection of a right.

For the record, my view is that (L2) gives a correct analysis of law.

Chapter 8 Laws as narrow-scope assurances: the triviality problem

My position is that laws are governmental assurances (not necessarily upheld) of protections of rights. There is an important objection (due to Thomas Holden) to this analysis. The objection is best conveyed by way of the following story.

Suppose that, as a prank, a legislative aide puts a phony bill (i.e. proposed law) on the desk of a senator. To simplify discussion, let us suppose that the aid has no desire to see the bill become law – he only wants the senator to have a good laugh. The proposed law would make it illegal for people to wear green shirts indoors. The Senator doesn't read the bill. But he absent-mindedly signs it. It then goes to Congress. Nobody in Congress reads it, but a sequence of absent-mindedly tendered signatures and votes convert it into law. Let GS be that law.

No human being ever wanted GS to be a law. Further, GS doesn't protect any right and it doesn't assure anyone that any right will be protected. Nonetheless, it does appear to be a law.

GS thus seems to be an example of a law that doesn't give anybody any

assurances of any protections of any kind. It doesn't even protect, or promise to protect, a potentate's ability to convert his strange whims into law. We thus seem to have a case of a law that doesn't protect anything, and doesn't constitute an assurance that any protection of any kind will be granted. So GS thus seems to be a counter-example to our analysis of law.

There is a second reason why the existence of laws like GS is, at least on the face of it, inconsistent with our analysis. GS was made in an impersonal, almost mechanical fashion. Human intentionality was involved only marginally. (In fact, we can further marginalize the role of intentionality by supposing that it was, not a legislative aide, but a random quantum event that deposited the "bill" on the senator's desk.) This is problematic because, ordinarily, we think of promises as being backed by mental states. If I say to you "I promise to pick you up at the airport tomorrow", my words issue a promise because (inter alia) they are produced by somebody who knows what they mean. (If an inanimate object – e.g. a set of wind-chimes – or insufficiently intelligent object – e.g. a parrot -- had made the same sounds, no promise would have been made, even though the auditor might mistakenly think otherwise.) GS is *not* produced in a comparable manner, making it *prima facie* hard to see how it would qualify as an assurance of any sort. So given that we have identified laws with assurances or promises of sorts, the existence of laws like GS seems inconsistent with our analysis.

§ How am I to respond to this? First of all, I do not want to take the revisionist approach of saying that GS is not a law. It seems to me a datum that GS is a law, and thus constitutes a real problem for our analysis.

I am going to give two responses to this. We will see that, despite first appearances, these responses are equivalent. But it will be more illuminating to give both responses separately, rather than attempt to telescope them into a

single one. Right now, let me give a very general outline of each response, before going on to give those responses in full.

As we've seen, governments are supposed to protect those under their authority (and they necessarily *succeed* in doing so, at least up to a point). The laws they issue are supposed to help them do that. It thus seems clear that laws are supposed to serve the interests of those who are subject to them, and are thus supposed to protect moral goods. Of course, there is a trivial sense in which that is true of everything. Even assassins and criminals are, in some sense, supposed to do good. But an assassin *qua* assassin is not supposed to do good, whereas a law *qua* law is supposed to do good. A law that fails to protect a moral good is less good *as a law* than one that fails to do so. (Later we will give argumentative grounds for this assertion, so as to head off the charge that it begs questions against the positivist.) But an assassin that kills the innocent is not, on that account, less good *as an assassin* than one that kills the wicked. So there is a non-trivial sense in which laws are, by their very nature, supposed to do good.

In actuality, GS doesn't do any good. But given that it is a law, what is it *supposed* to do? It is supposed to serve the interests of the people subject to the legal system of which it is a part. How is it supposed to help those people? Let us answer this question by asking another: when laws *do* help people, how do they do so? The law that prohibits murder does so by protecting the rights that people have (especially the right not to be murdered) that would be violated by the commission of murder. The laws preventing various forms of racial and sexual discrimination protect people's rights not to be subject to iniquities. It would thus seem that GS is *supposed* to protect the rights that would be violated by the wearing of green shirts indoors. Thus, GS can be seen as an assurance, on the part of the thing which created it (namely, the government) to protect the rights of those whose rights would be

violated by the wearing of green shirts indoors.

Here is another way of putting it. Under what circumstances would GS *not* be a failure? If the wearing of green-shirts indoors led to people's being harmed (because, say, the wearing of green shirts indoors released some toxic gas into the atmosphere), *then*, if enforced, GS would *not* be a failure. Under such a circumstance, GS would protect those rights that would be violated by the wearing of green shirts indoors – those rights being (inter alia) the rights to health and life. So GS can be seen as the government's saying to those people whose rights would be violated by the wearing of green shirts indoors: "fear not! Anyone wearing a green shirt indoors will be punished – so those rights of yours that would be violated by such an act will be protected."

Of course, there probably *is* no person whose rights would be violated by the wearing of green shirts indoors. Further, whereas most *assurances* are given by people who have (or, at least, pretend to have) certain intentions, GS was made rather mechanically and impersonally. In any case, it was certainly *not* made by anyone who had the intention of assuring anyone of anything. On the face of it, these facts seem to bear against our analysis.

But these facts, I will argue, have no such bearing; and we see this, when we consider the second (version of the) response I will give to the problem posed by laws like GS. Here is an outline of that other response.

Promises are not always made by *individuals.* Institutions can make promises. Suppose that a chain of fast-food restaurants (call it McJack's) uses low-quality beef, resulting in fatal cases of food-poisoning. It seems to me that the McJack's corporation *as a whole* can promise to make reparations to the victims, and can promise future customers that they needn't fear food-poisoning. It is thus possible for promises to be made in a (relatively) *impersonal* fashion. (I say "relatively" because, as we will see, human agency and intentionality must have *some* constitutive involvement in the making of

institutional promises. This prevents such promises from being *entirely* impersonal.)

A related point is that these promises can be made in a (relatively) *mechanical* fashion. (I say "relatively" because, as I said a moment ago, intentionality has to have *some* constitutive involvement. But, as we will see, the *proximate* causes of an institutional promise can be entirely lacking in any sort of human sentience.)

Given these points, the fact that GS is made in a relatively mechanical and impersonal fashion does not, by itself, mean that GS is a counter-example to our analysis of law.

But, if I am not mistaken, the threat posed to that analysis doesn't lie so much in the fact that GS is made in a relatively mechanical and impersonal manner as it does in the fact that GS doesn't assure anyone of any protection at all. This brings us to the second main component of our response to the objector's important challenge.

Let us begin with a short piece of fiction. Smith is the U.S. President and he says the following:

(1) "I hereby promise that, as of this day, every American citizen who buys a Plymouth Adobe[114] [a kind of car] will be given a robot that can cook any dish whatsoever."

To conclude our story, suppose that nobody ever buys a Plymouth Adobe, and that there are no robot chefs (since the relevant technology hasn't been developed).

Let us start with a simple question. Has Smith issued some kind of a promise? I am not asking whether he issued an intelligent or credible or otherwise meritorious promise, but only whether he did in fact issue a

promise. The answer is yes.

But there are two problems. First, a promise is presumably a promise *to* somebody or other. I cannot *just* promise. I must make a promise *to* Bob or Mary or Frank. To whom has President Smith issued a promise? To purchasers of the Plymouth Adobe? There is no such person.

Second, a promise is a promise *of* something. I cannot *just* make a promise to Bob. I must promise him tickets to the Laker's game or a brand new T.V. set. [115] Given this, supposing that somebody x did buy a Plymouth Adobe, *what* did President Smith promise x? The obvious answer is "a robot chef." But there is no object y such that y is a robot chef. So, it would seem, there is no object y such that y is what Smith promised to give x. So there is no x, and no y, such that Smith promised y to x. So Smith hasn't promised *anything to anyone.* Therefore, presumably, he *hasn't* made a promise. At the same time, it seems clear that he *has* done so.

But, of course, there is no real problem here. Given some old stand-bys of the philosophy of language[116] [117], we know how to resolve this situation. The meaning of (1) is *not*:

(2) Given any person x who buys a Plymouth Adobe, there is some robot chef y such that I hereby promise to give y to x. [118] [119]

Rather, the content of (1) is given by:

(3) I hereby promise that, for any person x who purchases a Plymouth Adobe, there will be some robot chef y such that y will be given to x.

(3) does not presuppose the existence of any robot chefs. (3) says (or promises) that *if* somebody buys a Plymouth Adobe, *then* there will be a

robot chef that is bequeathed to that person. (3) doesn't need there to be any *actual* robot chefs or purchasers of Plymouth Adobes. But it is still a promise to Adobe-purchasers that they will receive robot-chefs.[120]

If I say "I promise to give $100 to anyone who beats me at arm wrestling" I have made a promise or commitment, even if nobody ever does beat me at arm wrestling. My promise has the form: I promise that *if* there is anyone x who beats me at arm wrestling, *then* I will give $100 to x. It is a narrow-scope promise.

So *given only* that there is no x and no y such that GS is a promise of y to x, it doesn't follow that GS isn't an assurance or promise. It could be that GS is a *narrow-scope* promise. (It will turn out, in fact, that *all* laws are narrow-scope promises – even those laws that give actual protections to actual people.)

Before we can close our outline of the argument to be given, we must make one more point – one that, on the face of it, appears to have nothing to do with the concept of law. Suppose that Jones calls you up and says "I will be there tomorrow." At the level of literal meaning, this sentence-token expresses the innocuous proposition that somebody is going to be in a certain place at a certain time. Given only its literal meaning, that speech-act could be any number of different things. It could be a promise to save your life – to give you a much needed, and otherwise unobtainable, blood-transfusion. Or it could be a threat to kill you. Or it could be neither. It all depends on what the relevant background facts are. Given the right situational or historical background, any sentence-token, no matter what its literal meaning, could have any possible performative status – it could be a threat *or* a promise *or* a warning *or*…

Further, that same historical background determines, at least in part, *what it is* exactly that is being threatened or warned against or promised…Given

only that Smith's utterance of "I will be there tomorrow" is a promise, it is entirely open what it is exactly that is being promised: it could be a ride to the airport or a kidney or a ticket to the opera.

These points have obvious corollaries. First, *whether* a sentence-token has a certain performative status –*whether*, for example, it is a promise (as opposed to a threat or a simple indicative) -- is a function, in part, of the situational background.[121] Second, supposing that it is a promise, *what it is* exactly that is being promised is also a function of the background. Given only that Smith's utterance of "I will be there tomorrow" doesn't contain the words "I promise" or the word "kidney", it doesn't follow that Smith hasn't just promised to give you a kidney.

In conclusion, both the performative status and specific content of a speech-act are a function, in large part, of the background understandings brought to the table by both speaker and auditor.

Given these points, here is the general viewpoint for which I shall now argue. (I warn that, given what little we've said so far, the thesis about to be stated may have little intuitive appeal.) GS *is* an assurance of a protection. It is an assurance that, for any citizen x, and any right R had by x, x's possession of R is assured protection against any threats to its existence posed by the wearing of green-shirts indoors.

Of course, there is no actual person whose actual right is assured protection by GS. But in light of our points concerning narrow-scope assurances generally, those facts are not grounds for holding that GS is *not* an assurance. Nor are such grounds to be found in the fact that GS has an *impersonal* character, and was made in a mechanical manner. There is independent reason to believe that institutional promises often have these properties. Finally, nothing of consequence follows from the fact that the verbiage by which GS is expressed is not anything remotely like "it is hereby

assured that any one x such that…" *Most* promises don't involve expressions like "I promise" or "I assure"; and *most* promises are given by sentence-tokens whose literal meanings have little or nothing to do with what is being promised. In what remains of this chapter, we will defend the position just outlined.

§ According to the theory we have put forth, laws are assurances or promises of sorts. The making of a promise is an inherently social act. For this reason, no psychological state is *sufficient* for the making of a promise.

In fact, it isn't even clear whether there is any psychological state whose presence is *necessary*.[122] Suppose I say to you: "I promise to pick you up at the airport tomorrow." As I am uttering those word, I may have no intention of actually picking you up. Even so, I would still have made a promise – no less than if, while uttering those words, I had every intention of picking you up. My words are a promise not solely because they arise from this or that sort of mental state, but because those words occur in a certain social context.

This is not to say that mental states have *no* role in promise-making. Obviously they do. If an inanimate robot produced the sounds "I promise to pick you up at the airport tomorrow", no promise would have been made, *even if* the person hearing those words thought that the thing producing them was an animate object that really did intend to show up at the airport the next day. So, to qualify as a promise, those words would presumably have to have mental causes of *some kind* or other.[123] (That said, we will soon encounter some reasons to qualify, though probably not reject, even this weaker thesis.)

§ Promises are, at least typically, made with the use of language. So to deepen our understanding of the nature of promises, and therewith of laws,

we must take a brief detour through some commonplaces of the philosophy of language.

A given sound could mean anything. An utterance of “snow is white” could mean *grass is green.* By themselves, the intrinsic properties of a sound (or ink-mark) do nothing to fix its meaning. Its meaning is fixed by psychological, social, and historical factors – in conjunction with facts about orthography or acoustics.

Since promises are linguistic acts (or typically constitutively involve such acts[124]), everything we just said applies to them. What counts as a given promise is, in large part, a function of the social and historical background. The sound “I promise to pick you up tomorrow” could mean anything. It expresses (or constitutes) a certain promise only because it occurs in a certain psycho-socio-historical context.[125]

A corollary is that, given the right background, practically any act could constitute the making of a promise. I could develop a kind of private code with my neighbor Jim whereby, when he leaves his front-light on, he is promising to drive me to work the next day.

This brings us to a point that has an important place in our way of dealing with GS. There can be *impersonal* promises: promises that are not made by individuals and that do not correspond to the intentions or wishes of any one person.

The Union Carbide Corporation could, as a whole, promise to make certain amends to the people of Bhopal India.[126] But given only this fact, it doesn’t follow that any particular individual made such a promise. Indeed, nothing follows as to the nature of the intentions of any specific employee of that corporation. In fact, the Union Carbide Corporation could make a promise even though, at least for a while, no employee of that corporation even knows of the existence of that promise.

We may illustrate this last point with a bit of fiction. Union Carbide has a computer which scans all the major newspapers. If it finds an article with certain key-phrases, then it automatically transmits a message to some wire service or to the teleprompter read by some news anchor. So if the computer finds an article containing the words "Union Carbide guilty of atrocities in connection with…", the computer mindlessly transmits the following message to Bloomberg: "We at Union Carbide are profusely sorry for the incident at___ and wish to give the residents of that town $1,000,000 for their trouble." Under these circumstances, Union Carbide has made a promise, even though, at the time the message was transmitted, every single employee of Union Carbide may have been fast asleep.

We mustn't overstate things. The computer in question was implemented by people who had certain intentions. The Union Carbide executives who asked that such a computer be built had certain intentions: in effect, those executives wanted future victims of Union Carbide's malfeasance to be given an assurance they would receive compensation for the wrongs done to them. So it would be false to say that the aforementioned promise had *no* psychological basis. It was really only the final phases of the issuing of that promise – the part where the computer was mechanically carrying out the intentions of those kindly executives -- that could be so described.

§ This last point – that the executive who built the computer didn't have any *specific* individuals in mind – brings us to a distinction that is crucial to our way of dealing with GS and that, independently of that, is crucial to an understanding of law. The distinction in question is that between what I will call "wide-scope" and "narrow-scope" assurances.[127]

Let Smith be the executive who came up with the idea of having such a computer. It seems to me that, in implementing that computer (along with

the company policies relating thereto), Smith was making, or attempting to make, a kind of promise to future victims of Union Carbide. But in the statement that reports that assurance, the existential-operator must be given narrow-scope. That statement does *not* have the form:

(*) There is some x such that x has been wronged by Union Carbide and such that, in designing the computer, Smith wished to assure x that x will be given compensation.

Rather, that statement has the form:

(**) In implementing the computer, Smith wished it to be the case that, for any x, such that x should experience some misfortune as a result of Union Carbide's malfeasance, x have an assurance that x will be paid damages.

(*) describes what we might call a *wide-scope* (or *de re*) promise. (**) describes a *narrow-scope* (or *de dicto*) promise. If (*) is correct, that means that there is some specific entity to whom Smith is making a promise. He is making a promise to Mary or Fred or Bob.

If (**) is correct, that means that Smith is making an assurance not to any particular individual – not to Fred or Mary or Bob -- but to *anything* satisfying a certain description. In this case, that description would be: *thing x such that x has been harmed by Union Carbide's malfeasance.*

In order for (*) to be true, it is necessary that there be some *actual* entity that has been wronged by Union Carbide. But for (**) to be true, this is not necessary. (**) could be true even if Union Carbide ends up not harming a

single individual throughout its existence. In general, what I am calling "narrow-scope" assurances don't require that there be any *actual* person who is being assured of anything.

§ Another illustration may help show the relevance of these points. Suppose that Brown is the founding father of a country. Brown writes a constitution for that country. In that constitution, he writes: "It is hereby promised that anyone who is found to have been unjustly imprisoned by the government will never have to pay taxes again." As it happens, the government in question never does wrongfully imprison anyone.

Presumably a promise is a promise *to* somebody. Does it follow that Brown didn't make a promise? No. Obviously Brown did make *some* kind of promise. But when we are identifying the content of Brown's promise, we must take care to give the relevant operators appropriately wide-scope. The content of Brown's promise is given *not* by the proposition:

(*) there is some x such that x is (or has been or will be) wrongly imprisoned and Brown promises x that x will be exempted from having to pay taxes,

but rather by:

(**) Brown promises that, for any x such that x is wrongly imprisoned by the government, x will be exempted from having to pay taxes.

A statement of the form "x promised y to z" may be false on a wide-scope reading, but true on a narrow-scope reading. I will argue that, given a wide-scope reading, GS does indeed fail to qualify as an assurance, but that,

given a narrow-scope reading, it *does* so qualify.[128] I will also argue that *all* laws must be given a narrow-scope reading – even those that *do* provide actual protections to actual people – and that, consequently, our analysis of law doesn't have to be subjected to *ad hoc* changes to deal with flukes like GS.

§ Before we consider a completely "flukey" law like GS, let us see how these points apply to a law that is less flukey than GS, but more flukey than, say, the law prohibiting insider-trading. This will help put GS into perspective.

Suppose Congress passes a law that makes it illegal for employers to refuse to hire or promote someone x solely on the grounds that x has a certain sexual orientation. Let SO be this law.

SO does seem to be a promise: it assures people of the relevant sexual orientation that they needn't fear certain forms of discrimination. Suppose that SO comes about by a fluke. A random quantum event deposited a certain pattern of ink on a certain piece of paper, which then randomly landed on a certain Senator's desk. The Senator absent-mindedly signs it, thinking that it is some innocuous bill concerning a detail of the tax-code. It then goes to Congress, where as a result of a series of similar flukes, it is converted into law without any one person's even knowing of its existence.

Despite the random mode of its inception, SO still seems to be an assurance. It assures people of a certain sexual orientation that they needn't fear a certain kind of unfair treatment.

At the same time, SO is an *institutional* assurance. There was no one person who was making any assurances to people of that sexual orientation. In fact, none of the individuals composing congress even knew that any assurance of that sort was being provided to anybody.[129]

SO neatly conforms to (L5). At the same time, SO is a very different sort of thing from GS. But through extensions of the points we've made, we can show that none of these differences is relevant to the truth of our thesis.

§ Let us begin by re-stating the main points from the last few pages. First, something qualifies as a promise in virtue, in large part, of the socio-historical context. Second, institutional promises can be made in a mechanical, even mindless fashion. Third, statements of the form "x promised y to z" are ambiguous between wide- and narrow-scope readings.

When we synthesize these points, we can explain why GS is not an exception to (L5). First of all, so far as its being a promise is concerned, it is of no consequence that GS was made in an entirely automatic fashion, i.e. with little or no (immediate) assistance from human intentionality. We've seen that institutional promises can be made in a completely impersonal manner.

Further, it is of no consequence that there is no specific person, or group of persons, x such that GS is a promise to x of anything. After all, GS could be a narrow-scope promise. It could have the form "I promise that anything x satisfying description D be assured that…", as opposed to the form "there is some x such that x satisfies D and such that I promise x that…".

Finally, it is of no consequence that GS doesn't explicitly say who is being given the assurance in question or even what that assurance is. Where promises are concerned – indeed, where almost any speech-act is concerned -- implicature does quite as much as explicature. What it is that is being promised is as much a function of the socio-historical context as it is of what is explicitly encoded in the words involved in the act of promise making.

§ To find our bearings, let us continue, for the moment, to leave aside

flukes, like GS, and to focus on paradigmatic laws. The law prohibiting insider-trading assures the public that speculation will be conducted on a level playing-field, and that corporate insiders won't have an unfair advantage over others. Laws forbidding murder assure people that they will not be killed by those who wish them ill. Traffic laws assure citizens that their right to travel safely will be protected. Paradigm-cases of laws are assurances of the kind described by (L5).

The letter of the law doesn't usually say *who* is being protected by the law in question. That is because it doesn't have to be said. Who benefits from the laws prohibiting murder? People whose rights would be violated by the commission of murder. Who benefits from the laws forbidding insider-trading? People whose rights would be violated by the occurrence of insider-trading.

As we discussed earlier, not all laws prohibit things. The laws that give special privileges to those with more than five children don't *forbid* anything. But, of course, it is perfectly clear whose rights are protected by that law – those with more than five children.

We mustn't be led astray by psychological facts about the people who make the laws. When a law is enacted, the law-makers may have in mind a specific group or even individual. Suppose that Jones is a member of religion X and that, for this reason, he is unable to buy a house in a posh neighborhood despite the fact that he is an upstanding citizen who is willing to pay a more than reasonable amount. Fair-minded lawmakers respond to this incident by making it illegal to refuse to sell a house to somebody on the basis of that person's religious views.

But the resulting law is not just meant to protect Jones; it is meant to protect *anyone* whose rights would be violated by the practice of refusing to sell real estate to someone on the basis of religious considerations. That law

is a narrow-scope assurance: it assures that, for any x, x's ability to buy property will not be limited by x's religious views.

The corresponding wide-scope assurance would be this: there is some x such that the law assures that x will not be denied the right to buy property on the basis of x's religious views. If *this* proposition were the one giving the content of the law, then that law wouldn't be of much use to future victims of religious intolerance. (Soon we will discuss other reasons why that law is given by a narrow-, as opposed to wide-scope, assurance.)

§ As we noted, whether a given act constitutes a certain kind of speech act is a function in large part of the social context. The words "I'm coming over" could be a promise, a threat, or a simple statement of fact. It depends on the context.

Given this, let us look at the cultural context in which laws, including GS, are made. What does Congress *usually* do?[130] It *usually* makes laws like those just described. It *usually* provides explicit governmental assurances of protections. Given this socio-historical fact, there is a presumption that Congress is in the business of providing such assurances.

The basis of this presumption is not purely statistical. It is not just that Congress *happens* to provide such assurances on a regular basis. It is pretty clear that it is the job of Congress to do so, and that the frequency with which Congress provides such assurances is rooted in this systemic fact.

Let me explain this last point. Given how one becomes a member of Congress (one makes various speeches in which one makes promises to one's future constituency); given the expectations that the electorate has of those who are in Congress (they are expected to live up to those promises, and to serve the interests of the public); given the benchmarks in terms of which a congressperson's performance is assessed (her performance is evaluated in

terms of how adequately she has fulfilled the expectations just described); given the laws regulating the behavior of members of Congress (they cannot pass laws that make it impossible for them to be voted out of office by a disappointed electorate, they cannot pass laws giving themselves special exemptions from the law); and, finally, given what members of congress usually do (where the term "usually" corresponds not to a purely statistical regularity, but to one that is causally grounded in the truths described in the preceding clauses), it seems fair to say that Congress is *supposed* to serve the public good by providing protections. In other words, Congress doesn't just *happen* to provide assurances of protections: it is in the business of providing such assurances.[131]

As we noted a moment ago, not all promises use the words "I promise" (or their equivalent in another language). Sometimes context is enough to know that "I will be there tomorrow" is to be taken as a promise. Given only the verbiage found in the books containing the laws passed by Congress, it isn't always clear that some assurance is provided by the law. Further, given only that verbiage, it isn't clear *who* is being given an assurance, or what that assurance is. But all of those lacunae are filled by the contextual information described in the previous paragraph. I will now try to say *how* exactly those lacunae are filled.

§ Let us now close the argument. When a fluke like GS occurs, it cannot be looked at in a vacuum. It must be interpreted in light of the contextual information just described. Given this information, here is what, I propose, must be said: GS is an assurance, to anyone x whose rights would be violated by the wearing (on somebody or other's part) of green shirts indoors, that x will not have to fear such a violation.

It may be true that there is no person whose rights would be violated by a

case of green-shirt-wearing. But that is irrelevant. As is the case with many promises, there is no person x such that GS assures x of anything. Rather, GS assures that anyone x satisfying a certain description will be given a certain kind of protection.

It may be true that there is no *actual* right whose protection is assured by GS. Human nature and social institutions being what they are, it may well be that there is no person x, and no right R, such that the wearing of green-shirts indoors would violate x's possession of R. But that is irrelevant. GS is not given by the statement:

(*) there is some person x, and some right R had by x, such that GS promises x that her possession of R will be respected.

Rather, GS is given by the statement:

(**) GS is an assurance that, for any person x[132] and any right R had by x that is violated by the occurrence of green-shirt-wearing, x's possession of R will be protected.

So even though no *actual* person's *actual* right is assured by GS, that doesn't bear against our thesis that all laws, including GS, are assurances of protections of rights.

The apparent problem posed for our thesis by GS vanishes when we take into account the scope-based ambiguities brought to our attention in 1905 by Russell.[133] In any case, I have provided *some* reason for thinking as much. I would now like to provide additional reasons.

§ Let us now consider a set of circumstances under which GS would *not* be

a fluke. This will show that, despite its use of what might appear to be arcane technical points having nothing to do with the concept of law, our analysis has independent corroboration.

Let us start with a bit of fiction. It turns out that, because of some strange, hitherto unknown law of physics, if anyone anywhere is wearing a green shirt indoors, that causes massive amounts of some toxic gas to be secreted into the atmosphere. The presence of that gas in the atmosphere leads to death and injury on a large scale. The facts just stated become well known. But there are a few people who, for whatever reason, *still* insist on wearing green shirts indoors. Under pressure from their constituents and also from their own senses of right and wrong, members of congress appropriately make it illegal to wear a green shirt indoors.

Under *these* circumstances, GS is plainly an assurance. GS assures that, for anyone x whose rights are violated by the wearing of green-shirts indoors, those rights of x's will be protected.

Here we must be very careful. Under the circumstances just described, *there is* some x such that the wearing of green-shirts violates some right R had by x. After all, supposing that Bob is a U.S. citizen who has been harmed by the gas, the wearing of green shirts indoors does violate *Bob's* right to health and life. Nonetheless, GS is *not* given by the statement:

(#) there is some individual x and some right R had by x such that x's possession of R is hereby assured.

Even though, under the circumstances described, there are many individuals x such that GS protects x's rights, there is not any one person x such that the proposition giving the content of GS mentions x specifically. GS would be quite useless, or much *less* useful than it would otherwise be, if

it were given by (#).

A brief detour may clarify my meaning. As Frege[134] pointed out, there is no *one* individual x such that the statement "all whales are mammals" says that *x specifically* is a mammal.[135] If you say "all whales are mammals", you leave it open whether Shamu is a mammal, unless you add the statement "Shamu is a whale." By itself, "all whales are mammals" says only that *if* Shamu is a whale, *then* Shamu is a mammal.[136]

It is true, of course, that "all whales are mammals" gives us information that is crucial to understanding Shamu's physical structure, and that it doesn't give us any such information in connection with the planet Neptune or Socrates – or anything else that is not a whale. So that statement about whales certainly has *an* important connection to Shamu that it does not have to non-whales like the planet Neptune. But the connection is not internal to the statement itself; it is mediated through the additional empirical proposition that Shamu is a whale. By itself, the *content* of "all whales are mammals" leaves it quite open whether Shamu is a mammal. That is why "all whales are mammals, but Shamu is not a mammal", though false, is not self-contradictory.[137] What we must say is that, *given* the auxiliary empirical information that Shamu is a whale, "all whales are mammals" tells us that Shamu is a mammal. So the connection is empirical[138], and thus not rooted solely in content.

Similarly, even though, under the circumstances described a moment ago, there are many individuals x such that GS protects x's right to life and health, there is no individual x such that GS, or the proposition giving its content, says anything about x specifically. The statement "GS does not assure Rutherford that Rutherford's right to life and health will be upheld" is not self-contradictory. It might turn out that Rutherford is a cat or a rock or a person who does not live under U.S. law (because he lives in New Zealand).

So even if Rutherford is *in fact* a U.S. citizen who is protected by GS, that fact is not internal to the proposition that gives the content of GS. What we just said about Rutherford is true of *any* individual x. So it is *false* to say that there is some x such that the proposition given GS's content has the form: *x is hereby assured that x's rights will be protected from the menace posed to them by the wearing of green shirts.*

Of course, Congress may well have been *motivated* to create GS by the fact that specific people – Bob, Mary, Fred…-- were harmed by the toxic gas. But, for the reasons just considered, none of these individuals has the property of being a thing x such that the content of GS is given by a statement mentioning x specifically.

§ A different example may be necessary to clarify the relevance that these points have to the nature of law. In the U.S. there is a law[139] that forbids non-handicapped people from parking in certain specially marked parking spaces. Let HC be that law.

There are obviously many individuals x such that HC assures protection of some actual right R had by x. Given this, suppose we said that HC is given by the following statement:

(A) There are various individuals x such that, for various actual rights R had by x, it is assured that R will be protected.

Now imagine a world W with the following properties. By sheer good luck, as opposed to nomic necessity, there are no handicapped people in W. Also by sheer good luck, there are never going to be any handicapped people in W. At the same time, there *have* been handicapped people in W, and the

law-makers in W very reasonably (though falsely) believe that there will be such people in the future – they don't think that this run of good luck is going to last forever. Finally, W is as much like our world as the conditions just described permit it to be. So in W, the United States is governed by the same laws as here, one of those laws being HC.[140]

In W, (A) would be a false statement. In W, there is no person any of whose rights are assured protection by HC – no person whose access to supermarkets, clothing stores, and medical facilities is assured protection by HC.[141] Thus, there is no person x whose right to food, clothing, etc. is protected by HC. Of course, in W, people have the right to have easy access to food-stores and the like. But, in W, HC doesn't protect anyone's possession of that right.[142] More precisely, in W, there is no person x, and no right R, such that HC protects x's possession of R. Thus, if we say that, in our world, HC is given by (A), then we are committed to the false counter-factual that HC doesn't exist in W.

To avoid this, we must say that HC is given by the statement:

(B) It is hereby assured that, for anyone x who is handicapped, and for any right R had by x that is violated by the shortage of parking spaces that are convenient for handicapped people, x's right to R is protected.

It is readily seen that (B) is true in W, even though there are no handicapped people in that world.

In general, if they are to have the right counterfactual properties, laws must be seen as narrow-scope assurances, even in cases of laws where there are actual individuals whose actual rights are assured protection.

§ If Smith is seeking something non-existent (e.g. a person who has never

told a single lie), he is still engaged in a search. If Smith is endeavoring to do something that cannot be done (e.g. to prove that arithmetic is complete), he is still engaged in an endeavor. If Smith promises to protect a non-existent moral right had by a non-existent person, he is still giving a promise that a moral right will be protected.

Let us make this last point more concrete. Smith goes on national television, and promises to give the next person who sets foot on the moon without a space-suit a medication that cures that person of the ravages that occur to his body as a result of oxygen-deprivation and exposure to cosmic rays. As it happens, no one ever sets foot on the moon again. Further, there is no possible medication that would undo the effects of oxygen-deprivation or of exposure to cosmic rays.

Despite all this, Smith has still made a promise. He has promised that, for any person x, such that x steps on the moon without a space-suit, he (Smith) will make it the case x have a medication y such that y reverses the damage done to x due to oxygen-deprivation and exposure to cosmic rays. Even if no person ever *does* set foot on the moon again (let alone without wearing a space-suit), Smith has still made a promise or assurance. (If somebody were to set foot on the moon without a space-suit, and was *not* subsequently given any special medication, that person could rightly accuse Smith of breaking his promise. In her deliberations with herself, somebody deciding whether or not to go to the moon could correctly tell herself: “there is the fact that Smith has promised that…”)

There is nothing strange about a real promise or assurance made to somebody non-existent of something non-existent. Such a promise is dealt with in the same way as the fact that John is engaged in a search, even though what he is seeking is non-existent. It is dealt with *not* by positing things that don’t exist, but merely by putting the relevant operators in the correct order.

§ Let us sum up. Assurances can be made impersonally, even mechanically. Institutional promises sometimes have this property. Further, whether something counts as an assurance or promise is often a function, in large part, of the context – it is a function of what people *expect* under the circumstances. So given that *Congress* does such and such, and given the functional and historical role that Congress has in our society, such and such will be seen in a certain light. Those acts will register as promises.[143] Finally, laws are *narrow-scope* assurances. Given these points, the existence of flukes, like GS, is no threat to (L5). [144]

A different way of dealing with GS

There is a different way of dealing with laws like GS – laws that have a "flukey" character and that, in connection with this, appear not to conform to L5. It will help if we take a brief detour through a branch of philosophy that might appear unrelated to the law.

A car that doesn't run is still a car. Your car doesn't cease to exist when it stops running. If we say that something *must* run to be a car, then we render tautologically true the falsehood that all cars work.[145]

This raises a problem. The obvious definition of "car" is something like: "four-wheeled motor-powered vehicle."[146] Given what we just said, this definition is false. After all, a car's motor might be defective, and thus not power anything.[147]

The way to deal with this is by noting that the concept associated with the term "car" has a functional or teleological component. Something is a car not solely in virtue of what it actually does, but in virtue, at least in part, of what it is *supposed* to do.[148] A defective car that rolls off the assembly line

is *supposed* to perform certain functions. Given the objectives of the people who built the factory that produced the object in question; given the fact that the factory usually produces things that are consistent with those objectives; given that this last fact is not a purely statistical phenomenon, but is, to some degree, guaranteed by the physical conditions that obtain in the factory; and, finally, given the expectations that the community has of the things produced by the factory, there is clearly a sense in which the products of that factory are cars, even those that are non-operational.[149]

§ Let us reconsider GS in light of these points. What does Congress generally do? It *usually* makes laws that protect, or are clearly supposed to protect, rights. It makes laws like: "employers cannot make hiring- and promotion-decisions on the basis of race or gender", "one cannot buy or sell stock on the basis of information that isn't available to the public", "one must be read one's Miranda rights if anything one says to police is to be admissible in court."

Given this, when a legal "lemon" like GS is produced, there is clearly a sense in which it is *supposed* to give somebody some sort of protection. We might thus decide to reject (L5), and replace it with something like:

(L6) A law is something that (given its mode of origination, the intentions of those behind it, the expectations of those subject to it…) is *supposed* to be an explicit governmental assurance of the protection of a right.

In other words, a law is *supposed* to be the sort of thing described by (L5).

§ On the face of it, (L6) appears to have some advantages over (L5). For

one thing, (L6) might appear easier to reconcile with the fact that some laws seem less like assurances of protection than assurances of harm, as the following illustrates.

Consider the law that Stalin created whereby people must denounce their own family members for speaking ill of Stalin's leadership. That law was *supposed* to assure the Soviet people protection from the corrupting influence of counter-revolutionary thought.[150] If enforced, that law wouldn't have given any actual protection to anybody and, in fact, would do a great deal of harm to those subject to it. For the most part, the people of the Soviet Union probably knew this. So that law seems less like an assurance than a threat, making it unclear how (L5) could possibly be true of this law.[151]

(L6) seems not to have this problem, at least not as acutely. There is clearly some sense in which Stalin's evil law was *supposed* to be an assurance of a protection. Given the cultural background – given the propaganda in which Stalin's subjects were inundated, given the (supposed) objectives of the revolution, given the values that were promoted by Stalin's regime – there is clearly a sense in which that evil law was *supposed* to be an assurance of a protection. Relative to the propaganda issued by Stalin's regime, it is clear how that law was *supposed* to protect Soviet citizens.

§ Nonetheless, I do not think that there is any legitimate reason to prefer (L6) to (L5), or *vice versa*. Whenever we have a law that, in the relevant sense, is *supposed* to be an assurance of a protection of a right, it is easily seen to be an *actual* assurance of a right – albeit a "narrow-scope" assurance. I am inclined to think, in fact, that (L6) and (L5) are equivalent.

Consider the law just discussed. It is true that[152] there was no actual person x and no actual right R such that x's possession of R was assured by

protection by that law. But that is irrelevant. Stalin's evil law was given by the claim:

(a) It is assured that, given any person x and any right R had by x that is imperiled by the sort of counter-revolutionary menace posed by those who criticize Stalin's leadership, x's possession of R is hereby guaranteed protection.

That evil law is not given by the claim:

(b) There is some actual person x and some actual right had by x such that it is hereby assured that x's possession of R is protected from the menace posed by those counter-revolutionary thought.

Even if – as party ideologists would have had the Soviet public believe – there were actual people whose actual rights were imperiled by the existence of those critical of Stalin's regime, the law in question would *still* not be given by (b). What we said about the laws providing parking spaces to the handicapped applies here.

The fact that Stalin's law is *supposed* to protect people's rights is dealt with in the way that false suppositions are always dealt with in philosophy: by making sure that they are safely to the right of all the relevant operators.

§ It may be useful to take a brief look back at some of the points discussed earlier relating to the philosophy of language. The fact that there is no fountain of youth doesn't falsify, or otherwise impugn, our assertion that Smith believes the fountain of youth to be in Florida. It only means that when

we report Smith's belief, we must make sure that the false proposition believed (i.e. supposed true) by Smith is given narrow-scope with respect to the epistemic operator. We say: "Smith believes that there is a (unique) fountain of youth x, and that x is in Florida", and not "there is a (unique) fountain of youth x, and Smith believes that x is Florida."[153] [154]

With these points in mind, let us revisit Stalin's evil law. The following is a true statement:

(a) "Stalin's law is supposed to protect the Soviet people from the grave dangers to their welfare posed by counter-revolutionary thinking."

(a) does *not* mean:

(b) There are various grave dangers x to the welfare of the Soviet people such that x is posed by counter-revolutionary thinking; moreover, it is supposed to be the case Stalin's law protects the Soviet people from x.

(b) is false. In any case, there are possible circumstances where (a) is true and (b) is false; so the two don't mean the same thing. The real meaning of (a) is given by the statement:

(c) It is supposed to be the case that, for any grave danger x that counter-revolutionary thought poses to the welfare of the Soviet people, Stalin's law protects the Soviet people from x.

The supposition that Stalin's law assures the protection of some actual good is indeed false. But that means that, in the statement giving the content

of that law, we must give narrow-scope to the proposition describing the good supposedly protected by that law.

There is nothing *ad hoc* about this approach. We have seen that where *all* laws are concerned – even those that clearly protect the actual rights of actual people – the propositions describing the goods supposedly (and, in some cases, actually) assured protection must be given narrow-scope.

§ *Prima facie* GS seems inconsistent with a moralistic conception of law. But the existence of laws like GS forces us to confront an important fact about the concept of law. Once this fact is brought to the surface, a positivist approach to law loses a great deal of its initial appeal.

Just as a non-flying plane is worse *as a plane* than one that flies, so an unjust law is worse *as a law* than one that is just. A law that serves the interests of those subject to it is better *as a law* than one that thwarts those interests. From the viewpoint of a dictator's machinations, an unjust law may be preferable to a just one – just as, from an evil doctor's viewpoint, a botched surgery may be preferable to a successful one. But when we think of laws, the standards of evaluation that most naturally come to mind concern justice and justice alone. (It could be argued that this last fact has to do only with our erroneous psychological associations, and not with anything essential to the concept of law. We will soon address this position.)

The question that naturally arises in connection with any law is: what good does it do? GS does no good, and is therefore a failure. This suggests that, even though GS is morally hollow, it is not *supposed* to be that way. In general, something succeeds (fails) *as a law* to the extent that it does good (bad), and *only* to that extent. This is consistent with our position that laws are promises or assurances on the part of the government that moral goods will be protected. We regard any promise that isn't kept as a failure (*qua*

promise), and any promise that is kept as a success (*qua* promise). While being consistent with the fact that there are morally hollow and even evil laws, it also suggests, if it doesn't outright entail, that there is a principled relationship between law and morality.

GS is a law that doesn't protect *any* moral good. But it also seems to be a *defective* law. (To say that GS is defective is not to say that it isn't really a law. A defective car really is a car.) For this reason, GS's moral hollowness rather vindicates a moralistic conception of law. Given its genetic or formal properties, we have certain expectations of GS. It is made by the same mechanism that issues prohibitions of murder and theft, thereby protecting the individual's right not to be murdered or stolen from. We naturally see such morally pregnant laws as being the benchmarks in terms of which GS is to be understood. It is thus natural to see GS as an *attempt to do good*, albeit a failed attempt. This last fact is inconsistent with Hart's position that laws ultimately have only a circumstantial connection to morality.

§ A positivist might respond to this argument by saying the following:

You insist that an unjust (or otherwise harmful) law is, on that account, less successful *as a law* than one that is just (or otherwise does good). I admit that, if that is granted, then a case can be made for the sort of analysis of law that you propose. But your insistence on that point is entirely question-begging. A positivist's position is, in effect, that laws are *not* to be understood in terms of their relation to morality. A positivist would say that the benchmarks in terms of which *you*, apparently, are inclined to measure the "success" of a law have no basis in the concept of law, or at least nothing like the basis you think them to have. So I don't really see that you've made any progress here.

I would like to argue that, contrary to what the objector says, my argument does *not* beg the question. For the moment, let us set aside talk about laws, and instead discuss the concept of government. It would be hard to deny that a government fails *as* a government in so far as it thwarts the interests of those it governs. It seems a datum, even a tautology, that government x is better than y exactly if, all other things being equal, x does a better job than y of protecting the interests of the governed. It is hard to see what other benchmark could apply.

Of course, the people *in* government may not have any particular desire to serve the interests of their constituents. Their objectives may be entirely selfish (power, wealth, prestige). But that means that they don't care whether they are good governors. It does not mean that, in their case, being a good governor does not involve protecting the interests of those governed. Measured in terms of how well they *used* their governmental authority to achieve non-governmental objectives, such evil-doers may be great successes. But measured strictly in terms of their performance *as* governors, it would be rather hard to make the case they were successful.

I don't think that this intuition is an artifact of my own position. I have never heard a historian, or even a cynical political philosopher[155], say that *governmental* success could be understood in terms of anything other than the welfare of the governed. Stalin and Pol Pot were great successes – when it came to the acquisition and retention of power. But surely neither was a success *as* a head of government.

A doctor is a failure *as a doctor* in so far as he fails to heal his patients. Of course, we can imagine circumstances where the *personal* success of a doctor depends on his harming his patients. For example, a doctor may be in the employ of an evil dictator who pays him handsomely for each patient he

kills. But, under those circumstances, the success that such killing brings the doctor is not success *in* the practice of medicine; it is some kind of success that merely involves the practice (or mal-practice) of medicine. It is a tautology that a doctor *qua* doctor is supposed to heal. It seems scarcely less a tautology that a government *qua* government is supposed to protect the interests of the governed.

Once it is granted that governmental success is to be measured in terms of the welfare of the governed, it is not, I think, so hard to make a case that the successfulness of a legal system is to be measured in similar terms. There are no laws where there is no government. (That seems to be independently plausible, even a truism, and not merely a consequence of the peculiar system advocated in this work. It would be strange to say, "there are laws in place X, but there is no government there.") Laws are among the instruments of government. Given that governmental success is to be measured in terms of the welfare of the governed, the same must presumably be true of the success of a legal system.

Let us now close the argument. Given that a legal system is a success precisely to the extent that it serves the interests of those subject to it, it is not unreasonable to say that the same is true of the legal system deployed by a government. As we've seen, once it is granted that the success of a legal system is to be measured in terms of how much good it does, the truth of our moralistic conception of law straightforwardly follows.

§ In light of these points, we have a plausible answer to the question "why doesn't the existence of laws like GS warrant the rejection of our analysis of law (i.e. of our point that laws are explicit governmental assurances of protections of moral goods)?"

The law prohibiting murder protects the rights that would be violated by

the commission of murder. But we cannot identify laws with *actual* protections of rights, since there could be (and are) *unenforced* laws prohibiting acts that would violate people's rights. The *unenforced* laws (in certain countries) that prohibit spousal abuse are *supposed* to protect the rights of those whose rights would be violated by the perpetration of such crimes. In effect, those laws are the government's way of saying (falsely) to those (would-be) victims: "Fear not! Anyone who commits spousal abuse will be punished; so such abuse will be uncommon." Laws are thus naturally seen as *assurances* of protections of rights.

For the reasons we've discussed, laws are *supposed* to serve the interests of those subject to them. GS is *supposed* to help those subject to the relevant legal system. How *would* it do so if it did do so? It would protect those people from the harm, and consequent violations of their rights, done to them by cases of indoor green-shirt wearing. The reason we regard GS as such a failure is precisely that, in prohibiting the wearing of green shirts indoors, GS doesn't protect anybody's rights. The one circumstance under which GS would *not* be a failure is one in which the wearing of green shirts indoors *does* harm people and where, thanks to GS, that harm is now prevented. This makes it clear that what GS is *supposed* to do – what it would have to do to *not* be a failure as a law – is to protect people from violations of their rights caused by the wearing of green shirts indoors. Given this, we may naturally regard GS as the government's way of saying: "those of you who would be harmed by the wearing of green-shirts indoors – fear not! Any case of wearing a green shirt indoors will be henceforth punished and, consequently, the wearing of green shirts indoors will be much less likely than it would otherwise be."

Some may resist this analysis on the grounds that no individual *in* government had any intention of assuring anyone of anything. But we have

already seen the difficulties associated with such a viewpoint.

§ In light of the points just made, a number of grounds for resistance to our analysis of law prove less cogent than they might initially seem to be. Consider the U.S. law that requires adult males to enlist with the selective service. How can this law be regarded as an *assurance* of protection of a moral good? It just seems to be a kind of rule – one entity (the government) telling various other entities (eight-teen year old males) what to do.

Let us deal with this problem by asking a question: what is that law *supposed* to do? What is the raison d'être for it? It is supposed to make it easier to draft people, and thus for the government to supply the army with soldiers in times of war. So that law is supposed to help maintain the army. In its turn, the army is supposed to protect the citizens of the U.S. (And that is exactly what it does do -- whether by deterring foreign adversaries through its mere existence or through actual military action.) Whose rights would risk being violated if adult males did *not* enlist with the selective service? Everyone whose rights would be violated if there weren't a standing army. In other words, everyone (or almost everyone) in the U.S.

So the law in question is meant to protect (almost) everyone in the U.S. That law can thus be seen as the government's saying: "those of you whose rights would be violated were there no standing army, or were that army inadequately manned – fear not! It will be adequately manned, because if adult males don't enlist with the selective service, they will be penalized." We cannot say that the law *per se* is such a protection. After all, we could imagine that law's existing but not being enforced. But in such a case what we would have is an assurance of the sort just described that was *not* kept.

In conclusion, given the wide-scope/narrow-scope distinction, and given the fact that institutional promises don't always (at least not obviously) have

the properties that we associate with promise-making at the individual level, the existence of laws like GS appears not to warrant the rejection of our analysis of law.

Chapter 9 How are a brutal dictator's laws different from a gunman's threats?

In the last chapter, we considered a powerful objection to our analysis of law. In this chapter, I would like to consider another powerful objection to it.

In a recent article, a distinguished philosopher of law writes: "Surely the legal systems of brutal dictators would fall off the scale as being in the final analysis nothing more than the gunman in Hart's example."[156]

Obviously a gunman's demands don't necessarily have any basis in morality. The gunman isn't attempting to do any good, and he isn't promising (even falsely) that he will do any good. If this philosopher's plausible point is correct, then presumably the same is true of the so-called laws issued by a brutal dictator. In that case, my analysis of law is simply false.

Consequently, if I am to hold onto that analysis, I must *deny* that those so-called laws really are laws. But in that case, it looks as though I am guilty of the sort of revisionism of which, in chapters past, I have accused my opponents. After all, it seems a datum (and, I believe, it *is* a datum) that Stalin, and other gunman-like dictators, issued laws.[157] In any case, I do not want to weaken my analysis by making some revisionist view be a pre-condition for its success – by, for example, making my analysis implausibly require that Stalin and Mao never issued laws.

§ There are a number of points to be made in response to this criticism and in response to the aforementioned philosopher's point.

We saw earlier that, at bottom, a government is a protective entity. We must look at a dictator's laws in this light. The relationship of a brutal dictator to his subjects is, to some degree, symbiotic. The dictator cannot *single-handedly* keep an entire nation hostage. He needs the co-operation of military, law-enforcement, and administrative personnel. This means that the dictator must make certain concessions to these people; he must give them incentives to stay loyal. That by itself means that the dictator's power has *some* roots in the community he governs.

But the story doesn't end there. For each person x who directly benefits from the favors of the brutal dictator, there are people who, in their turn, directly benefit from x. Further, where those second-order beneficiaries are concerned, there are often third-order beneficiaries – and so on.

So even in the most one-sided of dictatorships, the dictator has *relatively* deep roots in the community. Otherwise, he wouldn't be able to hold onto power. A brutal dictator's legal system consolidates this symbiotic, though deeply inequitable, relationship.

By contrast, a gunman provides nothing for his hostages. There is no symbiosis. This is not to say that brutal dictators are themselves morally better than gunmen. But it is a fact that dictatorships presuppose a certain give and take, though an inequitable one, between the ruler and the ruled. A brutal dictator's legal system thus has *some* basis, though an insufficient one, in the moral needs of the community. A gunman's demands, on the other hand, have no basis at all in the moral needs of his hostages. So, to that extent at least, the legal system of a brutal dictator is not comparable with a gunman's demands.[158]

§ But these points refers us to what I believe may be a far deeper difference between a dictator's laws and a gunman's demands.

When a government official acts, he does so within a certain framework. I submit that this not only causally, but *constitutively*, puts limits on what kinds of morality and immorality are available to the official in question.

In the case of somebody who is operating in a vacuum, there is a fairly direct link between the wickedness or goodness of his motives, on the one hand, and the moral merit of the act itself. But in the case of a government official – whether he is a dictator or a duly elected congressman -- the meritoriousness of an act cannot be assessed along such strictly internalist lines. Facts about the moral character of the relevant institutions must also be taken into account. From some viewpoint, the institution is acting *through* the official. (I admit that this is a very vague statement. In a moment, I hope to give it at least a little more precision with the help of some insights due to Pettit and Jackson.) To put it more innocuously, when an official (e.g. a dictator) acts, the institution within which, and through which, he acts is implicated in that act. Consequently, if the institution embodies a certain morality, the official's act will as well, even though that official's personal motives may be entirely immoral or amoral.

§ Suppose that Smith is an elected senator within a country that has a morally wholesome political order, i.e. a political structure whose behavior is responsive to the needs and wishes of the people it governs. As a human being, Smith is entirely unscrupulous. He became a senator not because he wants to help people, but simply because he wants power. At the same time, given the laws governing the conduct of Senators in Smith's country, it would be next to impossible for him to pursue his self-interest without, to

some degree or other, doing good for his constituency. When Smith chooses to do X, he is in fact motivated by a desire for power – a desire to dominate, even belittle, others. But when he does X, the wholesomeness of the relevant political structure is expressed in that act.

There is a sense in which the relevant political structure is acting through Smith – a sense in which Smith, and his despicable personal motivations, are a mere vehicle for a larger, morally wholesome order. (When, in a moment, we discuss Jackson and Pettit's work, we will make it somewhat clearer what that sense is.) The morally wholesome framework within which Smith is acting so constrains his behavior that Smith's act ends up embodying a certain morality, even though Smith's psychological motivations are entirely immoral. The *institutional* basis of Smith's act is morally wholesome, even though the same cannot be said of its psychological basis, i.e. of Smith's motives. Since Smith's act is not just an act of an individual, but also of an institution, it follows (so it could be argued) that Smith's act *is* moral [159], at least in some respect, granting that Smith himself may have a gunman's psychology.

§ Earlier we briefly discussed how the moral character of a social structure cannot always be read off of the motives of those composing it. Jones, Brown, and Rutherford may all be self-interested scoundrels. But the political structure they create may yet embody a certain morality.

It seems, in fact, that for a social or political system to be good is *precisely* for its moral character to have a certain independence of those of the people composing it. Suppose that we have a political system that concentrates all power in the hands of a single dictator. It happens that, at the moment, the dictator is a supremely good human being and that, under his directorship, the government consequently does much good.

What we have here is not a good political order. We have a bad political order that happens to be manned by a good person, so that its badness is momentarily masked. That badness is expressed in the fact that when a less scrupulous person takes over, the system does nothing to shield its constituents from the immorality of their leader. A good political order is one that does good even when the people composing it are not good.[160]

A related point is that an *institutional* act may be morally wholesome, even though the psychological motivations of the relevant individuals may be immoral. Senator Smith privately wants to harm Jones. But Smith helps Jones, because the political structure within which Smith operates requires him to do so. In this case, an institution did something morally good *through* somebody whose personal motivations were immoral.

Let us generalize this line of thought. Wholesome institutional acts are typically realized by, and thus hard to distinguish from, acts of individuals. The motivations of those individuals may be immoral. This suggests that, when we assess the moral caliber of actions performed by those operating within an institution, we cannot use a wholly individualistic, psychological frame of reference. As I would now like to argue, this means that, even in cases where the dictator's psychological profile is indistinguishable from that of the gunman, the acts of the former are not likely to fall to quite as low a moral level as those of the latter.[161]

§ As we've seen, a government cannot be *wholly* antagonistic towards the people under its dominion. A government is, by definition, something that provides certain protections. The earlier described omnipotent ogre who does nothing to help the people of society S is not the government of S, or even a part of that government.

When a dictator acts, he does so within the framework defined by a

government – by something that necessarily has *some* moral value. For the reasons described a moment ago, so long as a dictator is operating within a structure that embodies *some* morality, a certain morality will tend to accrue to his actions, even though he himself may be motivated by pure malice.

The operative word is "tends." It seems that, if the morality inherent in the relevant political structure is to accrue to the dictator's acts, then the system must impose significant constraints on what the dictator does. The less determinative the system is of the dictator's conduct, the less the morality inherent in the former will be transmitted to the latter. So the morality inherent in the system *tends* to accrue to the dictator's actions, depending on the extent to which the system constrains what the dictator can do, and depending also on the extent to which that system is itself morally wholesome. A dictator is, almost by definition, given considerable latitude by the institution within which he operates. So the morality inherent in that institution, which may be small to begin with, may do tragically little in the way of imposing moral constraints on the dictator's behavior.

Nonetheless, there is a categorical moral difference between the acts of the dictator and those of the gunman. The gunman's act is not that of a benevolent institution: when he acts, nothing benevolent acts *through* him. His act is only as meritorious as his own motivations. When a dictator acts, an at least minimally benevolent institution does act *through* him. So, as a rule, his act is *not* as morally empty as his motivations.

§ I believe that a striking insight due to Phillip Pettit[162], and developed by Frank Jackson in collaboration with Pettit[163], is relevant to this topic, and will also give a fairly precise meaning to my grand and vague statement that "when a dictator acts, a benevolent institution acts through him." For reasons to be discussed later, I also believe that it would be very hard to

understand the nature of our obligation to obey the law without taking into account this insight.

We often hear assertions like the following: “the high-level of crime in the inner-cities is caused by the poverty prevalent in such places”; “sky-rocketing inflation, combined with a theocratic political system, led nation X to go to war with nation Y”; “a combination of patriotic fervor and increasing economic dependence on antagonistic foreign actors led the citizens of state X to elect a jingoistic war-like zealot.”

Such statements seem to allege the existence of distinctively *social forces*. The cause of some social condition is *poverty*, *inflation,* or *a theocratic system of government*. It seems a datum that such statements are not *categorically* false. Surely it is at least potentially true that, had economic conditions in Germany during the thirties been different, a totalitarian government would not have come into power.

But there is a problem with explanations of this type. The cause of any state of affairs is always *specific.* Suppose that nation X goes to war with nation Y. The cause of that situation consists of specific acts on the parts of specific people. Dictator Smith picks up the phone and makes certain noises to Minister Jones, who then puts a specific pin in a specific part of a specific map, which is then shown to deputy Minister Brown. And so on – until Private Rutherford makes a specific motion with a certain finger causing a certain bullet to penetrate a specific part of the thoracic cavity of Sergeant Murray.

When we take into account all of these specifics, the causal structure of the situation has been completely defined. There seems to be no room in that structure for grand forces like *inflation, poverty,* and the like.

Given this, some might say that those grand social forces don’t exist and that statements involving them are categorically false. But this seems unduly

nihilistic. If we purged our lexicon of terms like “inflation”, “poverty” and “democracy”, it would be hard, perhaps impossible, to express many known truths.

§ Pettit and Jackson give us a plausible and independently corroborated alternative to this nihilism. The system they propose involves the concept of what they call a *program-cause.* In connection with this, it involves a distinction between what they call causal *efficacy* and causal *relevance*.

Consider some true causal statement:

(*) The flame heated the pan.

The truth is that what really heated the pan was not *the flame.* That is much too encompassing an entity. The state of the pan is the result of the activity of certain specific mass-energy displacements at, or vanishingly near, the underside of the pan. The flame as a whole is innocuous, except in so far as it comprises those specific displacements. Had those specific mass-energy displacements occurred, the pan would have heated, even if the rest of the flame had been absent. Jackson and Pettit describe those specific mass-energy displacements as being “causally efficacious.”

Even though the flame as a whole is not causally efficacious, (*) is not exactly a *false* causal statement. It is not in the same category as a statement like “Smith’s coughing heated the pan” or “the fact that Mozart’s Jupiter Symphony is more than two minutes long is what heated the pan.” So in what sense is (*) true?

The presence of the flame in the relevant place *guarantees* the occurrence of the aforementioned causally efficacious events. *Ceteris paribus,* in any counterfactual scenario where that flame exists, there will be

causally efficacious events of the kind just described (though not necessarily those specific ones). So, as Jackson and Pettit put it, the presence of the flame *programs for* the right kind of causally efficacious events, and is therefore the "program cause" of the heating of the pan. The flame, say Jackson and Pettit, is thus "causally relevant", even though it is not causally efficacious. (If an event is causally efficacious in producing x, I will say that it is an "efficient" cause of x. My use of this expression is not intended to conform to Aristotle's, or his translators', use of that expression. If it does so conform, that is accidental.)

When we distinguish between causal relevance and causal efficacy, and recognize the concept of a program cause, we can hold onto causal statements, like (*), that we intuitively know to be true.[164]

Pettit and Jackson show how this distinction enables us to hold onto statements referring to general social forces. Nation X wages war with nation Y. What is *causally efficacious* in bringing this about is not some grand economic or political force, but rather a collection of specific actions. (The Generalissimo picks up a certain phone, pushes certain buttons, makes certain noises into the phone…) But the existence of large social forces *programs for* such specific events. The fact that nation X is in such dire economic straits, along with the fact that it has a system of government of such and such a type, *programs for* the sorts of specific events that will be causally efficacious in bringing about war between X and Y – just as the existence of a flame in the vicinity of the pan programs for the occurrence of specific micro-events that actually do heat the pan.

§ The Jackson-Pettit analysis suggests that there are two different things to bear in mind when evaluating the moral character of an act. We must bear in mind not only the forces that are causally efficacious in bringing about that

act, but also those that are causally *relevant*.

With these points in mind, let us turn our attention back to the case of the dictator. Dictator Smith's personal motives are causally efficacious in bringing about act X.[165] Those motives are, we may suppose, no better than those of a gunman. But the program-cause of Smith's act is a political order that necessarily embodies at least a flicker of morality. So unless we categorically exclude all consideration of program causes from our assessments of the dictator's act, that act will tend to embody a certain morality.

§ One might respond to this by saying: "Very well then. When we asses the moral caliber of dictator's act, let us leave out the institutional program-causes, and focus only on the causally efficacious mental states."

This is by no means an unreasonable proposal. I even think that, in certain contexts, it may be appropriate. But before accepting it wholesale, we must heed the following caveat. When we morally assess "the dictator's act", we want to study the very thing that reporters, historians, and economists denote with that expression. To echo a point made earlier, we don't want to study an artifact of our own philosophical analyses – of, in this case, our *a priori* belief that the dictator is in the same moral category as the gunman. As Jackson and Pettit stress, in most explanatory contexts, we are concerned with the causally relevant, and not merely with the causally efficacious. Events are typically individuated and understood in terms of program-causes: causally *efficacious* states of affairs are seldom discussed. If I want to know why the house burned down, I don't want to be told about micro-events e_1… e_n occurring in some vanishingly small region of space-time R. I want to be told something more pregnant, for example: "your former-friend Jake torched the place in retaliation for your negative review of his book" or "the occupant

of the house needed money and was trying to defraud the insurance company." The kind of micro-information in terms of which causally efficacious forces are described is explanatorily inert except in the arena of micro-physics.[166]

So if we insist *a priori* that our conception of event-individuation exclude consideration of program-causes, we run the risk of severing the explanatory links that matter the most to us. To be sure, there clearly are contexts where we are concerned with the dictator's personal motives for a certain act, and not with the social conditions that programmed for that act. But in other contexts we may wish to view the dictator's act from a less individualistic perspective. In any case, we cannot say *a priori* that we categorically never wish to view the dictator's act in that light.

§ There is another standpoint from which to compare and contrast a government's laws with a gunman's threats.

A law is a relatively enduring feature of a political system; it is not an event. What is an event is the *issuing* of a law. The U.S. law prohibiting insider-trading was *issued* by an act of Congress in 1933. But the law itself is not an occurrence; it is an enduring, though mutable, structural fact about the American legal system. (Even if a law exists only for one minute, and even if the event of making a law takes ten years, the law is, whereas the event of issuing the law is not, an enduring condition. The criterion by which something qualifies as an event as opposed to a condition, or *vice versa*, is not to be understood strictly in terms of how much time that thing occupies. [167])

So not everything that is true of the *act* of issuing a law is true of the law itself. In particular, I will now argue, even if the *act* of issuing the law is motivated by pure malice, it doesn't follow that the law itself is comparably

immoral. Indeed, where certain kinds of *institutional* acts are concerned, there will *systematically* be such a moral divergence; and where *non*-institutional acts are concerned, there will systematically *fail* to be such a divergence.

The law prohibiting insider-trading was (let us pretend, for the sake of argument) issued by legislators who (mistakenly) thought that such a law would harm the U.S. economy. But they still chose to enact it, simply because they were jealous of the financial success of some of their old college buddies, and knew that the law in question would bring that success to a speedy halt. Here a morally good law is created by an act which (setting aside our contention that, in some contexts, such an act must be looked at from an institutional, and not a strictly psychological, viewpoint) was lacking in moral value.

Up to a point, it may be a fluke that the acts of those petty congressmen led to a good law. But the structure of government *programs for* the fact that the law they create is something that meets at least certain minimal moral requirements. Indeed, that fact is programmed for by the very fact that those congressmen are operating within *a* government, let alone a relatively good one. So it is not *entirely* a fluke that the law in question is a good one.

A comparison may be appropriate. Suppose that a professional basketball player successfully makes a basket from 90 ft away from the hoop. The fact that the ball went *exactly* where it did is probably a matter of luck. But the fact that the ball went into that general area is a matter of skill. The basketball player cannot guarantee that the ball will go in *exactly* the right place. But given his level of skill, he can guarantee that the ball will go within (say) one foot of where it must go. So the fact that he scores a basket under those circumstances is, as we might put it, a relative, but not an absolute, fluke. If an amateur attempts the same thing, his success is a total,

and not merely a relative, fluke.

It is a fluke that the law those petty congressmen enacted was *as* good as it was. But it is not a fluke that it met certain minimal (and, in absolute terms, fairly high) moral requirements. The congressmen knew that they couldn't create a law that was *completely* immoral without dealing with blowback from their constituents. Also, those laws must meet requirements of Constitutionality. Petty though they may be, the congressmen are hemmed in by requirements that keep the laws they make in some kind of alignment, however minimal, with the needs of the citizenry. The structure of the government – the very fact that they are operating within *a* government -- *programs for* their making laws that are not entirely lacking in moral value. When malicious congressmen make a law that, despite their best efforts, happens to be good, that is going to be only a relative, and not a total, fluke.

Let us look at the gunman-scenario in light of these points. Just as we must distinguish the *act* of making a law from the law itself, so we must distinguish the gunman's act of *issuing* a threat from the threat itself. (Here language tends to fail us, since the term "threat", unlike the term "law", seems to be ambiguous between a certain kind of act and the ensuing structural fact. So a certain amount of verbal awkwardness is inevitable here. [168]) The issuing of the threat occurs at some specific time – the moment at which he says "Don't move! This is a stick-up…" The act of saying those things creates a certain condition – a state of terror, in which the people inside the bank rightly believe that they are not at liberty to do things which they would otherwise be able to do. Up to this point, the gunman-scenario is analogous with the law-maker's scenario. In both cases there is a specific act which creates a subsequent condition.

But now the analogy breaks down. In the case of the law that is issued

by ill-intentioned congressmen, the program-causes of that law ensure, or render probable, that the law will not be *as* morally empty as the act which created it. As a rule, laws created by malicious law-makers are not as bad as the acts (or individuals) that create them. But, as a rule, the condition created by the gunman's threat *is* going to fall to the moral level of that threat. (Remember that "as a rule" doesn't mean "100% of the time.") Of course, we could hypothesize situations where this is not the case. (The gunman issues a threat. Thanks to some strange causal mechanism, the sound-waves created by the gunman's words cause gumdrops to fall from the sky.) But such an event would be a total, and not merely a relative, fluke.

A brutal dictator is operating within a system that is less wholesome than the democratic framework within which U.S. congressmen act. But, for reasons we've discussed, such a framework will, as a rule, be more wholesome than that of the lone-gunman. For this reason – setting aside the point previously defended, to the effect that the dictator's acts must sometimes be looked at from an institutional, as opposed to strictly psychological, viewpoint -- the law created by a malicious dictator will, as a rule, have more moral value than the act of creating or issuing that law.

Let us sum up. We must distinguish the act of issuing a law from the law itself. Given the sort of thing that government is, laws made by malicious law-makers (including brutal dictators) will, as a rule, be morally better than the acts by which they were issued. So, as a rule, the legal systems of such scoundrels will not be as morally empty as the scoundrels themselves. But the gunman operates in the *absence* of the sorts conditions that create this moral disparity between the act and the ensuing condition. So, as a rule, the conditions created by his actions are not morally better than those actions themselves.

Summary of chapter 9

Some dictators are psychologically in the same category as gunmen. But when a dictator acts, a relatively wholesome institution is implicated in, and to some degree constitutive of, his act. The same cannot be said in connection with the gunman. For this reason, a dictator's act, unlike that of a gunman, necessarily has a certain dimension of morality.

We mustn't over-state things. There almost certainly will be contexts where we wish to view the dictator's act in a strictly internal, non-institutional fashion. From such a viewpoint, the dictator's act of issuing a law may sink to the low moral level of the gunman's threat. But it is also possible that, in other contexts, we do not wish to view the dictator's act from a strictly psychological standpoint, preferring to view it in broader institutional terms. We've seen some reason to believe that, thus viewed, the dictator's behavior is not comparable to that of the gunman, even though the two actors may be psychologically indistinguishable. [169] [170] [171]

But even if we set aside this last point, there is a difference between a gunman's threats and a malicious dictator's laws. Because the latter is operating within a government, the laws he creates will, as a rule, be better than the acts which created them. Because the former is *not* operating within such a framework, the conditions *he* creates will tend *not* to be better than the acts that create them.

Chapter 10 The moral structure of legal obligation

We have discussed the moral basis of law. Now we must discuss the moral basis of legal obligation. The connection between the two is not as straightforward as one might think.

Our position is that a law is a certain kind of assurance, on the part of the government, to protect some moral good. Even if this is correct, it does not follow in any obvious way that there is any moral imperative to *obey* the law. Consider Stalin's law requiring people to denounce anyone whom they suspect of anti-Stalinist sentiment. According to our analysis, that law *is* a promise to protect good, and it therefore has a clear moral basis (even though Stalin's intentions in issuing that law were immoral). But it is quite unclear what moral imperative there might be to obey that law. It would seem, in fact, that nothing would be more *im*moral than obeying it.

Even if we leave aside atrocities like the law just considered, there doesn't seem to any obvious reason why someone's *assurance* to do moral good – even if that assurance is plainly well-intentioned and thus has moral merit – should necessarily deserve a response any different from that of the most vile threat. Suppose that Smith is a public official who sincerely believes that consumption of large quantities of arsenic cures all physical ailments. Smith goes on television, and publicly "assures" all of his constituents that he will henceforth make sure that the food-supplies in his district are liberally supplemented with arsenic. Here Smith is providing his constituents with an assurance (a *narrow*-scope assurance) that he will do them good. (He is giving them an assurance that, for anyone x who is one of his constituents, x will be cured of x's ailments thanks to the pains Smith will take to make sure that x's food is replete with arsenic.) But it is obvious that our moral obligation is to thwart Smith's proposed efforts, not to comply with them.

Nonetheless, I will argue that legal obligation is *a* kind of moral obligation. In many cases, the moral obligation to obey the law is small compared with the moral obligation to break it. But the statement "you are not *supposed* to drive 66 mph if the posted limit is 65 mph" is not morally empty, and neither is the statement "you are *supposed* to denounce your best friend because he was critical of the Generalissimo." In such sentences, I will argue, the word "supposed" has a meaning identical with, and not merely analogous or similar to, the meaning it has in "you are supposed to help those in need." The general belief to the contrary involves a number of confusions that I now hope to expose.

Two erroneous reasons to deny that legal obligation has moral content

Let us start by undoing one confusion which, I think, lies behind certain conceptions of law. It is true that laws (in the legal sense) are not moral laws. It is not a moral law that one mustn't drive 70 mph. Because moral laws don't coincide with laws in the legal sense, many (e.g. Austin, Bentham) have been quick to conclude that moral and legal obligation are entirely unrelated.

Such a reaction would be too extreme. Given the physical conditions that obtain in my kitchen, the light turns on whenever the switch is flipped. That constant conjunction is *guaranteed* or *backed* by the laws of nature that hold. So the proposition *the light in J.-M.'s kitchen turns on whenever the switch is flipped* is backed by laws of nature, even though that proposition does not, by itself, express a law of nature.[172] Similarly, even though the proposition *one has a legal obligation to drive no more than 65 mph on certain highways* does not itself express a moral law, it doesn't follow that it is not *backed* by some moral law.[173]

§ This brings us to another confusion which, I believe, has given credibility to non-moralistic conceptions of law. Some legal obligations are down-right evil – e.g. a Roman citizen's obligation to report escaped slaves to the authorities. It would be immoral to uphold such obligations. This naturally induces people to think that a law can have *no* positive moral content.

Such a reaction involves a muddle that we are now quite familiar with: it falsely equates the concept *the right thing to do* with the concept *the sole thing that, under the circumstances, would have any moral merit at all*. In connection with this, it confuses the concept of a comparative good with the concept of an absolute good.

Some reasons to suspect that legal obligation has moral content

Let us begin by stating some general reasons to suspect that there is *a* moral component to legal obligation. What I say in this section is not meant to be probative, but only to induce an initial receptiveness to the possibility that, despite atrocious laws such as those associated with the political machinations of a Stalin or a Hitler, there is always a (defeasible) *moral* obligation to obey the law. It goes without saying that the moral obligation to obey the law is sometimes outweighed by a contrary moral obligation.[174]

As we've seen, a government is something that protects fundamental rights and provides basic services.[175] A government does this, in large part, by issuing laws. This suggests that laws have a dimension of morality and that legal obligation is not morally hollow. [176]

Here is a related line of thought. As we saw, there is no legal recognition, and hence no legal obligation, where there isn't governmental

protection. This suggests that there is more than a contingent, circumstantial connection between legal and moral obligation.

Here is a different line of thought. Some laws and legal systems do great harm. But, as we've discussed, it is possible that, when this is the case, we are dealing with some kind of pathological deviation that obscures the fundamentally moral nature of law and legal obligation. An over-grown or otherwise diseased liver leads to death. But it would be absurd to deny that *as a rule* livers are vital for our well-being. Given only that legal systems may do great harm, it would be equally absurd to deny that *as a rule* legal systems are fundamentally good things and to deny that *as a rule* there is a moral component to legal obligation.

§ There is one other reason to suspect that legal obligation has moral content. Consider the statement

(1) "the gunman has ordered you to hand over your money."

If the gunman falls asleep, and you escape without giving him your money, nothing that was *supposed* to happen has not happened. (1) doesn't have a *normative* component.

But the story is different with the statement:

(2) "the law prohibits you from driving 66 mph."

If you drive 66 mph, something that wasn't supposed to happen has happened. (2) thus has a normative component. This isn't to say that, if you drive 66 mph, anything regrettable has occurred. But it is a datum that there is a sense, however anemic, in which one is not *supposed* to drive 66 mph if

the posted limit is 65 mph.

It is hard to believe that the difference between (1) and (2) – that one of them is normatively pregnant while the other is not – has nothing to do with the differences in the moral characteristics of the relevant situations. A gunman isn't a part of a system that does anything positive for anybody. The law prohibiting you from driving 66 mph *is* part of such a system.

Let us briefly develop this last point. Many deaths and injuries are avoided, thanks to the fact that the speed limit is 65 mph as opposed to 95 mph or even 68 mph. So even though there isn't any obvious moral imperative to drive 65 mph as opposed to 66 mph, it is hard to believe that the difference between (1) and (2) – the fact that one of them is normatively-pregnant while the other is not – is completely unrelated to the background differences just discussed. It is hard to believe that this difference has no roots at all in the fact that traffic laws do good, or at least are a part of system that does good, whereas a gunman's orders are not a part of such a system.

As we've discussed, one has legal obligations where, and only where, one is dealing with a system that protects at least some of one's moral rights. Typically, if a legal system protects any of one's rights, it is likely to be the thing that protects most of them. If you fall within the jurisdiction of a legal system, it is more often than not the one thing which keeps you from being killed, enslaved, or otherwise dehumanized. It is hard to believe that this morally pregnant fact is unrelated to the fact that laws do, whereas a gunman's orders do not, correspond to *obligations* and that statements describing laws have normative content.

Of course, not all laws do as much good as the traffic law discussed a moment ago. So not everything that we just said about that law applies to laws in general. But what *is* true of laws in general, even wicked ones, is that they presuppose the existence of government and thus of a social structure

that is *fundamentally* moral. Given that statements describing laws are normatively pregnant, and that such statements are true only when the background situation satisfies rather high moral requirements, it is hard to believe that the normative component in legal statements is morally anti-septic.[177]

Though it may not establish it conclusively, this line of thought *suggests* that legal normativity is to be understood in terms of moral normativity, even after we allow for the existence of immoral or morally empty laws.

Some erroneous theories as the moral basis of legal obligation

In this section I wish to consider some plausible but *erroneous* theories as to the moral basis of legal obligation.

Governments protect their constituents. They do so, in large part, by issuing laws. Laws are effective only in so far as they are obeyed. Given this, it seems reasonable to say the following:

(*) *legal obligation is the moral obligation one has towards those whom one's government protects.*

But (*) is wrong. There are many cases where fulfillment of one's legal obligations would do nothing but harm. As a rule, I do nothing but harm if I obey the evil laws of a Stalin or a Hitler. There may be cases where, thanks to circuitous causal mechanisms, obeying such a law does good. But it would be hard to make a case that obedience to such laws *categorically* does a non-trivial amount of good.

Even if we leave aside immoral laws, (*) is still clearly wrong. I have a legal obligation not drive 66 mph if the posted limit is 65 mph. But it doesn't seem true that, if I drive 66 mph, I will hurt any of the people whom the government protects. Nor does it seem[178] that I would necessarily weaken the authority of the government and, therewith, its ability to protect people.

Of course, one can imagine circumstances where driving 66 mph (in a place where the posted limit is 65 mph) weakens the authority of the government. But one can just as easily imagine circumstances where driving 64 mph is legal and yet undermines government authority. (Not everything that undermines governmental authority is *ipso facto* illegal. Otherwise governments wouldn't have *make* such acts illegal by passing special laws.) So it doesn't appear that one's *legal* obligation not to drive 66 mph coincides with one's moral obligation not to undermine the authority of government.

§ But there is a much deeper problem with (*). Consider the *moral* obligation that I have not to kill other human beings (except in self-defense). *By itself* that obligation doesn't constitute any sort of legal obligation. Hitler's executioners had a moral, but not a legal, obligation not to kill the inmates of the concentration camps.

One typically has many moral obligations that are not matched by legal obligations. So even when the moral obligation is matched by a legal obligation, the existence of the former does not *by itself* suffice for the existence of the latter. Given only that I have a moral obligation not to kill Smith, it doesn't follow that I have any legal obligation not to do so. Smith and I might both live on a desert island, outside the jurisdiction of any legal system. In such a situation, I would have a moral, but not a legal, obligation not to kill Smith. So in the circumstances where I *do* have a legal obligation not to kill Smith, that obligation is not *identical with* the corresponding moral

obligation. What we have here are two distinct obligations that are discharged in the same way.

A bit of fiction may clarify the distinction I am trying to make. Smith fixes my car and charges me the reasonable fee of $100. I thus have a moral (and perhaps a legal) obligation to pay him that money. At the same time, my giving Smith $100 is necessary to prevent a nuclear war. Smith is an unstable person who, for some reason, has access to nuclear technology. Unless I soothe Smith's fevered mind with a crisp new $100 bill, he will deploy said technology

Here I have two different obligations. I have an obligation to pay Smith for his services. I also have an obligation to prevent a nuclear holocaust. Even though there is a single form of behavior that would discharge both of these obligations (namely, my giving Smith $100), these obligations are obviously quite distinct. Similarly, even though there is one mode of conduct that would discharge both my legal and moral obligation not to kill Jones, those obligations are distinct.

§ We may extend this line of thought. Given any law, no matter how trivial or wicked, we can imagine circumstances where the act of breaking it encourages a general disrespect for law and government, thus eroding or even undermining the government's ability to protect the public. Given this, it might seem reasonable to say that one's obligation to obey trivial or wicked laws coincides with one's moral obligation not to undermine the government's authority and, therewith, its ability to protect those under its dominion.

But such a thesis would be quite false.[179] Given practically any act that *would* weaken governmental authority, we can imagine circumstances where there is no *legal* obligation to refrain from that act; and given any legal

obligation, we can imagine circumstances where violating that obligation has no tendency to weaken government.[180] So a mode of conduct is never illegal *in virtue of* the fact that it weakens governmental authority, and is never legally required *in virtue of* the fact that governmental authority depends on it.

Let us sum up. In some cases, legal prohibitions seem to coincide with moral prohibitions. I have a moral *and* a legal obligation not to kill Fred. But I do not have that legal obligation *in virtue of* the fact that I have the corresponding moral obligation. In some cases, acting a certain way is necessary to prevent an erosion of governmental authority. But no form of conduct is ever illegal *in virtue of* the fact that it conduces to such an erosion. So my legal obligation not to kill Fred is not identical with my moral obligation not to do so.

§ But given only that my legal obligation not to kill Fred isn't identical with my moral obligation not to do so, it doesn't follow that my legal obligation not to kill Fred isn't identical with *some* moral obligation or other. (I will argue that it *is* identical with *a* moral obligation.)

Let me give an illustration of the general principle at work here. I have extra food that I'm not going to eat. In my neighborhood, there is a stray dog, named Rusty, that desperately needs food. Under this circumstance, I have a moral obligation to give Rusty some of my food. Let M_1 be this obligation. A few years ago, in exchange for a large sum of money, I promised dog-lover Bob that, for the next five years, I would never let any stray dogs in my neighborhood go hungry. Let M_2 be *this* obligation.

M_1 and M_2 are both moral obligations and, under the circumstances, they are discharged in the same way. But they are distinct obligations. M_2

will cease to exist after the five years elapse, but M_1 will persist. M_1 is an obligation to feed Rusty. M_2 is an obligation the discharging of which *involves* my feeding Rusty, but it is not itself an obligation to feed Rusty. Rather, it is an obligation to comply with a certain agreement between myself and Fred. At the moment, compliance with that agreement involves my feeding Rusty. But we could easily imagine circumstances where I could comply with that agreement and *not* feed Rusty. This would be the case if Rusty moves to another neighborhood, but I am still feeding the strays that are in my neighborhood. We could also imagine circumstances where I could fee Rusty and yet fail to comply with that agreement. This would be the case if I continue to feed Rusty, even though he isn't living in my neighborhood, but I don't feed the stray animals that *are* in my neighborhood. Thus M_1 and M_2 are distinct obligations, even though, under the circumstances, they are discharged in the same way.

For similar reasons, given only that my moral obligation not to kill Fred is not identical with my legal obligation not to kill him, it doesn't follow that the former isn't identical with *any* moral obligation.

The moral basis of legal obligation

I would now like to propose a different theory as to the moral basis of legal obligation. I warn that the next section will consist largely of dreary, but ultimately helpful, technicalities.

Some technical preliminaries

Let us begin with a background point. Some laws (e.g. *don't write*

checks unless you have sufficient funds in your checking account) prohibit certain kinds of behavior. Others (e.g. *married people are entitled to such and such tax-exemptions, handicapped people are entitled to such and such privileges, pedestrians have the right of way*) do not prohibit anything, and merely specify the legal consequences of certain conditions (being married or handicapped) or of modes of conduct (walking in certain places).

We might say that every law is "associated with" a certain kind of conduct or condition. In this context, the term "associated with" is being used in a purely technical sense to be defined as follows. C_i is the manner of conduct "associated with" law L_i exactly if, in virtue of behaving in manner C_i, one is not subject to any penalties provided for by L_i and one *is* eligible for any privileges thereby provided for. For example, let L_{78} be the law prohibiting the writing of bad checks, i.e. of checks that are not backed by sufficient funds. Obviously L_{78} makes provisions for the punishment of those who write bad checks, and does not make such provisions for those who do not do so. Thus, if you write a bad check, your conduct is *not* of the kind that is associated with (in our technical sense) the law prohibiting the writing of such checks; otherwise, your conduct is of that kind. Let L_{92} be the law giving special parking privileges to handicapped people. If you are handicapped, your condition is of the sort that is associated with the law giving special privileges to the handicapped; otherwise, your condition is not of that kind. (Henceforth, to avoid verbosity, the expressions "conduct", and "mode of conduct" will be elliptical for "conduct *or* condition" and "mode of conduct *or* condition.")

Often the mode of conduct associated with a law is to be understood in entirely privative terms. For example, there is no positive characterization of the mode of conduct associated with the law prohibiting the writing of bad

checks. In some dialectical contexts, it might be inappropriate to characterize *not* writing a bad check as a "mode of conduct." But this is not one of those contexts. From the viewpoint of a legal system, what is often important is precisely whether someone has *not* written a bad check, *not* driven drunk, *not* killed anyone. Where the law is concerned, the categories in terms of which conduct is to be understood are often purely privative. So our proposed use of the term "mode of conduct" appears to be appropriate in this context.

Before dealing with matters of substance, we must make a couple more terminological points. Let L_1, L_2…L_n be a complete list of the laws constitutive of the American legal system, and let C_1, C_2…C_n be the kinds of conduct associated with those laws. Let us say that your conduct *conforms to* C_i iff it is of the mode of conduct associated with L_i. So if L_{24} is *don't kill (except in self-defense)*, then your behavior conforms to C_{24} exactly if you don't kill (except in self-defense). If L_{56} is *married people are given such and such privileges*, then your condition conforms to C_{56} iff you are married.

As a human being who lives on a certain part of the Earth's surface, you have some kind of assurance or guarantee that, so far as your conduct conforms to C_1, C_2…C_n, you will have certain freedoms and privileges. You also have an assurance that, so far as your conduct fails to conform to C_1, C_2…C_n, you will *not* have those freedoms or privileges. So if you write bad checks, you will *ceteris paribus* forfeit some privilege that you would otherwise have. If you *don't* write bad checks, then *ceteris paribus* you will not lose that privilege. If you are *not* married, then you will *ceteris paribus* not have such and such tax-exemptions. If you are married, then *ceteris paribus* you will have such exemptions.

I should explain the presence of the words "ceteris paribus" in the last

paragraph. It would be an overstatement to say that, if you don't write bad checks, you will *definitely* avoid loss of some important privilege. If you don't write any bad checks, but are convicted of murder, you obviously will lose such a privilege. Our legal system assures you that, holding everything else constant (i.e. *ceteris paribus*), you will not lose any privilege so long as you don't write bad checks. Similarly, it would be an overstatement to say that, if you are married, then you definitely will enjoy some tax-exemption that you wouldn't otherwise have. Perhaps in addition to being married, you fall into some category that involves your paying taxes additional to those you would ordinarily pay. Our legal system assures you that, holding everything else constant, you will be given a tax-exemption if you are married. It does not assure you that you will *absolutely* be given a tax-exemption if married.

Given these points, it is not unreasonable to regard our legal system as corresponding to a certain kind of conditional assurance, one whose content – or, at least, the general spirit of whose content -- is given by the proposition:

(*) There is an assurance that, for anyone x satisfying certain conditions relating to place of residence, biology, age …, x will receive such and such protections iff x's conduct conforms to C_1, C_2…C_n. [181]

Let me make one last technical point before we go into an uninterrupted discussion of substantive matters. (*) assures that certain protections will be given to anything x *that satisfies certain conditions relating to place of residence, biology*, and so forth. (*) doesn't give an assurance to people who live in New Zealand or Pakistan; they are in the wrong place. (*) doesn't give an assurance to wildlife in North-America: they

have the wrong biology (they are not people). So what we need in (*) is not the unrestricted quantifier "for any x", but rather the restricted "for any x satisfying such and such requirements relating to place of residence, biology…"

But, henceforth, to avoid tedious repetitions, I will use the *un*restricted quantifier, "for any x", as an abbreviation for the restricted quantifier: "for any x such that x is a human who lives on such and such part of the earth's surface…"

The moral structure of legal obligation

Now, at long last, we can stop dealing with technical minutiae, and may concern ourselves with matters of substance.

Notice that (*) is a *narrow*-scope assurance. That this is so reflects an essential fact about the concept of a legal system. A legal system cannot be given by a proposition of the form:

(**) For any individual x, x has an assurance that x will receive such and such protections iff x's conduct conforms to C_1, C_2…C_n.

(**) makes a completely different statement from (*). Let x_1…x_m be the citizens of the United States. (**) holds in virtue of the fact that x_1 specifically has been given a conditional assurance of a certain kind; that x_2 has been given an assurance similar to, but numerically distinct from, that given to x_1; that x_3 has been given an assurance similar to, but numerically distinct from, each of the two assurances just mentioned; and so on, for each

of the hundreds of millions of people who are protected by the American legal system. So supposing that Smith is a U.S. citizen, (**) is true in virtue of the fact that Smith has been given some assurance that concerns Smith specifically, that Jones (another U.S. citizen) has been given some assurance that concerns *Jones* specifically – and so on, *mutatis mutandis*, for each of the remaining 300 million people in this country.

While it is *true* that Smith has been given such an assurance, no assurance concerning Smith specifically is any part of the *content* of any U.S. law or of the U.S. system of law. Under the American legal system, Smith does indeed have an assurance that he won't be killed. But the law is not: *murdering Smith is punishable by life imprisonment*. Rather, the law is: *for any citizen x, murdering x is punishable by life imprisonment*. So the content of the law is given *not* by the proposition:

(a) Smith has an assurance that he won't be killed,

but rather by the proposition:

(b) there is an assurance that, for anyone x, x won't be killed.

The connection between Smith and that law is not one of content, but of empirical fact. *Given* the empirical truth that Smith satisfies such and such conditions – that he is a human being (not a dog or a tree) who lives in area R (as opposed to the area associated with, say, New Zealand or Pakistan)…-- it follows that, thanks to the existence of assurance (b), proposition (a) is true. Given any individual x who is *in fact* protected by that law, an exactly similar argument shows that the protection is not one that is guaranteed strictly by the *content* of the law, but rather by that content *conjoined with* the empirical

proposition that x is an individual satisfying certain requirements (x is a human living in R…)

What we just said about the prohibition against murder is true of our legal system as a whole. There is no individual x such that the content of our legal system mentions, or is otherwise internally related to, x specifically. Pick an arbitrary individual – say, Donald Rumsfeld. There are Rumsfeld-free possible worlds where our legal system exists in exactly the form in which it exists in our world. So even though our legal system protects Rumsfeld, that protection is not guaranteed strictly by the *content* of our system, but rather by that content *conjoined with* the empirical proposition that Rumsfeld is an individual satisfying certain requirements.

Thus, even though (**) is *true*, it is not the proposition that gives the *content* of our (or, by an exactly similar argument, any other) legal system. Even though our legal system provides assurances to each of millions of individuals, *none* of these individuals is mentioned in any proposition that gives the *content* of our legal system.

§ In order to proceed, we must once again make a detour through some subject-matter that, at first, seems totally unrelated to the topic of discussion. This detour involves our revisiting some points about causality made by Jackson and Pettit.

These authors make it clear that program causes are typically more important than efficient causes where event-*explanation* is concerned. But the same is true of event-*manipulation*.[182] The efficient cause of the pan's heating is that micro-events $e_1 \ldots e_n$ occur. We can create such events only by creating situations that program for them. Given our limitations of knowledge and technology, we cannot *directly* create events of that kind. Such a feat might even be nomically impossible, given the interference-effects that would

have to be dealt with. Consequently, if we want to heat the pan, we must create the right *program* causes. We must, for example, create a flame.

Even if we could directly (i.e. not by way of a program-cause) create *exactly* the right efficient-causes, doing so would almost never be the right course of action. Suppose that our objective is to heat the pan and that we choose to do so, not by creating the right program-cause, but by creating *exactly* the right efficient-cause. First of all, the resources – cognitive, emotional, technological, economic…-- needed to pull off of such a feat would be astronomical compared with those required to create the right program-cause.

But even if we managed to circumvent those problems, others would take their place. Efficient-causes are *too* efficient. They don't have the elasticity of program-causes or, therefore, their subsequent amenability to adjustments in our objectives. Suppose that, while in the process of making dinner, we decide to add a few mushrooms to the omelet we are cooking in the aforementioned pan. This will appreciably affect the amount of heat needed to cook the food-stuffs in the pan. Although $e_1 \ldots e_n$ were appropriate for our needs *prior* to our adding the mushrooms, they will probably not be appropriate to our needs after that change. So long as we direct our efforts towards efficient-causes, our new culinary objectives will involve a complete over-haul of our previous methodology. But this is not so if we think in terms of program-causes. It is easy to increase or decrease the intensity of *the flame*.

In general, the re-parameterizations demanded by changes in our objectives can often be accomplished without jettisoning our chosen *program*-causes. But efficient causes, almost by definition, have zero parametric elasticity. A consequence is that the occurrence of anything unforeseen, no matter how slight, demands a complete abandonment of our

choice of method.

§ Points similar to those just made apply to a government's efforts to help those under its authority. For reasons of economics and logistics, it would be totally out of the question for the government to deal with every single civil problem on a case-by-case basis. There must be standing policies and fixed procedures; there must be a relatively fixed legal and administrative apparatus. By itself, this apparatus does nothing. But its existence *programs for* the sorts of events that efficiently bring about the resolution of problems that arise. Given a specific problem, that apparatus guarantees that the relevant personnel (police officers, rescue-workers…) will do what must be done to resolve the situation.

In large part, governments serve and protect by creating the right program-causes – by creating general legal and administrative conditions that guarantee the occurrences of the needed specific events. There is no other way. Unless a government were responsible for a trivial number of people, it would not be practically possible for it to deal with problems on a strictly individualized basis.

Of course, *some* problems, e.g. Pearl Harbor or 9-11, receive individualized attention. But the government couldn't possibly give such attention to every problem that arises, or even more than a tiny percentage of them.[183]

Laws are not events.[184] They are enduring features of the political system that help[185] program for responses of a certain kind to events of a certain kind. In light of these points, suppose that the American legal system were given by an assurance of the following form:

**(***) There is an assurance that, for anyone x, x will receive such and

such protections iff x's conduct conforms to C_1, C_2…C_n *except* in cases where conformity to C_i (for an arbitrary i) is patently unnecessary from the standpoint of the public weal. Where such exceptional cases are concerned, conformity to the relevant standard is at the individual's discretion, and non-conformity does not jeopardize the individual's assurance that he will receive the aforementioned protections.

(***) says that there is *no* penalty for driving 66 mph, or even 76 mph, in places where the posted limit is 65 mph, so long as doing so is obviously without adverse consequences. In such a case, (***) guarantees that the driver will not be ticketed or otherwise subjected to governmental censure. In general, given any situation where conformity to C_1, C_2…C_n is *patently* unnecessary, (***) says that the individual *doesn't* have to conform to those standards.

If our system of laws were given by a proposition like (***), that fact would *program for* disastrous consequences. For the sake of argument, suppose that speeding tickets were given *only* in cases where doing so was needed to head off some accident. In that case, practically everybody would add 10 mph to their average driving-speed. (I know I would.) But then our highways really would become death-traps.

Suppose that, in cases where it was obviously harmless to do so, there would be *no penalty* for taking somebody else's car for a joy-ride. That fact would program for innumerable non-trivial violations of people's rights – cases where x takes y's car for a joy-ride, and then destroys it, or does minor damage to it (e.g. bends the fender) and then destroys it rather than return it to the owner and risk being found liable for the damage. It is obvious how analogues of these examples can be constructed in connection with any law that can *almost* always be harmlessly flouted.

Of course, (***) does not give any kind of immunity to the individual who mistakenly *thinks* that non-conformity to C_1, C_2…C_n is harmless. If Smith drives ten mph over the posted-limit and ends up killing someone as a result, (***) holds him responsible, since (***) says that the non-conformity must *actually* be harmless. In Smith's case it was not. So (***) is by no means innocuous: it does make provisions for penalties in this sort of situation, and thus has a non-trivial amount of deterrence-value.

Nonetheless, the civil situation programmed for by (***) would be quite unsustainable. Suppose the law allowed you to commit an act with legal impunity, as long as it was *obviously* harmless to do so. Consider what a difference that legal change would make to your behavior. Now consider what a difference it would make to the behavior of each of the remaining three hundred million people in this country. The situation would be catastrophic.

For this reason, governmental assurances of protections *must* be given through a proposition like (*). If that assurance is given through a proposition like (***) – a proposition with an escape-clause that permits any sort of behavior, so long as it is harmless – the result will be a situation where protections have *not* been adequately provided. The government will have failed to do what, ethically speaking, it must do.

Here it is important to distinguish between two possible situations. Let S be a situation where the legal system is given by (*), and where, for the normal reasons, people's conduct deviates from C_1, C_2…C_n. Let S* be a situation where (***) is the law and where, for the normal reasons *plus* the extra latitude granted by the escape-clause, people's conduct deviates from C_1, C_2…C_n. The second deviation will dwarf the first. In S*, it will not necessarily be *illegal* to go 85 mph in a place where the posted limit is 65 mph. It will be illegal only if it is not obvious that doing so would be without

negative consequences. In S, it will necessarily be illegal. In S, but not S*, you cannot (legally) be given a ticket for going 85 mph if, in fact, it is obviously safe to do so.

If we were dealing with omniscient actors – people who were flawless in their ability to assess the risks involved in exceeding the speed limit – there would be no difference between S and S*. But so long as we are not dealing with omniscient actors, there will be dramatic differences. S is a situation where people don't drive 85 mph even if they think that (leaving aside problems relating to law enforcement) they are completely convinced that it would be harmless to do so. S* describes a situation where people *do* drive 85 under those circumstances. In S*, the highways are deathtraps. In S, they are not. What we just said about traffic-laws applies (mutatis mutandis) to *every* kind of law. S is, whereas S* is not, a viable civil situation. The law must be given by (*), not by (***) – not by anything with an escape-clause in it.

§ A legal system would be useless if its laws were too specific. If laws were relativized to hyper-specific situations, then (unless the law-makers were superhumanly capable of predicting the future) most laws would be quite useless. Laws must typically be general.

A corollary is that there will be *inefficiencies* in any legal system that is not constructed by creatures of super-human predictive ability. There will always be laws that can, or even must, be broken on some, or even all, occasions. It is not possible to reduce the "co-efficient of inefficiency" to anything remotely close to zero. But the consequences would be disastrous if laws ceased to be binding the moment they became inefficient. As we discussed, if the law prohibiting driving 90 mph became inapplicable to contexts where driving that speed were harmless, our highways would

become deathtraps. This is because people sometimes would misjudge when it was harmless to drive 90 mph. (People even misjudge when it is "patently obviously" harmless to do so.)

(***) would be preferable to (*) if, and only if, people were infallible. But they are not; so it isn't.[186] There are indeed inefficiencies in the law. But there are also inefficiencies in people's judgments about when it is harmless to break the law. The inefficiencies in the former protect us from the catastrophes that would otherwise be brought about by the inefficiencies in the latter.

(***) says that the law doesn't apply when it is inefficient: if there is no reason not do so, it is legal to drive 90 mph, to sell heroin, and to use hand-grenades to kill rodents. We've already discussed why (***) is untenable and why, in general, a legal system cannot be given by a proposition containing the sort of escape-clause found in (***). [187]

§ We are now, I believe, in a position to see the moral basis of legal obligation. Every government has a moral obligation to serve the welfare of its constituents. There is no way to discharge this obligation unless these constituents abide by certain standards of conduct. So a given government has a moral obligation, and therefore a right, to demand that its constituents conform to certain standards of conduct. This demand must be given by a proposition that does not include an escape-clause of the kind earlier discussed. That government has a moral obligation, and therefore a right, to demand of each individual[188] that her conduct conform to the standards described in that proposition. It follows that each individual has an obligation to conform to those standards. If I have the right to demand that Jones refund my money, then Jones has an obligation to do so.

§ It may help if we make it clear exactly what is meant by the term "demand" in this context. Given only that I have the right to *ask* Jones to let me sleep in his guest bedroom, it doesn't follow that he has an obligation to let me do so. But the government doesn't just have the right to *ask* citizens to conform to certain standards. The government must *exact* such conformity. A government would be useless if it merely *asked* its constituents to comply with it – it if merely *asked* people not to kill and steal, but didn't *require* them to refrain from such acts. So when we say that the government has the right to *demand* conformity to certain standards of conduct, we mean that it has the right to enforce compliance with those standards.

For the reasons we've discussed, a government cannot demand that a given person obey *only* those laws that are not trivial or wrong. Every legal system, no matter how well designed, will contain inefficiencies. If a government asks each individual to obey the legal system *minus* the inefficiencies, that will program for an untenable civil situation. The government owes it to its constituents to demand of each of them that she obey *all* the rules, inefficiencies and all. Consequently, each of them has an obligation to obey all of those rules, including trifles like GS.

Of course, the government could waive the requirement to obey this or that *particular* inefficient law. If the government declared that GS could be disobeyed with impunity, there would be no negative consequences. But there would remain innumerable other inefficiencies. There would be no way to *list* all of these, and no way to lift the corresponding obligations on a one by one basis. For reasons of logistics and epistemology, it would not be possible to *enumerate* the actual cases where it would be harmless (or beneficial) to exceed the speed limit. The only way to lift the obligations associated with these inefficiencies would be through some general escape-clause, such as that found in (***). But we've already seen what the

consequences of such a clause would be.

Two kinds of legal injustice

There remains a delicate problem to be dealt with. Before we can deal with this problem head-on, we must make one more point about the concept of legal injustice.

Given practically[189] any law, no matter how wholesome, there will inevitably be *some* circumstances where breaking it is necessary. Let (#) be the law that prohibits killing another human being unless there is a "clear and present" threat to one's own life. This is a reasonable and morally wholesome law; it is not comparable to the suffrage laws of the 19th century or the racist laws of the Jim Crow South.

At the same time, there are circumstances where (#) ought to be broken. A battered wife may not be able to kill her husband when he *does* pose a clear and present threat to her safety: in such a situation, he will be on his guard and thus maximally *hard* to kill. But (let us suppose) unless she does kill him, she can be sure that the quality of her life, as well as those of her children, will be intolerably reduced. It seems to me that, under such a circumstance, the woman has every moral right to kill her husband. At the same time, it would probably be unwise to broaden the law to allow for such killings. At the very least, such a broadening would have to be made with extreme caution. Suppose the law became:

(##) It is illegal to kill another human being unless there is a clear and present threat to one's safety *or* unless there is a clear long-term threat to one's quality of life.

(##) is too permissive, and would program for civil disaster. It would allow people to kill their pointy-headed supervisors, and their unruly and overly demanding children. This is not to say that (#) is completely incapable of being appropriately broadened or refined. But it would be hard to broaden it *enough* to permit legitimate killings (the battered wife who wants her children and herself to be free from terror) without also permitting unjust killings (the youth who is growing increasingly tired of his father's stodgy unwillingness to bankroll a summer in Tibet so that he can "find himself"). Even if such a balance were struck, it would not be *completely* faultless; situations like those just described, in connection with the battered wife, would still occur (though less frequently). Every law, no matter how good, has *some* co-efficient of inefficiency.

We must distinguish the kind of legal injustice just discussed from a very different kind. Consider the U.S. suffrage laws of the 19th century. It wasn't only in extreme or hard to foresee situations that those laws were unjust. On the contrary, it would only be in extreme situations that they *would* be just. Suppose that, because of some strange disorder, Mary becomes violently anti-social whenever she gets within a certain distance of a polling station. In that case, a sexist suffrage law might help stave off a disaster. But such a law would *typically* be unjust. Here we are dealing with legal injustice that is in a different category from the sort arising in the just discussed case of the battered wife. Not all injustices associated with a legal system have to do with a margin of inefficiency. Some such injustices are (we might say) *inherent* in the law.

How can we have a moral obligation to obey unjust laws?

I now wish to discuss a problem with my analysis of law that arises in

connection with laws that are *inherently* unjust. As before, it will be helpful to begin with a bit of fiction.

G is a government and L is its legal system. In some respects, L is effective. It keeps the peace, and ensures the presence of much needed infra-structure. At the same time, L is grossly unfair to many people under its dominion. The extent of the unfairness doesn't fall within the margin of inefficiency that *inevitably* characterizes a legal system. The unfairness has to do with laws that, from a moral perspective, are *patently* sub-standard. So the unfairness is of the kind associated not with (#), but with the U.S. suffrage laws of the 19th century.

Jones is a female citizen in the nation governed by G. Some of L's laws are unfair to Jones. It is not a datum – though it may be true - that she has any *moral* obligation to obey all of these laws. But it does seem to be a fact that she has a *legal* obligation to obey those laws, even those that are not fair to her. Suppose that L includes a sexist law prohibiting Jones from checking books out of public libraries. One day Jones dresses up as a man and manages to check out a book. It is a datum that Jones has broken the law. If Jones were not subject to the sexist law just mentioned, then Jones could not be *in violation of it*: you cannot be in violation of a law unless it defines a legal obligation for you.

I have argued that a government has an obligation, even a right, to demand that its constituents conform to certain modes of conduct. I have attempted to find in this fact a categorical moral basis for legal obligation. But it would be hard to maintain that a government has the right to demand that its constituents conform to an *unfair* set of standards – for example, to sexist or racist standards. Therefore, it would be hard to maintain that, for example, a woman has an obligation to comply with a demand that she not vote, not check out books from the library, and so on. Given that a woman is

under no obligation to conform to such standards, it seems that my analysis can find no moral basis for her obvious *legal* obligation to conform to them. This would seem to scuttle my attempt to identify legal obligation with a certain kind of moral obligation.

§ It is fitting that we end by discussing this particular objection to our analysis, since our response to it will put into play the conceits most central to this work: the non-binary nature of morality, the close relationship between someone's being protected by a legal system and their having obligations under it, and the related fact that there is no government where there isn't a modicum of moral protection.

As we've said many times before, rights and obligations have a "dimension of weight", to use Dworkin's expression. I have *more* of an obligation to save the life of my children than I do to refrain from stealing medication from the pharmacy. But it doesn't follow that I have *no* obligation to refrain from such a theft. That obligation has been outweighed, not cancelled. A corollary is that the pharmacy-owner has *a* right to stop me, even though his right to spare himself a $1.00 loss in profits is outweighed by my right to save my children.

In imposing L on the citizenry, G is not doing something wholly bad. We must remember that x doesn't have any obligations under a given legal system unless x receives vital protections from it. Failure to extend considerable protections to x amounts to a failure to recognize x; and this amounts to x's having no obligations under that system. So even though L is unfair to Jones, it would be wrong to say that L does *no* good, or even anything less than a significant amount of good, for Jones.

Further, G has an obligation to protect not just to Jones, but all of her fellow citizens. While it is unfair in important respects to Jones, L is fair in

important respects to many others -- and also, as we just saw, to Jones herself. (There is no contradiction in saying that L is fair to Jones in some respects, and unfair in others.[190]) So in imposing L on the citizenry, G is doing something that is *on balance* good.[191] It would thus be false to say that G has *no* right to impose L on the citizenry – even though its right to do so is less strong, and thus more easily outweighed, than its right to impose other, more enlightened legal systems.

For reasons, we've discussed, Jones' compliance with L is necessary for, or at least conducive to, L's being successfully imposed on the citizenry.[192] (In any case, her non-compliance would certainly be a hindrance.) So Jones' compliance is necessary for, or at least conducive to, the occurrence of a non-trivial good.

It would thus be false to say that Jones has *no* obligation to comply – even though her obligation to comply is outweighed by her right to flourish and is also weaker than the obligation she would have to comply with a more fair legal system.

There is another dimension to Jones' obligation to L. Jones is recognized by L, and therefore has legal obligations under that system, just in case she receives vital protections from it. Her compliance with L conduces to the effective functioning of that system, and her non-compliance conduces to its malfunctioning. So L is not just *a* good thing; it is a good thing on which she personally depends. Her obligation to L is thus of a dual nature. (If Smith does good for me, then presumably my moral obligations to him are *ceteris paribus* greater than they would be if he just did good for someone or other.) This is consistent with our hypothesis that her legal obligations under L are moral obligations. At the same time, it is consistent with the fact that Jones is acting morally if she breaks the laws of L that are unfair to her.

A refutation of some seeming truisms concerning law and morality

Some of the points we've made warrant a reconsideration of some views that would otherwise appear quite reasonable.

We are inclined to say that, within a given legal system, some legal obligations are *more* moral than others. Some legal obligations are morally wholesome (e.g. my legal obligation not to kill people, except in self-defense); others are morally neutral (e.g. my legal obligation not to drive 66 mph if the posted limit is 65 mph); and others are downright immoral (e.g. my legal obligation to report an illegal alien who is heroically feeding his family of twelve).

This point of view is false, and it involves the fallacy of supposing that obligations are identical if they are discharged in the same way. My *legal* obligation not to kill people is neither more, nor less, morally wholesome than my *legal* obligation not to drive 66 mph; and each is neither more nor less morally wholesome than my *legal* obligation to turn in the heroic illegal alien. Each of these legal obligations is *a* moral obligation. But even though it is *a* moral obligation, my *legal* obligation not to kill others is entirely distinct from the corresponding moral obligation – even though there is some one (privative) mode of conduct that discharges both obligations. The same point *mutatis mutandis* applies to the other two legal obligations. Permit me to substantiate this perhaps alarming point of view.

Let us recall a story that we told earlier. I have an agreement with Smith that I will feed stray dogs in my neighborhood. Each of us entered into this agreement freely, and it is morally wholesome in all other respects. I thus have a moral obligation to honor it. Because Rusty is a stray dog in my neighborhood, honoring that agreement involves feeding Rusty. But even if I hadn't entered into any such agreement, I would still have a moral obligation

to feed Rusty. I have extra food; Rusty needs it; so it would be immoral of me not to feed Rusty.

Here one might be tempted to say that I have only *one* obligation -- albeit one that can be *described* in different ways (either as an obligation to feed Rusty or as an obligation to honor an agreement with Smith). But this is exactly what we must not say. There are counterfactuals scenarios where Rusty is well fed but where Poppy (some other dog) is a stray in my neighborhood that needs food. Under those circumstances, honoring my agreement with Smith would involve my feeding Poppy, not Rusty. So my moral obligation to feed Rusty doesn't have the same counterfactual properties as my moral obligation to honor my agreement with Smith. Therefore, by Leibniz's law, they are distinct obligations, even though, under the present circumstances, some one mode of conduct discharges both of them. In general, entirely distinct moral obligations can sometimes make the same *behavioral* demands of us.

Let us take the next step in our argument. There are no *degrees* of legal obligation. One either has a legal obligation or one doesn't. Legally, I am 100% obliged not to kill and 100% obliged not to litter. If we said that these legal obligations *coincided* with the corresponding moral obligations, then there would be degrees of legal obligation. But there aren't; so they don't.

I have a purely moral obligation not to kill people. Let M_1 be that obligation. I also have a *legal* obligation not to kill people. Let M_2 be *that* obligation. M_2 is *a* moral obligation – as we've seen, all legal obligations are moral obligations. But M_2 is entirely distinct from M_1. We've already seen why. There are contexts where I have a *moral*, but not a legal, obligation not to kill. So even when I have *both* a moral and a legal obligation not to kill, the one obligation is not *identical* with the other.

At the same time, my *legal* obligation not to kill is *a* moral obligation. It isn't identical with my moral obligation not to kill, but rather with my moral obligation to comply with a system of laws that does much good (for myself and others).

I have a legal obligation not to drive 66 mph if the posted limit is 65 mph. Let M_3 be that obligation. For reasons that we've discussed, M_3 is also *a* moral obligation. M_3 is not identical with my moral obligation to drive at safe speeds, but rather with my moral obligation to comply with a certain system of laws.

M_1 is much stronger than M_3. But M_2 and M_3 are of equal strength. They are in fact the same obligation: the obligation to comply with a certain system of law.

Here is what complicates the situation. M_1 is an extremely strong moral obligation. The behavior that discharges M_2 *also* discharges M_1. The behavior that discharges M_3 does *not* discharge anything having M_1's weight. This makes it seem as though my strictly *legal* obligation not to kill people has much more moral weight than my strictly legal obligation not to speed. But this is not the case. The one legal obligation is reinforced by a distinct, but similar, moral obligation of great weight. The other legal obligation is not comparably reinforced. But the two strictly *legal* obligations have equal moral weight.

My strictly legal obligation to turn in the valiant illegal alien is *a* moral obligation. But it is a moral obligation, *not* to turn in a heroic soul, but rather to comply with a certain system of law. I have a strong *moral* obligation *not* to turn in this fine individual. For this reason, it *seems* as though my legal obligation to turn in the heroic illegal alien has no moral content. But that is not the case. We are dealing with two distinct moral obligations such that

there is some mode of behavior that discharges the one moral obligation and violates the other. For this reason, it falsely *looks* as though my legal obligation has a purely negative moral value.

§ There is one final loose end to be tied up. As we just noted, legal obligation doesn't come in degrees. You are either legally required to do x or you aren't. But moral obligation does come in degrees. This appears to conflict with our thesis that legal obligation is a kind of moral obligation.

This conflict is illusory. Given an obligation to do x, that obligation may be had to varying degrees. But whether you have *any* degree of moral obligation to do x is a binary affair. Some things are more intelligent than others, and some things (e.g. rocks) have no intelligence at all. Whether a thing has *any* intelligence is a binary affair, even though intelligence can be had to varying degrees. You have legal obligations under system L exactly if your moral obligations to L rise above a certain threshold. The degree to which you have those moral obligations is indeed a matter of degree. What is not a matter of degree is *whether* they rise above that threshold.

You have a legal obligation under L exactly if you have a moral obligation that rises above a certain threshold. This is consistent with the idea that legal obligation is a kind of moral obligation; and it is also consistent with the fact that legal obligation is binary, whereas (generic) moral obligation is not.

Why "ought" does not imply "can"

"Ought implies can." This is an old saying among ethicists.[193] But recent work in moral philosophy has produced a compelling argument against

this old saw.[194] I think that the slogan “ought implies can” embodies some of the distortions about morality that have derailed attempts to see the moral structure of legal obligation. The best way to undo those distortions is to attend to the aforementioned argument.

That argument can be expressed through a story. I borrow a million dollars from Smith. For whatever reason, I spend all the money and, despite my best efforts, I am utterly incapable of paying Smith back. Surely I still owe Smith the money. The fact that I don’t have the financial wherewithal to do so does not obliterate my debt to Smith. Given the physical structure of my brain – given the low ratio of neurons to glial cells in the part of my brain mediating financial intelligence – it may be a nomic impossibility that I should ever find a way to repay Smith. But surely I still *owe* him the money. Given that I owe him the money, it presumably follows that I *ought* to pay him the money. It would be absurd to say that I *owe* somebody money but have no obligation to give him that money. So even though I cannot repay Smith, I still owe him money and thus ought to give him money. I ought to do it but I cannot. Thus “ought” does *not* imply “can.”[195]

Does this argument show that “ought implies can” is wrong *tout court*? I think the term “ought” is ambiguous. Sometimes that word denotes the available course of action that is, on balance, the most moral (or least immoral) one. Trivially, this kind of “ought” would be inapplicable where there was no possibility of action – where there was no “can.”

But, I believe, this is not the only sense of the word “ought”, or even the prevalent one. Even though *on balance* the psychotherapist ought to tell the authorities about his patient’s murderous plan, there is still a sense in which he ought to keep his patient’s confidence.[196] The therapist would certainly feel *bad* about betraying his patient. I don’t think that this feeling is a sheer delusion, comparable to a hallucination. The therapist realizes that, in

betraying the trust of his patient, he has violated *some* kind of obligation. Given this, it is not a leap to say that the therapist has put himself on the wrong side of *some* kind of "ought." If we don't say this, it becomes hard to explain why the therapist – though not *regretting* his course of action – may yet feel *remorse* about it.

A possible response would be to say that the therapist *shouldn't* feel remorse about his decision: he did the right thing, and should therefore be entirely guilt-free. (In fact, this is probably what people trying to console the therapist would say.) In so far as the therapist feels badly about betraying his patient's confidence, it is because his emotions are not rightly aligned with the Moral Law.

But there is an easy way to cast doubt on the veracity of such a view. Suppose that the therapist felt absolutely no remorse at all about betraying his patient's confidence; suppose that he no more felt guilty about it than he did about throwing away a used shampoo bottle. Surely we would regard him as an unfeeling monster. We *want* the therapist to feel a certain remorse –not because we think he has done wrong and therefore deserves punishment, but because we want the therapist to care about the patient. Under the circumstances, a feeling of remorse is an inevitable expression of such concern, and its absence would express an utter lack of it. So it is not plausible to say that the therapist's feeling of remorse corresponds to some kind of moral hallucination on his part – to a failure on the part of his emotions to track the moral realities. On the contrary, a *lack* of remorse on his part would be appropriately described in those terms.[197]

Undercover detectives often feel guilty about betraying the friendships that they must develop, and then break, to make a case.[198] They don't wish that they *hadn't* betrayed those friendships: they don't wish that they had actually let the drug dealer go scot-free. But they still feel a certain remorse

about having to violate the trust that was placed in them. Surely such feelings are not sheer delusions, but rather spring from an authentic connectedness to deep moral truths[199], e.g. the principle that it is wrong to betray a friend. It is not inconsistent with actual verbal usage to say that, on one delineation of the term "ought", such feelings of remorse spring from a knowledge that one has done something that one ought not to do (even though not doing it would have been worse).

In this second sense of the word, "ought" does not always imply "can". This second kind of "ought" is a comparative, not a binary, notion. The therapist ought to keep the confidences of his patients. But he also ought to prevent the killing of innocent people. He cannot do both. So the therapist – even if he is as moral as it is logically possible to be – cannot do everything he ought to do. So, on this delineation, "ought" does not imply "can".

I think that, on this delineation of the term "ought", whenever one is subject to a law, one ought to obey it. At the same time, one sometimes ought to disobey a law that one ought also to obey. If my children are dying, and my only hope of saving them is to steal some penicillin from a pharmacy, then I *ought* to do so. The loss to Walgreen's, and to the infrastructure of the United States, is negligible compared to the potential loss to my children (and myself). But the obligation not to steal has been outweighed, not cancelled. Circumstances being what they were, something that *ought* not to be done *had* to be done. One "ought" was outweighed by another. In any case, this seems a truer description of the moral structure of the situation than the statement that my obligation not to steal had been completely cancelled.[200]

The thesis that "ought" implies "can" has the consequence that, since I cannot *both* acquire the penicillin *and* avoid stealing from Walgreen's, *one* of those obligations has been cancelled: either I don't have to help my children or I *don't* have any obligation not to steal from Walgreen's. More generally,

"ought implies can" is incompatible with the idea that there can be degrees of obligation and, what is related, that there can be conflicting obligations. It entails that if X isn't what you *most* have to do, then there is *no* moral value to doing X. So if obeying the traffic laws isn't the *best* thing to do, then it is without moral value at all.

"Ought implies can" thus embodies the sort of spurious dichotomization that obscures the moral structure of legal obligation. Since you can't make it to the hospital on time *and* drive under the speed limit, then there is absolutely *no* sense in which you ought to drive the speed limit: the law prohibiting you from driving 120 mph in a residential neighborhood ends up corresponding to *no* moral obligation and is thus drained of any moral content. Whatever is not the *best* thing to do ends up having *no* moral merit at all. The other side of the coin is that, if something has *any* moral merit, then it becomes absolutely *de rigueur*. If you have *any* moral obligation to drive the speed limit, then you *must* do so: it nearly enough becomes a *moral law* that you cannot speed. So "ought implies can" ends up entailing either that laws are moral absolutes (Augustine, Aquinas) or that laws have at best a highly circumstantial connection with morality (Austin, Hart).[201] Both of these views are false. Laws are not as morally empty as a gunman's threats or the army regulations concerning the kind of shoe-polish one can use. Nor are laws in the same category as moral verities.[202]

The triviality-objection

There is an important objection (due to Thomas Holden) to our analysis of law that we must consider (as before, I will put it in the voice of a hypothetical interlocutor):

Practically anything that anybody does, no matter how morally bankrupt, can be seen as having *some* basis in morality. Suppose that Smith shoots Jones for absolutely no reason. In shooting Jones, Smith was exercising his finger, thereby preventing atrophy. He was also employing synaptic pathways that he probably seldom uses, thus preserving neural plasticity. *By itself* it is good that Smith's musculature isn't atrophying. *By itself* it is good that Smith's cerebral cortex is being stimulated. What is bad is that these goods occur in the context of such a heinous act.

This poses a serious problem for your analysis. Since any act has some connection with morality, however peripheral, your claim that laws are protections of moral goods becomes quite trivial. Suppose we have some law that forces innocent people to spend their whole lives digging ditches. That law ensures that people exercise regularly, develop their back-muscles, spend time outdoors, and so on. Of course, *by themselves*, regular exercise, a strong back, and so on, are goods. But obviously any connection between that law and morality is insufficient to warrant a moralistic conception of law, such as the one you advocate.

We can go further in this direction. The more unfair a law is, the more it forces those subject to it to develop strength of character, patience, and other moral excellences.[203] But surely we don't want to say that the more unfair a law is, the deeper is its connection with morality. A positivist would feel vindicated, not threatened, by the discovery of laws that made such demands on our character.

So given only that laws have *a* basis in morality, it doesn't follow that your moralistic conception of law has been vindicated, or even that your position is significantly different from that of the positivist. Any positivist

will admit that the most heinous laws could, in the circuitous manner just described, have a rather deep connection with morality.

According to the position I am advocating, a law is a certain kind of governmental *assurance* of protection of a right. A law is therefore a promise, on the part of the government, to protect something of moral value. This promise may not be kept (not all laws are enforced). For this reason, a law may fail to ensure any good. In other cases, the right whose existence is assured protection is entirely bogus. If enforced, that sort of law doesn't do any good, and may do great harm. But, I have argued, a law is in all cases a promise, on the part of the government, to protect something of moral value.

This *does* make my position meaningfully different from that of the positivist, and it does render entirely non-accidental the connection between law and morality. But that connection is not captured by saying that all laws *do* good. As the objector points out, the good done by many laws is far too trivial to establish any significant connection between law and morality. Let us now try to make clear what the nature of that connection actually is.

§ The crimes of a gunman may do him a world of good. Thanks to them, he has money, freedom, and security. Those same crimes may even improve the lives of innocent people. Thanks to them, the gunman's children are well-fed, and the gunman's ailing mother now has proper medical care.

The gunman's threats are not assurances of protections of rights. The gunman is not saying (or otherwise conveying): "I am hereby assuring protection of such and such rights had by such and such people." But that is, I have argued, precisely what a government is saying when it issues a law. So my analysis regards the gunman's acts as being very different from the

issuing of a law. It is irrelevant both how much good the gunman's acts do and how much bad that law does.

§ Given that laws are promises or assurances of protections of rights, the connection between law and morality becomes clear. Once again, it will help if we make a few general points about language. Speech-acts – assertions, promises, questions, orders – have *conditions of satisfaction.*[204] These are the conditions that must be met if an utterance is to accomplish what, at the level of semantics, it is supposed to accomplish. If I say "Fred is at the store", the relevant satisfaction-condition is that Fred is at the store. If Fred is not at the store, then my utterance has not accomplished what, *at the level of semantics*[205], it is supposed to accomplish. (Truth-conditions are a special type of satisfaction-condition.) If I say "go to the store!", the relevant satisfaction-condition is that you go to the store. If you do not do so, then my utterance will not have accomplished what, at the level of semantics, it is supposed to accomplish. If I ask you "are you going to the store?", the satisfaction-condition is that you provide me with a true statement as to whether you are going to the store.

Here we must be careful. If Fred is not at the store, then an utterance of "Fred is at the store" has not accomplished what, *at the level of semantics*, it is supposed to accomplish. But oftentimes our own psychological objectives are at variance with the satisfaction-conditions of our utterances. *My* objective may be to mislead you – to make you believe something that isn't true. But that doesn't change the fact that, *qua* assertion, my utterance is *supposed* to be true. This point (mutatis mutandis) applies to every other type of speech-act. When I ask you "is Fred at the store?", *my* objective may not be to elicit a truthful answer from you. *My* objective may be to make you believe falsely that I do not know where Fred is. So if you *do* provide a

truthful answer to my question, and thus make it clear to yourself that I *do* know Fred's whereabouts, that may foil my plans. But that doesn't change the fact that, *qua* question, my utterance is supposed to elicit a truthful answer from you as regards Fred's whereabouts.

Promises have satisfaction-conditions. Suppose I promise to pick you up at the airport tomorrow. The satisfaction-condition of my promise is fulfilled exactly if the proposition *I pick you up at the airport tomorrow* is true. Of course, *my* objective may be to *not* pick you up and leave you in the lurch. But that doesn't change the fact that, *qua* promise, my utterance is supposed to be followed, at the appropriate time, by the event of my picking you up.

When a government issues a law, it is assuring protection of a right. Suppose that the government makes it a law that employers must allow employees who are single-parents extra-time off with pay. That law assures the truth of some proposition more or less like the following:

(^) given any single parent-employee x of a company C, x's need for extra paid time-off is protected from any attempts that C might make to thwart x's satisfaction of said need.

The objectives of the people *making* that promise may *not* be to provide single-parent employees with any such protections. Perhaps the law-makers have no intention of enforcing the law. We could even imagine a situation where, paradoxically, they issue that very promise to guarantee that single-parent employees would *not* be given the aforementioned protections. The law-makers may know that, the cultural context being what it is, employers will lash out against that law by giving single-parent employees even *less* time-off than they had previously.

None of this changes the fact that, *qua* assurance of a protection of a

right, the law-maker's act is supposed to lead to a situation where some right that is currently not protected *is* protected. Under the circumstances just described, the law-maker's act, *qua* promise, is a success exactly if (^), or some similar proposition, is true.

So whatever the psychological motivations of the law-makers may be, there is a very clear sense in which a law is *supposed* to lead to protection of a right. There are many reasons why that promise might not *in fact* lead to any such protection. It may be that the promise is not kept. It may be that, given the nature of the assurance, the promise *cannot be kept* or, at any rate, cannot be kept in a non-vacuous sense. Remember Stalin's law that critics of his regime are to be denounced to the authorities at once. As we discussed, that law was an assurance – one to the effect that *given anyone x whose rights are threatened by the sort of counter-revolutionary menace posed by those critical of Stalin's regime, x's rights would be protected from said menace.* Since[206] *there is* no person x satisfying the conditions in question, the italicized proposition is incapable of being non-vacuously true. So Stalin's law is an assurance that cannot be met in a non-vacuous fashion. Nonetheless, Stalin's law, *qua* assurance, is *supposed* to lead to a situation where the italicized proposition is (non-vacuously) true. In so far as no such situation comes into existence, that law, *qua* law (i.e. *qua* a certain kind of governmental assurance), is a failure. The fact that Stalin's own objectives were malicious in no way bears against this last fact.

According to our analysis, a law is a certain kind of assurance of protections of a right. The satisfaction-condition for any law, *qua* law (as opposed to *qua* machination on the part of the law-maker), is to create a situation where some moral good that is not currently being protected *is* protected. This establishes an internal, non-accidental relationship between law and morality.

Of course, some laws fail to *do* any good. But this doesn't threaten our analysis. Just as assertions don't always do what, *qua* assertions, they are supposed to do (namely affirm a true proposition), so promises don't always do what, *qua* promises, they are supposed to do. The fact that not all assertions are true, and that not all of them are even *intended* to be true by those uttering them, doesn't warrant our rejecting the obvious truth that, at some fundamental level, assertions are *supposed* to be true. Similarly, given only that not all governmental assurances of protections of rights lead to the actuality of such protections, and that not all of them are even *intended* to do so, it would be absurd to say that such assurances are not, in some significant sense, *supposed* to do just that. We may therefore conclude that our analysis posits a non-trivial connection between law and morality and that, at the same time, our analysis– unlike other moralistic conceptions of law – is not vulnerable to the obvious truth that there are evil laws. Our analysis – unlike, say, the extreme form of anti-positivism advocated by Augustine -- does not put us in the awkward position of having to deny this (apparent) datum.[207] [208]

We must also remember the fact that there is no legal obligation where there isn't a high degree of moral protection. This means that one doesn't have any legal obligations to a legal system that doesn't, in absolute terms, do a fair amount of good for one. This shows that, even though there are individual laws that may harm one a great deal, there is a sense in which *as a rule* the laws to which one is subject do good for one. Nothing comparable can be said in regards to a gunman's demands.

§ Let us conclude this section by addressing the question: Is legal positivism right or wrong? There are different definitions of the term "positivism". It is wrongly taken for granted that these definitions are all equivalent.

Dworkin says that "positivism" is the conjunction of (i)-(iii). Thus defined, positivism is entirely correct. We saw why. Legality is binary; there are no degrees of legality. By contrast, principles are not binary. By itself, this makes it hard to see how principles could have any place in legality (outside of being incorporated into specific rules). Furthermore, legality is about rules. If a police officer accuses you of doing something illegal, he must be able to cite some rule that you have violated. If instead he cites some *principle*, we must assume that he is a joker or ironist. So if principles are to be part of the law, then law and legality must be separated. But this cannot be done. The law therefore consists of the rules and rulings. So far as principles are involved, it is because the rules and rulings incorporate them. Given this, the truth of (i)-(iii) straightforwardly follows.

Given only that (i)-(iii) are true, it does *not* follow that laws have no moral basis. For the sake of argument, suppose that all laws are rules and rulings. It doesn't follow that every rule – every order, every demand issued by a gunman -- is a law. To think otherwise is nearly enough to infer *all G's are F's* from *all F's are G's.* To qualify as a law, a rule (or ruling) has to be a protection of a moral right. (Even that condition, though necessary, is not nearly sufficient, as we have seen.) So even if all laws are rules or rulings, any law must protect some moral good.[209]

Sometimes "positivism" denotes the view that laws and legal systems don't necessarily have any basis in morality.[210] (This is the sense in which we have used this term throughout most of the present work.) Thus defined, positivism is false, as we have seen.

The metaphysics of law

According to Dworkin, the law incorporates morality by virtue of the fact

that moral principles are *constitutive* of the law. For example, the principle that business should be conducted fairly is one of the components of the U.S. legal system.

There is a purely *a priori* point to be made in response to this. A legal system is a man-made, spatiotemporally-bounded entity. The U.S. legal system came into existence a few hundred years ago, and it will eventually go out of existence. The states of affairs that sustain the existence of that system have spatiotemporal boundaries. So a legal system is a spatiotemporal entity. Of course, a legal system is a much more encompassing entity than, say, a car or a fork, and its boundaries are much more nebulous. But that is consistent with the obvious fact that legal systems have temporal (and, to some degree, spatial) co-ordinates.

A legal system is thus *not* in the same category as platonic *abstracta*, such as ethical or mathematical principles. Dworkin's position requires that abstract objects (namely, ethical principles) be constituents – veritable parts – of non-abstract object (namely, legal systems). On the face of it, it isn't clear how this is possible: surely an ethical principle cannot be a constituent of a car or a fork, or any other paradigmatic spatiotemporal entity. So here we have a kind of *a priori* objection to Dworkin's view.

Of course, this objection is less than decisive. First of all, spatiotemporal entities obviously *instantiate* abstracta. Second, certain spatiotemporal entities – e.g. mental states (beliefs, desires) and linguistic expressions – have abstract entities (propositions, principles) for their *objects.* So we mustn't over-state the gulf between the concrete and the abstract.

But it remains a fact that nothing platonic can be a veritable *constituent* of anything spatiotemporal. Nothing concrete can have something abstract as a *component.* Admittedly, there are some apparent exceptions to this. Suppose that mental state (or condition) M is a grasping of principle (or

proposition) P. In that case, one might be tempted to say that a spatiotemporal entity (namely, M) has a non-spatiotemporal entity (namely, P) as a component. It might seem that meanings are veritable constituents of *languages* and that, at the same time, languages are spatiotemporal entities. [211] (The English language came into existence at a certain time, and will eventually cease to exist.)

But these exceptions to our principle are clearly spurious. I grasp the proposition that 1+1=2. If that proposition were a *part* of my grasping, and thus of some mental content of mine, then that proposition would have spatiotemporal boundaries; it would, in fact, exist within a sub-region of the region occupied by my mental state. But that is obviously not the case. So even though my belief is a spatiotemporal object, and even though there is an internal relation between my belief and that proposition, it is still clearly false to say that my belief has that proposition as one of its *constituents.* The constituents of the proposition believed in (or, by considerations analogous to those just put forth, desired or regretted or intended…) are not constituents of a belief *in* that proposition.[212]

Admittedly, this does leave us with a mystery. Beliefs are individuated by their objects, so that if x is a belief in P, then x is *necessarily* a belief in P. [213] Being propositions, the objects of belief are abstract. Hume plausibly argued that any internal relation between x and y must be understood in terms of spatial *containment.* Given these points, it seems to follow that, contrary to what we just argued, propositions are constituents of mental entities. What has gone wrong?

What has gone wrong is that Hume's principle applies to necessary relations *between spatiotemporal entities.* The property of triangularity is individuated by its liaisons to other abstracta: for x to be that property *is* for it to be the case that (inter alia) anything satisfying x is a planar three-sided

object such that…But the property of triangularity doesn't *contain* anything: only spatiotemporal entities can contain anything, since containment is a spatial relation. So necessary relations are mediated by containment relations only where *both* parties of the relation are spatiotemporal. This requirement is not satisfied in the case of beliefs. Hence our original point – that propositions are not actual parts of beliefs – prevails.

An analogous line of thought shows that, even though linguistic expressions, and other man-made methods of communication, express *abstracta*, nonetheless no expression, and no system of expressions, i.e. no language, has anything non-spatiotemporal as a veritable constituent.

Of course, many authors hold that languages are *not* man-made – that they are platonic objects (function-theoretic pairings of physical-types with meanings, or some such) that people *use*.[214] According to this viewpoint, what came into existence in 1200 A.D. is not the English language *per se*, but is rather humankind's *use* of it. Supposing this is true, it doesn't confute our position. *If* languages are platonic and thus non-spatiotemporal entities, then the fact that platonic abstracta are components of them is irrelevant to the principle that nothing *non*-platonic and spatiotemporal can have something platonic for a part.

So, ultimately, nothing spatiotemporal can have anything non-spatiotemporal as a component. In so far as Dworkin says otherwise, he is simply wrong. More importantly, in so far as Dworkin's position is that judges' (and legislators') decisions are what connect morality and the law, his position coincides with that of the positivist, as we saw earlier (pp. 26-42).

To sum up, purely *a priori* considerations of metaphysics – in particular, the fact that nothing concrete can have something abstract as a veritable constituent – warrant the rejection of Dworkin's particular brand of anti-positivism. Any legal system necessarily embodies a certain morality.

But, *pace* Dworkin, an ethical principle cannot possibly be a *constituent* of a legal system -- except, as we saw (pp. 26-42), in a sense that reduces his position to that of the positivist.

§ Our analysis does not have this problem. An assurance is a temporally-bounded entity. Assurances are created at certain times, and they cease to be effective, and thus to exist, at certain times.

Remember that we must distinguish the *making* of an assurance from the assurance itself. The former is an event, not a condition; the latter is a condition, not an event. It seems that events tend to have better-defined spatiotemporal boundaries than conditions. So assurances *per se* tend to have rather more nebulous spatiotemporal co-ordinates than the issuing of assurances. Nonetheless, an assurance is a spatiotemporal entity. If the McJack's corporation issues an assurance, that assurance comes into and goes out of existence, and is thus not a platonic *abstractum.*

On our view, a law is an assurance of a moral protection, and a legal system is a system of such assurances – a pattern of issuing, and carrying out, such assurances. So our view, unlike Dworkin's, is consistent with the fact that laws and legal systems are non-platonic entities.

At the same time, our view is consistent with the fact that laws have a basis in morality. After all, an assurance to protect a moral good is an inherently moral act or, at any rate, is one that must be understood in moral terms. Given our view, there is no need to embrace the Dworkinian absurdity that ethical principles are actual *parts* of temporally-bounded objects such as legal systems.

Chapter 11 International law

It is obvious that individuals can be subject to laws. But is the same true of *whole* nations? The system put forth in this book suggests a certain answer to that question, and the purpose of this chapter is to state that answer.

Of course, the question "can nations have legal obligations?" is to be distinguished from the question "do nations have *moral* obligations to one another?" It is obvious that, morally, certain countries *should* behave in certain ways towards other countries. But it is not obvious whether such behavior is, or even can be, regulated by laws.

By way of anticipation, here is what we will find. First, there are not currently international laws. Second, international laws can nonetheless exist *in principle*, even though the conditions that would permit their existence would bear little resemblance to current international conditions. So there are possible scenarios where there is international law, but at this point in time we are not in one of those scenarios or in anything that bears even a remote resemblance to one of those scenarios.

In the next few paragraphs, I will outline the argument to be given in this chapter. Doing so will involve repeating some of the main contentions of this book.

There is law only where there is government. The government in a given area is the main guarantor of rights in that area. Given this, let G be the government in some area, and let x be a denizen of that area. x has legal obligations under G iff x depends on, and receives, vital protections from G. There is legal obligation only where there is moral protection. For G to issue a law is for it to issue an assurance to the inhabitants of the relevant area that it will protect some right of theirs.

For there to be international law, there would have to be an international government – a government of nations. After all, since there is no law without government, ungoverned nations are no more subject to laws than are ungoverned individuals. Just as an *intra*national government is something which is responsible for protecting the rights of individuals, so an *inter*national government would be something which is responsible for protecting the rights of nations. There would be an international government, and hence international law, only if there were something on which nations depended for protections of their rights and from which they received such protections.

As a matter of empirical fact, there is no international body on which nations can (or do) depend for their rights. So far as a nation's rights are honored, it doesn't have anything to do with assurances issued by some international body. (It has to do with that nation's strategic position or its military prowess or with *ad hoc* alliances that it has formed with other nations. To a lesser extent, it has to do with the cultures and moral outlooks of its would-be adversaries.) There is nothing to which whole nations are related in a manner comparable to that in which individuals are related to governments. There is nothing that gives to whole nations the protections that governments give to individuals and that form the basis of the legal obligations that individuals have. For this reason there is no international law.

At the same time, there is no reason in principle why there couldn't be something which protected whole nations in the way in which governments protect individuals. So there isn't any reason why whole nations couldn't receive protections analogous to those given to individuals, and which form the basis of individuals' legal obligations.

There are two caveats. First, the existence of the kind of thing capable of

providing such protections – in other words, the existence of a genuine government of nations -- would be possible only in conditions that were dramatically different from those which characterize today's world. Second, and more importantly, the existence of a government of nations, and thus of international law, wouldn't *necessarily* be desirable. This is because, as we will discuss, a nation's rights are not the rights of the individuals composing it, and it may consequently be necessary to violate the rights of a nation to ensure that those of its citizens are given an acceptable degree of protection. The idea that the rights of nations are *inherently* worth honoring embodies a failure to recognize that the rights of nations are not in any significant sense identical with, or even capable of being understood in terms of, the rights of individuals. Permit me now to give a more complete statement of the argument just outlined.

Do nations have rights?

We must start by asking a very basic question. Do whole nations have rights? Here is the relevance of that question. A law is an assurance that a right will be protected. So if whole nations don't have rights, then it immediately follows that there can be no international laws.

Of course, the people who compose nations have rights. The people of North Korea have rights. But does North Korea itself have rights? Given only that each individual North Korean has rights, it doesn't follow that the *nation* of North Korea has rights. To think otherwise is to commit the "fallacy of composition." Given only that each particle composing my body weighs less than a millionth of an ounce, it doesn't follow that I weigh less than a millionth of an ounce. Here is a more apt illustration. Every individual *in* the Mafia has legal rights, but it doesn't follow that the Mafia itself has legal

rights. So there is no straightforward inference from “every North Korean has rights” to “North Korea itself has rights.”

Nonetheless, it is clear, I believe, that whole nations *do* have rights. Let me give an outline of an argument for this view. Ultimately, a right is always a right to flourish, and anything that is capable of flourishing *ipso facto* has rights.[215] Rocks don’t have rights, because the concept of flourishing doesn’t apply to them. But people have rights, because the concept of flourishing *does* apply to them.

For x to be the sort of thing to which the concept of flourishing is applicable, it is sufficient that x have aspirations. Any thwarting of aspirations is *ipso facto* a bad thing (even though, to recall a point that we have stressed, such a bad is often outweighed by some good) and any fulfillment of an aspiration is *ipso facto* a good thing (even though, to echo the last parenthetical remark, such a good is often outweighed by some bad). [216]

It seems plain that nations have aspirations and that, consequently, the concept of flourishing applies to them. Any nation has certain state-level aspirations or objectives – objectives that cannot be identified with those of any one individual, even in cases of the most despotic forms of government.

In some cases, it is tempting to believe that a whole nation’s objectives can be *identified* with those of some one person. But we mustn’t succumb to this temptation. The objectives of North Korea are far more closely allied with Kim Jong Il’s personal objectives than the objectives of the U.S. are allied with those of any one individual. But it would still not be correct to say that North Korea’s objectives are categorically identical with those of Kim Jong Il. The reasons for *not* identifying North Korea’s objectives with those of Kim Jong Il are to be found in obvious applications of the points made earlier (in Chapter 9) regarding the differences between institutional and

psychological explanation.

In any case, it is clear that nations have objectives that cannot be identified with the objectives of any one individual, and that, consequently, there is a sense in which a nation *as a whole* can flourish and therefore has rights.

Can nations be given assurances?

We've seen that whole nations have rights. But can whole nations be assured that their rights will be protected? In other words, given what nations are, and given how *groups* of nations must therefore be interrelated, would it ever be possible for whole nations to be assured that their rights would be protected? And supposing the answer to this last question to be "yes", could nations ever depend for protection of their rights on such assurances? My view is that both of these questions are to be answered affirmatively. But first we must make it clear what would be involved in an affirmative answer to these questions.

Suppose that we are unarmed and are surrounded by gunmen who wish to take our money. I am in no position to assure you that your money will not be taken. Of course, I could *say* "I assure you, my friend, that your money will not be taken." But that "assurance" would be hollow; it would be an assurance in name only and would be an example of what J.L. Austin described as an "infelicitous utterance."

Suppose I randomly go up to a man and a woman and say "I pronounce you man and wife". Because I am not a clergyman, and because these words are not uttered in the context of the relevant sort of ceremony, I haven't *actually* performed the sort of speech-act that is ordinarily associated with those words. My utterance is "infelicitous", since it occurs in the absence of

the sorts of social conditions that would permit my utterance to discharge the relevant performative function – namely, the function of establishing a marital union between two people.

The same principle applies in the case where we are surrounded by the gunmen. I am not in a position to give you a meaningful assurance, and am thus not really in a position to give you *any* assurance. "Meaningful assurance" is a pleonasm, like "real gold." Just as there aren't two kinds of gold – fake and real – so there aren't two kinds of assurances – meaningful and meaningless. My utterance of "I assure you that your money won't be taken" was meaningless, and thus wasn't really an assurance at all.

What condition must be met if x is to be able to give y an assurance? Given what we said in the last few paragraphs, the condition would *seem* to be this: x must be capable of *delivering* on that assurance. But this answer is not quite correct. Suppose that, unbeknownst to you, I have super powers and that, through a mere effort of my will, I can make all of the gunmen vanish. Because you don't *know* that I can make the gunmen vanish, I am not capable of assuring you that they will vanish. Because you don't know that I have this power, I cannot instill in you a rational feeling of confidence that I can make the gunmen vanish, even though I *can* in fact make them vanish.

The other side of the coin is that, if x has a reasonable expectation that y can bring about such and such, then y can *assure* x of such and such *even if y is incapable of bringing it about.* Once again, imagine a scenario where you and I are surrounded by gunmen who want our money. As before, suppose that, unbeknownst to everybody involved, I have the power to make the gunmen vanish through a mere act of thought. The head gunmen says out loud that, after you and I have handed over our money to him, he will divide it equally among his mates. The head-gunman is not *in fact* able to keep this promise, since (unbeknownst to him) I can and will make him vanish before

he can take our money. But he has nonetheless made a genuine assurance. The reason for this is that, even though he cannot keep his promise, there is still a reasonable *expectation* (on the part of the people to whom he is issuing his promise) that he can do so.

Here is what we may conclude from this discussion: x can assure y that P only if, under the circumstances, y has a reasonable expectation that x's statement that P will be the case is true.

This doesn't mean that y has to be *completely* convinced that x's statement will come true. Somebody who is believed to be reasonably, but not infallibly, reliable can make meaningful assurances. If I believe that Smith *generally* keeps his word, Smith can make meaningful assurances to me, even though I also believe that, under extraordinary conditions, he might break his word.

There is another point to make in this same vein. Suppose that x can make it be the case that P, and can thus correctly predict that P will be true, but that y gives no credence to x at all. y firmly believes that x is a pathological liar who cannot be relied upon to keep his word. In that case, for obvious extensions of the reasons considered in the preceding paragraphs, x isn't capable of *assuring* y that P.[217]

Why there are currently no international laws

Let us map these points onto the topic of international law. Laws are assurances of protections. If they exist, international laws are assurances of protections to whole nations. For an entity to be capable of issuing such assurances, there must be a reasonable expectation that it can and will live up to those assurances. So if there are to be international laws, then there must be some entity that (meaningfully) assures whole nations of protections and

that can (therefore) reasonably be expected to live up to those assurances.

Within a nation, a government can provide such assurances because it has a great deal of control over what its subjects do. Because it has such control, a government can *make* the assurances that it provides come true. The U.S. government has an enormous amount of power over its citizens. It thus has the power to *make* it be the case that its citizens typically not kill one another. That is why the government can meaningfully assure its citizens that they will not be killed. That, in turn, is why it can create a *law* prohibiting murder.

If there are to be international laws, there must be an entity that has a comparable amount of control over what whole nations do. Here we must ask two questions. First, as a matter of empirical fact, is there such an entity at present? Second, if not, could there be?

The answer to the first question is "no." If two people are brawling on a street corner, the police put an end to it in a matter of seconds. But there have been no comparable efforts to stop many (or, arguably, any) of the extraordinarily vicious international conflicts that have been raging for the last forty years. A brawl between nations may go on indefinitely; and when such a brawl ends, it is because somebody won, and not because an international peace-keeping body intervened.

There is thus no entity that is capable of (meaningfully) *assuring* nations that their rights will be protected. In any case, whether there is any entity *capable* of issuing such assurances, there is no entity that *does* so. Given our analysis of law, it straightforwardly follows that there is currently no international law.

Here is a different argument. As we saw in chapter 7, for entity x to be subject to the laws issued by entity G, x must depend on G for vital protections. So if x doesn't need G's protections, then x simply isn't subject

to G's laws. It seems pretty clear that, as a matter of empirical fact, the U.S. doesn't depend on the protections of the U.N. At any rate, it would be an extreme stretch to say that the U.S. depends on the U.N. in a manner comparable to that in which an individual depends on his government. The relationship between an individual and his government, whereby the former has legal obligations under the latter, is not in any significant sense paralleled by the relationship between the U.S. and the U.N.. Given our analysis of law, it follows that the U.S. isn't subject to any laws that the U.N. might issue. (Of course, what we just said about the U.S.'s relationship to the U.N. is true of its relationship to any other international body one might consider.)

Of course, the U.S. does (at least sometimes) benefit from the assistance of an international body, such as the U.N.. But it doesn't depend on such assistance in the deep and systematic manner in which an individual depends on his government to protect his rights. So the U.S. doesn't have the sort of dependence on the U.N., or on any other international body, that an individual has on his government and that grounds the legal obligations of the former under the latter. It follows that the U.S. doesn't have legal obligations under any international body.

Let us extend these points. Laws presuppose government. International laws presuppose an international government – a government of governments. We've just seen that, even if such a body exists, certain national behemoths (e.g. the U.S.) are exempt from any laws that such a body might issue. It follows that those nations are not subject to any laws at all, and have no legal obligations. To sum up, if a nation is strong enough not to need protections, then it cannot be subject to international law.

Of course, there are many nations that *do* need such protections. Given current international conditions, are *those* nations subject to international law? No. Even though they need such protections, they typically don't

receive it. And when they do receive it, it isn't from an international body. Given our analysis of law, it follows that such nations don't have any legal obligations towards such a body and, therefore, don't have any legal obligations at all.

Consistent protection necessary for legal obligation

One point should be made explicit. There is legal obligation only where there is *consistent* protection. Occasional protection is not enough. This is because laws are *assurances* of protections, and there can be such assurances in a context where protections are sporadic.

Within a nation, one *consistently* receives certain basic protections from one's government. That is why the government can (meaningfully) assure people that it will protect their rights. That, in turn, is why people have legal obligations under their governments.[218]

But, at this point in time, there is nothing that provides comparably consistent protections to whole nations. We mustn't overstate matters. There *is* international intervention, and international bodies are sometimes partially responsible for this. But such intervention doesn't happen at all consistently. Consequently, nations have no *assurance* that it will happen – no assurance that an international body will protect their rights. It follows that whole nations have no legal obligations under any such body and, consequently, that they have no legal obligations at all.

Consistent protection not sufficient for legal obligation

But even if a given nation – Guyana, for example -- consistently received protections from some international entity, it *still* wouldn't

necessarily follow that it had legal obligations under that entity. This is because there are different kinds of consistency, and the kind of consistency characteristic of the protection that Guyana receives, according to our supposition, might not be of the relevant sort.

How does the U.S. government protect its citizens? It doesn't do so by intervening every time one of them is about to be attacked. It isn't as though, every time citizen Jones is about to be mugged, the U.S. government thwarts the muggers through an *ad hoc* intervention. The government protects Jones by creating general conditions that make it inadvisable, as a rule, to violate certain basic rights of others. It creates conditions that *program for* Jones' protection and that render unnecessary constant *ad hoc* interventions on its part. (Here I am using the expression "program for" in the technical sense given to it by Jackson and Pettit.) The consistency with which Jones' rights are protected thus has a systemic basis.

What is true of Jones is *universally* true. Governmental protection is systemic protection. The statement "no *system* of protections is operative in area A, but there is government in that area" is tautologically false. Since there is no law without government, it follows that there is no law where there is no system of protections. The protections that underlie legal obligation must therefore be consistent in a systemic, and not a purely statistical, sense.

Let us relate these points to the topic of international law. For Guyana to have legal obligations under some international entity X, it isn't enough that X provide Guyana with protections. It isn't even enough that X *consistently* do so. The protections given by X must have a systemic basis. X must protect Guyana by creating general conditions that make it inadvisable for other nations to violate Guyana's rights. Guyana has no legal obligations towards X unless X has so structured international conditions that it is generally

inadvisable for nations to violate Guyana's rights.

Here is what I am led to believe by my extremely limited empirical knowledge of current international conditions. So far as Guyana's rights are not violated, that doesn't have to do with the presence of systemic facts of the sort just described. It has to do with specifics concerning the interests and capabilities of other nations. To a lesser extent, it may even be thanks to the protections issued by some international body, such as the U.N. But it is very hard to believe that the protections that Guyana enjoys have anything comparable to the systemic basis had by the protections that I enjoy, and that underlie my legal obligations.[219]

Epistemological versus political a priority

I would like to develop the general line of thought presented in the last ten or so pages. In order to do so, I must first make a point about how I am going to use a certain expression – the expression "a priori."

According to Kant, my knowledge that 1+1=2 is part of my innate cognitive endowment, and is thus *a priori*. My knowledge that I can expect certain basic protections from the U.S. government is obviously not *a priori* in *that* sense. But it is *a priori* in a less strict, but still significant sense.

I need *special* information to know that Smith's car broke down yesterday. But I don't need special information to know that, as a rule, I can expect certain basic protections from my government. I don't need *special* information to know that, if I walk into a McDonalds in a town that I happen to be passing through, I will probably live to tell people about my subsequent dining experience. I don't need to have done special advance reconnaissance about the psychopathologies of the people working or eating at that

establishment. The reason I don't need such special knowledge is that I can take certain governmental protections for granted: I know that each of the people in that McDonalds knows that he will be severely punished if he tries to kill me and that, consequently, he (probably) won't attempt to do so.

Of course, violent incidents occur in restaurants. But such occurrences presuppose extraordinary conditions, and are thus in the nature of singularities. If somebody *does* try to kill me while I am in that McDonald's, it is because that person is extraordinarily deranged or is otherwise in some kind of extraordinary situation (e.g. he is a member of the Mafia and knows that he can silence any potential witnesses to his crime). But within the extremely broad limits set by such singularities, I can reasonably assume that I will be safe in the McDonalds. Indeed, given the highly irregular nature of the circumstances under which my safety would be threatened, I would need special information to reasonably believe that I would *not* be safe to enter the McDonalds, and would not need special information to reasonably believe that I *would* be safe to do so. As we said a moment ago, the reason I don't need such special knowledge is that I can take certain governmental protections for granted.

My knowledge that I enjoy such protections is a concomitant of my being a cognitively normal person past the age of four, who has more or less normal social relations with others, and who lives on a certain part of the Earth's surface. Such knowledge is therefore coeval with my being able to participate meaningfully in social and political life. Thus, considered as a socio-political entity, though not as a biological or cognitive subject, I am (so to speak) born with such knowledge. So even though it is *a posteriori* in a strictly epistemological sense, that knowledge is *a priori* in a socio-political sense. By contrast, my knowledge that Smith's car just broke down is not coeval with my acquiring a socio-political identity, and is thus not *a priori* in

the sense just described (or *a fortiori* in the strictly epistemological sense). Henceforth, I will use the term “a priori” to denote socio-political, not epistemological, a priority. Having settled this terminological point, let us move on.

Let Smith be an arbitrary Finnish citizen. Smith doesn’t need to do special reconnaissance to know where he can go (in Finland) without being killed or otherwise violated. Nor does he need to be accompanied by a special security detachment. As long as he remains in Finland, Smith can take it for granted that, as a rule, he will be safe.

Of course, there will be *specific* places where it is ill-advised for him to go. For example, there may be some particular biker-bar where his life would be in jeopardy. But such a place would be singularity – an exceptions to the principles that govern the area. (A singularity is a situation where conditions are so extreme that ordinary governing principles don’t apply.[220])

Let Smith* and Finland* be the counterparts of Smith and Finland in a counterfactual world W, satisfying the following conditions. In a purely *statistical* sense, Smith* is no more likely to be violated in Finland* than Smith is likely to be violated in Finland. In fact, we may suppose, Smith* is *never* violated by any of his fellow Finns*.

At the same time, Smith* *does* have to do special reconnaissance in order to know where he can safely go in Finland*. This is because, so far as Smith* has any guarantee of safety, that has to do with random facts about the emotional interests of Smith*’s fellow Finns*, and not with any general, *a priori* expectation of safety that Finns* enjoy.

This is not to say that it is just Smith*’s good luck that he is never violated. Smith* is a cautious and intelligent person. Before he dares to venture out of his house, he does extensive research (through his computer) about the psychologies of his fellow Finns*. On that basis, he correctly

deduces that he can travel safely anywhere in Finland*.

Let us continue our description of W. What we just said about Smith* is true of all of Smith*'s fellow Finns*. They are all very unlikely to be violated anywhere in Finland*, and each of them knows this. But they cannot take this for granted. Though it is completely justified, each Finn*'s feeling of security is based on special *ad hoc* researches on his part, and no Finn* has a legitimate purely *a priori* expectation that his fellow Finns* will respect his rights.

Given how we've described W, it seems an obvious truth, even a tautology, that there is no *government* in Finland*. We have a situation where people freely act benevolently towards their fellow Finns*. In any case, we have a situation where, so far as Finns* behave well towards their fellows, it isn't due to governmental coercion, and is due instead to the idiosyncrasies of their personalities and situations. We have the kind of situation that Proudhon and Bakunin dreamt of: a situation of benign anarchy.

The occupants of Finland* don't have a rational *a priori* expectation that their rights will be protected. Their knowledge that they will be thus protected is based on *ad hoc* investigations. By contrast, Smith (no asterix) *does* have exactly that sort of *a priori* expectation. This is inseparable from the fact that Finland (no asterix) is not in a condition of anarchy, whether benign or not.

Let us now close the argument. There is government only where there is an *a priori* assurance that one's rights will be protected. There is no such *a priori* assurance where there is no *system* of protections. There is no legal obligation where there is no government. So there is no legal obligation where there isn't a *systemic* basis for the protections that one receives. Thus even if international entity X sincerely and veridically assures nation Y that it will protect it in perpetuity, it doesn't necessarily follow Y has any legal

obligations under X. It is necessary that these protections be consistent in a systemic, and not a purely chronological or statistical, sense.

It is safe to say that, given current international conditions, there is no entity that is capable of making systemically based assurances of the kind described. Guyana (or Iran or the U.S) obviously doesn't have any *a priori* expectation that the U.N. will give it such and such protections. It may well be that the U.N. *will* give it such protections. But this could be established only on the basis of special information. Guyana thus doesn't have the kind of *a priori* assurance that *I* have that I can walk the streets safely and that grounds my legal obligations under the U.S. government.

Systemic versus causally determined

I would now like to extend the argument just given. Doing so will involve distinguishing between two concepts that tend to be conflated.

"Systemic" and "causally determined" are not the same thing. It could be that X consistently helps Y, and also that it is deterministically necessary that X forever continue to do so. But it doesn't follow that X's assistance has any *systemic* basis or, therefore, that Y has any legal obligations towards X. Even if its occurrence is deterministically necessary, a sequence of protections wouldn't necessarily amount to protection in the relevantly systemic sense.

Some fiction will help to illuminate the statements just made. Nations A, B, and C provide important protections to Guyana and, moreover, are causally compelled to continue to do so in perpetuity. At the same time, *if* there were some miniscule change in the distribution of mass-energy, then A, B, and C would stop giving this help to Guyana. *If*, at time t, particle P were at some location so much as one trillionth of an inch from location L, then A, B, and C *would* stop protecting Guyana.

So far as the citizens of Guyana *don't* realize that their fate depends on P's exact location at t, they are unaware of the relevant facts; they lack the information needed to have an informed and rational feeling of certainty that A, B, and C will continue to protect their nation. But there is no way that anyone (including the citizens of Guyana) could possibly *know* that P is exactly at L at time t. Such high-resolution knowledge is a pragmatic, and probably also a theoretical, impossibility. So to the extent that anyone firmly believes that A, B, and C will continue to help Guyana, that belief is not rooted in the relevant information, and any feeling of confidence associated with it is irrational. Under these circumstances, there can thus be no genuine *assurance* that A, B, and C will continue to help Guyana, even though A, B, and C are causally determined to do so in perpetuity.

Meaningful assurances can be tendered only in contexts where predictions can be made with some confidence. I cannot *assure* you that you will win the lottery – even though, given the conditions that obtain, it may be deterministically necessary that you win. But I can assure you that I will pick you up at the airport – even though, given the conditions that obtain, it may be deterministically necessary that I *fail* to pick you up. (Given those conditions, the laws of physics may dictate that my car break down before I make it to the airport.)

If one is to make a prediction with any rational confidence, one must know of some structure that *programs for* the predicted event. (Here, of course, I am using the expression "programs for" in the technical sense given to it by Jackson and Pettit.) If something programs for the occurrence of E in world w, then there cannot be just a *single* difference between w and some world where E doesn't occur. After all, if E is programmed for in w, there is *ipso facto* a system or structure of events in w that demand E's occurrence. Consequently, there will be vast, and therefore easily known, differences

between w and any world where E doesn't occur.

Now suppose that, at time t, E occurs in w, but that its occurrence is not programmed for in that world. In that case, there is no limit to how much, prior to time t, w may resemble a world where E *doesn't* occur. After all, a world where E isn't programmed for is one where there is no redundancy at all, so far as E's generation is concerned. This means that, given such a world, *if* things had been even minimally different, E wouldn't have occurred. This in turn means that, in such a world, there can be no assurance that E will occur, since the information that would warrant a rational prediction of E's occurrence would be much too fine-grained to acquire.

Whether an event can be rationally predicted isn't a function of whether it is causally pre-destined to occur. Given the current positions and momenta of the various molecules composing a certain cloud, its exact shape in five minutes may be causally predetermined. But given what we know about the cloud, and given what, in light of interference-effects, it is nomically *possible* to know, there aren't even conceivable grounds for making rational predictions as to the cloud's exact shape five minutes hence.[221]

Predictability doesn't track efficient causality. (I am using this term in the sense in which it is used by Jackson and Pettit.) Efficient causality is a function of what the *exact* initial conditions in question are. Since, for reasons of physical law, and also of human psychology, these cannot typically be known, causal predetermination is, by itself, little guarantee of rational predictability.

But so far as events occur within systems, rational prediction doesn't presuppose such prohibitively minute and extensive knowledge. Suppose that you have a blueprint of some washing machine, and thus know its internal structure. When you put a quarter into that machine, you know within an extremely small margin of error what is going to happen to the quarter.

Because that quarter's fate, so to speak, is now in the hands of the washing machine's structure, a number of things that would otherwise be relevant to predictions about its behavior suddenly become irrelevant. It becomes irrelevant *exactly* how much pressure you put on the quarter when you initially insert it. It becomes irrelevant *exactly* what the quarter's mass is. It becomes irrelevant whether it is tarnished. Outside of a structure, like the washing-machine, the quarter's behavior may turn on what the quarter's *exact* mass is or on *exactly* how forcefully it was pushed.

The exact dimensions of specific events are virtually unknowable. But the structures that program for events are (relatively) easy to know. For this reason, predictability tracks program-causality, not efficient-causality.

A corollary is that the ability to issue meaningful assurances is a function of what the relevant *systemic* facts are. A consequence of this corollary is that meaningful assurances cannot be made where such systemic facts are absent. Given our analysis of law, yet a further consequence is that there can be no laws where events don't happen *within* a structure.

A national government has the power to coerce people. So it has the ability to *make* predictions about human behavior come true, and thus to create an environment that is highly predictable, at least in certain respects. This predictability is what makes it possible for governments to issue meaningful assurances of protections, and is thus what makes it possible for there to be laws.

At this point in time, international events don't occur within a system. There is thus no systemic basis for rational prediction in the international sphere. Of course, such predictions can sometimes be made. But they must be made on an *ad hoc* basis. To know what nation X will do, one has to have quite specific knowledge about its military capacities or about the dictator's thought-process. One has to have something comparable to the microphysical

knowledge needed to predict exactly what shape a given cloud will have in a few minutes. And just as various barriers make it impossible to have knowledge of the latter sort, so various barriers typically (though not always) make it impossible to have knowledge of the former sort.

A related point is that, given current international conditions, there isn't any entity that could possibly coerce nations into acting in a way that would make their behavior amenable to the sort of predictability needed for the meaningful issuing of assurances or, therefore, for the existence of law.

In light of these points, let us suppose, for argument's sake, that the following three conditions are met:

(1) Guyana currently has some kind of *ad hoc* understanding with nations A, B, and C, whereby it consistently receives certain benefits from them.
(2) A, B, and C are *causally* bound to continue to act in this way for all time.
(3) A, B, and C sincerely *tell* Guyana that they will continue to give it these protections.

Given only that (1)-(3) have met, it *still* doesn't follow that Guyana has the kind of assurance that it would need to have in order for it to have any legal obligations towards A, B, and C. So far as Guyana can predict that it will continue to receive this help, this prediction must be made on the basis of knowledge concerning the specifics of these nations – on the basis of knowledge of what, in this context, must be described as efficient-causes as opposed to program-causes. This means, as we saw earlier, that a possible world where A, B, and C *don't* help Guyana after a given time can be arbitrarily similar to our world *before* that time. This in turn means that, even though A, B, and C may be causally bound to help Guyana in perpetuity, Guyana doesn't have any real assurance that they will do so. More precisely,

it means that Guyana doesn't have the kind of systemically rooted assurance that I have that the U.S. government will continue to protect me and that grounds my legal obligations under that government. Consequently, Guyana has no legal obligations towards A, B, and C.

What we just said about Guyana could equally be said of any nation at the present time. Given the analysis of law put forth in this book, it follows that, at present, there are no international laws.

Let us sum up. Predictability tracks program-causality, not efficient-causality. Where there is no system, there is no program-causality. As a matter of empirical fact, there is no international system.[222] So unlike *intra*national events, *inter*national events are not to any degree pre-programmed. They may be causally pre-determined; and in some cases, some people might know might make the right predictions on the basis of a knowledge of these pre-determining factors. But the possibility of such predictions is the exception, not the rule. There could be assurances to whole nations only if such predictions were the rule, not the exception. There could be international laws only if there were such assurances. So there is no international law.

A variant of the argument just given

Let us begin by remembering a point made in Chapter 10 (pp. 186-187). Two entirely distinct moral obligations can be *discharged* in the same way. Suppose that I enter into the following agreement with Smith. Smith will build me a swimming pool iff I behave in a humane fashion towards any homeless people that might happen to be in my neighborhood. (Both parties enter into this agreement freely, and it is wholesome in all other respects.) *Independently* of this agreement, I have a moral obligation to be humane

towards Willy, the local homeless person. But now, because of my agreement, there is a *second* reason why I am morally obligated to be kind towards Willy. If I were to be mean to Willy, I would be violating my agreement with Smith.

Here one might be tempted to say that I have but one moral obligation: an obligation to be humane towards Willy. But this is not so. My moral obligation to honor my agreement with Smith is distinct from my moral obligation to be humane towards Willy. Let W be a counterfactual world satisfying the following conditions. I have entered into the aforementioned agreement with Smith. Willy doesn't exist, and Buddy is the one homeless person in my neighborhood. In W, I have no moral obligation to be humane towards Willy, since Willy doesn't exist in that world. But I still have the very same obligation towards Smith that I have in this world. But in W, that moral obligation is discharged by being humane towards Buddy, not towards Willy. Thus, my moral obligation to be humane towards Willy does not have the same counterfactual properties as my moral obligation to honor my agreement with Smith. So, by Leibniz's Law, those obligations are distinct. It just happens that, in this world, they are fulfilled in the same way.

Let us map these points onto the concept of legal obligation. Independently of any legal system, I have a moral obligation not to kill Fred. Further, my legal obligation not to kill Fred is *a* moral obligation. But it is not identical with my independently existing moral obligation not to do so. Rather, my legal obligation not to kill Fred is identical with my moral obligation to comply with a social system that protects my fellow citizens and myself. Up to a point, my moral obligation to comply with that system is *discharged* in the same way as my independently existing moral obligation not to kill Fred. In other words, there is some one mode of conduct (namely, my abstaining from killing Fred) that discharges my moral obligation not to

kill Fred and *also*, up to a point, discharges my moral obligation to comply with the just mentioned social order. (Remember that, where the law is concerned, behavior is often to be characterized in *privative* terms. More often than not, one honors one's legal obligations by *not* acting in a certain way – by *not* killing, *not* stealing, and so on.) But, as we saw a moment ago, two entirely distinct moral obligations can be discharged in the same way; and my moral obligation not to kill Fred is quite distinct from my moral obligation to comply with that social system. Consequently, even though my legal obligation not to kill Fred is *a* moral obligation, it is entirely distinct from my *independently* existing moral obligation not to kill Fred.

In general, a legal obligation to do X is never identical with an independently existing moral obligation to do so. At the same time, one's legal obligation to do X is always identical with *a* moral obligation. It is identical with one's moral obligation to comply with a social system that protects oneself and one's peers. So legal obligation presupposes a system of protections, as opposed to a series of *ad hoc* protections. For reasons that we've seen, it seems unlikely that, in the international arena, such a system exists. It therefore seems unlikely that there are international laws at present.

international law: a second perspective

In this section, we will put forth an argument that is completely different from those just given. The argument given here is far from airtight, since it involves repeated use of a questionable empirical premise. (Every time I use this premise, or some variant thereof, I will indicate that I am doing so by parenthetically inserting the letters "QEP", which are short for "questionable empirical premise.") Also even if that premise is granted, what follows is only a probabilistic claim. Finally, that probabilistic claim is not specifically

concerned with our own period of history. Given that premise, what follows is that, for any time t, it is highly probable that international law does not exist at t.

Let us start with a terminological point. In this context, when we talk about a "group" of people, what is meant is a reasonably large group – a group large enough to constitute a political entity. The points about to be made are not meant to be true of small groups.

In previous chapters, we saw that legal obligation presupposes moral obligation. In connection with this, we saw that X has legal obligations towards Y only if Y rises above a certain threshold in the way of protecting X's rights. A consequence is that a system of norms or coercions defines legal obligations only in contexts where the absence of such a system would lead to a greater degree of rights-violation than its presence.

At this point in history, there are things that are *referred* to as "international laws." Of course, given only that something is referred to in that way, it doesn't follow that it really is a law. Given the point made in the last paragraph, there is a simple way of determining whether those so-called international laws are *actual* laws. Suppose that those so-called laws were absent, and also that nothing comparable to them took their place. In that case, would nations' rights be more severely violated than they are now? If the answer is "no", then there are no international laws, and the things that are *referred to* as such are laws in name only. On the other hand, if the answer is "yes", it doesn't necessarily follow that those things are laws; all that follows is that *one* condition necessary for their being laws has been met. Whether the answer is "yes" or "no" is a purely empirical question. But I will now argue that that the probability of the answer being "yes" is quite low.

Given a group of people, that group either dissolves (either because its

members kill one another or they leave the group) or it ends up imposing some system of order on itself. There is simply no possibility of a situation of benign anarchy obtaining in a group of individuals. So where groups of individuals are concerned, the presence of law generally does more good than its absence; the benefits of law outweigh its encumbrances. Given our analysis of law, this means that groups of people typically satisfy at least one of the conditions necessary for the existence of law.

Obviously, a situation of anarchy among nations often leads to conflict and, as a result, to violations of rights (both of nations and of individuals). In *absolute* terms, the quantity of rights-violation is enormous. But, as we will now see, it is far from clear whether that quantity is high in the relevant *relative* terms.

Nations often violate one another's rights. But the violated countries seldom cease to *exist* in consequence of such attacks; there are very few cases where one nation actually *destroys* another (QEP). Consider the nation of Georgia. In the 1700's, it was invaded and appropriated by Russia. The latter imposed its political structure on Georgia, and involved itself in the most intimate aspects of the day to day affairs of the Georgian people. Russian control over Georgia intensified during the Soviet Era.[223]

There can be doubt that this was a rather extreme form of international aggression. But surely the nation of Georgia *survived* during all of this (at least on *one* natural delineation of the term "nation"[224]). This is evidenced by the fact that the Georgian people became an autonomous state at practically the instant at which the Soviet Union fell. What we just said about Georgia is true of *all* of the former Soviet Republics; and it applies, to some degree at least, of many other nations that have been the objects of colonization and other forms of international aggression. As a rule, such aggression very rarely leads to the wholesale *extermination* of nations (QEP).

It leads to subjugation and humiliation – but not, typically, to extermination.

In light of these points, let us do an exercise. Pick a period of history when, in your view, there were *clearly* no international laws. Make sure to pick a period during which there were actually nations or, at least, political organizations roughly comparable to nations. (So, for example, don't pick the year 85,000 B.C.) Let t* be the period of time that you choose. Now consider a situation where, instead of the (let us suppose) 134 nations that existed during t*, there are 134 individual people. Further, suppose that those individuals interact with one another in a way that parallels the way in which the world's nations at t* interact. (I must reiterate that those individuals are *not* interacting in the way that 134 *actual* individuals would interact if in a context of lawlessness.)

Under these circumstances, there would be many cases where somebody was kicked in the shins or punched in the nose, and there might even be a few cases where somebody held somebody else in captivity. But for the reasons given two paragraphs ago, there would be few or no *killings* among those 134 individuals (QEP). Though there would be a reasonable amount of discontentment, there would be nothing even remotely comparable to the kind of situation that would obtain if you took 134 *actual* individuals and put them in a context where they were governed by no rules at all. As we saw earlier, a group of *that* kind would either disband or kill itself off, supposing that it didn't quickly impose some kind of order on itself.

As we noted, the members of our hypothetical group are not unscathed. But would they be better off if they were subject to *laws*? Would those rules do more good than harm? Would the injustices they thwart exceed the inconveniences they create?

Not necessarily. The rights-violations experienced by those 134 individuals, in their condition of lawlessness, are not much worse than the

rights-violations typically experienced by actual individuals *within* nations (QEP). To see this, pick 134 *actual* individuals at random. (So imagine your best friend Gabe, his cousin Sally, her boyfriend Russell, and so on.) Of course, all of these people are subject to law and are probably law-abiding for the most part. Practically any given one of those individuals will be disgruntled to some degree, and will have to put up with serious encroachments on his ability to have a happy life.

But – and this is the important point – the rights-abuses experienced by these people would not even be remotely comparable to the corresponding abuses experienced by individuals in a condition of lawlessness. In the one case, we are almost never dealing with anything more than hurt feelings and broken dreams.[225] In the other situation, we are dealing with wholesale physical slaughter. Clearly the rights-abuses experienced by the 134 people whom you chose at random (your friend Gabe, his cousin Sally….) would be much more comparable to the rights-abuses experienced by the 134 people whose behavior and experiences modeled those of the world's nations (at time t*) than they would be to the rights-abuses experienced by a group of 134 individuals who were stuck together in a condition of lawlessness. This suggests that, in a state of anarchy, the rights-abuses that whole *nations* experience are comparable to the rights-abuses experienced by individuals who are governed by law.

In its turn, this last point makes it unclear whether the benefits that nations would receive from a system of international law would exceed the general inconvenience and loss of freedom of being under such system. As we've seen, there is no possibility of law in a context where law would do more harm than good. It follows that it is unclear, at best, whether law would even be possible in the context defined by the 134 nations that we've been discussing.

In absolute terms, a state of anarchy among nations may be extraordinarily malignant. But when we wish to assess whether law is possible in a given context, the relevant question is not "in absolute terms, how much damage follows from a condition of lawlessness?", but rather "how much *better* would the presence of law be than its absence?" If the answer is "not at all", then there is no possibility of law in that context. As we've seen, a state of anarchy among nations is likely to define a situation where the answer to that question is "not at all – the presence of law would make things worse" (QEP). In an international context, the encumbrances of law are likely to outweigh its benefits (QEP). Given this, and given that law is possible only where its presence would do more good than its absence, there is only a low probability that international conditions during an arbitrarily chosen period permit the existence of international law. This suggests that the question "are there currently international laws?" is rather like the question "did Smith roll double sixes five times in a row?" Both questions are empirical; but the probability of an affirmative answer in either case is low.

§ So far as this last (putative) fact has not been seen, and so far issues concerning international law generally remain obscure, that is largely due to a failure to distinguish facts about individuals from facts about institutions.

A nation's state-level objectives are not to be identified with those of its citizens. Even in cases where a nation's state-level objectives coincide with those of some one person – e.g. a dictator or legislator – we are, strictly speaking, dealing with two numerically distinct objectives that happen to converge. This is a direct consequence of the points made in Chapter 9 concerning the differences between psychological and institutional explanation.

Remember our story (pp. 154-156) about the petty congressmen who create the bill forbidding insider-trading. What *they* are doing, as individuals, is one thing – they are attempting to thwart an old college buddy who has surpassed them financially. But what *Congress* is doing is very different -- it is discharging its function of creating laws that program for the welfare of the nation. So the objectives of those congressmen are different from the objectives of Congress as a whole, even though in this case some one set of behaviors realizes the attainment of both of those objectives. Similar remarks show why *Stalin's* decision to exterminate the kulaks is very different from the *Soviet Union's* to do the same thing – even though there is some one set of actions that leads to the attainment of both objectives. So a nation's objectives and aspirations are *categorically* distinct from those of the individuals composing it.

As we discussed previously, a right is always a right to flourish. Flourishing is to be understood in terms of the attainment of objectives.[226] So given that a nation's objectives are categorically distinct from those of its citizens, it follows that the same is true of a nation's rights. A nation's right to develop nuclear weapons, or to invade country X, isn't identical with the rights of any of its citizens.

A corollary is that the rights of many of the individuals composing a nation may be horribly violated even as the rights of the nation as a whole are honored. A related corollary is that the rights of the individuals composing that nation may be honored even as the rights of the nation as a whole are violated. So a minor violation of a *nation's* rights might involve an extreme violation of the rights of many of the individuals composing it; and a major violation of the *nation's* rights might lead to little or no violation of the rights of the individuals composing it.

In order to illustrate the principles under discussion, let us a consider a

case where a nation's political institutions make little or no allowance the welfare of the individuals composing those states. Of course, it may be objected that any political institution *ipso facto* makes at least *some* non-trivial allowance for the welfare of the people composing it. (In fact, this is what I myself have argued.) But even if this is true, actual political institutions are still fruitfully understood in terms of how well they approximate to that scenario – just actual physical objects are fruitfully understood in terms of how well they approximate to impossibilities such as frictionless planes. So let us suppose that N is a totalitarian state, and that the *only* value recognized by its political institutions is fidelity to the fuehrer. Thus, N's objectives *as a whole* are met so long as the fuehrer is doing well.

If enough of N's citizens are killed, that may make it harder for the fuehrer to maintain his quality of life. But whether the fuehrer is doing well is not a direct function of the *quantity* of N-citizens that are killed. If the fuehrer becomes violently upset because his valet died, the objectives of N *as a whole* fail to be met. By the same token, if a town of a million people is vaporized, that may not have any effect on the fuehrer's welfare or, therefore, on that of N as a whole. Because of the values embodied in its political institutions, N's welfare is not a function of how many of its citizens are hurt, but of the relationship of those who are hurt to the fuehrer.

As previously noted, even though no actual political situation is as extreme as the one just describe, any such situation approximates to it to some degree or other, and can thus be understood in terms of it. It is a distinct possibility that, given the values currently embodied in its political institutions, North Korea *as a whole* is flourishing, even though a quarter of its citizens are starving. (In a moment, I will deal with the obvious reservations that many will have about this claim.) Facts about institutions do not correspond in any simple way with facts about individuals; and facts

about a nation's rights don't correspond in any neat way with facts about the rights of the individuals composing them.

Some will take it as obvious that, because a quarter of its citizens are starving, North Korea is a "failed state"; and some will take this supposed fact, in its turn, as proof that I am wrong to see a nation's welfare as being so disjoint from that its citizens.

But this usage of the term "failed state" embodies the very conflation I am trying to expose – the conflation of individual with national facts. When people describe North Korea as a "failed state", they mean that it has failed to serve the interests of its citizens, not that it has failed to serve *its own* interests. But in this context, what is relevant is how well North Korea serves its own interests. International law governs nations, not individuals. So when we wish to know whether international law is possible, we must think in terms of the welfare of *nations*, as opposed to that of individuals. We must look *not* at how well nations are serving the interests of their citizens, but at how well the interests of nations *as wholes* are being served. From *that* viewpoint, North Korea may not be a failed state, notwithstanding the wretched condition of its citizenry.

When we look at nations, we tend to identify *their* welfare with those of their citizens. As we've seen, we are wrong to make this identification. To the extent that we do so, we cannot assess whether *nations* (as opposed to the individuals composing them) would be better off for the presence of law or, consequently, whether international law is desirable or even possible.

Of course, violations of national rights typically involve horrible violations of human rights. That is surely one reason why we tend to identify national with individual rights; and that, in turn, is probably why many people see it as self-evident that a nation's rights are as inviolable as those of a person. But so far as the welfare of today's nations is confluent with that of

their citizens, that is a historical accident – a contingent fact that has no *a priori* basis. International law protects nations, not people (except in so far as those nations, in their turn, protect people). This suggests that international law wouldn't *necessarily* be desirable, even if circumstances where it were possible. We will pursue this matter in a moment.

Why there could be international laws

We have seen a few different reasons to believe that, as a matter of empirical fact, there are no international laws. But *could* there be such laws? Yes.

To see why, let us review the essentials of our analysis of law. G is a government in area A iff G is the main guarantor of rights in A. For G to issue a law is for it to assure the people in A that it will protect some right of theirs. G's legal system, supposing that it has one, is a collection of such assurances. G is capable of meaningfully making such assurances only in so far as it has the power to *coerce* people into behaving in certain ways and, therefore, into respecting one another's rights. G can make meaningful assurances, and can thus issue laws, only to the extent that it has the power to coerce people into making its assurances come true.

For there to be international law, there would have to be a government of nations: a metagovernment. For something to qualify as such a metagovernment, it would have to be something on which individual nations depended for fundamental protections, and which also provided those protections. Supposing that there were such a thing, for it to issue laws would be for it to assure nations that it would protect some right of theirs. For it to be capable of (meaningfully) issuing such assurances, it would have to have

the power to coerce nations into acting according to its will.[227] It would have to have the kind of power over nations that actual *intra*national governments have over individuals.

If those conditions were met, then there would be international law, and nations would have legal obligations to one another. As we discussed earlier, not all of those conditions are currently met. In fact, none of them are met. But there doesn't seem to be any reason in principle why they couldn't be met. A situation satisfying those conditions would not be a surd – it would be in the same category as a square circle or a disembodied grin.

§ But there are two caveats. The first caveat involves our making a distinction between "universal" and "non-universal" international law. There would be universal law iff there were some system of laws to which *all* nations were subject. There would be non-universal law iff there were some system of laws to which some, but not all nations, were subject.

Now we can state the first of our two caveats: any situation satisfying the conditions necessary for universal international law would bear little resemblance to current international conditions. For there to be universal international law, there would have to exist something that was at least approximately as powerful with respect to each of the world's nations as national governments are with respect to each of the individuals composing them. Otherwise, there wouldn't be anything which had the coercive abilities to make meaningful assurances of the sort discussed earlier (pp. 217-219). Further, that thing would have to be something that guaranteed the rights of all nations, as opposed to something which merely dominated them.

The first of these conditions couldn't be met in any scenario that bore even a remote resemblance to actual international conditions. As we saw earlier, behemoths that don't need the protection of government cannot be

subject to law. And in so far as there is *one* thing that isn't governed by law, there is no possibility of law at all, as we may show through a bit of fiction. [228]

In nation N, the number seven is sacred, and anything associated with it is seen as being of divine origin. Consequently, each person whose social security number ends with a "7" has complete immunity from prosecution, and can simply do as she wishes: she can kill, loot, and so on. Such people are referred to as "Lucky Sevens", and people who are not Lucky Sevens are referred to as "Citizens." The existence of Lucky Sevens would render null and void any assurances which N's government might otherwise be able to provide. At any minute, a Lucky Seven can strip a Citizen of his life, property, or dignity. In any case, so far as a Citizen holds onto his life, money, or dignity, it is only by the good graces of Lucky Sevens; and a Citizen thus no *assurance* (of the relevantly systemic kind) that his rights to these things will be protected. So far as a Citizen *does* have such an assurance, it is of the irrelevant *ad hoc* kind that we discussed earlier, in connection with Guyana, and is not of the kind that underlies legal obligation. So there wouldn't be any significant sense in which Citizens enjoyed the protections of *law*.

Given only what we've said, it is possible that there can be law *within* the group of Lucky Sevens. But there obviously couldn't be laws *common* to Citizens and Lucky Sevens. This proves our point that *in so far as* there are ungoverned agents, there is no law.

Just as the existence of Lucky Sevens renders impossible the existence of laws that apply to *both* Citizens *and* Lucky Sevens, so the existence of behemoths, like China or the U.S., renders impossible the existence of laws that apply to *both* behemoth nations *and* to smaller nations. So there is, at present, no possibility of *universal* international law – of a system of

international law that applies to *all* nations. At the same time, given what only what I've said, it is still a possibility that nations of approximately equal strength, e.g. Denmark and Albania, are subject to common laws.

Of course, smaller nations, e.g. Barbados, have a reasonable assurance that the U.S. or China won't invade them tomorrow. But those assurances are not of the relevantly systemic kind. They involve *ad hoc* knowledge of the strategic and cultural dispositions of these other countries. Consequently, those assurances don't underlie any legal obligations on their part.

So far as there are behemoths, there is no possibility of universal international law. Consequently, for there to be universal international law, all nations would have to be of approximately equal strength.[229] It is quite obvious that any world where this condition were satisfied would be radically different from ours. So the first of our two pre-conditions for universal international law ("there would have to exist something that was at least approximately as powerful with respect to whole nations as individual governments are with respect to individuals") couldn't be met in any world that was even approximately like ours. The second of our two conditions ("that thing would have to be something that guaranteed the rights of nations, as opposed to something which merely dominated them") can be met only if the first is met. So even though universal international law is theoretically possible, it couldn't possibly exist in any scenario that, in respect of international relations, bore even a remote resemblance to reality. But, to echo a point made a moment ago, this leaves open the possibility that there could be law *within* a group of nations that were all of approximately equal strength; in other words, it leaves open the possibility that, at present, there is non-universal international law.[230]

The second caveat is, I believe, much more important. Given that there

could in principle be (universal) international laws, would their existence *necessarily* be desirable? As we saw a moment ago, the answer to this question is “no.”

A quick look at modern history will help us to draw one of the more important consequences of this answer. It is safe to assume that, for the most part, the citizens of the U.S.S.R under Stalin were not happy. At the same time, the U.S.S.R. under Stalin had various state level-objectives. It aspired to exterminate the kulaks, to industrialize large sectors of its economy, and so forth. It is by no means incoherent, and is probably factually accurate, to suppose that the fulfillment of these state-level objectives was not compatible with the happiness of most Soviet citizens. It is safe to assume, in fact, that the happiness of those individuals wasn’t even compatible with the existence of the system of government associated with those objectives. It follows that, if the rights of the Soviet people were to be upheld, it would have been necessary to violate the rights of the Soviet nation as a whole. In general, it may be morally right to violate a nation’s rights.

§ But given this last point, it might seem that there is an incoherence in our analysis of law. Throughout this work, I have defended the principle that national governments, and hence nations themselves, are at their very core things which protect the rights of individuals. But in the last section, I said that the rights of nations might *conflict* to an excessive degree with those of individuals. Haven’t I contradicted myself here?

No. As we discussed, the Soviet government under Stalin was *at bottom* a protector of rights. (That government was, of course, a dismal failure. But it failed *by comparison* with good governments, e.g. the present day Icelandic government, and not in absolute terms.) But that is perfectly consistent with the proposition that the rights of the Soviet Union as a whole would have had

to have been violated in order for the rights of individual Soviet citizens to be protected to anything even approximating the extent to which they *would* have been protected under a different form of government. A violation of the Soviet Union's rights would thus have been a pre-requisite to giving an *acceptable* degree of rights-protection to the citizens of the Soviet Union.

Here is a fictional illustration of the principle behind these remarks. Molly is extremely abusive to her daughter Sally and subjects Sally to horrible psychological and physical torment. But these abuses happen *within* the context of a parental relationship. Molly provides basic protections for Sally: she gives her food and shelter; she drives her to school; and she pays her medical bills. Also, when Molly is abusive to Sally, it is done under the banner of good parenting and even, to some degree, with the *intention* of being a good parent. When Molly locks Sally in the closet for five days without food or water, it is "for her own good." After all, she is just trying to keep Sally on the straight and narrow path of virtue, and sometimes stern measures are what is necessary to do that. (Molly only wishes that somebody had been so strict with *her*. That way she wouldn't have ended up being an unwed mother at age fifteen.)

Molly's despicable abuses are thus cases of *mis*parenting and thus of *mis*nurtering. It isn't that Molly is *failing* to be a mother, and thus a protector. On the contrary, Molly's wickedness is expressed in a distinctively maternal way. That wickedness is expressed not in *not* mothering, but in *mis*mothering and *mis*protecting. Molly is failing *as* a mother; but she is not failing to *be* a mother. When Molly beats Sally for watching a movie with Tom Cruise in it, on the grounds that any exposure to such sexually charged filth would destroy Sally's chances for eternal salvation, Molly is doing in a perverse and extreme way what a good mother would be doing in grounding her daughter for coming home drunk after a date with a member of the Hell's Angels.

Molly's wickedness consists *in her doing her job as a mother and nurturer*. That is, it consists in her doing that job *in a bad way*. It does not consist in her *not* doing that job.

The way in which Molly is being evil to Sally is thus different in kind from the way in which a mugger or a burglar might be evil to Sally. Molly's abuses presupposes a parent-child context; they are evil articulations of that kind of protective relationship, and are thus very different from the evils to which Sally would be subjected by somebody on whom she didn't depend for protections. (Even if – as is probably the case, given her conduct – Molly is to some degree being *deliberately* evil towards Sally, Molly is still being evil towards Sally *in a maternal way*. Molly's maternal responsibilities – giving advice, disciplining, nurturing – are the vehicles through which her ill-will is expressed. So regardless of what the intentions behind it are, Molly's destructive behavior is a form of *mis*mothering, not of *not* mothering.)

Social Services would surely be justified taking Sally away from Molly, and in finding more caring parents for her. So while it is true that Molly is Sally's protector (even though she is also her destroyer), that is perfectly consistent with the proposition that it would be morally *de rigueur* to pull Sally out of that relationship, and embed her in a better one. Similarly, while it is true that governments are *ipso facto* providers of protections, that is perfectly consistent with the proposition that, in some cases, it might be morally necessary to destroy a nation's political structure, and then embed its citizens in a better one.

This last point mustn't be taken to mean that nation's rights are not typically to be respected. It must be taken to mean that the conditions under which a nation's rights can be violated are not comparable to those under which an individual's rights can be violated. It is *categorically* wrong to violate a person's rights (even though, in some cases, not doing so would be

more wrong). But it is only circumstantially wrong to violate a nation's rights. It is wrong to do so only if the amount of overlap between a nation's welfare and that of its citizens exceeds a certain threshold. It may well be that, in every actual country, this threshold is exceeded. But, as we've seen, a brief survey of historical and hypothetical fact shows there to be no reason why this must the case.

A paradox concerning political institutions and their relation to individual rights

I would now like to discuss a topic that doesn't have any special relationship to international law, or even to law in general, and that concerns political institutions generally. But it is a point that must be addressed; and I think that, given the analysis put forth in this work, we may be in a good position to weigh in on it.

We've already seen that a nation's rights *can* conflict with those of its citizens. But it is easily seen that, as a matter of analytic necessity, such conflict is inevitable. In fact, what is true of nations is true of political institutions generally: there is no *conceivable* situation where there are political institutions but where nobody's aspirations are to any degree thwarted. A situation where nobody's aspirations are thwarted is a situation where everybody does as she wants; and such a situation is one where nobody's rights are protected and where, consequently, there are no political institutions. Thus, a political institution is *ipso facto* one where people's freedoms are abridged and where, consequently, their rights are violated to some degree. So the rights of political institutions, including nations, *necessarily* conflict to some degree with those of individuals. (Some would

say that not every abridgement of a freedom is a violation of rights. I would disagree, for reasons that we have given many times and that we will revisit in a moment.)

At the same time, it is obvious that the rights of individuals would be trampled upon if it weren't for the protections given to individuals by political institutions. If it weren't for the law, my quality of life would plummet.

So political institutions necessarily conflict with individual rights and, at the same time, the rights of individuals wouldn't be honored unless there were political institutions.

But isn't the very last sentence paradoxical? Isn't it self-contradictory or, at least, incoherent? And doesn't this last fact show that one of my premises must be wrong? Supposing that a contradiction follows from my premises, it follows that one of those premises must be wrong. So *either* political institutions don't necessarily conflict with individual rights *or* individuals don't depend for protections of their rights on such institutions. Given this choice, one must reject the first disjunct – one must say that political institutions *don't* necessarily involve any violation of individuals' rights. This is, in fact, what most theorists would say.

But there is no paradox, and there is no need to deny either of our two premises. On the one hand, there is no need to deny the obvious fact that the existence of any political institution necessarily involves at least some thwarting of individual-objectives; and there is no need to deny the concomitant fact that, in so doing, such institutions are violating individual rights. On the other hand, there is no need to deny the obvious fact that, but for the existence of political institutions, the rights of individuals would be trampled upon.

The only reason there *seems* to be a paradox here is that issues relating

to *causal* necessity are being conflated with issues relating to *conceptual* (or analytic) necessity.[231] As a matter of *conceptual* necessity, a political institution is one where individuals are constrained – where people cannot do exactly as they wish. So the existence of political institutions is a *conceptually* or *analytically* sufficient condition for the existence of a scenario in which people are not allowed to do as they wish, and in which they consequently cannot flourish to a maximal degree. After all, in a situation where everybody is allowed to do exactly as she pleases, there is *by definition* an absence of any sort of political structure.

At the same time, political institutions are necessary for the protection of the individual's rights. But this necessity is causal, not analytic. Hart (1977) is very clear about this, and illustrates it with a piece of science fiction similar to the following. Suppose that we had impenetrable shells, that we didn't need food, or shelter, or the cooperation of others. In that case, we wouldn't need political institutions: we could flourish perfectly well without them. So while it is true that, in actuality, we depend on political institutions for protections of our rights, the nature of that dependence is causal, and not conceptual.

It is thus in no way a paradox to say that the protection of individual rights presupposes the existence of political institutions and *also* that the existence of political institutions presupposes a violation of individuals' rights. For there is no incoherence in saying that the existence of political institutions is *causally* necessary for the protections of individual rights, while also saying that the existence of such institutions *entails* the abridgments of individual rights.

Here is an illustration of the principle at work. Let X, Y, and Z be three straight and rigid objects.[232] For X, Y, and Z to form a triangle, it is necessary that they be arranged in a certain way. (They must be arranged in

such a way that not all three of them intersect but such that any two of them intersect.) Given the laws of physics, and given the materials of which X, Y, and Z are made, it could well be that, if those objects are put in that arrangement, they eventually cease to be straight. In other words, arranging X, Y, and Z in the form of a triangle may make it *causally* necessary that they not be straight. At the same time, it is clear that for X, Y, and Z to form a triangle, it is *analytically* necessary that they be straight. The following statement is thus in no way incoherent: "If X,Y, and Z are arranged in the form of a triangle, that is *analytically* sufficient for their being straight, and it is also *causally* sufficient for their soon becoming bent." So the proposition "if X,Y, and Z are arranged to make a triangle, they will necessarily stop being straight" is compatible with the proposition "if X, Y, and Z form a triangle, they are necessarily straight."

Given this, let us reconsider the two premises that lead to the political paradox we've been discussing. For clarity's sake, let me again say what they are:

(1) political institutions demand the abridgment of certain individual liberties and, therefore, necessarily conflict to some degree with the rights of individuals;

(2) the rights of individuals would be completely and totally dishonored if it weren't for the existence of political institutions.

For reasons exactly analogous to those stated in connection with rigid bodies X, Y, and Z, the truth of (1) is in no way inconsistent with the truth of (2), the reason being that the occurrence of "necessarily" in the one case refers to causal necessity, while the other occurrence refers to analytic

necessity. Thus one is guilty of a version of the fallacy of equivocation so far as one sees any paradox here.

There is another, rather more pedestrian reason why (1) and (2) don't conflict with each other. The rights that political institutions protect may be *different* from the rights that they thwart. The existence of law may thwart certain aspirations that I have, while being necessary for the fulfillment of other aspirations of mine. The existence of law is obviously necessary (in a causal sense) for the existence of plumbing, mass-media, books, and thus for my continuing to have a life at all like the one I am used to. So the existence of law is causally necessary for my continuing to write books, play the piano, and so forth. In fact, the existence of law is probably causally necessary for my continuing to live at all.

At the same time, the existence of law is *analytically* incompatible with at least *somebody's* being able to fulfill at least *some* of his aspirations. But the aspirations that law thwarts are not always the aspirations that law protects. For example, because of the law, my life-long dream of painting a giant replication of the Mona Lisa on Mount Rushmore is going to be thwarted. But the fact that the law thwarts *this* aspiration of mine, and thus to that extent thwarts my right to flourish, is obviously compatible with the fact that the law enables me to fulfill my desire to learn to play all of the Beethoven sonatas, and thus protects my right to flourish to *that* extent. The law has a dual nature: it simultaneously thwarts and protects aspirations and, therefore, rights. But the rights protected are not always the rights that are thwarted. So, in those cases, the appearance of paradox generated by the dual nature of law is easily dissolved.

But, in other cases, the appearance of paradox is not always so superficial, and cannot be dissolved by saying that the rights that are thwarted are distinct from those that are protected. And because of such cases, we must

invoke the points made earlier – concerning the difference between causal and analytic necessity – to dissolve that appearance.

Here is an illustration. People generally aspire to be free – to live as they want to live, and not as others want them to. Because there is law, it is impossible for people to live *exactly* as they want to. But to the limited extent to which they can do so, it is thanks to the law. So here we have a case where the law both thwarts, and renders possible, the fulfillment of a single aspiration. Obviously we cannot dissolve this paradox in the manner discussed in the last paragraph, since we were there concerned with cases where the aspirations protected by law were *distinct* from those thwarted by it.

But given the points made earlier, concerning the difference between causal and analytic sufficiency, there is no trouble dissolving the paradox. The law is *causally* sufficient for the thwarting of certain aspirations and, therefore, rights. But the law is *analytically* sufficient for the protecting of those same rights. After all, where there is law, there is *by definition* a protecting of rights. (Compare: "If X,Y, and Z are arranged in the form of a triangle, that is *analytically* sufficient for X, Y, and Z being straight, and is also *causally* sufficient for their becoming bent in due course.") .

Most importantly, both (1) and (2) must be accepted. One is guilty of empirical error if one denies that political institutions are necessary for the protections of people's rights. One is guilty of conceptual incoherence if one denies that political institutions always involve an abridgment of individual liberties and of the individual's concomitant ability to flourish (at least in certain directions). Finally, one is guilty of a revisionist and implausible morality if one denies that an abridgment of liberties, and of one's concomitant right to flourish, is a violation of one's moral rights.[233]

In any case, a nation's rights are not identical with those of its citizens

and, moreover, the former can easily conflict to a prohibitively high degree with the latter. This point is easy to corroborate through examination of both historical and hypothetical fact. This warrants a negative answer to the question "would the existence of international law *necessarily* be desirable?"

Suspensions of rights versus positive violations of rights

There is an objection that many might have to our analysis of law. The objection doesn't have any special connection with our analysis of *international* law (even though, if correct, it nullifies that analysis). But it is an important objection, and this is as good a place as any to deal with it.

I have said that x has legal obligations under government G iff x depends on G for vital protections and also receives them. But there is an obvious problem with this view. Suppose that, for no reason at all, the police arrest and torture me. After being released, I still have legal obligations under the U.S. government. It isn't as though, just because I was tortured, I no longer have any legal obligation to obey traffic laws.

Governments often torture and otherwise violate their own citizens. But those citizens continue to have legal obligations under those governments. In any case, being tortured or otherwise violated by one's own government doesn't categorically exempt one from legal obligations under that government. So it seems that, contrary to what we have maintained, one can have legal obligations towards something that *fails* to protect one.

I believe that this point embodies a confusion -- a confusion of a piece with the binary moral thinking earlier descried -- and that this confusion has been one of the many factors inhibiting acceptance of a moralistic conception of law. There is, I will now argue, a difference between positively *violating* someone's rights, on the one hand, and *suspending* protection of a person's

rights, on the other. Consequently, a person's rights can continue to be protected even while that person is undergoing positive rights-violations.

As before, fiction may be of use. N is a totalitarian state. The government controls every aspect of life – what one does for a living, who one marries, what one reads, and so on. Because the government has such a tight lock on what people do, there is virtually no crime in N. Citizens of N never rob, beat, or kill one another (at least not in their capacity as private citizens). At the same time, the moral rights of N's citizens are constantly being violated: at every moment, the government is violating the rights of its citizens by stripping them of their freedom.

The rights-violations that N's citizens undergo are obviously not due to *suspensions* of the protections they receive. On the contrary, those violations are inseparable from those protections: it is precisely *because* people in N have so little freedom that their right not to be killed, beaten, and robbed are protected so well. Some of the rights of those people (e.g. the right not to be robbed by a fellow citizen) are protected exceptionally well, and the reason for this is specifically that other rights (e.g. the right to choose what to do with one's life) are so badly violated. So the rights-violations that N's citizens experience have nothing to do with a suspension of protections.

There is no paradox here. A rights-violation is no more an *absence* of rights-protection than a pain is an *absence* of a pleasure. Pains and rights-violations are positive realities (in the ontological, not the moral, sense of "positive"); and there is thus no barrier to rights-violations co-existing with rights-protections, even if one and the same entity is responsible for both – just as there is no paradox in saying that so and so brings me great sorrow and also great joy. (So and so might be a spouse who is attractive and stimulating, but also manipulative and deceitful.)

Suppose that one day, for no good reason, officer Brown beats me. What

we have here is not a situation where the protections that I usually receive from the government were *suspended.* Rather, what we have is a situation where I continue to receive certain protections but where, in addition, I also receive positive mistreatment. The fact that these protections continue to exist is expressed in the fact that, when people other than Brown try to beat me or steal from me, the police arrest them; in the fact that, if my insurance doesn't pay on some claim which it is contractually obliged to pay, I have recourse to the judicial system; in the fact that, if (for some reason) I am walking in the middle of the street, drivers swerve so as not to hit me; and so forth. It would be extremely unnatural to say that, while Brown was beating me, the protections of rights that I typically have were *suspended* and that, when Brown stopped his malfeasance, those rights were suddenly restored.

In fact, the right counterfactuals make it clear that such a view is simply false. While Brown is beating me, it is still the case that *if*, at that moment, somebody burgled my house, that person would be arrested (if seen by the police). So I continue to receive the protections that I had up until the time of the beating.

We must therefore regard that beating, not as a suspension of governmental protections, but as a positive evil that is *added* to those protections. Given the distinction between a positive evil, on the one hand, and a suspension of protections, on the other, it is clear why our analysis of law ("no legal obligation without governmental protection") is consistent with the fact people often continue to have legal obligations under governments even after being horribly violated by them.

It may be true (though for reasons we will give in a moment, it is by no means obvious) that if such rights-violations are sufficiently pervasive, *then* legal obligation may be suspended. Suppose that government G not only beats Smith, but doesn't arrest people who steal his belongings, or who don't

honor contracts with him. Depending on how we develop this scenario, Smith may indeed stop having legal obligations under G. But that is because the government has failed to give to Smith the protections that, according to our theory, are a pre-requisite to his having any legal obligations. Owing to its pervasiveness, this rights-violation collapses into a rights-suspension. Whenever there is a case that abuses are *so* pervasive that legal obligation is suspended, it is because there has been a suspension of protections.

§ But, I will now argue, it is much harder than one might initially think for a government to *suspend* the rights of some specific individual. An obvious consequence of this is that governmental violations of the rights of its own citizens tend to consist in the perpetration of evils *additional to*, and not *in lieu of*, those protections.

We must begin by remembering the fact that the protections a given citizen enjoys are not issued through assurances that concern *that citizen specifically*. (See pp. 174-175.) Even though Donald Rumsfeld is protected by the law, there is no law of the form: *you may not steal from Donald Rumsfeld.* The law has the form: given *any* entity x meeting certain requirements of geography, biology…you cannot steal from x (without incurring such and such penalties).

So laws are not given *directly* to the individuals who benefit from them. To borrow the terminology of contemporary semantics, laws are not *de re* with respect to their beneficiaries. Those benefits are conferred not by targeting *specific* individuals – e.g. Donald Rumsfeld and Allan Greenspan – for protection, but by creating *general* conditions of a certain kind, i.e. by creating civil conditions that program for a general respect for (certain) rights. The government doesn't protect (say) Allan Greenspan through a series of repeated *ad hoc* interventions specifically on Allan Greenspan's

behalf, but by creating conditions that make such intervention unnecessary and that apply to whole classes of individuals.

Because the governmental protections associated with law are so indirect, it would be rather hard for the government to *suspend* them. The idea that they could be readily *suspended* presupposes that they are doled out on a person-by-person basis and can thus be withdrawn on such a basis. That idea, in its turn, assumes that such protections are like the allowances one gives to one's children. For little Timmy to stop receiving his allowance, his parents need only stop giving it to him. Similarly, so the idea under criticism goes, for Allan Greenspan's rights to be suspended, the government just has to *not* give them to him anymore.

But the governmental protections that underlie legal obligation are not given out one person at a time, and therefore cannot be taken away one person at a time. This is not to say that such protections cannot be taken away. But because those protections are not given to any individual *directly*, and are provided through the creations of general conditions from which the individual benefits, a government would have to take special measures to suspend those rights. To suspend Smith's rights, the government in question would have to make sure that shop-owners stopped honoring his currency, that drivers stopped swerving to avoid running over him, that emergency room doctors stopped treating him, and so on. So suspending his rights would actually require extraordinary and quite *ad hoc* measures on the government's part.

This suggests that, if the government violates Smith's rights (by, for example, beating him), it is probably not by *suspending* the protections that it gives to him, but rather by superimposing positive abuses on those protections.

Historically, there have been cases where existing legal rights were

actually suspended. In Nazi Germany, Jews actually *stopped* enjoying governmental protection. But it isn't as though the Nazi government withdrew those protections on a person-by-person basis. It did so through general measures. It did so by making *systemic* changes. This is obviously consistent with the idea that laws provide general protections – protections that involve the creation of general conditions favoring the honoring of people's rights -- as opposed to protections given out on a person by person basis. It also confirms our point that suspending the protections that ground legal obligation would involve *systemic* changes. This, in its turn, confirms our point that, when an individual's rights are violated, it probably isn't through a *suspension* of the protections that ground his legal obligations towards the government in question, but is rather through perpetration of positive evils that co-exist with those protections.

There is thus no reason to believe that our principle *no legal obligation without governmental protection* is threatened by the fact that a person often continues to have legal obligations towards a government after that government has horribly violated her rights.[234]

Chapter 12 The concept of legal interpretation

Earlier we put forth some views on law and legal interpretation that might seem quite nihilistic. If the judge rules one way, that ruling is the law; if he rules the other way, *that* ruling is the law.[235] Of course, a later judge may overturn a ruling on the part of a previous judge. But until that happens, the first ruling stands. As we discussed, the fact that a later judicial ruling is

needed to neutralize the first only confirms the view that judicial rulings constitute, and therefore do not interpret, the law.

Without *rejecting* any of these points, I now wish to show that they don't have the nihilistic consequences that they might seem to have and that they don't warrant anything like the now widely rejected "legal realism" advocated at the beginning of the last century. (Legal realism is the view that judicial interpretation isn't answerable to *any* objective-standards and that it is therefore *mere* legislation. The law is whatever the judge's indigestion makes him feel like *saying* that it is.)

Here, more exactly, is what I would like to show. Even though judicial rulings create, and do not identify, the law, some rulings are still better than others. I am not making the trivial and obvious point that some rulings are *morally* better than others or do more social good than others. I mean that some judicial decisions are truer to existing law than others.

In fact, judicial decision x can be *more* evil, and can do much *more* social harm, than decision y, even though x is truer to existing law than y. I will further argue that this last fact is consistent with the fact that judicial decisions define, and do not interpret, the law and also with the fact that laws are assurances of moral protections. The next few pages will show why we should give no credence to any appearance of paradox that might be generated by the preceding claims.

Why legal interpretation is not mere legislation

What is legal interpretation? There is a tendency to give one of two answers to this question. First, there is the idealistic answer (held by Blackstone): to interpret a law is simply to identify what the law already is; judges do not create, but merely discover, law. Then there is the cynical

answer (held by the legal realists): since whatever the judge says is *ipso facto* the law, it follows that so-called legal interpretation is mere legislation and thus isn't interpretation at all.

We've seen why the idealistic view must be rejected. We've seen that the judge's ruling is *ipso facto* what the law demands and that, consequently, legal interpretation doesn't consist in *identifying* existing law. So the legal realist is right to this extent.

But it would be unwise to embrace the view that judicial interpretation is *mere* legislation and isn't answerable to objective standards. It seems a plain fact that some judicial rulings embody a correct understanding of the verbiage found in books, while other rulings embody spurious understandings of that same verbiage. This is not to say that *all* judicial rulings are either completely right or completely wrong. But it is a fact of every trial lawyer's life that some judges are more likely than others to rule in a manner consistent with the material found in the relevant books.

So there is reason to believe *both* that Blackstone's idealism is false *and* that the same is true of the nihilistic approach taken by the (so-called) legal realists.[236] I would now like to propose a third view.

§ Let N be a nation whose population is composed of two ethnic groups, the X's and the Y's. The X's dominate the political activity in N. In that nation, there is an evil law L that allows only X's to vote. Many of the Y's have X-blood (inevitably, a certain amount of inter-breeding has occurred). At the same time, for various cultural reasons, X's are very dismissive of those who are not *entirely* of X-ancestry. This leads those who are part X and part Y to identify with the Y's. In other words, it leads those who, in terms of genetics, are part X and part Y to be 100% Y in terms of psychology and cultural affiliations.

Of course, L exists because the X's wish to retain their hegemony. L can thus be seen as an assurance to the X's that their right to retain power, and their consequent expectation of a certain quality of life, will be protected. (L thus conforms to our analysis of law.)

One day, a person named Smith who is genetically half X and half Y attempts to vote. Culturally and psychologically, Smith is a Y, not an X. The wording of L leaves it entirely unclear whether Smith ought to be allowed to vote, since L was written at a time when X's and Y's had not yet mated and, therefore, when the line between the two groups was sharply drawn. (L speaks without qualification of "X's" and "Y's", and makes no allowance for those who are of mixed heritage.) The decision goes to Judge Jones. Being a moral and forward thinking person, Jones says that Smith can vote. His reasoning is that it is that L is ambiguous on this particular issue, and that whenever we are dealing with legal ambiguity, we must opt for the morally wholesome disambiguation.

But the official government attorney appeals this decision, and the case goes to Judge Brown, who sits on a higher court. Brown rules that Smith cannot vote. Brown's reasoning is this:

The letter of L leaves it open whether Smith can vote or not. So in a narrow, technical sense, L is open on the question of whether Smith may vote or note. But the purpose of L – the reason it was brought into existence -- is to ensure that X's, and X's only, may vote. Genetically, Smith is only half Y. But *culturally* he is 100% Y. For this reason, permitting him to vote would be contrary to the spirit of L. So denying Smith the right to vote is, whereas giving him the right to vote is not, consistent with the spirit (if not the letter) of L.

§ Let us evaluate this situation. It is obvious that Jones' decision is *morally* better than Brown's. But which of these two decisions is better *as a legal interpretation*? Answer: Brown's. But why?

Sadly, Brown's reasoning is on the mark. L is an assurance of a certain kind. It assures the X's that some of their freedoms will not be abridged. Since L was first issued, society has changed in a relevant respect – the line between the X's and the Y's is now genetically blurry. Given those changes, Brown's decision is a natural adaptation of *that* particular assurance to contemporary society. By contrast, Jones' decision is not a natural adaptation or extension of that original assurance.

Let me explain what I mean by this last point. Suppose I ask you to buy me a pack of cigarettes. You go to the store. They are out of cigarettes, but they do have cigars. (Let us assume that there is no other place to purchase tobacco-products of any kind.) So you buy me cigars.

Technically, you have failed to comply with my request. From the standpoint of formal logic and literal meaning, you no more succeeded in complying with my request than you would have if you had brought me a candy-bar. But you were right from a pragmatic or psychological viewpoint to bring me cigars; and, from that same viewpoint, you would have been wrong to bring me a candy-bar. People who want cigarettes typically have a violent craving for nicotine. For this reason, given a choice between having *no* tobacco, on the one hand, and having a cigar, on the other, a smoker will not hesitate to choose the latter. So when we consider the psychological and social realities in which my request is embedded, it is clear that your decision to bring me cigars was the right one, even though, from a purely logical or semantic standpoint, you no more succeeded in complying with my request than if you had brought me a soft-drink.

In general, human speech-acts – assertions, promises, threats,

assurances, blandishments – must not be understood in narrowly logical or semantic terms, but rather in broader psychological and social terms.

A corollary is that, in so far as such speech-acts are vague or incomplete, they must be completed in a way that befits not only their literal meanings or their narrowly logical properties, but also their psychological and social underpinnings.

Since laws are issued by, and embodied in, speech-acts (usually, but not necessarily, identical with the depositing of verbiage in certain books), everything we just said applies to them. This is what I meant when, a few paragraphs ago, I said that "in this context, consistency is to be understood in a cultural or psychological, as opposed to narrowly logical, sense."

Let us return to the Jones-Brown scenario. As we noted, Jones' decision is morally wholesome, whereas Brown's is evil. But this rather confirms our point that it is Brown's decision that is the *legally* correct one. After all, the original assurance was evil. It is therefore to be expected that the same would be true of a correct (so-called) interpretation of it – one that is true to its spirit.[237]

None of this is discrepant with our analysis of law. L is a governmental assurance of protection of a moral good. It assures the X's that their freedom will not be abridged in a certain way. As we've seen, any protection of freedom is *a* good, albeit one that may be incompatible with a far greater good.

We may conclude that legal interpretation is a principled endeavor, and is answerable to objective facts – facts about the history and culture of the relevant political organization. The points made earlier (pp. 26-42) do not warrant any kind of nihilism about legal interpretation.

§ But there is one other point to make in connection with the Jones-Brown

scenario. We saw that Brown's decision is *legally* better than Jones'. But let us now suppose that Jones' decision had prevailed. In other words, let us suppose that the government didn't appeal Jones' ruling. Would that have created some kind of *contradiction* within N's legal system? More precisely, would it have resulted in a situation where that legal system comprised two laws x and y such that x permitted what y forbad?

No. We'd have a system that didn't give the right to vote to those who were *entirely* of Y-ancestry, but did give that right to those who were partially of Y-ancestry. There is obviously nothing contradictory about such system.

Of course, as we saw, Jones' ruling is inconsistent with the spirit of the law that he is interpreting. But that ruling didn't result in there being some one mode of conduct that is both legal and illegal; it didn't result in there being two laws such that the one permitted what the other forbad. Given only that Jones ruled as he did, there is no reason why it isn't logically possible to comply with all of the laws in the system in question. So, *pace* Fuller, contradictions are not, or at least not necessarily, created by judicial rulings that embody false understandings of the relevant law.

Legal interpretation as precisification

But this last point raises an important question. How can a decision that embodies an erroneous legal interpretation *not* result in a contradictory legal system? Remember what we said in connection with *Riggs.* If the judge rules that the young man is not entitled to the money, it is illegal for the bank-manager to open the safety-deposit box. If the judge rules that the young man *is* entitled to the money, then it is illegal for the bank-manager *not* to open the safety-deposit box. The judge's ruling defines legality.

Judicial rulings create legal obligations. So, in effect, such rulings create laws. Given a law L, an erroneous judicial interpretation of L is presumably one that forbids what L permits or permits what L forbids. So an erroneous judicial decision is one that creates a law L* that forbids what L permits or that permits what L forbids. In the last section, we said that an erroneous legal interpretation does *not* create such an inconsistency. Something must give way here.

There is yet another problem. As we just noted, the law is whatever the judge says it is. How is this to be reconciled with the point, defended a moment ago, that legal interpretations have an objective basis? As we ourselves argued earlier (pp. 26-42), the judge cannot be interpreting the law if he is creating it. So, in this work, I have evidently asserted both that the judge is, and is not, interpreting the law.

I would now like to argue that, in actuality, I have not contradicted myself. I would also like to explain why, *even though* legal interpretation is not mere legislation – even though some interpretations are truer to existing law than others – an erroneous legal interpretation does not result in a contradiction within the relevant system. The essence of the matter lies in the following simple formula:

(LI) Legal interpretation is delineation.

In this context, I am using the word “delineation” in the technical sense in which logicians use it. I will now explain what that sense is, and also what (LI) is supposed to mean.

Legal interpretation as delineation

If you say “Harry is bald”, that is consistent with Harry’s having one hair, or two hairs, or three hairs….but not with his having a million hairs. But for many numbers n between (say) a hundred and ten-thousand, it is unclear whether *Harry has n hairs* is consistent with “Harry is bald.” (The asterices are meant to be quasi-quotes.) There are many values of n such that there doesn’t seem to be any fact as to whether a person with n hairs can truly be described as “bald.”[238] So, it would seem, the predicate “bald” doesn’t have precise truth-conditions.

A “delineation” (or “precisification”) of the sentence “Harry is bald” would be an assignment of *precise* truth-conditions to it. So a delineation of that sentence would be given by a proposition like: *“Harry is bald” is true iff Harry has exactly 1,200 hairs or less.*[239] A logician who assigns precise truth-conditions to “Harry is bald” is not exactly saying what that expression currently means. Rather, he is *giving* that expression a precise meaning that complements, but doesn’t replicate, the meaning that it already has. The logician is “precisifying” or “delineating” the meaning of that sentence. Being an act of meaning-stipulation, rather than meaning-identification, a delineation cannot, strictly speaking, be correct or incorrect.

At the same time, there is *a* clear sense in which delineations are answerable to objective standards and in which, consequently, one delineation can be objectively better than another. A bit of fiction can make this clear. According to logician A’s delineation of the term “bald”, a person is “bald” iff he has less than one hundred hairs. According to logician B’s delineation of that same term, a person is “bald” iff he has less than a million hairs. Even though neither delineation corresponds *precisely* to existing usage of the term “bald”, it is obvious that A’s delineation is (as we might put it) *more* correct than B’s or, in any case, has a better objective basis than B’s. In general, even though delineations are not true or false, there is a clear sense in

which they are answerable to objective facts and are thus not *purely* stipulative (or legislative, as we might also put it).

§ To take the next step in our argument, we must register a few pedestrian facts about linguistic meaning. They are best introduced through a short story.

You are afraid that the neighborhood bully is going to beat you up. Because I am a better pugilist than the bully, I am able to protect you from him. Hoping to convey this to you, I say "I assure you that I will protect you from the neighborhood bully: you have nothing to fear." As it happens, you don't speak English, and thus have no idea what I am saying.

Have I assured you of anything? No. I have *attempted*, but failed, to give you an assurance. So far as an "assurance" isn't understood, it isn't an assurance at all. Of course, my words have a definite meaning. But because they don't succeed in assuring you of anything, no assurance has been given. It is irrelevant that you *would* have been given an assurance *if* you had had the relevant information.

In general, if A is an attempted assurance, then to the extent that it is unclear what it is that A assures, or who it is that is on the receiving end of that assurance, A is *not* an assurance. This is a consequence of the fact that "assure" (like "know") is a "success"-verb. If you are not *in fact* assured by my words, then I have given you no assurance. It is irrelevant whether it is my fault or yours that my attempt to assure you failed.

It should be pointed out that "threaten" is also a success-verb. I say to you "if I see you in this neighborhood again, I'm going to break your legs." Unbeknownst to me, you don't speak a word of English. Under these circumstances, I have *tried* to threaten you. But, it seems to me, I haven't succeeded in doing so.

Given any speech-act – whether it is an assertion, a promise, a threat…-- so far as it is unclear what is meant by it, that speech-act is a failure. It could be that the speech-act in question is entirely free from vagueness, and that, so far as it is not understood, that is entirely due to incompetence on the auditor's part. Nonetheless, a speech-act is a failure so far as it isn't understood, it being irrelevant why it isn't understood.

Laws are assurances. Those assurances are typically (though not necessarily) made by depositing certain ink-marks in certain books. As we've seen, to the extent that an "assurance" is unclear, there has been no assurance.

§ Some science-fiction will help us continue with our argument. In the year 2206, the U.S. legal system is not significantly different from how it is now. (Obviously there have been *some* changes. But the basic framework has remained intact.) In that year, there come into existence robots that are psychologically very similar to humans. Their intelligence-level is the same as ours, and they have more or less the same emotions and values. These robots come into existence very suddenly, and there is a period of a few years during which their status under U.S. law is unresolved.

During that interim period, robot Gloxo is arrested one day for some petty crime. Of course, any human being in Gloxo's circumstances would have various legal protections. He'd have the right to remain silent, the right to an attorney, and so forth. The legal system assures every human being of such protections. This is not to mention that, under that system, any human being would have the right not to be beaten or killed.

But does the U.S. legal system give *Gloxo* any such assurance? Have robots been guaranteed these protections? No. As we saw, whenever there is any *question* as to whether somebody has been given a certain assurance, the answer is *ipso facto* "no." So given that it even needs to be asked whether

Gloxo enjoys legal protection, it *ipso facto* follows that he does not.

This is not hard to corroborate. By hypothesis, there is no verbiage concerning robots, even those with human emotions, in the relevant law-books. One will say that laws concerning robots are *implicit* in such verbiage and that, consequently, there *are* existing laws concerning robots, albeit hidden ones. But we've already seen why that view is a non-starter. If a judge rules that robots do not enjoy protections, then robots *ipso facto* don't enjoy them; and if a judge rules that they do, then they *ipso facto* do enjoy them. In neither case will be there a *contradiction* within U.S. law, meaning that, prior to the judge's ruling, there were no laws concerning robots. We will develop this point in a moment.

Of course, Gloxo is *morally* entitled to such protections. Further, there is *a* sense – to be explicated in a moment -- in which a judge who ruled that Gloxo was *legally* entitled to such protections would be ruling correctly. Such a ruling would probably be more in keeping with the spirit of U.S. law than a contrary a ruling.

But it is a fact that, as of yet, Gloxo has no such assurance. It is easy to demonstrate this by developing a point made a moment ago. Suppose that the police beat Gloxo to death. Thanks to pressure from the ACLU, those police officers are brought to trial. The case goes to Judge Green. Green decides that, legally, robots aren't human beings and that the police therefore didn't commit homicide.

It is pretty clear that Green's decision is evil and, further, that it is inconsistent with the moral outlook embodied in the U.S. legal system as a whole. (Remember that Gloxo has human emotions, even though he is biologically non-human.) But it is also clear that Green's decision hasn't created a legal *contradiction*. His decision did not result in a situation where some law prohibited what some other law permitted If Gloxo had *already*

had some legal assurance that he had such rights, then Green's decision *would* have resulted in a situation where Gloxo both did, and did not, have certain rights. Since Green's decision did not have that result, it follows that Gloxo never had those assurances to begin with. For exactly similar reasons, if Green had ruled that, legally, Gloxo *was* a human being, *that* ruling wouldn't have created any contradictions within U.S. law.

Let us continue with our story. Once again, suppose that Green rules that, legally, robots are not people and thus don't enjoy legal protections. (So the police who killed him are found not guilty of committing any crimes against a *person*.) As we've just seen, Green was not *identifying* what the law *already* said in connection with Gloxo or robots generally. So what was he doing, or at least attempting to do?

Prior to Green's ruling, human beings did have certain assurances under U.S. law. They were assured that they have the right to remain silent, the right to an attorney – not to mention the right not to be beaten or killed. These assurances are laws. (They are, respectively, the law that the police may not interrogate you without mirandizing you, the law that others may not beat you, and so on.) Let $L_1 \ldots L_n$ be these laws. There is no doubt about what *these* laws are. The question that Green must answer is not "what do $L_1 \ldots L_n$ say about robots?" The answer to that question, as we've seen, is "nothing." (That is why a judge was needed in the first place.) The question that Green must in fact deal with is (at least approximately) this:

Green's question (GQ): **Let $L^*_1 \ldots L^*_n$ be a set of laws that are just like $L_1 \ldots L_n$ *except* that, unlike $L_1 \ldots L_n$, $L^*_1 \ldots L^*_n$ give *robots* the assurances that $L_1 \ldots L_n$ give to humans. And let $L\hat{}_1 \ldots L\hat{}_n$ be a set of laws that are just like $L_1 \ldots L_n$ *except* that, unlike $L_1 \ldots L_n$, $L\hat{}_1 \ldots L\hat{}_n$ *deny* robots the**

assurances that $L_1…L_n$ give to humans. Should I replace $L_1…L_n$ with $L^*_1…L^*_n$ or should I replace $L_1…L_n$ with $L^\wedge_1…L^\wedge_n$?

Green's "interpretation" of the law – his "interpretation" of $L_1…L_n$ – doesn't consist in his *identifying* those laws, but rather in his *extending* them in one way or another. Green is not *identifying* what the law already says in connection with robots. Green is replacing existing U.S. law with a body of law that, while not conflicting with hitherto existing U.S. law at any point, is not characterized by the same incompleteness as that law. Green is producing a *delineation* (or *precisification*) of U.S. law.

§ I would like to give further corroboration for the view just presented. In any literal and non-extended sense of the word "meaning", every English-speaker knows the meaning of expressions like: "you have the right to an attorney", "you have the right to remain silent", and so on. In fact, the laws that lie at the center of the most bitter legal controversies are typically very easy to understand ("abortion is legal", "people of the same gender cannot marry one another", "schools cannot teach creationism instead of evolutionary theory").

Of course, *some* laws are expressed in language that is hard to understand. In such cases, an expert may be needed to know what is literally meant by some statute, and *bona fide* interpretation would be necessary. But *legal* controversies generally don't involve difficulties relating to *that* sort of interpretation. The language of the Constitution, even the most controversial parts of it, isn't technical or otherwise hard to understand. [240]

§ Of course, some may insist that, in the scenario described a moment ago,

Green *is* interpreting the old laws – that it is a paradigm-case of what we call "interpretation."

Very well. In that case, our position amounts to this:

(LI*) Legal interpretation is delineation. To interpret the laws is *not* to identify existing law. To interpret the law is to replace law that is incomplete in some respect with law that is not thus incomplete but otherwise coincides with the old law.

Given that *this* is what it is to interpret the law, it is clear why, no matter how Green rules, his ruling won't create a legal contradiction, i.e. a pair of laws such that one permits what the other forbids. In general, it is clear why a judge can "interpret" a law in either of two opposed ways without creating any kind of contradiction within the system in question.[241]

At the same time, our analysis makes it clear why this last fact is consistent with the presumption that some legal interpretations are better than others. Recall our cigarette-story. There was no way for you to comply with my request to buy cigarettes. Under the circumstances, the best that could be done was to comply with a request that was similar to it, namely: *bring me cigarettes or, failing that, some other nicotine delivery-system.*

The position that Judge Green is in with respect to existing law is similar to the situation that you are in with respect to my request that you buy me cigarettes. Green cannot *apply* existing law to the Gloxo-case. There is no law to apply, since all the existing laws are silent with respect to robots. The best that can be done under the circumstances is to apply new laws that are not comparably silent but that are otherwise identical with the old.

In *some* sense of the word "consistent", some new laws satisfying that condition are going to be consistent with the old laws, and other new laws

satisfying that same condition will not be. But here the word "consistent" doesn't denote logical consistency. Since the old laws are silent regarding Gloxo, *any* laws that are not silent regarding Gloxo, but are otherwise identical with the old laws, will be *logically* or formally consistent with the latter. But some such laws will be inconsistent with the old laws in an extended psychological sense – in the sense in which, in the cigarette-story, your bringing me cigars is consistent with my request and your bringing me a candy-bar is not. Given what $L_1 \ldots L_n$ are, and given the attitudes of the people who made those laws, along with the mores of the culture in which they are embedded, Green would be guilty of *a* kind of inconsistency if he ruled that Gloxo had no rights. But in this context, the term "inconsistency" is not to be understood in narrowly logical or semantic terms, but in terms of broader psychological or cultural norms. Of course, that sort of psychological or cultural inconsistency isn't sufficient for the existence of any sort of *legal* contradiction, since it doesn't result in there being two laws such that one forbids what the other permits.

Given these points, it is clear why Fuller is wrong to say that judicial mis-interpretations (so-called) of the law categorically create *contradictions* within the law. The law needs to be "interpreted" in situations where it is unclear who has been assured of what. But, as we've seen, for it to be "unclear" whether the law assures Gloxo of the right to counsel, and so on, is simply for Gloxo *not* to have such an assurance. It isn't for Gloxo to be assured that he *doesn't* have that right. In other words, it isn't for Gloxo to have some kind of negative assurance. It is for Gloxo not to be given any assurance, whether positive or negative. So legal interpretation is needed precisely when the system in question *fails* to issue an assurance and when it therefore fails to issue a law.

Of course, there *are* plenty of laws prior to the Gloxo-case, and these

play a crucial role in the resolution of that case. Green's decision isn't made out of whole cloth. It consists in issuing a set of laws that is, in some respect, more complete than, *but otherwise coincident with,* those laws already in existence. The italicized clause shows to what a large extent his decision, though not identification of existing law, is constrained by objective standards and is thus not an expression of mere caprice on his part.[242] [243]

Why the Constitution is a part of U.S. law (despite what we said on pp. 26-42)

As we've seen, it is simply not an option to say that legality can ever be distinguished from what is permitted by law or that illegality can ever be distinguished from what is prohibited by the law.

Given this fact, an argument can be made that the Constitution is *not* actually a part of U.S. law. Earlier (pp. 26-42) we put forth this argument, and we concluded that, indeed, the Constitution is *not* a part of U.S. law.

What I now want to show is that, in light of our points regarding legal interpretation, we *can* hold onto the presumption that the Constitution is a part of U.S. law while also holding onto the obvious truth what is legal must coincide with what the law permits. (As we've seen, that obvious truth is simply inconsistent with the Blackstone-Fuller view that legal interpretations *identify* existing law.)

Right now let us recap what we said earlier concerning the place of the Constitution in U.S. law:

The law is fixed the moment it is fixed what is legal and illegal. The

judge's word defines legality and illegality. (Legally, the money is Smith's if the judge says it is, and it isn't if the judge says it isn't.) So far as the judge's word is discrepant with the Constitution, the latter is inert in terms of defining legal reality, and is thus no part of the law. By the same token, so far as the Constitution is *not* thus inert, it is thanks only to some judge's (or legislator's) words. So in the final analysis, the Constitution is a part of U.S. law only in so far as it *influences* the decisions of judges and legislators. This means that the relation between the Constitution and the law is one of *causality* or *influence*, and not one of identity or part-hood. So the Constitution is a cultural or social paradigm that has a great deal of influence on the law, while not actually being a part of it.

The conclusion of this argument – the Constitution isn't a part of U.S. law – is obviously a dubious one. In connection with this, I would now like to argue for two points. (1) That conclusion isn't *necessarily* wrong, since there is independent corroboration for it. (2) Given the points just made concerning legal interpretation, we can accept all of the premises of the argument given in the last paragraph *while also* holding onto the presumption that the Constitution *is* a part of U.S. law. I believe that (2) is much more important than (1), and will thus provide a commensurately more developed argument for it.

My argument for (1)

Any legal system is embedded in a larger culture. Given this fact, there will inevitably be institutions that mediate between the law and the host-culture, while not themselves being veritable components of the legal system.

Where such mediating institutions are concerned, it will sometimes be ambiguous whether they are veritable components of the legal system or are merely influences on it. The Constitution could be seen in such terms. This would accommodate the *apparent* fact that what the Constitution demands can be so discrepant with what is legal and illegal, and thus with what the law permits and forbids. (I use the word "apparent" because, if the argument to be given in the next section is cogent, those discrepancies are not actual.)

Some corroboration for this view is found in the following line of thought. In the U.S., legal reality (what people can and cannot legally do) typically has a certain consonance with the Constitution. Of course, legal reality is not *fixed* by the Constitution. (This is a point that we've stressed.) But, at least arguably, it is typically not hard to establish a reasonably solid connection between the one and the other. This gives credibility to the view that Constitution is a veritable constituent of American law.

But there are other countries that have constitutions that don't have such a good fit with legal reality. When we look at the constitutions of other nations – e.g. the constitution of the U.S.S.R. or Nigeria or Pakistan[244] – we find that they don't have remotely the kind of role in the legal systems of those countries the U.S. Constitution has in our legal system. In fact, the U.S. seems to be more of the exception than the rule in this respect: in most countries that have a constitution, the constitution has *at most* a certain influence on legal reality, and obviously isn't constitutive of it.

Further, there are counterfactual scenarios in which, in the U.S., the Constitution comes to diverge more and more from legal reality, even though the basic political structure of this country hasn't changed. We could imagine a series of politically extreme Supreme Court decisions resulting in a progressive alienation of the Constitution from legal reality.

So while some of us may perceive it as a truism or tautology that the

Constitution is a part of U.S. law, that perception may only reflect idiosyncratic and contingent facts about American history. That "truism" is, at best, a contingent empirical truth.

My argument for (2)

Let us begin by reviewing a now familiar story, along with some now familiar philosophical consequences of that story.

In *Riggs*, the judge had to decide whether the young man was legally entitled to the money. The judge decided that he was not. For the sake of argument, let us suppose that the young man appeals the case, and that it eventually finds its way to the Supreme Court. So the Supreme Court must decide whether the first judge's decision was Constitutional, i.e. whether the young man was *Constitutionally* entitled to the money.

For reasons that we've seen, the Constitution doesn't say, even obscurely or implicitly, whether the young man is entitled to the money or not. (Incidentally, some might say that if we could look inside the minds of the framers of the Constitution, we would, or at least might, find that the Constitution *did* contain a definite position on cases like Riggs. In the last footnote of the present work, I discuss why this sort of psychological approach to the concept of legal interpretation is untenable.) So the question before the Supreme Court is not "what does the Constitution *actually* say about this case?" Rather, the question is, at least approximately, as follows.

***The Supreme Court's Question (SCQ)*: Let $C_1 \ldots C_n$ be the laws that are *actually* issued by the Constitution. Let $C^*_1 \ldots C^*_n$ be laws that are just like $C_1 \ldots C_n$ *except* that, unlike $C_1 \ldots C_n$, $C^*_1 \ldots C^*_n$ say that the young man**

is entitled to the money. And let $C^_1 \ldots C^_n$ be laws that are just like $C_1 \ldots C_n$ *except* that, unlike $C_1 \ldots C_n$, $C^_1 \ldots C^_n$ say that the young man is not entitled to the money. Should we replace $C_1 \ldots C_n$ with $C*_1 \ldots C*_n$ or with $C^_1 \ldots C^_n$?

Before we close the argument, let us take a brief look back. Our earlier argument involved the inference from:

"No matter how the judge interprets the Constitution, his word determines what is legal and what isn't, and thus what is law"

to:

"The Constitution is a part of U.S. law only in so far as the judge's word allows it to be – so it isn't really a part of U.S. law at all (it is merely something which *influences* U.S. law)."

But that inference involves a non-sequitur if, as we have argued, SCQ correctly describes what the Supreme Court's task of "interpreting" the Constitution consists in. If the Supreme Court's interpretations (so-called) of the Constitution are cases of filling in blanks in the laws actually present in the Constitution, while otherwise leaving those laws the same, then two things immediately follow. First, no matter *how* the Supreme Court rules, its ruling will be right: if, as SCQ suggests, Constitutional interpretation is *a* kind of legislation, then such interpretations cannot fail to constitute the law. At the same time, the Constitution, and the laws found therein, *are* a part of

U.S. law. Contrary to what we argued earlier, the reason that the Supreme Court's decision is *ipso facto* the one that defines U.S. law is not that the Constitution is *not* part of the law, but is rather that, by definition, the Court's interpretations (so-called) consist in filling in blanks in the law, and thus cannot come into conflict with existing law. The Court fills in the blanks in the laws *in* the Constitution: and *that* is why, even though the Constitution *is* a part of U.S. law, the Court's decisions do not (in any strict, logical sense) contradict the laws present in the Constitution. (The Court's decisions *may* contradict those laws in the extended sense discussed earlier – the sense in which, had you brought me a candy-bar instead of cigars, you would have been utterly failing to comply with my request.)

In conclusion, given our analysis of legal interpretation, we can hold onto the presumption that the Constitution is a part of U.S. law *without* doing so by embracing the absurd view, implicitly accepted by Fuller, that law and legal reality are distinct.

The logical foundations of our analysis of legal interpretation

I would now like to give reasons of a strictly logical kind in support of our view that legal interpretation is a species of delineation, i.e. the legal "interpretations" extend, and do not identify, existing realities.

There is a debate in philosophical logic concerning whether there is objective vagueness, i.e. whether there is any proposition P such that there really is no fact as to whether P is true or not. Here is one widely accepted, and seemingly very reasonable, point of view on the matter:

There are many values of x satisfying the following conditions. We

have all the empirical information there is to have regarding the state of x's coiffure, i.e. we know exactly how many hairs x has, how thick they are, how much hair-free skin there is on x's head, and so on. Further, we have no relevant linguistic deficits, i.e. we have as good an understanding as there is to have of the sentence *x is bald*. (The asterices are meant to be quasi-quotes.) It presumably follows that we know the identities of the concepts corresponding to the constituent expressions, in particular the word "bald." But, despite those facts, we feel that there is no non-arbitrary answer to the question *is x bald?*

There is a related fact. Obviously some people are bald and others are not. At the same time, there doesn't seem to be any number n such that a person with n hairs is bald whereas a person with n+1 hairs is not bald. We cannot accommodate this apparent datum by saying that, for some values of n, a person with n hairs is definitely neither bald nor non-bald. That position would introduce an artificially sharp barrier between being determinately bald and being determinately neither bald nor non-bald (and also between being determinately neither bald nor non-bald and being determinately *not* bald). So it seems to follow that, for some numbers n, a person with n hairs is neither definitely bald nor definitely non-bald nor definitely neither bald nor non-bald. In a word, the vagueness of the predicate "bald" can no more be dealt with by seeing it as generating *three* categories – bald, not bald, and neither bald nor not bald – than it can by seeing it as generating just two categories (bald and not bald). For it is just no less artificial or arbitrary to posit a sharp line between bald and neither-bald-nor-not-bald than it is to posit a sharp line between bald and not-bald.

Given this, it seems to follow that, for some values of n, there is no fact as to whether a person with n hairs is bald or not bald. It isn't that

such people are definitely in a no-man's-land between baldness and non-baldness – we've just seen that such a position involves positing unrealistically sharp barriers between those who have and lack the property of baldness. It is that there is no fact as to whether such people are bald: it is objectively indeterminate.

The argument just given involves two muddles. One of these is a conflation of sentences with propositions and, more generally, of expressions with meanings. The other muddle concerns the relationship of a proposition -- or more generally, a piece of information -- to its logical consequences. Let us begin by discussing this last point.

As Wittgenstein (1974, 1983) stressed, entailments correspond to *internal* relations. (As I am using the word "entailment" and "entails", P entails Q exactly if, given P, it is analytic or, at least, *a priori* that Q is the case. So *x=2* entails *x is greater than one*, whereas *x is a college professor in the U.S.* does not entail, though it may otherwise confirm, *x has an annual salary of over $10,000.*) For any x, the proposition *x is a triangle* entails *x has three sides.* This obviously reflects something about the structures of the concepts involved -- *triangle*, *three*, and so on. (In this context, the term "concept" is meant to denote a non-psychological entity – the meaning of a predicate.)

Entailment relations are not circumstance-dependent. This is because such relations are grounded entirely in facts about the structures of the concepts involved. The very structures of those concepts guarantee the truth of the transition in question. Entailment relations, and analytic relations generally, are those that hold in virtue of facts about content alone. So they are content-internal and are therefore *individuative* of content. From this it

follows – so I will argue -- that, so far as there is indeterminacy, there is a failure to *identify* content, rather than a failure on the part of content to have a definite truth-value. I hope to show that the arguments presented earlier (to the effect that *there is* objective indeterminacy) overlook the content-internal nature of entailment relations; I hope to show that those arguments implausibly require that what a proposition tells you can be distinguished from what you can infer from its truth.

Propositions are paradigm-cases of contents. Given what we just said, it follows that propositions are individuated, at least in part, by their entailment relations: different entailments, different propositions. Of course, distinct propositions may entail, and be entailed by, the same propositions. This is the case with any two analytically equivalent proportions, e.g. *1+1=2* and *triangles have three sides.* But there is no doubt that, for P and Q to be different propositions, it is *sufficient* that, for some proposition R, R is entailed by the one proposition but not the other.

More importantly, the idea that propositions are individuated by their entailment-relations *is* consistent with the fact that *distinct* propositions may be analytically equivalent. Let P* be the proposition *for some n, 1+n=2*. Each of *1+1=2* and *triangles have three sides* entails P*. But the *way* in which *1+1=2* entails P* is very different from the *way* in which *triangles have three sides* does so. In the one case, the series of inferences linking the two propositions is short and direct. (One needs only the principle of existential generalization to make the connection.) In the other case, the series of inferences is long and circuitous.

So the fact that *1+1=2* and *triangles have three sides* are different propositions *is* reflected in facts about their entailment relations. It isn't reflected in facts about *what* they entail, but it is reflected in facts about *how* they entail what they entail.

Given these remarks, it is easily seen that if, for any proposition p_1, each of p_2 and p_3 entails/is entailed by p_1 in the *same way*, then p_2 and p_3 would not only be analytically, but also *cognitively*, equivalent. In other words, they would be the same proposition. So even after we take into account the fact that analytically equivalent propositions can be distinct, there is still no denying that propositions are individuated by their entailment relations.

A corollary is that if P's identity is fixed, it *is* determinate, for any proposition P*, whether P entails/is entailed by P* or not. Thus, in so far as it is indeterminate (in the sense of there being no fact of the matter, not in the sense of being unknown) *whether* p entails q, it is indeterminate *which* proposition p is. (Permit me to clarify the parenthetical remark in the previous sentence. It can be indeterminate, in the sense of *unknown*, whether p entails q without there being any question as to p's identity. Until the 1930's, it was *unknown* whether the proposition *1+1=2* entailed *arithmetic is incomplete*. But there has never been any doubt as to the *identity* of the proposition expressed by "1+1=2.")

So when it seems as though we are dealing with a case where there is no fact as to whether one proposition entails another, what is really going on is that we haven't settled the identity of at least one of those propositions.

At the same time, it does seem absurd to say that, for some number n, "Smith is bald" is true if Smith has n hairs and not true if Smith has n+1 hairs.

§ Given these points, it is natural – and, I believe, necessary -- to take the following position. What we might take to be cases of objective indeterminacy are really cases where language *leaves it open* which exact proposition is meant by a given sentence. This is not to say that cases of non-linguistic indeterminacy turn out to be cases of linguistic indeterminacy, i.e. it

is not to say that there is sometimes no fact as to what a given sentence means. (That would itself be a form of objective indeterminacy. For it would mean that some metalinguistic proposition of the form *sentence S has proposition P for its meaning* would be such that there was no fact of the matter as to whether it was. This, in turn, would mean that there was no fact as to whether S did, or did not, mean P. But this, it is easily seen, would be a straightforward case of objective indeterminacy.) It is to say that, in a case where we believe there to be no fact as to whether some proposition is true, language has determinately *failed* to assign a unique proposition to a given sentence.

Suppose that English semantics, taken in conjunction with a complete knowledge of the state of the top of Harry's head, leaves it open whether the sentence "Harry is bald" is true or not. So we know exactly how many hairs Harry has, and we also know exactly how thick they are, how large Harry's head is, and so on. (So we know precisely the ratio of hairy to non-hairy skin.) Let P be a proposition in which all of this information is embodied. So P is some proposition like *Harry has 3,876 hairs on his head, and the average thickness of these hairs is…and (consequently) the ratio of hair-covered skin on Harry's head to hair-free skin is such and such…*So even though we have all of this information – even though we know the truth of P -- the rules of English semantics simply don't yield a verdict as to whether the sentence "Harry is bald" is true or not.

Here is the view that is generally, though not universally[245], held in connection with this sort of scenario. (It is the view I reject.) We say that there is a proposition – call it Q – such that Q is what is meant by "Harry is bald", but such that P does not entail the Q and P also doesn't entail not Q. Remember that P is a proposition giving the exact physical make-up of Harry's head – a proposition saying exactly how many hairs he has, how

thick they are, and so on. P gives all the relevant information, and there is thus no relevant empirical fact that is not covered by P. Short of including the proposition *Harry is bald* or *Harry is not bald* (or some trivial variant of one of these, e.g. *either Harry is bald or there are square circles* or *either Harry is not bald or there are non-flying birds that fly*) , P gives you every piece of empirical information that could conceivably be relevant.[246]

We can now close the argument. Suppose we say that "Harry is bald" *does* encode a determinate proposition Q, but that, because Harry is a borderline case, there is no fact as to whether Harry is bald. (It is objectively "indeterminate.") In that case, given what we've said in the last few paragraphs, we must say there is no fact as to whether P entails Q, or as to whether P does *not* entail Q. But this conflicts with the fact that propositions are individuated by their entailment relations. What a proposition entails, along with how it entails it, is completely determinative of the identity of a proposition. So far as it is open what is entailed by a sentence S, it is open what is said by S. So far as it is open what is said by S, it is open which proposition it expresses.

Of course, in some cases there is a determinate proposition whose consequences we may have trouble identifying. There is no doubt as to the identity of the proposition meant by the sentence "there are infinitely many primes", even though that proposition has infinitely many consequences that will never be known. But it is one thing to say that we don't *know* whether a given proposition entails some other proposition, and it is quite another to say that there is no fact of the matter. If we say that there is no *fact* as to whether proposition P entails proposition Q, then we are saying that a proposition's identity can be determinate even though facts about what it entails are not determinate. But such a position is incoherent, given that a proposition just *is*

a piece of information and that pieces of information are individuated by their entailment-relations.

§ Here, then, is what we must say. Under the circumstances described, i.e. given that P is true, English semantics doesn't warrant the truth of "Harry is bald" or of "Harry is not bald." It isn't that there is no fact as to what the rules of English semantics demand. There is a fact. They don't demand that "Harry is bald" is true, and they don't demand that "Harry is bald" is false. In this context, they don't demand anything.

There is nothing mysterious about this, and it doesn't require us to reject the law of excluded middle. If I say "I am going to the store", my words leave it open whether I am going by car or by bicycle. They determinately *don't* entail that I will go by car and they determinately *don't* entail that I will go by bicycle. They also determinately *don't* entail the negations of either of those propositions. We are dealing with garden-variety compatibility-relations, not with counter-examples to classical logic.

In connection with this, we mustn't say that the *concept* of baldness is such that, where some people are concerned, there is no fact as to whether it applies to them. (Here we must be careful to distinguish that concept from the words used to express it, e.g. "bald.") For the sake of argument, suppose that the concept of baldness is genuinely indeterminate, i.e. that there are, at least potentially, some people such that there is no fact as to whether or not they fall in its extension. In that case, it follows straightaway that, for some possible object A, the proposition *A is bald* is such that there is no fact as to whether it is true. Thus, the idea that concepts may be indeterminate is equivalent with, or at least entails, the idea that propositions may be indeterminate. Since propositions cannot be indeterminate, as we've just seen, it follows that the same is true of concepts. So for any concept C, and

any object x, the proposition *x falls in the extension of C* is either true or false. There is a fact of the matter.

This is not to say that the predicate “bald” is without meaning. It is replete with meaning. For any x, if you know that *x is bald* is a true sentence, you know of infinitely many possible scenarios that they are non-actual. (Among those scenarios is one where Harry has a million hairs, one where he has a million and one hairs, one where he has a million and two hairs…) So, in virtue of knowing that *x is bald* is true, you have a non-trivial amount of information about the state of affairs that actually obtains.

But concepts and, therefore, propositions are *individuated* by their entailment-relations. Therefore, in so far as English semantic rules leave it open whether the predicate “bald” can be truly predicated of someone like Harry; and in so far as there is some proposition P such that, even after all the relevant empirical data is in, those same rules leave it open whether P entails, or is entailed by, *x is bald* (for some possible value of x); to that extent, English semantics has not assigned a *single* concept to the word “bald”, but rather a family of concepts, any two of which overlap considerably (no two of them apply to a person who has a billion hairs), but no two of which are strictly identical.

This doesn’t mean that there is no fact as to what our semantic rules demand. There is a fact. Supposing that Harry is a borderline case, it is a fact that these rules *don’t* demand that the expression “bald” be predicated of Harry, and it is also a fact that those rules *don’t* demand that the expression “non-bald” be predicated of him. It isn’t that these rules demand that “bald” *not* be predicated of Harry. It is rather that they *fail* to demand that it be predicated of him (or that it not be predicated of him).

There is no need to reject bivalence to accommodate this. The rules of cricket don’t demand that I brush my teeth five minutes from now or that I

not do so; those rules are compatible with each alternative. We are dealing with garden-variety cases of compatibility, i.e. of cases where a proposition P doesn't entail some other proposition Q and where P also doesn't entail not-Q.

§ Let us now take the next (and final) step in our argument. After we've taken this step, it will be clear what bearing these purely logical points have on the nature of law and legal interpretation.

Given *any* assurance a, there is some proposition p(a) such that a is kept (or honored) exactly if p(a) is true. Let us say that p(a) is the *validator of a.* Suppose I say to Jones "I promise to pick you up at the airport tomorrow." (Let *t* be the date of the day following my utterance.) My words are not true or false. (Promises are not true or false, though they may be sincere or insincere.) But there is some proposition C such that my promise is kept exactly if C is true. C is the proposition: *At t, JM picks up Jones at the airport.*

Let us generalize this point. Since not all speech acts are assertions, not all speech-acts are true or false. But all speech-acts have *satisfaction*-conditions. (Imperatives must be obeyed; questions must be answered truthfully; promises must be kept; and so on.) And given any speech act *a*, there is some proposition *V(a)*, such that *V(a)* is true or false exactly if *a*'s satisfaction-conditions are met. Of course, *V(a)* is the validator of *a.*

Since the validator of any assurance is a proposition, it follows straightforwardly from we what said a moment ago that there are no indeterminate validators. In its turn, this straightforwardly entails that there are no indeterminate assurances. So given any assurance A, given any individual x, and any privilege p, there is a fact as to whether A assures x of p.

Of course, there are cases where a so-called "assurance" leaves it open whether x may do p (for some values of x and p). Let us illustrate this with a short story. At a department-meeting, I say "I assure all of you that you are welcome at my party." But, for reasons to be described forthwith, my words leave it open whether I have assured department-member Roderick of anything. Roderick knows I dislike him, and thus has some reason to believe that, given the subtle and largely tacit understandings that govern verbal exchanges, Roderick didn't fall within the scope of my assurance. But Roderick also knows that I am reasonably forgiving, and also that I prefer to communicate directly rather than through innuendo. So Roderick has reason to believe that he *did* fall within the scope of my assurance, i.e. that he *may* be welcome at my party.

Here our tendency is to say that I gave an assurance that is *indeterminate* with respect to the question whether Roderick falls under it. But this is exactly what we must *not* say, as it embodies the absurd idea that propositions, and pieces of information generally, are *not* individuated by their entailment-relations. For the sake of argument, suppose that I *have* given an assurance that is indeterminate on the question whether Roderick is welcome at my party. Let A be that assurance, and let R(A) be the validator of that assurance. For the reasons given a moment ago, there is a fact as to whether R(A) entails *Roderick is welcome at JM's party.* R(A) either entails that proposition or it doesn't.

Given that R(A) is determinate in respect of whether it entails *Roderick is welcome at John's party*, it follows that A is correspondingly determinate and, consequently, that Roderick either has, or has not, been given an assurance as regards whether he would be welcome at my party.

So the idea that Roderick *has* been given an assurance, albeit an "indeterminate" one, is inconsistent with the most fundamental facts about

proposition-individuation (and, more generally, information-individuation). What we must therefore say is that I have *not* given Roderick an assurance, whether positive or negative. I have not assured him that he will be welcome at my part, and I have not assured him that he will not.

Of course, my words are not as informative as they should be. But this isn't because I have produced an "indeterminate assurance." It is because I have produced a determinate assurance that was *compatible* with Roderick's being welcome and also with his not being welcome. Under the circumstances, so as not to leave Roderick hanging, I *should* have provided an assurance that was *in*compatible with one of those possibilities. But the fact that I did not do so has nothing to do with my providing Roderick with an assurance that falls outside the scope of classical logic.

§ Now let us close the argument. Since laws are assurances, everything we just said in connection with the Roderick-situation applies to them. Since the validator of any assurance is either true or false, it follows that laws, being assurances, are determinate in every respect.[247] More formally, given any law L, and any person x, and any assurance A, either L determinately assures x of A or L determinately *fails* to do so.

In cases where we have a tendency to say that it is "indeterminate" whether L assures person x of A, L *doesn't* assure x of A. This doesn't mean that L gives some kind of *negative* assurance to x; it doesn't mean, in particular, that L assures x that he will *not* be given A. It means that L *fails* to give A to x.

Of course, not everybody will agree with our view that laws are assurances of rights. (For example, Hart would disagree with this.) But in the present context that doesn't matter. Everybody agrees that laws either prohibit or permit things. But given any speech-act that expresses a

prohibition or a permission, the validator of that speech act is a proposition, and is thus determinately true or false. Given this, it follows, as we've seen, that the corresponding prohibition or permission is determinate in every respect.[248]

§ Legal interpretations time-period-specific

I would like to end this book by defending two bold claims.

First, it is possible, at least in principle, for the Supreme Court (or any body tasked with the interpretation, so-called, of law) to rule one way at time t, and in a directly opposed way at time t*, *while being right on both occasions.*

My argument will take it for granted that legal interpretation is a principled endeavor, i.e. that some legal interpretations (so-called) are objectively better than others. It is *vacuously* true that, if legal interpretation isn't answerable to objective standards, then all of the Court's decisions are correct, even if they contradict one another. And, of course, I am not interested in defending that vacuity.

Further, my argument for this does *not* involve an abandonment of any of the laws of classical logic. I am taking it for granted that, for any proposition P, *P and not P* is false. Given the right kind of deviant logical system, it is trivially true that self-contradictions (on the part of a court or anything else) can be true. And, to echo what I said a moment ago, my purpose is not to defend such a triviality.

Finally, my argument doesn't involve any commitment to the position that U.S. law is inconsistent. Obviously *if* U.S. law is inconsistent, then it is trivially true that the Supreme Court could be right to rule in opposed ways at different times. And, to echo the last echo, there is no reason, or need, to defend such a triviality.

My point is a strictly logical one and isn't intended to have any historical content. My point is only to show that -- given classical logic, given a consistent body of laws, and given the view that legal interpretation is answerable to objective standards -- a body tasked with interpreting the law *could*, in theory, rule one way at a given time and in an opposed way at some other time, and be right on both occasions. So, for example, given only that the Supreme Court rules at time t that corporal punishment

in school is Constitutional, and rules at time t* that it is unconstitutional, it doesn't follow that the Supreme Court was wrong on either occasion.

The second claim concerns a doctrine known as "originalism." Originalism says that the Constitution should be interpreted in light of the intentions of the framers. On the face of it, this is obviously a plausible claim. But I will argue that originalism implicitly *defines* the expression "intentions of the framers" in terms of a certain conception of how the Constitution should be interpreted, and that, consequently, what originalism represents as historically based insights into Constitutionality are in fact artifacts of own definitions.

Why a judicial body, operating within a coherent system of law, can contradict itself and be right on both occasions

Let us start with a bit of fiction. In nation N (where English is spoken) the following verbiage is deposited in the law-books in the year 2006:

(*) "All bishops must wear green shirts."

Let us suppose that the linguistic meaning of (*) remains the same for next several hundred years. So the *intension* of the word "bishop" doesn't change. (Obviously its *extension* will change during that period.)

Before proceeding, we must make some general points that don't have any special relevance to the concept of law. (Readers familiar with Putnam (1975) and Burge (1979) will recognize the influence of these authors' work on the present discussion.) Whether one is a bishop is not a function solely, or even primarily, of one's intrinsic properties, but also of one's social liaisons. Given only that Bob and Twin-Bob are molecule for molecule twins, and that Bob is a bishop, it doesn't follow that Twin-Bob is also a Bishop. It could be that Twin-Bob lacks the relevant social liaisons.

There are some important corollaries. You can have a concept of what a bishop is – you can grasp the concept of bishophood, and can understand sentences like "there are fewer bishops than priests" --

without knowing the exact criteria that someone must satisfy to become a bishop. I myself can understand the sentence like “Smith is a bishop”, at least in some significant sense of the word “understand”, even though I have little or no knowledge of the Church-internal apparatus that bestows bishophood on someone. A consequence is that if I perform some speech-act concerning bishops – if, for example, I say “find out the name of the oldest Bishop in Alaska” – there will be, at most, an inadequate representation in my own psychology of the conditions that must be satisfied if the success-conditions associated with my speech-act are to be fulfilled. (If a speech-act is an assertion, its success-condition is that it be true; if a command, its success-condition is that it be carried out; if a question, its success-condition is that it be answered truthfully.) There is nothing in my psychology that tells you in virtue of what, exactly, somebody is a bishop – nothing in my psychology that distinguishes real bishops from bishop-impersonators.

There is another corollary. Because the property of bishophood is, in part, relational, it follows that the non-relational components of that property may vary from time to time. To be a bishop is to be embedded in a certain social practice in a certain way. The *non*-relational, intrinsic properties associated with one’s being thus embedded may vary from time to time. At time t, the Church may decide that only those with IQ’s of over 140 can be bishops. At time t*, it may be decided that even those of sub-standard intelligence can be bishops, provided that they are sufficiently pious. What is going on here is not that the property of bishophood *per se* has changed. There is some significant property such that, in virtue of being a bishop, a bishop of today has that property in common with a bishop of 1,200 years ago, even though the exact conditions that one must satisfy to become a bishop may well have changed over the centuries. (It is a matter of historical fact that the exact requirements for priesthood have changed over the centuries. It is not unreasonable to suppose that the same is true of the corresponding requirements for bishophood.)

There is no paradox here. To accommodate these facts we don’t have to take the absurd position of saying that the *property* of bishophood has changed. (Although *instances* of properties are spatiotemporal, at least in some cases, properties themselves are non-spatiotemporal objects and therefore don’t change.) Nor do we have to take the revisionist position of saying that the so-called bishops of the 1200’s were not bishops in the same sense as the bishops of the 21st century. (So we

don’t have to say that, in the last sentence, the first occurrence of the word “bishop” means one thing and that the other occurrence of that word means something else.) Rather, what we say is this. Like many other properties, the property of being a bishop is one whose instances can be realized in different ways, and the conditions that realize possession of that property at time t may be different from the corresponding conditions at time t*. So even though the property of bishophood *per se* cannot change, a person’s having an IQ of 95 may be incompatible with his having that property at time t, and compatible with his having it at time t*.

Someone is a bishop in virtue of his social-functional role. The conditions that one must fulfill to have that role may change from time to time. One possible reason for this is that, as the structure of society changes, so do the conditions that one must fulfill to have a certain role *in* society.

A bit of fiction will help us complete our argument. SC is the supreme court of nation N. In the year 2008, SC rules that people with IQ’s of less than 140 are *ipso facto* not bishops, and are therefore not legally required to wear green shirts. In the year 2025, SC rules that Hobson’s having an IQ of less than 140 does *not* disqualify him from being a bishop and that he *is* legally required to wear a green shirt. Given the points made in the preceding paragraphs, it follows SC *might* be ruling correctly on both occasions – even though (we may suppose) SC is operating within a coherent legal system and even though, as we have seen, judicial decisions have an objective basis. (I reiterate that, throughout this discussion, it is taken for granted that certain basic principles of classic logic hold – in particular, the law of non-contradiction.) In general, given only that some judicial body (that is a part of a coherent system of laws) rules one way at one time, and in a directly opposed way at some other time, it doesn’t follow it is ruling incorrectly on either occasion.

Given what we said, another non-trivial point follows. Even if the members of the SC in 2025 know exactly what was in the minds of the people who wrote (*), it doesn’t follow they will know the actual content of (*). This follows from the fact – noted a few paragraphs ago – that one can understand, and meaningfully utter, sentences like “Smith is a bishop” while having only a dim idea of the exact conditions that one must satisfy to be a bishop -- while having, in particular, little or no knowledge of the Church-internal regulations conditioning the bestowal of bishophood on someone. In general, knowing the psychology of the lawmaker doesn’t necessarily settle what the law in question is

or how it is to be applied or delineated.

There is no *a priori* reason why what we just said about SC needn't be true of the Supreme Court. So – and this is a purely logical point, and may have no historical consequences at all – it is *theoretically possible* that, on some of the occasions where the Supreme Court reversed itself, its ruling was correct both before and after the reversal. It seems to me that this point is at least worth bearing in mind in discussions of legal interpretation, even though it is an open question to what extent it fits the behavior of any actual judicial body.

§ Originalism not a coherent doctrine

A moment ago we saw *one* reason to believe that facts about the law-maker's psychology don't always have the relevance to questions of legal interpretation that one would expect them to have. In this section, I would like to present some other reasons for thinking this and, therewith, for doubting the veracity of a doctrine known as "originalism." This is the position that the Constitution must be interpreted in light of the intentions of the framers.

First of all, the moment we speak of "the intentions of the framers", we are already committed to some sort of idealization. (The need to address this point was brought to my attention by Peter J. Rudinskas.) There were various different framers – Adams, Jefferson, Washington…These people had numerically and qualitatively different minds. Surely Jefferson's intentions weren't exactly the same as Washington's. When we talk about the intentions that *they* had in connection with the future governance of the nation, we are already speaking in institutional, non-psychological terms (or, if you prefer, we are speaking in terms of institutional rather than personal psychology).

Whenever it is part of an expression that is meant to denote a *collective* attitude, the word "they" is in the same category as "the average man" or "frictionless planes." In other words, in such a context, the word "they" denotes an idealization, not an empirical reality. From a psychologist's viewpoint, there is no *they*. There are only various individuals whose numerically and qualitatively distinct psychologies may converge in certain respects at certain junctures.

When we talk about "the intentions of the framers", we obviously don't mean the intentions of the framers taken severally. Originalism is thus exceedingly dubious if the term "intentions" is understood in purely individualistic, psychological terms.

So if originalism is to hold water, we must take the expression "the intentions of the framers" to refer to the intentions of some kind of collective entity. But how are the intentions of that entity to be understood? Since, as we have seen, we cannot understand those intentions strictly in terms of the psychologies of the framers, it seems not out of the question that those intentions must be understood in terms of their social, and perhaps even their historical, role. But the socio-historical role of those intentions is inseparable from the socio-historical role of the Constitution. This suggests (though, I admit, it hardly establishes) that the originalist's position collapses into circularity (or, if you prefer, triviality): Constitutional interpretations are to be evaluated in terms of the intentions of the framers, but the intentions of the framers are to be understood in terms of the historical and social role of Constitutional interpretation.

Here is a related argument. For the reasons noted a moment ago, the expression "the intentions of the framers" denotes an idealization. An idealization is, by definition, something that has no precise counterpart in the spatiotemporal world. So when we talk about "the intentions of the framers", we are not talking about an empirical reality, but about some kind of theoretical construct. Of course, the theory embodied in one person's use of that expression may be different from that embodied in some other person's use of it. But it seems likely that, as a rule, the theories in question will *themselves* concern the role that the Constitution should have. If this is right, then it follows that originalism is guilty of vicious circularity or, if you prefer, of triviality.

A concrete example may help make it clear what this last argument amounts to. It is reasonable to predict that, at some point in the future, the Supreme Court will have to decide whether gay marriage (marriage between two people of the same gender) is protected by the Constitution. Suppose that, in connection with this issue, one of the members of the Court reasons thus:

(*) ***If Jefferson had been asked whether the Constitution protected gay marriage, he would have said "no"; and the same is true of Washington and Adams…***

(*) is surely correct. I very much doubt that, during his lifetime, Jefferson would have condoned homosexual marriage. He might not even have acknowledged the coherence of the concept of such a marriage. But this proves nothing that is of relevance to Constitutional interpretation. For if, during their life-times, we had asked Jefferson or Washington whether African Americans should be allowed to vote or hold higher office, those luminaries might well have given us answers that wouldn't correspond even remotely with later Supreme Court decisions – with decisions that even originalists regard as paradigms of sound jurisprudence. The same point holds in connection with the question whether the government should have the power to break monopolies, whether corporal punishment should be permitted in schools, and so on.

So when we talk about what Jefferson (or Adams or Washington…) *would* have said about such an issue, we are not talking about the actual Jefferson – the person who lived at a time when there was slavery and when it would have been unheard of for women to hold higher office. We are talking about what would be said by some idealization of Jefferson -- by somebody who had Jefferson's *basic* values, but whose own interpretation of those values wasn't distorted by the temporal and cultural parochialism to which Jefferson, like all other mortals, was inevitably subject. So when legal theorists talk about "what Jefferson would have said", they are not making an empirical statement about what the *actual* Jefferson would have said, but rather an *a priori* proposition about what he *should* have said, given his values (or what we believe to be his values).

But the originalist's intention in running counterfactuals like (*) is to *derive* or expose Jefferson's values, so as to generate an appropriate conception of Constitutionality. The originalist is thus caught in a vicious circle. He must use a counterfactual to identify the values on which his conception of Constitutionality is based. At the same time, he cannot run the counterfactual in question *until* he has identified those values.

There is a possible response to this:

"As you point out, the idiosyncrasies of Jefferson's psyche are irrelevant to issues relating to

Constitutional interpretation. For this reason, I agree with you that, when discussing Constitutional issues, we must idealize away from those idiosyncrasies and must therefore speak of what Jefferson *should* have said, given his values, as opposed to what he actually might have said, given his peculiarities and limitations as a human being. But I disagree with your point that this exposes some kind of circularity in the originalist's position. Jefferson's values are definable *independently* of any views of Constitutional interpretation. Jefferson believed in certain things. For example, he believed that, with some obvious qualifications, people should be allowed to do as they wish and that, consequently, government should have a minimal regulatory role. While such views have obvious consequences as regards Constitutional interpretation, they can (and probably must) be understood *independently* of any such views. After all, the concept of small government is one that a person can fully grasp even if she has never heard of the U.S. Constitution. So your charge of circularity is under-argued."

This line of thought simply displaces the problem onto an analogous problem. This, I will now argue, is because the expression "Jefferson's values" *itself* denotes an idealization, as opposed to a psychological reality. Consequently, talk of "Jefferson's values" in this context involves a circularity parallel to that discussed a moment ago.

Let us tell a story. We go back in time and ask Jefferson: "Should women be allowed to vote?" Jefferson says "no." Being creatures of the 21st century, we disagree with Jefferson. At the same time, we are (let us suppose) originalists; so we wish to reconcile his response with our belief that the Constitution should be interpreted in light of his intentions (along with those of Washington, Adams… who, we may suppose, gave the same answer as Jefferson in response to our question). We effect this reconciliation by telling ourselves that, in answering that way, Jefferson was not being true to his core values. It wasn't the *true* Jefferson speaking. Rather, it was a Jefferson whose inner goodness was marred by cultural baggage that was extrinsic to his character. (Compare: "The Smith I know wouldn't ever say such a thing -- that was the bottle talking.")

But, like many complex human beings, Jefferson had many values; and while some of them may

well have been fair-minded and egalitarian, others may have been retrograde and misogynist. So when we say that, in answering as he did, Jefferson was not being "true to his values", the word "his" doesn't denote the Thomas Jefferson the human being. Sadly, Jefferson's answer may have reflected who he really was. Rather, the word "his" denotes a Thomas Jefferson who has been purged of his ugly 18th century prejudices. But in that case, we are using the term "Jefferson" not to refer to the actual Jefferson, but to some idealization of Jefferson that corresponds to a certain conception of how the country should be run and therefore, presumably, of how Constitutional issues are to be settled. For analogues of the reasons given a moment ago, this entails that the originalist's position collapses into a vicious circularity (or an innocuous tautology).

Let us sum up. In discussions of Constitutional interpretation, one often hears claims about what Jefferson (or Adams or Washington…) *would* have said about such and such, e.g. the issue of gay marriage. But it is irrelevant what Jefferson *the actual person* would have said. Even the finest minds are marred by ugly biases. For all we know, Jefferson might have thought that any expression of homosexuality should be illegal or, on the contrary, he might have secretly thought that men should be legally required to abstain from relations with women and to marry other men. So when we talk about what Jefferson *would* have said, we are talking about what some kind of idealization of Jefferson would have said – some Jefferson who didn't suffer from the cultural limitations and random neuroses that mar even the best of minds, including his. But it is more likely than not that the relevant idealization is going to incorporate a certain conception of Constitutionality, and this condemns to vicious circularity any attempt to derive a correct conception of Constitutionality from such counterfactuals. Once we start dealing with an idealization of Jefferson's psyche, as opposed to his actual psyche, c*priori* propositions about values. The result is, at least plausibly, that such counterfactuals presuppose certain conceptions of Constitutionality, and thus cannot substantiate any such conception.

References

Austin, John. 1832. *The Province of Jurisprudence Determined.* London:

Weidenfeld and Nicolson.
--1890. *Lectures on Jurisprudence.* London: John Murray.

Baker, Mark. 1985. *Cops.* New York: Simon and Schuster.

Baker, Jean H. 2005. *Sisters: The Lives of America's Suffragists.* New York: Hill and Wang.

Bell, J.S. 1987. *Speakable and Unspeakable in Quantum Mechanics*. Cambridge: Cambridge University Press.

Bohm, David. 1957. *Causality and Chance in Modern Physics*. Philadelphia: University of Pennsylvania Press.

Bork, Robert. 1989/90. "The Limits of 'International Law'." *The National Interest* 18: 3-10.
-- 1990. *The Tempting of America.* New York: Simon and Schuster.

Burge, Tyler. 1982. “Other Bodies.” In Woodfield, Andrew, ed., *Thought and Object*. New York: Oxford.
-- 1977. “Belief de re” in *Journal of Philosophy*. (1977: 339-362).
-- 1979. “Individualism and the Mental.” In French, Uehling, and Wettstein (eds.). *Midwest Studies in Philosophy* IV, Minneapolis: University of Minnesota Press, pp. 73-121.
-- 1982. “Other Bodies.” In Woodfield (ed.). *Thought and Object.* Oxford: Oxford University Press.

Cappelen, Herman and Lepore, Ernie. 2005. *Insensitive Semantics: A Defense*

of Semantic Minimalism and Speech Act Pluralism. Oxford: Blackwell.

Carnap, Rudolph. 1934. *The Unity of Science.* Bristol: Thoemmes Press.

Chisholm, Roderick. 1961. "Responsibility and Avoidability." In *Determinism and Freedom in the Age of Science.* Sydney Hook ed. New York: Collier Books.

Coleman, Jules L. and Murphy, Jeffrie G. 1990. *Philosophy of Law.* San Francisco, CA.: Westview Press.

Davidson, Donald. 1984. "Mental Events". In *Essays on Actions and Events.* New York: Oxford University Press.

Dummett, Michael. 1978. *Truth and Other Enigmas.* Cambridge, MA.: Harvard University Press.

Dworkin, Ronald M. 1977. "Is the Law a System of Rules". In *The Philosophy of Law,* edited by Ronald Dworkin. Oxford: Oxford University Press.
-- 1986. *Law's Empire.* Cambridge MA.: The Belknap Press of Harvard University Press.

Fletcher, Connie. 1990. *What Cops Know.* New York: Simon and Schuster.

Fodor, Jerry. 1981. "Special Sciences". In *RePresentations.* Cambridge: The MIT Press.

Frege, Gottlob. 1884 (1953). *The Foundations of Arithmetic.* Translated by J.L. Austin. Oxford: Basil Blackwell. Reprinted in 1980 by Northwestern University Press.

Fuller, Lon. 1964. *The Morality of Law.* New Haven: Yale University Press.

Geach, Peter. 1972. *Logic Matters.* Berkeley: University of California Press.

Golding, Martin P. *Philosophy of Law.* 1975. Englewood Cliffs, New Jersey: *Prentice Hall, Inc.*

Grice, H.P. 1975. "Logic and Conversation." Reprinted in Martinich (1990).

Hart, H.L.A. 1961. *The Concept of Law.* Oxford: The Clarendon Press of Oxford University Press.
--1977. "Positivism and the Separation of Laws and Morals." In *The Philosophy of Law*, edited by Ronald Dworkin. Oxford: Oxford University Press.

Hegel, G.W.F. 1956. *The Philosophy of History.* Translated by C.J. Friedrich. New York: Dover Inc.

Hempel, Carl Gustav. 1952. *Fundamentals of Concept Formation in the Empirical Sciences.* Vol. 7 of the *International Encyclopedia of Unified Science.* Chicago: University of Chicago Press.
-- 1966. *The Philosophy of Natural Science.* Englewood Cliffs, NJ: Dover.

Hintikka, Jaakko. 1969. "Deontic Logic and its Philosophical Morals." In

Models for Modalities. Dordrecht, Holland: D. Reidel Publishing Company.

Hook, Sydney (ed.). 1961. *Determinism and Freedom in the Age of Science.* New York: Collier Books.

Jackson, Frank and Pettit, Philip. 2004a. *Mind, Morality, and Explanation.* Oxford: Oxford University Press.

Jackson, Frank and Pettit, Philip. 2004b. "Causation in the Philosophy of Mind". In Jackson, Frank; Pettit Philip; Smith, Michael 2004a.

Jackson, Frank and Pettit, Philip. 2004c. "Functionalism and Broad Content". In Jackson, Frank; Pettit Philip; Smith, Michael. 2004a.

Jackson, Frank and Pettit, Philip. 2004d. "Program Explanation: A General Perspective". In Jackson, Frank; Pettit Philip; Smith, Michael. 2004a.

Jackson, Frank and Pettit, Philip. 2004e. "Structural Explanation in Social Theory". In Jackson, Frank; Pettit Philip; Smith, Michael. 2004a.

Jackson, Frank and Pettit, Philip. 2004f. "In Defense of Explanatory Ecumenism". In Jackson, Frank; Pettit Philip; Smith, Michael 2004a.

Kant, Immanuel. 1785 (2002). *Groundwork for the Metaphysics of Morals.* Translated by Allen W. Wood. New Haven, Connecticut: Yale University Press.

Kaplan, David. 1989. "Demonstratives". In *Themes from Kaplan.* Edited by

Joseph Almog, John Perry, and Howard Wettstein. Oxford: Oxford University Press.

Kelsen, Hans. 1945. *General Theory of Law and State.* Translated by Anders Wedberg. New York: Russell and Russell

Kim, Jaegwon. 1993. *Supervenience and Mind.* Cambridge: Cambridge University Press.

Kripke, Saul. 1972. *Naming and Necessity.* Cambridge: Harvard University Press.

Kuczynski, John-Michael. 2004. "A non-Russellian treatment of the referential-attributive distinction" *Pragmatics and Cognition* 12:2 (2004), 253-294
-- 2005. "Can One Grasp Propositions without Knowing a Language?" *Teorema.* Volume XXIV/2, pp. 43-65.
-- 2005b. "Counterfactuals: The Epistemic Approach." *Philosophia Scientiae.* 9 (1), pp. 1-45.
-- 2006. "Implicit Comparatives and the Sorites." *History and Philosophy of Logic*, 27 (February 2006), 1-8.
-- 2006b. "Formal Operations and Simulated Thought." *Philosophical Explorations*, vol. 9, No.2, June 2006 (pp. 221-234).
-- 2007. "Two Concepts of Sentential Form and the So-called Computational Theory of Mind." Forthcoming in *Philosophical Psychology.*

Lange, Marc. 2000. *Natural Laws in Scientific Practice.* Oxford: Oxford University Press.

Lewis, David. 1966. “An Argument for the Identity Theory.” Journal of Philosophy 63 (1966): 17-25.
-- 1973. *Counterfactuals.* Oxford: Blackwell.
-- 1980. “Mad Pain and Martian Pain,” in Ned Block, ed., *Readings in the Philosophy of Psychology*, Volume I (Harvard University Press) 216-222.
-- 1986. “Are We Free to Break the Laws?” In *Philosophical Papers Volume II.* New York: Oxford University Press.

Mackie, J.L. 1982. *The Miracle of Theism. Arguments for and against the Existence of God.* New York: Oxford University Press.

Martinich, A.P. e.d. 1990. *The Philosophy of Language.* Oxford: Oxford University Press.

McCawley, James D. 1981. *Everything that Linguists Have Always Wanted to Know about Logic.* Chicago: University of Chicago Press.

McDowell, John. 1998. *Meaning, Knowledge, and Reality.* Cambridge: Harvard University Press.
-- 1998b. “Singular Thought and the Extent of Inner Space”, in McDowell 1998 (228-259).
1998c. “Criteria, Defeasibility, and Knowledge”, in McDowell 1998 (369-394).
1998d. “Knowledge and the Internal”, in McDowell 1998 (395-413).

McGinn, Colin. 1988. *Mental Content.* Oxford: Oxford University Press.

Nagel, Ernest. 1962. *The Structure of Science.* Indianapolis: Hackett.

Orwell, George. 1949 (1983). New York: Penguin Putnam.

Pap, Arthur. 1961. "Determinism, Freedom, Moral Responsibility, and Causal Talk." In *Determinism and Freedom in the Age of Science.* Sydney Hook ed. New York: Collier Books.

Pettit, Phillip. 1975. *The Concept of Structuralism.* Berkeley, CA: University of California Press.

Plantinga, Alvin. 1974. *The Nature of Necessity.* Oxford: Clarendon.

Popper, Karl. 1957. *The Poverty of Historicism*. London: Routledge & Kegan Paul.

Putnam, Hilary. 1975. "The Meaning of 'Meaning'". In *Collected Papers: Volume II.* Cambridge: Harvard University Press.

Quine, Willard van Orman. 1960. *Word and Object.* Cambridge: The M.I.T. Press.

-- 1966. "Quantifiers and Propositional Attitudes." In *The Ways of Paradox.* Cambridge: Harvard University Press.

Rawls, John. 1971. *A Theory of Justice*. Cambridge, MA.: Harvard University Press.

Recanati, Francois. 2001. *Literal Meaning.* Oxford: Oxford University Press.

Rubin, Alfred P. 1990. "Correspondence." *The National Interest* 122-124 (Spring 1990).

Russell, Bertrand. 1905. "On Denoting". In A.P. Martinich (ed) *The Philosophy of Language* Oxford: Oxford University Press.
--1920. *Introduction to Mathematical Philosophy.* London: George Allen Unwin.
-- 1927. *The Analysis of Matter.* London: George Allen and Unwin.

Sabine, George H. 1937. *A History of Political Philosophy.* London: George G. Harrap and Co.

Searle, John R. 1965. "What is a Speech Act". In *Philosophy of Language*. A.P. Martinich (ed.). New York: Oxford University Press.
--1995. *The Construction of Social Reality.* New York: The Free Press.
-- *Speech Acts.* 1970. London: Cambridge University Press.

Strawson, Peter. 1950. 'On Referring'. In A.P. Martinich (ed.) *The Philosophy of Language.* Oxford: Oxford University Press.

Taylor, Richard. 1961. "Determinism and the Theory of Agency." In *Determinism and Freedom in the Age of Science.* Sydney Hook ed. New York: Collier Books.

Tucker, Robert C. 1981. *Politics as Leadership.* University of Missouri Press.

Van Inwagen, Peter. 1993. *An Essay on Free Will.* New York: Oxford University Press.

Weber, Max. 1947. "Politics as a Vocation." In *From Max Weber: Essays in Sociology.* Translated by H.H. Gerth and C. Wright Mills. London: Kegan

Paul, Trench, Trubner & Co., LTD.
-- 1966. *Writings on Law in Economy and Society.* Translated by Edward Shils and Max Rheinstein. Cambridge, MA.: Harvard University Press.

Wilkins, Burleigh Taylor. 1978. *Has History any Meaning? A Critique of Popper's Philosophy of History.* Ithaca, New York: Cornell University Press.
-- 2006. "Review Essay on *The Roots of Evil*." *The Journal of Ethics* (2006) 10: 193-199.

Wittgenstein, Ludwig. 1958. *The Philosophical Investigations.* Translated by G.E.M. Anscombe. Oxford: Basil Blackwell.
--1974. *Philosophical Grammar.* Rush Rhees (ed.). Anthony Kenny (transl.). Oxford: Blackwell.
-- 1975. *Philosophical Remarks.* 1975. Edited by Rush Rhees, Raymond Hargreaves. Translated by Maximilian A.E. Aue.
--1983. *Remarks on the Foundations of Mathematics.* Edited by G.H. von Wright, Rush Rhees, and G.E.M. Anscombe. Cambridge: M.I.T. Press.

---

[1] Of course, it might be said that the flourishing of a serial killer is not a good thing, so that, except subject to heavy qualifications, it is false to say that "human flourishing is a good thing."

Here there are two points to make. First, Smith's flourishing is that of a gifted musician, not a serial killer. Second, as we've already intimated, and will later discuss at greater length, the idea that a sadist's flourishing is bad *tout court* involves a failure to distinguish several very different notions of badness. Obviously the harm that a sadist causes others is bad *tout court.* But it doesn't follow that, taken by itself, a sadist's flourishing is bad.

[2]

And, of course, purely *a priori* or analytic domains (e.g. number theory).

Russell (1927) gives a clear explanation of why principled regularities are so unlikely to be given expression by an infallible constant conjunction, and also of the conditions under which such conjunctions are found.

[3]

This point was made by Mill in his *System of Logic.*

[4]

I have never seen any acknowledgement of this fact – that principled relations are given by *ceteris paribus* statements – anywhere in the positivist literature. As far as the anti-positivist literature is concerned, I believe that Dworkin's important analysis of the distinction between rules and principles does, at least a fashion, incorporate an acknowledgement of that point.

[5]

To echo the last note, I have never seen any acknowledgement of this fact – that where concepts having a teleological component are concerned, principled relations aren't typically given by perfect statistical correlations – anywhere in the positivist (or, for that matter, the anti-positivist) literature.

[6] Obviously there are doctors (e.g. Dr. Mengele) who do *deliberately* do harm. But usually when a doctor does harm, it is by *failing* to do what a better doctor would have done under the circumstances -- it is by *failing* to diagnose an illness correctly or by failing to prescribe the appropriate medicine. Surely Dr. Mengele was a singularity, and his behavior cannot be taken as a paradigm of most actual sub-standard medical care.

[7]

Hart (1977: 34-37).

[8]

Cf. Rawls (1971: 3): "Justice is the first virtue of social institutions, as truth

is of systems of thought." Of course, a system of laws is a paradigm case of a social institution.

[9]

The words "ceteris paribus" are crucial in this context. The conceit here is one that Rawls (1971) puts to important account.

[10]

See Dworkin (1977: 40). I believe that Freud (in *Group Psychology and the Analysis of the Ego* and in *Totem and Taboo*) put forth some views as to the nature of this (alleged) blunder or, at any rate, as to the nature of similar blunders. If I am not mistaken, in *the Genealogy of Morals*, Nietzsche puts forth a system of thought that is consistent with this view. Coleman and Murphy (1990: 52-53) briefly discuss the relevance of Nietzsche's ethical and cultural views for the philosophy of law.

[11]

Austin (1832), Hart (1977: 24-25). Austin (1832) described Blackstone's idea that judges "find" but never "make" law as a "childish fiction" (See Hart: 1977: 24-25).

[12]

Dworkin (1977: 38-39).

[13]

In this work, I am presupposing a certain realism about moral truth. I am taking it for granted that it really is morally wrong to do certain things (e.g. kill babies) and that it really is morally right to do other things (e.g. to give money to those in need).

[14]

Hart (1977: 27).

[15]

Hart (1977: 27).

[16]

Dworkin (1977: 1).

[17]

Dworkin (1977: 1-2).

[18]

Dworkin (1977: 44-46).

[19]

Dworkin (1977: 44-46).

[20]

Dworkin (1977: 46-49).

[21]

In chapter 10, we will discuss the exact nature of this vagueness, and will find that our initial theoretical impulses in connection with it are quite misguided.

[22]

Dworkin (1977: 46-49).

[23]

One mark of a principle seems to be that it is given a sentence that is preceded by the expression “all other things being equal.” All other things being equal, products should be safe, contracts should be fair, and so on. Rules are not hedged by *ceteris paribus* clauses; they are qualified by a finite number of discrete *exceptions.* As every philosopher of science knows, *ceteris paribus* clauses are not equivalent to lists of exceptions.

[24]

Dworkin (1977: 44-49).

[25]

This is not my own view. I think that there is a significant sense in which a rule may be *interpreted*. Here I am stating the argument to be evaluated, not

the position I wish to maintain.

[26]

Arguably, the term “rule” means one thing in the expression “mathematical rule” and quite another in the expression “legal rule”. The second has a normative meaning: it says what people are *supposed* to do. The first says what *is* the case: certain numbers *are* paired off with others. I am aware of this (possible) ambiguity. But I think that the argument which I am presenting is itself, at some level, guilty of that very equivocation. That is why I am not registering that ambiguity (if such it be) in my exposition.

[27]

Dworkin (1977: 52-60).

[28]

This is a slight simplification, but a harmless and useful one in this context.

[29]

This is not my own view. I believe that there is *always* a fact of the matter, and that apparently vague predicates are in fact *relational* predicates, whose non-vagueness is revealed when the unspoken, contextually-supplied benchmark is identified. See Kuczynski (2006). But there is no doubt that utterances of the form “it is indeterminate whether such and such” are sometimes intended to convey that there is no fact as to whether such and such (notwithstanding that, thus intended, all such utterances are, in my view, false). In Chapter 10, we will discuss this issue and also its bearing on the concept of legal interpretation.

[30]

In Kuczynski (2007), I argue that the kind of “discretion” that must be exercised in connection with judgments relating to historical or literary analysis must *always* be exercised, even in cases where the thinking consists in the so-called “mechanical” application of some algorithm. The term

“purely mechanical thought” is thus to be taken in a loose, or a purely comparative, sense.

[31]

Arguably, these values are projections of the fact that our country came into existence in consequence of an act of defiance.

[32]

Robert Bork (1990) has characterized a number of Supreme Court decisions as judicial usurpations of the law. Some of these cases are similar to *Henningsen*.

[33]

It seems that the debate between the positivist and the anti-positivist can be understood in terms of an analogue of Euthyphro’s question: is it legal because the judge approves, or does the judge approve because it’s legal?

[34]

In any case, this is what Kant said. I believe that Kant is right: it is *always* wrong to break promises (even though it is sometimes *more wrong* to keep them). This is an issue that (in some form) we will discuss at length.

[35]

What we’ve just said about the concept of legality closely parallels what Dworkin says about rules, and owes much to his insights.

[36]

The expression “the slaughter bench of history” is borrowed from Hegel’s *Philosophy of History.*

[37]

Arguably, there are special cases where the reality coincides with the information through which one learns about it. It could be argued, for example, that one’s knowledge of one’s sensations is entirely unmediated and that, consequently, the reality collapses into the information through which

one learns about it. (I take this position in my article "Are any of our beliefs about ourselves non-inferential or infallible?" (*Journal of Theoretical and Philosophical Psychology*, Volume 21, Number 1, Spring 2001 pp. 20-46). I now have doubts about the cogency of the argument provided there, believing that it might be guilty of a version of what Sellars referred to as the "myth of the given." In any case, when we are dealing with laws – in fact, when we are dealing with anything other than states of consciousness – it is not an option to say that the reality collapses into the information through which we learn about it.

[38]

Contrary to what Berkeley himself said.

[39]

See Dummett (1978).

[40]

For an illuminating discussion of the perils of slippery slope reasoning, see van Fraassen (1980: 14-19).

[41]

Granted, the Constitution itself made certain provisions for slavery. But, arguably, these were just concessions to the mores of the time, and contradicted the deeper intent of the Constitution (and perhaps also the secretly held moral convictions of the authors of the Constitution). There often is, I think, a fact as to whether a certain law is Constitutional or not; and what settles the matter is often the presence of some principle in the Constitution. We will revisit this question in chapter 9.

[42] This position is associated with Lon Fuller's work. See Fuller (1964). See also Wilkins (2006: 195).

[43]

The history of women's suffrage in the U.S. is more complicated than I am making it out to be. (The omitted complications are irrelevant.) As far back as 1790, some states (e.g. New Jersey) allowed women to vote. But it was not until 1920 that it became a federal law that women are to have the same voting privileges as men. See Baker (2005) for a more complete history. [44]

The exact change consisted in the addition of the 19th amendment to the Constitution. Notice that this amendment is a *rule*, not a principle:

"The right of citizens in the United States to vote shall not be denied or abridged by the United States or by any State on account of gender. Congress shall have the power to enforce this article by appropriate legislation."

This ammendment is completely binary. It doesn't say, even implicitly, that *ceteris paribus* the legal right to vote will not be subject to gender-based restrictions. It says simply that women have the right to vote. It thus isn't comparable to the principle that all citizens are entitled to life, liberty, and the pursuit of happiness. X's right to those things is to be balanced against other people's rights to those same things: there is an implicit *ceteris paribus* clause. So the principle that citizens are entitled to life, liberty and and the pursuit of happiness has a dimension of weight. By contrast, the 19th ammendment says that women are permitted to vote, not that their right to vote is a consideration to be weighed against others.

Of course, as with many rules, there are *exceptions* to the 19th ammendment. If a woman is a felon or is severely mentally imparied, then she cannot vote. But as we discussed earlier, a finite list of *exceptions* is very different from a *ceteris paribus* clause. Rules are characterized by finite lists of exceptions; principles are characterized by *ceteris paribus* clauses. So the

19th ammendment has none of the hallmarks of a principle (even though it was obviously motivated by principles), and all of the hallmarks of a rule.

[45] Strictly speaking, of course, it is not the mere presence of such principles in the Constitution that causes judges to make the decisions they in fact make. It is judges' awareness of those principles that does so.
[46]
I will end this section by discussing an objection, not to the substance of my views, but to my general methodology:

You have proposed two, opposed pictures of the relationship between U.S. law and the U.S. Constitution. On one those pictures, the Constitution is a part of U.S. law. On the other it is not. But it is not possible for both of those views to be correct.

The concept of law is not free from vagueness. When we subject a vague, lay-concept to theoretical precisification, we often find that there are junctures at which either one of two (or more) opposed pictures is *equally* capable of accommodating the relevant hard core of data. We alluded to such cases earlier, in connection with the concept of temperature. There is no shortage of such cases in the history of both science and philosophy. There are cases where the hard core of empirical data is adequately accommodated by saying that, collectively, molecules X, Y, and Z have a temperature of 10,000° or by saying that they have a temperature of 5°. This is not to advocate any kind of anti-realism (e.g. the view that there aren't *really* electrons or viruses). I am not advocating any kind of indeterminacy in connection with *ontology* – in connection with the existence of a certain kind of entity. I am merely registering the pedestrian fact that our lay-concepts are

sometimes confused, opaque, and not entirely free from contradiction and, further, that there are different ways of disciplining those concepts. It is a common-place that lay-concepts don't always yield determinate verdicts, or may yield contradictory verdicts, when they are forced to operate outside the narrow corridors of pre-theoretic intuition. A theoretician's task is to provide extensions or articulations of those concepts, so far as possible, that yield the same verdicts with respect to the paradigm-cases and that, unlike their pre-theoretical counterparts, yield clear verdicts with respect to non-paradigmatic cases. If a theoretician provides multiple, internally consistent articulations satisfying the conditions just described, that is not to her discredit. It is irrelevant whether such articulations are mutually contradictory, since the theoretician is asking that but *one* of them be accepted.

[47]

Hart (1977: 19).

[48]

Hart (1977: 27).

[49]

See the preface to *The Varieties of Reference.*

[50]

See the introduction to Fodor's *RePresentations*. Fodor makes it clear that he is, in general, a functionalist, He also says why, despite this fact, he cannot accept a functional analysis of phenomenal states (e.g. itches, tickles). David Lewis (1966, 1980) famously *does* accept a functional analysis of phenomenal states. With regard to those who wish to individuate qualia in terms of their phenomenological, as opposed to functional properties, he says flatly that he simply has no idea what this non-functional mystery-property could possibly be. While he grants that there is such a thing as phenomenology, Lewis maintains that phenomenology *is* functional role.

[51]
Fuller (1964).
[52]
In any case, the term “law” picks out such a category in so far as it picks out anything worthy of study. For the sake of argument, suppose that the condition a thing must meet to be picked out by the term “law” were irreducibly disjunctive. An example of such a disjunctive condition might be: *x is a law if x is an order issued by a gunman or x is issued by a democratically elected legislature or x is a form of conduct conformed to by most of the members of any healthy society or x is issued by a legislature appointed by the king or x is a sentence in a book issued by such and such a printing-press.*… In that case, we would be better of coming up with a different term for each disjunct, and then studying the concept expressed by that disjunct, leaving to one side the term “law” and the hyper-disjunctive concept denoted thereby. It is hard to believe that the concept *law* is thus disjunctive. In any case, it is worth attempting to produce a non-disjunctive analysis.

Wittgensteinians might say that the concept of law is a “family resemblance” concept. But, as I am aspiring to show, I don’t think it is necessary to take that route.

[53] Economists are often tasked with determining whether some law* has a good effect over all. Of course, the relevant conception of goodness is a narrowly economic one. But surely economic goodness is not categorically separable from moral goodness. *Ceteris paribus* a law that makes everyone poor is not a law that it is *morally* good to have, and *ceteris paribus* a law that makes everyone rich *is* a law that it is morally good to have.
[54]

Consider some of the various sentences that use the term “law”:

“immigrants who enter this country without a visa are in violation of the law”;

“it is against the law for a male not to register with the Selective Service at age eight-teen”;

“it is against the law to litter”;

“insider-trading is against the law”;

“the law requires sex-offenders to let everyone in their neighborhoods know of their past sex-crimes”;

“the law does not allow an illegal immigrant to have a driver’s license”;

“it is against the law for a doctor to re-use syringes”;

“it is against the law for drivers (in California) not to have car insurance.”

Unless we had some term that applied *both* to laws* that are morally wholesome and those that are not, the propositions expressed by these sentences would become practically inexpressible. Complex moral issues – issues that, in some cases, we may be certain will never be settled -- would have to be settled *in advance* in order to know how to express some obvious truth (e.g. the truth that we now express with sentences like “in California,

drivers are required by law to have car insurance"). So in order to make those propositions expressible, we'd have to find some term that was neutral between good and bad laws*. But then all of the issues now studied by the "philosophy of law" would be studied by the "philosophy of__", where the blank is to be filled in with this new term.
[55]

Consider a book on the economic consequences of certain kinds of laws*, e.g. Bork's book on anti-trust law*. Suppose we had to re-write that book, but this time without the use of a term that is neutral between moral and immoral laws*. (Note: when I talk about "immoral" laws*, I mean those that are *comparatively* immoral, since my view, of course, is that all laws* embody a certain morality.) Some sentences using the term "law" would then have to be translated into sentences where that term was replaced with a term that applied only to morally *wholesome* laws*. Other such statements would have to be translated into statements where the term "law" was replaced with a term that applied only to morally unwholesome laws*. So unless the person doing the translating were a supremely gifted moral philosopher, the result of that translation would almost inevitably turn a number of true statements into false ones.

Further, it would then become impossible for the author to state his viewpoints in a way that didn't prejudge the issues being investigated. After all, one of the objectives of that book is to *figure out* whether anti-trust laws were helpful or harmful. If the author had to use terms that presupposed such a categorization, what would otherwise be significant, non-trivial claims would crumble into tautologies and contradictions. The only way to undo this damage would be to do what Bork did in the first place – use some term that applies to both categories.

[56] See Hempel (1952) and Nagel (1962: chapter 11). Superficially, some of the points that Kripke (1972) makes in connection with heat might seem to effect a kind of re-identification of the lay and scientific concepts of heat. While I agree with Kripke's position, I don't think it warrants the rejection of the point made here. Given what Hempel and Nagel say, there are clearly some cases where – even if we allow for the points made by Kripke – the concept of temperature at work in statements like *such and such molecules have thus and such temperature* is connected to the common-sense concept of temperature only by way of a rather complicated theoretical apparatus. In such cases, it really does seem *de rigueur* to say that the concept of temperature at work is a scientific delineation or reconstruction of the common-sense concept, and not one and the same with that concept. There are some contexts where, from a *physicist's* point of view, the temperature in some part of space where your hand would quickly freeze is actually 10,000°. This will be the case if there are three molecules in that area that have the right mean kinetic energy. It is very hard to believe that the concept of heat operative here is one and the same with the concept of heat implicated in Little Timmy's statement that "it is hotter in Georgia than it is in New Hampshire." See Nagel (1962: chapter 11). The lay-concept has little application outside of extremely narrow horizons. It seems that what we are dealing with is, as Carnap puts it (though not in connection with this exact concept), a "rational reconstruction" of Little Timmy's concept.

Despite what we've said, some will *still* argue that the concept of heat implicated in Little Timmy's statement is identical with that implicated in the statements about the three molecules. Even so, it is fairly uncontroversial that *some* lay concepts must be rationally reconstructed – jettisoned in favor of similar, but distinct concepts – in the interests of scientific or philosophical endeavor. An example is the concept of probability. Even if we allow for the

fact that there are different *theories* about probability (e.g. the frequency theory), and even if we allow for the externalist theories of Putnam and Kripke, there is still (so it is widely held) no *one* concept corresponding to the lay-concept of probability.

Incidentally, the term "rational reconstruction", in the sense in which it is being used here, occurs in the introduction to Carnap's book on probability. I do not know whether that is the first occurrence of it in the literature.

[57]

Of course, physicists do use the term "temperature"; but the concept it picks out is (at least arguably: see the previous footnote) not exactly the concept picked out by the homonym of that term used by the layperson.

[58]

See Hempel (1952).

[59]

See Nagel (1962: 535-546) for a discussion of the feasibility of attempts to reduce the sociological to the psychological.

In *Group Psychology and the Analysis of the Ego* and *Totem and Taboo*, Freud gave psychological analyses of sociological concepts. I do not know whether Freud held that social concepts were *categorically* capable of analysis in strictly psychological terms.

[60]

Carnap (1934) holds this. See Fodor (1981) for some criticisms of this general viewpoint. See also the last chapter of Hempel (1966). The best discussion I know of concerning the conditions that have to be met for inter-disciplinary reductions is found in chapter 11 of Nagel (1962). These days – for the reasons given by Nagel, Hempel, and Fodor -- practically nobody would agree with the sort of extreme reductionism advocated in Carnap (1934).

[61]

This is not my own view.

[62]

In *Totem and Taboo*, Freud discusses the psychological basis of law. The distinction between good and bad laws has no place in that discussion. It could be argued that this is to be attributed not to anything intrinsic to the concept of law, but rather to the limitations of Freud's analysis. But it is still suggestive.

[63]

In any case, there is no such evidence that I know of that has been marshaled by any advocate of a moralistic theory of law.

[64]

Also, evil laws*, like those prohibiting women from voting, have much the same *psychological* reality as morally wholesome laws*. Suppose that Smith is a sane and well-informed person who lives in the U.S. in 1820. Obviously what Smith chooses to do, at least in certain contexts, will be a function of what the laws* are, it being irrelevant whether those laws* are morally wholesome or not. When Smith is to decide whether she can borrow her neighbor's horse without asking permission, she is going to think in terms of what "the law" is. (Of course, this will be but one of many considerations.) When Smith must decide whether she can vote, she is again going to think in terms of what "the law" is.

Of course, like many other psychological categories, this category will be divided into sub-categories. Smith will probably have different beliefs about the moral merits of the laws* she confronts, and these beliefs will be operative in many contexts. Since Smith sees the law* forbidding theft as just, she will either defer to it or feel guilty if she does break it. Since Smith sees the law* forbidding women from voting as wrong, she may choose to

break it (by going to the polling station disguised as a man), and will not feel guilty about doing so. So we mustn't deny the obvious truth that human psychology makes *some* allowance for the difference between moral and immoral laws*.

But the psychological differences just discussed seem to be internal to a unified psychological category. One sees it is morally acceptable to break *the law* in cases where it is not wholesome, as morally unacceptable to break *the law* in cases when it is wholesome.

In light of this suppose we decided that the laws* prohibiting women from voting were not laws at all, on the grounds that they lack the needed moral qualifications. In that case, to recover the psychological explanations just discussed, we'd have to come up with some new term – say "shmaw" – to denote the psychologically significant category currently associated with the term "law." But then the very problems now faced by the anti-positivist in connection with the term "law" would be reproduced within the philosophy of shmaw.

[65] Oliver Wendell Holmes famously held a "bad man" theory of law (Golding 1975: 37-38). If you ask a good person what the law is, he will give you a tendentious, revisionist answer. He will say, for example, that there aren't *really* laws prohibiting women from voting, since such prohibitions oughtn't be dignified as "laws." So if you want to know what the law is, you must ask a bad man. The bad man will give you a dispassionate answer, one that doesn't conflate the prescriptive with the descriptive.

As we've seen, the bad man's conception of law is the one used by historians, economists, and the layperson. The good man's conception applies to some concept that has little or no explanatory utility. So, to this extent, Holmes seems to be right (leaving it open whether the same can be said of the

conception of law that Holmes derived from these considerations).

[66] Fuller (1964).

[67]

This brings us to a loose end that must be tied up. I have argued that there is no legal obligation where there is no legal protection. Clearly slaves in the U.S. were not really given any protections by the U.S. government, and thus didn't have any legal obligations to it. But if they had no such obligations, then, it might be argued, it follows that they were not enslaved by any *laws* and that, contrary to our hypothesis, there were consequently no *laws* permitting slavery.

The argument just given involves a non-sequitur. Given that the government didn't protect slaves, at least not to a significant degree, it follows that *they* were not legally bound by the laws permitting slavery (or, arguably, by any other U.S. law). But it does not follow that such laws didn't exist. Those laws did impose duties on people who were recognized, and thus given protections, by the U.S. government. Here is an analogy. (The analogy must be interpreted as narrowly as possible: any invidious overtones it might be seen as having have no basis in the author's views, and are to be blamed entirely on the evil laws being discussed.) Cows obviously don't have obligations to the U.S. government. But even though I am not a cow, I have a legal obligation to abide by those laws: for example, I cannot take another person's cow without his permission. The laws permitting slavery imposed similar legal obligations on those whom the government did recognize.

[68]

Hart (1977: 35-36).

[69]

See Rawls (1971).

[70]

To simplify discussion, suppose you know that he will really do it, and also that you are totally unable to convince him not to do it. Suppose also that you know all fifty people at his work place to be morally excellent people.

[71]

Cf. Hintikka (1969: 203): "I give an honest promise (let this act be described by *p*) to bring it about that *q* (say, have a cup of coffee with you). Unknownst [sic] to me, however, my father has fallen ill, which creates an obligation to see him that overrules my earlier promise. It seems to me that moral philosophers have felt somewhat uncomfortable in discussing this kind of situation, and perhaps one can also see why. For the fact that the obligation created by my promise is overruled means that it is false to say *simpliciter* that *Oq*, i.e. that I am obliged to fulfill the promise…Yet it is clear that not everything is morally all right if I have to break my promise, however firmly this course of action may be prescribed to me by the norms I abide by. I have somehow done something wrong. This 'moral failure' is what easily makes on hesitant to say that in such a case there is no absolute duty to keep the promise."

[72] Hart (1977).

[73]

There is, I believe, a very big difference between saying that laws are *stated* vaguely, on the one hand, and saying that laws *themselves* are vague or "indeterminate", on the other. I believe that, unless we are to part with some truisms about the nature of propositional content and entailment, we must at all costs avoid saying that laws themselves (whether logical, legal, moral, causal…) are ever "vague." (See Kuczynski 2005.) This point has *a* relevance to the debate between the positivists and the anti-positivists – but

not one we can explore right here. We will discuss it at length in chapter 10.

[74]

Strictly speaking, “nothing” is not quite the right word. It would be more correct to say that, if we do not assent to Radbruch’s proposed change in nomenclature, then *much less* of moral consequence follows from the fact that the term “law” is applicable to some directive than would otherwise follow. In particular, it could not be inferred, on that basis alone, that the directive in question ought to be obeyed. But, strictly speaking, *something* of moral consequence always follows from the fact that a directive is correctly described as a “law.” For reasons that we’ve already considered, and for reasons that we have yet to consider, no law is ever *entirely* hollow morally (though many are more bad than good, and must therefore not be obeyed).

[75]

This is, of course, reminiscent of a famous argument found in Wittgenstein’s *Philosophical Investigations.*

[76]

Fuller (1964: 34) makes this point. See Coleman and Murphy (1990: 38-39) for a discussion of Fuller’s point.

[77]

Hart (1977: 36-37). This example is not Hart’s. Hart actually doesn’t give an illustration of his point.

There is a response to make to this point of Hart’s. Judge Jones might rationalize his draconian conduct by saying that society depends on rules, and that when people (including judges) make exceptions to rules – even in cases where the interests of justice seem to demand it – the rules begin to deteriorate, and society is led down a slippery slope to anarchy. Thus, in cases where it seems evil to stick to the letter of the law, there is in fact some long-term social justification for doing so. I do not myself think that this is a

particularly good reply (not every relaxation of a rule leads to anarchy – in some cases, bending or breaking the rules might be needed to prevent a descent into anarchy). But it is one that is often given by sticklers for the rules, and ought to be considered (if only to be rejected).

[78]

See Fuller (1964).

[79]

It is not entirely clear to me whether causal and natural necessity are to be identified with each other. But my reasons for this lie outside the scope of the present discussion.

[80]

Hart (1977: 35-36).

[81]

The need to consider this sort of law was brought to my attention by Thomas Holden.

[82]

In fact, we will find some reason to give credence to the seemingly paradoxical thesis that distinctively *governmental* evil consists in *mis*-protections: in an over-use or mis-use of mechanisms whose purpose, and whose standard mode of behavior, is to protect.

[83] Of course, appendices don't do any good when healthy, even though a diseased is fatal.

[84] But it is not hard to find that data that gives it at least a flicker of support. A gunman's demands are not *supposed* to do anyone any good (except for the gunman himself). This suggests that there is no law in a context where the *supposition* that good is being done is altogether absent. It

seems to me that we *don't* have laws in a context where the person giving orders doesn't even make a pretence of having *moral* authority – where he openly admits that any authority he has (if the term "authority" is even applicable in such a context) is based on nothing other than brute force. There are dictators who are psychologically on all fours with gunmen. Such dictators make it a point to cloak their orders under the banner of concern for the welfare of the people. At one level, this is obviously just a Machiavellian ploy. But, it seems to me, if this pathetically transparent pretence were absent -- if the dictator openly admitted that he had no concern at all for the welfare of his subjects – the result would no more register as a law than the gunman's threat. This suggests to me that we don't have a law *until* we have something that is supposed to do good and that, consequently, that something is a law *in virtue of* some expectation or supposition, perhaps false, that satisfies some kind of moral requirement. If correct, this would entail that laws are necessarily things that are supposed to do good. While being consistent with the fact that there are laws that *don't* do any good, this would suffice to establish a principled, internal relationship between law and morality

[85] Of course, the principles of which Dworkin speaks are prescriptive, not descriptive, whereas the principled relation between (for example) a match's being struck and its being ignited is descriptive, not prescriptive. Dworkin is talking about ethical principles, whereas we have been discussing nomic principles. But Dworkin's point seems to apply to principles in both the descriptive and prescriptive senses. We must also remember that prescriptive claims easily reinterpreted as descriptive claims. Consider "if you believe P, and you also believe that P implies Q, then you *should* infer that Q." This corresponds to the descriptive claim "if P, and if P implies Q, then Q."

Similarly, “if you borrow money from x, then *ceteris paribus* you *should* re-pay x” corresponds to the descriptive claim “if you borrow money from x, then *ceteris paribus* you have an obligation to re-pay x.”

[86] Aquinas writes:

“As Augustine says, that which is not just seems to be no law at all. Hence the force of a law depends on the extent of its justice…Every human law has just so much of the nature of law as it is derived from the law of nature. But if in any point it departs from the law of nature, it is no longer a law but a perversion of law (*Summa Theologica,* I-II, q. 95).” Quoted in Golding (1975: 30).

See also Sabine (1937: 221-225), and Murphy and Coleman (1990: 15).

[87]

Quoted in Murphy and Coleman (1990: 14). See also Murphy and Coleman (1990: 15-19) for documentation of philosophical dissatisfaction with the Aquinas’ view.

[88] See Hart (1977: 30).

[89]

Strictly speaking, this needs to be qualified. As we will soon discuss, terms like “immoral” and “immorality” are systematically ambiguous, complicating the assessment of Aquinas’ position. Aquinas is in fact wrong. But if we ignore the ambiguities just mentioned, we are likely to draw exaggerated conclusions as to the moral status of law from the falsity of Aquinas’ position. Indeed, the doctrines of Austin and Hart seem to embody just these sorts of inferences, as I hope to show.

[90]

See Austin (1832) and Austin (1890). See also Hart (1977: 19). Max Weber's sociological views embody this conception of government. See Weber (1966). As we might expect, Max Weber has a rather Austinian conception of law (though it is rather more sophisticated than Austin's: a matter that lies outside the scope of the present work).

[91]

Hart (1961), Dworkin (1977: 39-40).

[92]

See Murphy and Coleman (1990: 21).

[93]

Hart (1977: 20-21), Dworkin (1977: 39-40).

[94]

See Searle (1965) and (1995). Searle's important distinction between "constitutive" and "regulative" rules is relevant here. A "regulative" rule regulates some pre-existing activity. A "constitutive" rule *creates* or *defines* a new activity. The rules of chess are constitutive rules: they define the game of chess; there was no chess before they came along. By contrast, the laws forbidding insider trading are regulative rules: they regulate existing commercial activity (the selling of stocks).

The laws prohibiting murder are also rules in the "regulative" sense: they don't define a new activity; they regulate existing activity. By contrast, the laws conferring legal status on certain partnerships are rules in the "constitutive" sense; they define new modalities of activity.

The command-theory of law cannot handle the existence of laws of the latter kind. I believe that the distinction between constitutive and regulative rules has important implications for the nature of law. To my knowledge, these implications have not been delineated.

[95]

Hart (1961), Hart (1977: 18-19), Dworkin (1977: 40-41). A position that would meet with wider acceptance would be to say that, if the law requires me to do something, I have a *prima facie* obligation to do it, whereas I do not even have a *prima facie* obligation to do what the gunman says.

But it is still a matter of debate whether one even has a *prima facie* obligation to obey the law. It would seem, at least at first, to depend on the overall moral character of the legal system in question and on the consequences of weakening or obliterating that system for the protection of the relevant community's moral rights.

By way of anticipation, I will argue that, while one does always have an obligation -- not merely a *prima facie*, but an actual, obligation -- to obey the law, that obligation is often outweighed by an obligation to break the law.

[96]

Hart (1977: 19). See also Golding (1975: 27-29).

[97]

Some have said that the Mafia *was* the government in Sicily at certain times. Though we will not explicitly address this particular historical question, we will make it clear what such a claim amounts to.

[98]

For obvious reasons, use of Weber's original thesis in this context would not only *not* help us free our analysis of vicious circularity, but would *add* additional sources of circularity. Weber's original thesis is, itself, a paradigm of viciously circularity. See the next endnote.

[99]

Weber's *actual* thesis is that the government is the thing which has a "monopoly on the *authority* to use force" (my emphasis). See Weber (1947).

There is a rather obvious problem with this view. What is meant by the

term “authority”? If what is meant is *governmental* authority, then Weber’s analysis becomes trivial and circular: “the government is that thing which has a monopoly on the use of governmental authority to use force.”

If what is meant is *legal* authority, then Weber’s analysis once again becomes circular. A *legal* right is one that is granted by the government (though, as we will see, not all rights granted by the government are legal rights). In any case, the concept of legality presupposes that of government. So we cannot non-circularly analyze the concept of government in terms of that of legal authority.

If what is meant is *moral* authority, then Weber’s thesis is simply wrong. The government is not always, or even *ever*, the thing which has a monopoly on the *moral* authority to use force. To have such a monopoly would be to be the sole thing which has the right to use force. But in any state there are cases where a private citizen has the moral right to use force, e.g. to defend himself against other citizens or against unjust acts on the part of the government.

To sum up: if by “authority” Weber means governmental or legal or moral authority, then Weber’s analysis becomes trivial or false. In this context, I don’t see what other sort of authority might be relevant. So it seems that if Weber’s analysis is to stand any chance of being correct, then we must get rid of the term “authority”, so that it becomes simply: “the government is the thing which has a monopoly on the use of force.” This thesis is defensible (though false, as we will see). I believe that it was held by John Austin, among others. In fact, I believe that Weber*’s thesis permeates thought on government in one form or another.

[100]

Weber (1966), Tucker (1981). In some of his works, Robert Tucker references with approval Weber’s thesis that the government is the entity which has a “monopoly on the legitimate use of force within a given area”

(Tucker 1981: 9). But it seems to me that Tucker's approval is *actually* directed more towards what I have called Weber*'s thesis than towards Weber's actual thesis. The concept of a monopoly of power has an important role in Tucker's work; the concept of *legitimate* power, and of the *authority* to use force, have a more questionable involvement.

It is worth pointing out that much of Tucker's work as a political scientist concerns Stalin's regime. (Tucker has written two biographies of Stalin.) Given this, it is not hard to see why Tucker would tend to see governments as nothing other than monopolizers of force. Nonetheless, for reasons I am now about to discuss, I do not think that Weber*'s position is correct; and, for the reasons given in the previous footnote, I think that Weber's actual thesis is an empty tautology.

It might also be worth pointing out that Tucker's focus as a political scientist is Marxism, along with its political aftermath. Marx cynically thought of laws as manipulations on the part of the ruling class to consolidate existing class-divisions and thereby to retain power. Although Tucker himself is not a Marxist, and is in fact politically conservative, it may be that Marx's conception of law has rubbed off on him.

[101] At any rate, there is a strong argument to be made for this, even though we will discuss some ways of resisting that argument, and therewith of reviving (L2). But *regardless* of whether that argument is cogent or not, (L5) is the correct analysis of law. This is because, for reasons that we will soon discuss, (L2) is correct iff it is equivalent with (L5).

[102] Aficionados of the *Philosophical Investigations* will recognize the stamp of Wittgenstein's thought on this discussion.

[103]

Actually, so-called "legal realists" would say otherwise. These days, legal

realism is generally rejected – rightly, I believe. See Hart (1977: 22) for a brief discussion of the problems with legal realism.

[104]

Hart (1961) makes similar points.

[105]

See Fuller (1964: chapter 2) for a compelling defense of the position that laws must be public.

[106]

Fuller (1964: chapter 2) clearly and forcefully argues, on grounds similar to those here given, that laws must be public – that a private law is a contradiction in terms.

[107]

See Fuller (1964).

[108]

Fuller (1964) makes it clear that laws must be public, and that private resolutions do not qualify as laws.

[109]

Actually, it turned out that the “proof” was flawed, and wasn’t really a proof after all. See Bohm (1957) and Bell (1987).

[110]

There is one other nuance to attend to. Not every assurance publicly made by a politician, even an elected official, corresponds to a law. (When George Herbert Walker Bush said on television: “read my lips: no new taxes”, he obviously wasn’t enacting a law.) This may induce some to reject (L5). But such a reaction would involve a couple of confusions.

According to (L5), a law is an assurance of a protection made by the main guarantor of rights in the area. In any country that is not a pure autocracy, no *one* person is the main guarantor of rights. Even if the most powerful person

in the land makes such an assurance, unless he is speaking *on behalf of* the main guarantor of rights in the land, that assurance does not constitute a law. So given only that George W. Bush cannot make a law just by issuing a televised assurance of some protection, it doesn't follow that (L5) is false. George W. Bush is not the main guarantor of rights in this country. No one person has that honor.

But this might be seen as leading to vicious circularity in our account. Presumably so and so is speaking "on behalf of" the government iff so and so's statement has the force of law (or some such). Thus (L5) would seem to be viciously circular, as it collapses into:

(L6) L defines a legal obligation for X, i.e. is a law for X, exactly if the following conditions are met. L is an explicit and public *assurance* of a protection of somebody's right (not necessarily X himself), and the entity providing that assurance is the main guarantor of rights in the area. Furthermore, the assurance in question must be made by somebody (or something) who is speaking on behalf of said guarantor of rights and, therefore, *with the force of law*.

But (L5) does not collapse into (L6). The circularity in question is merely apparent.

As we saw, it is of the essence of law that everybody *know*, or readily be able to know, what the law is. So there must be ways of clearly distinguishing assurances that are laws from those that are not. The issuing laws must therefore be governed by rigidly followed, hyper-standardized conventions. Otherwise, it simply won't be clear whether the "law" in question is a counterfeit or the real McCoy. But those conventions are there only to *indicate* whether something has the property of being a genuine law: being

issued in accordance with such conventions is not what *makes* something be a law.

It may be helpful to pursue the analogy with currency. The government goes to great lengths to ensure that the bills it issues have distinctive markings – markings that counterfeiters would have difficulty replicating. But possession of such markings is not what *makes* a piece of paper be legal tender. If a counterfeiter were to perfectly replicate those markings, the resulting piece of paper would not be legal tender. If the official mint of the U.S. government were to go haywire one day, and fail to make those markings, the resulting pieces of paper *would* be legal tender. What makes something be legal tender is that it has the right pedigree. The markings in question are excellent *evidence* that a piece of paper has the right pedigree. But they obviously don't *make* a piece of paper have that pedigree or, consequently, qualify as legal tender.

Any assurance of a protection that *clearly* has the right pedigree is *ipso facto* a law. The aforementioned conventions are often needed to make it clear that an assurance has the right pedigree. But if the usual protocols were not needed to make that clear, then they would drop out. Those protocols do not *give* legal standing to an assurance: they simply indicate an *existing* legal standing. A law is an assurance that *in fact* is issued by the government. When an assurance is issued by an official who is speaking *on behalf* of the government, that merely *indicates* that the assurance *already* has legal standing: it is in no way part of what makes that assurance count as law. So (L5) is adequate as it stands; there is no need to replace it with the viciously circular (L6).

[111] As we've noted, one of the standard criticisms of the command-theory of law is that it doesn't fit with the constructive aspects of law. A

piece of money is a legal certificate. But it is unclear what sort of *command* corresponds to a dollar bill. Money lifts restrictions. Commands impose restrictions. Our theory can accommodate this. A dollar bill is a legal entitlement to certain things. Each person has *a* moral entitlement to everything she wants. That entitlement is often *outweighed* by the moral entitlements that other people have. But, for reasons we've seen, "outweighed" does not mean "obliterated." Possession of a dollar bill is a partial protection of one's moral entitlements. The same point holds with regard to the privileges afforded by the law to those who form business or marital partnerships.

It is worth ending this section by highlighting a linguistic fact known to every philosopher of law. In many languages, the words for "law" and "right" are identical. See Hart (1977). For example, the German "Recht" means both *law* and *right*, the same being true of the French "droit". Of course, French and German are both indo-European languages (though they are in different sub-branches: French is a Romance language, German is a Germanic language). So it is possible that we are dealing with a "verbal coincidence" which has been historically transmitted down the various tributaries of the indo-European family-tree. But it seems more likely that this linguistic fact embodies our recognition of an internal relation between these two concepts. It is worth asking whether, in any non-indo-European languages, there is one word for law and right. My guess is that this is indeed the case.

[112] In an essay called "Thought and Talk", Donald Davidson uses the expression "a definite but feeble intuition". I am borrowing his excellent turn of phrase. Davidson is discussing a topic that is completely unrelated to what we are discussing.

[113]

The term “explicit” in this context isn’t to be taken as strictly as it is taken in logico-semantic contexts. In such contexts, the term “explicit”, as applied to an expression x, means that what x is intended to convey is completely represented by x’s literal meaning, or that there is some kind of strict correspondence between x’s grammatical and syntactic forms. But in the context of a treatise on the concepts of law and government, that delineation of the term “explicit” is irrelevant, being much too strict. Here to characterize x as “explicit” means only that it satisfies a garden-variety, common-sense conception of straightforwardness and non-circuitousness.

[114]

There was a Saturday Night Live skit involving a fictitious car called an “Adobe.”

[115] I believe that, ultimately, what is promised is never an object, but rather the truth of a proposition. What I promise Bob is not simply *tickets to the game* but rather *that I will give him tickets to the game* or *that he will have tickets to the game* or some such. But we needn’t worry about this nuance here.

[116]

Russell (1905).

[117]

Quine (1966).

[118]

Permit me to use the standard jargon for a moment. Because the “promise”-operator (so to speak) in (2) has *narrow-scope*, any promise issued by somebody uttering (2) would be *de re* with respect to some actual robot chef and some actual Plymouth Adobe purchaser. In other words, the very existence of that promise would *presuppose* the existence of such individuals.

So given that there are no such individuals, no promise has been made. See Burge (1977) and Quine (1966) and Strawson (1950).

[119]

See Strawson (1950) for a closely related point.

[120]

Permit me to use the jargon once again. The "promise"-operator in (3) has *wide*-scope. Consequently, given only that there is no actual robot chef or Plymouth Adobe purchaser, it does *not* follow that (3) has failed to issue a promise. It is clear that (3) *is* a promise. It is what I will (somewhat inaccurately) refer to as a *wide-scope assurance* (or promise).

[121] Of course, certain expressions – e.g. "I promise", "I am warning you", "I urge you" -- marginalize the role that such background information has in the way of determining the performative status of a sentence-token. If I say "I promise to pick you up tomorrow", then context has to do less to make my words constitute a promise than it would if I had just said "I will be there tomorrow." But even when one's words are weighted with such locutions, the role of background information is by no means null. If I say "I promise…", it doesn't necessarily follow that I have made a promise. Given the context, it could be clear that I am telling a joke or play-acting.

[122]

In the *Philosophical Investigations,* Wittgenstein lays great stress on the difficulties facing attempts to locate the *psychological* basis of social acts, like promise-making.

[123] In fact, those words must have mental causes of a rather specific kind. Suppose that, because of some strange, Tourette's-like mechanism, I involuntarily produce those words whenever I have a certain kind of ticklish

sensation. In that case, my producing those sounds would not constitute a promise, even though they had a mental cause.

Nonetheless, it does seem fair to say that a promise must have *some* kind of mental basis. (It seems that it must be caused, at least in part, by a desire to make you think that I am going to pick you up. See Searle (1970) for an attempt to state with some precision what the promise-maker's mental contents must be.) At the same time, an utterance qualifies as a promise in virtue of *at least in part* the social characteristics of the situation in which it arises.

[124]

At least typically. I leave it open whether promises could be made without the mediation of some sort of language or system of symbolism – whether, for example, there could be telepathic promises.

I also leave it open whether promises that initially appear not to use any symbolism *do* in fact do so. For example, consider a promise made with a knowing look. Our initial reaction is to say that no symbolism was involved. But one could argue that, in actuality, the look does involve some kind of language. Perhaps the language involved is of a very primitive, pre-verbal sort. (I believe that this is held by some psychologists – that the looks through which we communicate are a pre-verbal form, but nonetheless linguistic, form of communication.) Or perhaps it is an *ad hoc* system of symbolism -- such as agreeing with one's crime-partner to cough loudly if one sees the security guard coming.

[125]

Actually, the fact that these words express a promise is *doubly* dependent on context. First, the fact that those words have for their literal meaning the proposition *I am going to pick you up tomorrow*, is obviously a function of the socio-historical context. But *given* that those words have that literal

meaning, the fact that they are taken as a promise – as opposed to, say, a threat or a simple statement of fact – depends on facts about a narrower context: a context *within* the aforementioned socio-historical context. So there is a semantic and also a pragmatic dependence on context.

[126]

I am, of course, alluding to the tragic incident in 1984.

[127]

In reality, such a decision would probably be made by a number of people. But it will simplify exposition if we telescope all of these individuals into some one individual.

[128]

There is another point I wish to highlight. Notice that the *wide-scope* promise made (in our fictitious scenario) by Union Carbide to the people of Bhopal was made in an entirely mechanical fashion. A computer mechanically responded to ink-marks on a piece of a paper, and it then mechanically sent a certain message to somebody at a wire-service. Virtually no mental activity was involved. (Suppose that the message was not read by a living news anchor, but by a mindless robot that looked and sounded like a news anchor. In that case, no mental activity at all was involved.)

But Smith's *narrow-scope* promise is not comparably mechanical. It was made by a person, or a group of people, who knew what it was they were promising. In general, when institutions mechanically make wide-scope promises, that is, it seems to me, parasitic on their *non*-mechanically making narrow-scope promises.

This is not to say that narrow-scope promises cannot be made mechanically. Suppose that the computer just described sends a message to the wire-service saying: "We at Union Carbide promise that anyone born on Bermuda after the year 2125 who is harmed by our corporation will be given

stock-options in our corporation." In that case, Union Carbide is promising that *anything* x satisfying a certain description (*thing x such that x is born on Bermuda after 2125 and is harmed by Union Carbide*) be given stock options. But suppose that nobody is born on Bermuda after 2125 -- because, let us suppose, it has been converted into a nuclear testing site. In that case, there is no x such that Union Carbide is promising anything to x.

In fact, even if somebody x is born in Bermuda after that time and is harmed by Union Carbide, it still isn't entirely clear whether Union Carbide is making any promises to x specifically.

[129] In fact, it is possible that no member of Congress even knew of the existence of that *kind* of assurance. The existence of that assurance seems to presuppose the existence of that sort of sexual orientation (or at least of the *property* of having that orientation, though not necessarily of the instantiatedness of that property). Given this, suppose that, by some fluke, none of the members of Congress even knew of the existence of people of that sexual orientation. In that case, Congress as a whole has made an assurance such that no member of Congress even knows of the existence of that *kind* of assurance, let alone that there has been a specific instance of that kind.

[130]

Of course, laws are not always enacted by the U.S. Congress. Sometimes they are issued by a brutal dictator, sometimes by a benevolent town-mayor. But the points we make in connection with Congress are easily transposed to these other contexts.

[131]

This connects with some points we were discussing earlier. According to Dworkin, one of the tenets of legal positivism is that laws have a certain

pedigree. The conceit seems to be that morality has nothing to do with pedigree – after all, the morally empty army regulations relating to boot-polish may have the right *pedigree.*

As we discussed, given only that something must have a certain pedigree to be a law, it doesn't follow that laws can be morally empty. The opposition described in the last paragraph is spurious. Our discussion of promises gives us new evidence of this. Suppose that I promise to pick you up at the airport tomorrow, and that I do this by saying "I will be at the airport tomorrow." (So the promise is conveyed through implicature: I *don't* say "I promise to…" or even "I will pick you up at the airport.") The reason that utterance constitutes a promise is precisely that it has a certain pedigree. First of all, its cause is a person who had certain mental states. The act of my making those noises, and the mental states involved in that act, are embedded in a certain situation – one defined by my friendship with you, and the understandings and expectations involved therein. That friendship, in its turn, is embedded in a certain culture – a culture characterized by a certain language, mores, and ethical precepts. So my words qualify as a promise precisely because they *originate* from this sort of situation – because they have a certain pedigree. Their having this pedigree is not only *not* incompatible with their having a moral dimension – with their imposing certain moral obligations on my person, and with their very occurrence having moral properties – but is precisely what *gives* them that moral dimension.

The fact that *Congress,* as opposed to an inanimate malfunctioning robot, performs a certain act is what *gives* that act a certain moral profile: it is what makes that act be an assurance of a certain kind, and what gives people the right to be disappointed if certain events don't follow that act, and so on. So the idea that laws are individuated by their pedigrees is in no way incompatible with the idea that laws have a moral dimension. What *gives* a

speech-act – or, for that matter, any act – its moral profile is, at least in part, its mode of origination.

It is true that nothing *is* a moral verity in virtue of its mode of origination. (Moral verities are given by propositions, which are abstract objects. So talk of "modes of origination" is out of the question here.) But it doesn't follow that nothing *embodies* a moral verity in virtue of its mode of origination. The reason that my showing up at the airport at a certain time is a *moral* act is precisely that it has a certain pedigree: it was caused by an intention to honor my promise, along with a general intention to do good. If that event had been caused by my intention to steal your luggage, then it would not have been moral. So the fact that my act is moral *does* consist, at least in part, in its having a certain mode of origination.

[132]

Of course, it isn't just *any*body x who is assured such a protection. It is anyone x meeting certain conditions relating to place of residence (she must live in the geographical area associated with the United States), biology (she must be a human being and not a tree or a squirrel); and there are probably other requirements that she must meet. But to avoid verbosity, I will omit these qualifications, and use the unrestricted quantifier "for anyone x" as a stand-in for the non-elliptical quantifier just described, this being: "for any one x, meeting such and such requirements regarding place of residence, species-membership…".

[133]

See also Quine (1966) for the classic modern discussion of the interactions between epistemic and other operators.

[134]

Frege (1884: 60).

[135]

Frege (1884: 60-61) cryptically expresses this by saying that "all whales are mammals" doesn't say anything about whales, but only about the *concept* whale. Understandably, this statement is likely to meet with resistance. But there is quite definitely *a* reading on which it is correct. As we will now see, there is no x such that "all whales are mammals" says that x is a mammal. *A fortiori* there are no whales $x_1 \ldots x_n$ such that, for any i, "all whales are mammals" says that x_i is a mammal. So on the *wide*-scope (de re) reading, "all whales are mammals" doesn't say anything about whales.

At the same time, "all whales are mammals" *does* say that, for any whale x, x is a mammal. So on the narrow scope (de dicto) reading, "all whales are mammals" says something about whales. But Frege's point seems to be that to say that a statement is *de dicto* about x is really just to say that it is about some *concept* that x falls under, and not, strictly speaking, about x itself.

And there is plainly a sense in which that is correct. The statement "Socrates is bald" is *about* Socrates in a far stricter sense than the statement "someone x was uniquely a great Athenian hemlock-drinking philosopher, and x was bald." The second, we are told, is *de dicto* "about" Socrates. Now the first is true in a world W exactly if, in W, Socrates is bald. The second can be true in a world W where Socrates is not bald and false in a world W where Socrates is bald. Now it seems uncontroversial that the first sentence really is *about* Socrates. It also seems that would almost render the term "about" ambiguous if we said that the second was also "about" Socrates: after all, the relationship that Socrates' baldness (or non-baldness) has to the truth of the one is entirely different from that that which has to the truth of the other. So it would, at the very least, by misleading to say that the second one was "about" Socrates.

To some extent, we head off any misleading suggestions by saying that

the sentence “someone x was uniquely a great Athenian hemlock-drinking philosopher, and x was bald” is *de dicto*, but not *de re*, “about” Socrates. But this only raises the question whether *de dicto* aboutness and *de re* aboutness are two different sub-species of the same species. There is no doubt that so-called *de re* aboutness really is a form of aboutness. But given how different so-called *de dicto* aboutness is from *de re* aboutness, it is to be doubted whether the former is but a different variety of the sort of thing that the latter is. Consequently, given only that “all whales are mammals” is *de dicto*, but not de re, about Shamu – and only *de dicto* about any other whale that exists – we can see the reasonableness of Frege’s point that it is not *really* about whales, but only about the concept of whale-hood.

[136]

But by that reasoning “all whales are mammals” is about Socrates no less than Shamu. After all, Socrates, no less than Shamu, satisfies the condition of being a thing x such that *if* x is a whale, *then* x is a mammal. In fact, *everything* satisfies that condition. So if we insist on saying that “all whales are mammals” is *about* Shamu specifically, then we must say that it is *about* Socrates and my desk and the planet Neptune. But in that case, the concept of “aboutness” has been stretched to nothingness.

[137]

Russell (1920: 168) argues on similar grounds that “a man” doesn’t refer to any particular man.

[138]

Even though it may *necessarily* be the case that Shamu is a whale – even though there is no (metaphysically possible) world where he is a non-whale – it is still an a posteriori, empirical fact that he is a whale. Though false, the statement “Shamu is a fish that looks like a whale” is not self-contradictory, and can only be refuted on empirical grounds. See Kripke (1972) and Putnam

(1975).

[139]

There are probably several such laws. But in this context, we may pretend that there is only one.

[140]

The fact that there are no handicapped people in W surely wouldn't make it impossible for lawmakers in W to create HC. They want HC to be ready to protect anyone who should have the misfortune of becoming handicapped, as well as the unborn handicapped. So our description of W – specifically, our statement that HC exists in that world – is *not* question-begging.

[141]

One might respond by saying that laws like HC protect the rights of everyone, even those who are not handicapped. But that seems to be a way of saying that, if a previously healthy person *does* become handicapped, *then* HC is going to protect his rights. It might also be a way of saying that HC protects the non-handicapped by protecting their handicapped loved ones. In this context, these considerations are irrelevant, given that nobody in W is handicapped or is going to become handicapped.

[142]

In fact, in W, HC encroaches on that right. If you desperately need food from store X, and the only free space is the handicapped spot, then HC thwarts your right to have access to food.

Of course, the same is true of HC in our world. But in our world, HC's encroachment on the rights of the non-handicapped is outweighed by its protection of the rights of the handicapped, so that its overall effect on human rights is positive. In W, HC's overall effect is negative. (Except, perhaps, in some extremely indirect sense. For example, the mere existence of HC in W might reinforce a general attitude of kindliness and tolerance, granting that it

doesn't provide any specific person with any concrete protection of anything.)

[143]

In response to this, one could *conceivably* say that, when an institution mechanically produces some so-called assurance, no assurance has *in actuality* been provided: it just *seems* otherwise. So conceivably one could say that the earlier discussed so-called assurance provided by Union Carbide (the one generated by the computer) wasn't really an assurance at all, just a simulated assurance. (And, of course, one might say the very same about GS.) To think otherwise, we will be told, is to be guilty of the "pathetic fallacy" – of thinking that if the wind-blows and makes the noise "I promise to give you prosperity", an *actual* promise has been made; of thinking that, when a computer turns on and displays a smiley face on its monitor, the computer is *actually* welcoming you.

I feel that this line of thought conflicts with the data and can be defended only by making the *a priori* assumption – earlier discussed and discredited -- that facts about institutions are in lock-step with facts about the individuals composing them. In connection with the topic of government, we've seen that an institution can be morally good even though the same is not true of the people composing it. The objectives of an institution can (fortunately) be very different from those of the individuals involved. So we've seen little reason to believe, and every reason to *dis*believe, that there is any straightforward reduction of institutional to individual facts.

I do not deny – in fact, I firmly believe – that facts about institutions *supervene* on facts about individuals. In other words, where there is no difference at the level of individual facts, there can be no difference at the level of institutional facts. But this doesn't warrant the conclusion that institutions can do only what the individuals composing them do, as I will

now argue.

Supose that X is an object having a temperature of 78°, and that X consists of particles $x_1 \ldots x_n$. Now let X* be some other object that consist of particles $x^*_1 \ldots x^*_n$, where $x^*_1 \ldots x^*_n$ are, as regards their interactions and also their intrinsic properties, qualitatively identical with $x_1 \ldots x_n$. It is true that X* must also have a temperature of 78°. So facts about the temperature of an object certainly *supervene* on facts about the particles composing it. But it does not follow, and is in fact categorically *false*, that objects have the same temperatures as the particles composing them. Temperature is an aggregative property; so the concept of temperature doesn't apply to individual particles.

Similarly, given only that the objectives of an institution supervene on facts about the individuals composing that institution, it doesn't follow that institutions always have *the same* objectives as those individuals. Not only does this not follow; it seems seldom to be the case where institutions of above a certain size are concerned.

Our strong pre-theoretic presumption is that, when the computer at Union Carbide mechanically hammers out a message (with the words "We at Union Carbide promise $1,000,000 to the people harmed by this latest tragedy at…") that it then transmits to the newscaster, Union Carbide has locked itself into some kind of commitment to the people that it has harmed – that it has made some kind of promise. Because the behavior of the computer was – though mechanical – adequately rooted in company policy, and wasn't some freak accident, we feel that Union Carbide *has made a promise,* even though every employee of that corporation was asleep at the time.

The situation would be different if the computer's behavior were not adequately rooted in the policy, and the standard operating procedures, of Union Carbide. Suppose that the computer's behavior was not grounded in any company policy. In fact, that the computer is not *supposed* to send

messages of the sort described earlier ("We at Union Carbide are profusely sorry about…"). Its *actual* purpose is to alphabetize company documents. But one day, because of a random quantum event, the computer transmits a message of that kind to Bloomberg. In *that* case, perhaps, Union Carbide *wasn't* making any sort of assurance – just as I wouldn't *really* be making a promise, if during an epileptic fit, I involuntarily made the sounds "I promise to drive you to the airport tomorrow."

But, of course, this is not what is going on in the scenario described earlier. In that context, the computer's behavior, though mindless (at least as regards its proximal causes), was sufficiently rooted in company policy and procedures that its behavior *does* register with our intuitions as being an act of Union Carbide's making a promise to those whose its malfeasance has harmed.

Now it was indeed through a freak accident that the "bill" corresponding to GS was placed on the Senator's desk. But the ensuing process –the congressional apparatus that turned GS into law -- was not a freak (granting that the members of Congress involved may have been on "auto-pilot" during the process, and thus failed to look at the bill with due care). The mechanism involved is a deeply entrenched one. In fact, it is precisely *because* that process was so deeply rooted in a stable mechanism that it was able to go through with so little assistance on the part of human intentionality. So leaving aside the very first instant of its existence, it was specifically the *non*-freakish nature of the processes that made it law that enabled GS to come into existence.

Let us close the argument. As we saw, there are good reasons to believe – and, it seems, only a dubious *a priori* reason *not* to believe – that Union Carbide's computer-based message *does* constitute a real promise on the part of that corporation. Even though that message was made mechanically, it was

still made in the right way: the computer's behavior in that instance was adequately rooted in the policy and general function of the company. I don't see how GS is relevantly different or, therefore, why GS doesn't qualify as an assurance.

Permit me to substantiate this last point – which, I grant, ought not to be accepted without further ado. Let RD be some law that is made very deliberately and non-mechanically by members of Congress. So RD is, we may suppose, a law that doesn't allow a certain form of racial discrimination to occur. All of the members of Congress feel that RD is, for reasons of morality and social policy, a necessary law to have. Accordingly, they consciously convert it into law.

It is very natural to regard RD as a kind of commitment, on the part of the government, to protect those people x whose rights would be violated by certain forms of racial discrimination. I don't think that there is too much trouble reconciling RD with (L5).

Let us look at GS in this light. When GS comes into existence, it does so through the very process that brought RD into existence. Of course, the process that *created* the bill that was converted into RD is different from the process that created that bill that was converted into GS. But in both cases, the conversion-process – the process *of turning the piece of paper into an actual law* – is, in the relevant respects, the same. Where the one is concerned, that process happened very mechanically; the relevant members of Congress were on "auto-pilot", as we put it. Where the other was concerned, those same members of Congress were very much aware of what was happening. But in both cases, the same process, the same mechanism, was involved. Of course, in the one case, but not the other, various members of Congress made themselves a part of that process. But there is no denying that the same legislative apparatus is at work in both cases.

Where GS is concerned, *the process* was not circumvented. Given that RD is a promise (an assurance of protection of a right), we would, I think, be guilty of theoretical arbitrariness if we were to say that the same wasn't true of GS.

Here is an example of a case *the process* has been circumvented – a case where GS really *isn't* an assurance and where it (also) isn't a law. Congress never passes GS. The so-called bill never lands on the Senator's desk (or if it does, the Senator looks at it and then throws it in the waste-basket). But one day the official printing press of the U.S. government, which publishes all the law books, goes haywire and randomly deposits GS (or the ink-marks corresponding thereto) in the law books. Consequently, many a judge and lawyer and police officer *thinks* that there is now some strange law forbidding people from wearing green shirts indoors.

Here, it seems, we really do have a freak accident; and no assurance was made by Congress. But, by the same token, I don't think that, in this case, anyone would say that GS was a law. Most would say that, because of printing-error, people mistakenly *thought* that it was a law. So even though, in the case of this sort of freak accident, we don't have an assurance, we also don't have a law. So this scenario doesn't give us any reason to believe that laws – even mindlessly made ones – are anything other than assurances. [144]

This line of thought gives us new insight into a problem that any moralistic analysis of law must face (a problem that we have already discussed from a few different angles): the existence of laws that seem to be mere orders issued by an immoral dictator.

We've noted many times that any government, even the most brutal and autocratic, are ultimately providers of basic protections. Of course, once something establishes itself as a government by providing certain minimum

protections, it can then exploit the dependence that people subsequently have on it to do evil towards those very people. This creates the illusion that the government in question is *fundamentally* antagonistic towards its own subjects.

But this is to misunderstand the peculiarly insidious nature of governmental evil. People depend on governments – for protection from their fellow citizens and foreign invaders, for roads, hospitals, and so forth. When a government is evil, it takes advantage of that dependence. So *governmental* evil presupposes the existence of a protective relationship between government and subject.

That is what makes governmental evil such an abomination. If someone on whom you depend -- a parent or best friend – uses that dependence to harm you, a worse evil has been committed than if the same harm is done to you by somebody on whom you do not depend. In such a case, there has been a betrayal, in addition to a harm.

Let us start with a historical fact -- or, rather, what might initially appear to be a historical fact but is actually, I will argue, a tautology. When governments are evil towards their subjects, it is always done under the pretext of legitimate governmental concern. There is always an attempt to assimilate acts of what are, in fact, political repression to exercises of legitimate governmental functions. The political rival becomes a “criminal” or a “traitor.” The legal code is often broadened to legitimate the arrest (and subsequent torture…) of this “criminal.” In general, when governments misbehave towards their own constituencies, that misconduct always involves a misuse of some protective mechanism: the legal system, the police, the military.

This is not, I believe, *solely* because it is a Machiavellian expedient for governments to cloak misconduct under the banner of morality and legality.

It is, I will now argue, because such pretexts are necessary for such misconduct to be acts of *government* at all.

Let us illustrate this with a story. Government G, which is run by despot Jones, commits some atrocity. Under Jones' leadership, G kills a much-beloved dissident, named Brown, whose only crime is attempting to make life better for his fellow citizens. Instead of attempting to legitimate his act, Jones admits on national television that he killed Brown simply because he didn't like him. There is no pretext of legality or of governmental concern for the welfare of populace. Jones simply says: "I did it because I felt like it. It was evil. It was illegal. It was not at all government-like. It isn't going to do any of you any good. But I still did it, and I'm glad I did so."

What we have here, it seems to me, is not a case of *governmental* evil. We have an evil act on the part of an individual. That individual happens to have control over the instruments of government, and he used those instruments to commit that act. But, in the end, we have a case where a person, not a government, is being evil.

For there to be an act of *governmental* evil, the act in question has to be done in the name of government: the protective functions characteristic of government must be woven into that act. This can be done through some ludicrous extension of the legal system or through some obviously false appeal to national security. There is no limit to how thin the pretext may be. But until there is *some* such pretext, we don't have a *governmental* act; we just have a private individual who has governmental power.

An act of governmental evil is necessarily one that involves a *mis-use of the protective mechanisms characteristic of government* (e.g. the legal system, the police, the bureaucracy).

Governmental evil is mis-protection. An evil government is not like a robber who barges into a house and kills its occupants. An evil government is

like a mother who, in her effort to keep her infant warm, puts so many blankets on it that it suffocates or like a nurse who kills her patients by giving them too much medication.

Here is a comparison. (Actually, it is a special case of the phenomenon we are discussing.) Smith is a police officer. One day, while not on duty, Smith beats Rutherford. Here we don't have a case of *police*-misconduct. We just have a case of misconduct on the part of somebody who happens to be a police officer.

But to be guilty of *police*-misconduct, it is not enough that Smith misbehave *while* on duty. His police-powers must be implicated in a certain way. If, while on duty, Smith steals a bicycle, he is not guilty of *police*-misconduct. But if he then tries to legitimate that theft by saying that he had to commandeer the bicycle to capture a fugitive, or if he tries to arrest somebody who tries to prevent his act of theft, *then* he is guilty of police-misconduct.

For there to be police misconduct, it is not enough that Smith misbehave *while* on duty: he must do so in *the name of duty.* Police are supposed to enforce laws. Police-misconduct is *mis-enforcement*. It is not garden-variety misconduct that happens to be performed by a police officer.

Similarly, for there to be a case of *governmental* evil, it is not enough that somebody who happens to work for the government (even if it is the dictator himself) do something evil. The act in question must be an *act* of government. Since the business of government is to protect, the act must be one of protection, i.e. of *mis*-protection. It must be an over-extension or misplacement of the protective functions constitutive of government.

With the possible exception of Pol Pot's Cambodia, the Soviet Union under Stalin came as close as any society in history to being a state whose government was a gunman, and whose citizens were the gunman's hostages.

On every single occasion that Stalin publicly made a decree of some kind, he justified it by saying that it was necessary to protect the citizens of the U.S.S.R.

Here is an example. A number of Soviet soldiers were captured by the Germans and subsequently released. One would think that they would be welcomed as heroes. Stalin ordered that they be sent to the Gulag. His justification was that, because these soldiers had spent so much time in the company of Westerners, their minds had been contaminated by western values and that *in order to protect the purity of Soviet people*, these lost souls would have to be separated from the rest of the population.

Of course, this justification was obviously a Machiavellian ploy on Stalin's part to consolidate his rule. But, as we discussed earlier, that is not the whole story. In order for this nefarious act to be an act of government, as opposed to the act of a free actor who happened to be an officer of the government, it had to be an act of protection. That act was an evil *of government* because it *was* a protection: a mis-protection – like (to modify our comparison) a mother's attempt to nourish her baby by feeding it poisoned baby-food.

A despot's orders are laws because they *are* assurances of protections. When Stalin made it a law that these soldiers be sent to the Gulag, he *was* providing an assurance of protection. He was assuring that, for any Soviet Citizen x, x's right to be free of the decadent Western values would be upheld. (Of course, in actuality, there is no *actual* right had by any *actual* person x such that Stalin's decree *actually* protected that right of x's. But we have already seen why that is irrelevant to its being the sort of assurance that our analysis identifies with laws.)

We noted a little while ago that, when a dictator cynically fabricates some legal pretext for an act of political repression, he thereby converts what

would ordinarily be an act of *personal* evil (on the part of somebody who happened to have governmental power) into one of *governmental* evil. Here we must keep in mind something that we discussed earlier: given somebody who is working within an institutional framework, we cannot understand that act along entirely internalist, psychological lines.

From the viewpoint of the Machiavellian dictator, that pretext *is* just a cynical ploy. But, as we have seen, the dictator's *psychology* is not always determinative of the true nature of the dictator's action. The dictator has no intention of doing anything remotely protective. But the institution for which he works is fundamentally protective (even though it may exploit that protective relationship to commit evils on a scale that would otherwise be impossible). Whether he knows it or not, the dictator commits harm *through* the protective mechanisms characteristic of government: his harms are mis-protections, even though his psychology may to no degree be that of a protector.

Marx famously distinguished between the concepts of "subjective" and "objective" value. (Given my extremely limited knowledge of this sector of philosophical history, I have some reason to believe that this distinction is found in Hegel and was first made by Rousseau in *the Social Contract*.) Suppose that I give money to the poor with the intention of alleviating their suffering. Marx says that *subjectively* I am doing the right thing – after all, I doing my best to help those in need -- but that *objectively* I am doing the wrong thing. By providing the poor with momentary relief, I am reconciling them with a system that enslaves them, and am postponing their making the revolutionary changes that are needed to give their lives true meaning. From a psychological or internalist standpoint, I am moral. But from the viewpoint of larger historical forces – of perhaps (though I may be collapsing two different concepts here) the relevant *program* causes – I am immoral (or, at

any rate, my act realizes some kind of immorality).

I believe that Marx's (or Hegel's or Rousseau's) distinction is of profound importance (though some of his applications of it seem to have been misguided). From a strictly psychological standpoint, the dictator's act is entirely nefarious. *Subjectively* he is entirely immoral. But from the standpoint of larger historical forces, his act is not *as* lacking in moral value. It is an idea worth considering that Marx was making the same distinction as Pettit and Jackson.

In a book called *Search for a Method*, Sartre discusses how historicist theories, like those of Marx and Hegel, can be reconciled with the fact that, at some level, it is always *individual actions* that drive history. His points have a family resemblance with those of Jackson and Pettit.

[145] In response to this, one might say that, strictly speaking, all real cars *do* work – a "car" that doesn't run is a car in name only. But such a proposal doesn't correspond very well to the concept associated with the term "car." It does seem to be a fact that some *cars* don't run. It seems to be a fact that, when a car stops running, it doesn't vanish, to be replaced by a car-like non-car. It seems to be a fact that, when a hitherto non-operational so-called car is fixed, a car doesn't come into existence *ex nihilo*. Of course, we could use the term "car" to denote only operational vehicles of a certain kind. But, in that case, by all appearances, the term "car" wouldn't mean what it currently does.

[146] This particular definition probably isn't adequate. It may be that "car" is indefinable. (That is probably what Jerry Fodor would say.) But in this context, this doesn't matter. All that matters here is that, in some cases, x falls into a category C in virtue of what x is *supposed* to be like, not in virtue of

what x is actually like.

[147]

We might circumvent this by saying that a car is a four-wheel vehicle that has *a* motor, not necessarily one that works. But then the problem we are discussing in connection with the term “car” would be displaced into the term “motor.” A motor cannot be defined as something that *actually* powers something else, since that would render tautologically true the falsehood that all motors work.

[148]

There may be limits to this. If the Ford-factory produces something that has no resemblance to a car – a heap of amorphous metal -- that thing is perhaps not a car. But, arguably, this only proves our point. In that case, the thing bears so little resemblance to a four-wheeled, motorized vehicle that there is no longer a supposition to the effect that it is a car.

[149]

There is a corollary. Given only that something has all the properties of a working car, it doesn’t follow that it is a car. Suppose that a sand-storm on Neptune produced something that is qualitatively just like a working Ford. That thing is not a car. (At the same time, if it were found by humans and put to the appropriate use, then in respect of its being categorizable as a car, its amenability to certain human endeavors would trump its faulty pedigree. Or so it seems to me.)

[150] This doesn’t mean that Stalin himself *intended* for that law to protect anybody from anything. But when we say that that law was “supposed” to protect the Soviet people from counter-revolutionary thought, we are not talking about what *Stalin* wanted that law to do. Given only that such and such is *supposed* to be the case, it does not, at least not obviously, follow that

any one person supposes such and such.

[151]

Except Stalin himself, along with his cadres, and also people who might have benefited purely accidentally (e.g. somebody who is given the promotion because his rival was conveniently denounced). But the fact that these people would have been given a protection does not, by itself, appear to be quite enough to regard the law in question as an assurance.

If carried out, the gunman's threat may benefit certain people – not just the gunman himself, but the people who happened to dislike, or be in competition with, the person whom the gunman kills. But it doesn't follow that the gunman's threat is an assurance of some kind. Obviously it is not.

[152]

Leaving aside the few, probably irrelevant exceptions discussed in the previous footnote.

[153]

See Quine (1962) and (1966). The point is due to Russell (1905), who makes it very clearly. Russell illustrates the distinction in question with an anecdote about a "touchy" yacht-owner. A guest of the yacht-owner says: "I thought your yacht was bigger than it is" – meaning "there is length L such that your yacht has L, and I thought that, for some length L* that is longer than L, your yacht had L*." This statement can be true without there having to be any incoherence on the part of the rude guest.

The yacht-owner responds by saying: "my yacht is not longer than it is" – meaning: "it is not the case that there is some length L and some length L* such that L* is longer than L and such that my yacht has both L* and L." So the yacht owner was cleverly exploiting the ambiguities of scope inherent in the guest's statement, thereby attributing an absurdity to him.

[154] Another example may help. Consider the statement:

(1) “John is supposed to write a novel by November.”

This doesn’t mean:

(2) There is some novel x such that John is supposed to write x by November.

(2) absurdly suggests that a novel exists before its author creates it. The meaning of (1) is:

(3) It is supposed to be the case that: by November, there is some novel x such that John writes (or has written) x.

[155]

Of course, here it will be said that Machiavelli is an exception to this. Maybe so. But if my fading memories of my not very careful reading of that philosopher serve me correctly, Machiavelli’s point is that a ruler must use deception, and the like, *in order* to serve the interests of the governed.

[156]

Wilkins (2006: 197). It should be pointed out that it is unclear whether Wilkins is speaking in his own voice or in the voice of a proponent of a position that he is merely considering. I believe the latter to be the case.

[157]

Here it is worth pointing out a fact about George Orwell’s novel *1984*. The society therein described is the most oppressive society conceivable. In many ways, that society is modeled on Soviet society under Stalin. The protagonist

of *1984*, Winston Smith, says (or thinks to himself) that in his society there are no laws (Orwell 1949: 6):

"The thing he was about to do was open a diary. This was not illegal (nothing was illegal, since there were no longer any laws), but if detected it was reasonably certain that it would be punished by death, or at least by twenty-five years in a forced-labor camp."

It is interesting that he says this, for it would seem that the laws in Oceania were *too* plentiful – that they had grown out of all control.

I believe that what Orwell is suggesting is that there is no law where there is *sheer* repression. For there to be laws, the repression must be but one aspect of a relationship that is, in some respects, protective of rights.

There is another aspect to Smith's (Orwell's) point. In Oceania, the rules that people *actually* had to follow (e.g. "believe all the government's lies") could never be *explicitly* stated by the government. Even though those rules applied to everybody, and had to be known to everybody, there was still a sense in which they were "under-the-table": they could never be openly stated by anyone (in fact, they were not even to be consciously *thought* by one, even though everyone was to have internalized them). The government's authority depended on this conspiracy of silence. In Oceania, the rules that were really operative could never be stated, or even thought, entirely explicitly. Therefore, from our viewpoint – and, I think, from Orwell's – they don't qualify as laws.

[158]

A comparison might help. A turtle's shell is organically connected with the rest of its body. The shell cannot be removed, at least not without harming the turtle. Imagine that a turtle was afflicted by a disease whereby its shell

became so heavy that the turtle could no longer move. The shell would *still* be organically connected to the turtle, and it would *still* perform important functions. At the same time, the shell would also be a terrible encumbrance.

Now imagine a perfectly healthy turtle on whose back some prankster has placed a lead weight so heavy that the turtle cannot move. The lead weight is not organically connected with the turtle: if it were to be removed, nothing but good would follow.

The case of the dictator is comparable to the case of the over-grown, cumbersome shell. The case of the gunman is comparable to the case of the lead weight. A dictatorship is a case of something that has organic roots in the community, and that has some utility, becoming over-developed and of negative utility. The case of a gunman is a case of an encumbrance that has no roots in the group of people thereby encumbered. The fact that there are legal obligations under dictatorships (albeit ones that are often outweighed by conflicting ethical obligations), whereas there are none in the gunman scenario, is an expression of the difference just described. So there are sociological and also ethical differences between the gunman scenario and the brutal-dictator scenario. Of course, the analogy certainly seems to have its limitations. Brutal dictators *can* be removed, whereas a turtle's shell cannot be removed without killing the turtle (so far as I know, given my extremely limited knowledge of veterinary science).

But this very line of thought, taken a step further, actually shows the validity of the analogy. It would be easy to remove a turtle from its shell, but it would be hard to do so without killing the turtle. It was easy to remove Saddam from power, but it is proving to be extremely difficult to create a viable political order to replace his dictatorship: we de-shelled the turtle, but we are finding it very hard to keep the turtle alive. By contrast, when a gunman is taken out by snipers, that is in no way to the detriment of his

hostages: the removal of a gunman doesn't create the sort of vacuum left by the removal of a dictator.

It should also be noted that, where even the most wicked dictatorships are concerned, there are usually *some* people (other than the dictator) who resist the dictator's removal. Once again, Iraq is a case in point. Even now that Saddam has been deposed, there are those would try to reinstate him. This shows that Saddam's regime, though wicked, nonetheless had *some* roots in the needs of the community – granting that those roots were too shallow to warrant describing that regime as morally legitimate.

[159] Not just in the trivial and irrelevant sense that it brings about good things, but in the non-trivial sense that the act *per se* – leaving aside its consequences – is good.

Obviously any act may *lead* to good. In this discussion, when I refer to the "moral" wholesomeness or unwholesomeness of an act, I am leaving aside consideration of what that act may *cause*. Given the right conditions, the most malicious intention of a Hitler or Stalin could in principle *lead to* good. In this context, all of my claims about "wholesomeness, or the lack thereof, must be understood in non-instrumental terms.

[160] This is one of the main themes of Rawls (1971).

[161]

In response to this whole line of thought, one might say something like the following:

"It is not *Smith's* act that is good. What is good is the institutional act which is realized by, but still *distinct* from, Smith's act. It seems to me, in fact, that your entire discussion ignores the distinction between realization

(constitution) and identity. The clay realizes (constitutes) the statue. But, as Bernard Wiggins observed, the clay is not *identical* with the statue, since not every predicate applicable to the one is applicable to the other. (For example, you can destroy the statue without destroying the clay.) So given only that the institutional act is realized by that of the dictator, it doesn't follow that every property of the former – in particular, its goodness – accrues to the latter.

We may go further in this direction. Actions are, tautologically, individuated by the intentions that lie behind them. My typing involves my deliberately making certain movements with my fingers. My making those movements creates many sub-atomic disturbances. But it would be incorrect, or at least misleading, to describe my action as one of creating certain sub-atomic disturbances. That action must be described in terms of my intentions: it is an act of writing a certain word. [See Davidson's "Actions, Reasons, and Causes" in Davidson (1984).] If my intention were to create certain sub-atomic disturbances – if, for example, I were a physicist and my creating such disturbances were part of an experiment I was conduct – then *those very movements* of my fingers would mediate an entirely distinct action from the one I am in fact performing.

When somebody asks "what is x doing?", the answer they wish to be given is couched in terms of x's intentions. If I ask "what is that infant doing?" I want to be told something like "he's trying to draw" or "he's trying to get his mother's attention", not "he's creating such and such sub-atomic occurrences." The fact that questions of the form "what is x doing?" are to be answered in this way suggests, I think, that what people do – i.e. what *people*, as opposed to what their bodies, do -- is to be understood in terms of their intentions.

The goodness of the institution may not have any representation in the dictator's intentions. Given only that the institutional basis of the dictator's

act is good, it is therefore as misleading to describe the dictator's act as good as it is to describe the baby's gesticulation as one of creating such and such sub-atomic disturbances."

I agree with most of this. Identity is not realization; and actions are individuated, at least in part, by their motivating intentions. So I agree that, *given only* that the dictator's act *realizes* some wholesome institutional act, and *given only* that the institutional act thereby realized embodies a certain morality, it doesn't follow that some morality is transferred to the dictator's act. But I still don't think that the relationship between the dictator and the host-institution is such as to block this transference.

The dictator's relationship to the host-institution is not comparable to the baby's relationship to the sub-atomic events that it unwittingly precipitates in the process of toddling about.

Let us start with a background point. Suppose that, with the intention of shooting Jones, I squeeze a certain trigger on a certain gun. As a result, the gun fires and kills Jones. Surely it can be said that *I* killed Jones. It would be wrong, or at least mis-leading, to say that all *I* did was cause the events that killed him. From some viewpoint, the physical events occurring within the gun are constitutive of what I have done, even though I may not have the foggiest understanding of the nature of those events.

Suppose that Stalin orders that so and so be killed, the result being that so and so *is* killed. It is quite possible that Stalin doesn't know all the intricacies of the bureaucratic apparatus that carried out that order. But it can surely be said that *Stalin* killed so and so. It would be highly misleading to say that Stalin caused the killing of so and so but did not himself do so. The cascade of administrative events mediating between the order and the termination of so and so's life are actually *constitutive* of some act of Stalin's.

When a dictator uses a system of government to execute his intentions, the systemic facts that execute his intention are not merely *caused* by him, but are *constitutive of* some act of his. For reasons we've discussed, the institutional facts that realize the dictator's act are, at least up to a point, inherently moral (even though, in this case, they are being put to immoral use). So even if we allow for the distinction between realization and identity, and for the fact that actions are individuated, at least in part, by the intentions that cause them, we must allow that inherently moral mechanisms have a constitutive role in the dictator's act. This does bear out our point that there is *a* dimension of morality to the dictator's act that is lacking to that of the gunman.

What we are finding is that, even though the dictator's motives may not be wholesome, it is unlikely that his act will altogether fail to implicate the wholesomeness inevitably present, to however small a degree, in the institution within which he is operating.

[162] Pettit (1975).

[163]
Jackson and Pettit (2004a).

[164]
See Jackson and Pettit (2004b) and Jackson and Pettit (2004d).

[165] It could be argued that, strictly speaking, this is false, and that what is *really* causally efficacious is the occurrence of various sub-atomic events at various neural contact-points. But certainly the dictator's psychological motivations are, in some intuitive sense, *much closer* to being what is causally efficacious than are any grand social forces. So they are, as we might put it, "relatively causally efficacious" – relative to those social forces, they

are causally efficacious: they are *almost* what is causally efficacious, comparatively speaking.

In any case, the supposed fact that the dictator's psychological states are not, strictly speaking, causally efficacious is irrelevant in this context. What is relevant, and what is uncontroversial, is that the *program*-causes of the dictator's act may be much more wholesome than his private motivations, thus complicating the moral assessment of the dictator's act. (I said "supposed" a few lines back because many, including Jaegwon Kim (1993), would say that the mental states *are* causally efficacious, and not merely causally relevant. In Kuczynski 2006b, I defend Kim's view.

[166]

This point is made very clearly, and compellingly argued for, in Jackson and Pettit (2004e) and Jackson and Pettit (2004f).

[167]

For reasons that I give elsewhere (in "The analogue-digital distinction and the cogency of Kant's transcendental arguments") the difference between a state and an event is the difference between, on the one hand, a change and, on the other hand, a change in the manner of change. So states are changes and events are hyper-changes.

Since this may sound strange, even tautologically false, some argumentation is in order. Consider the chair you are sitting on. There we have a classic example of a "steady state" – of a condition of comparative equilibrium – as opposed to an event. Let S be the just described equilibrium. Now kick the chair over. There, of course, we have an example of an event – of a disruption or cessation of a state. Let E be that event.

S comprises innumerable displacements of mass-energy. There is no a single molecule or atom or muon in the chair that is not, at every instant, being wrenched away from its previous location. So S *comprises* various

events. Nonetheless, S is not itself an event. Why is this? Because all of the events just described – all of the changes – happen *within* a certain structure. The *manner* of change remains constant. But when you kick the chair over, a change occurs that is *not* internal to some system or pattern. States are series of framework-internal changes; events are framework-external events.

The *making* of a law involves a *change* in a framework, that is why it is an event. (Of course, that change is itself internal to some *larger* framework. The framework consisting of the laws at a given time is subject to change. What is not subject to change is the framework consisting of a pattern of conduct that remains within certain bounds set by the Constitution and also our culture.)

[168] These remarks suggest that the very concept of an event is relative (and not merely in the sense demanded by relativistic conceptions of motion). X can be an event relative to one framework while failing to be an event relative to some other framework. In any case, the concepts *event* and *state* must be understood in dynamic, and not strictly chronological, terms. That is why a law could, in principle, last for one minute even though the creating of a law could, in principle, last for ten years.

I say “seems”, because my own feeling is that, strictly speaking, the expression “threat” unambiguously refers *not* to the issuing of the threat – not to the act of saying (e.g.) “your money or your life” – but to the condition subsequently created. But here it isn’t necessary, or possible, to argue for this.

[169]

Again, I do not wish to over-state things. There are evil dictators – dictators whose acts are entirely lacking in moral value. But for the reasons given, in order for the dictator’s act to collapse into that of a gunman, in respect of its

moral caliber, it is not enough that the dictator's motives be comparable to those of the gunman. The dictator's act is to be understood as a convergence of a number of different forces; he is embedded in a social, even a moral, order. Therefore, barring some kind of extreme internalism, his acts are not always entirely separable from the forces constitutive of that order. The moral caliber of the dictator's acts is therefore not entirely distinct from the moral caliber of those forces. It would only be in rather singular circumstances that those various forces so completely coalesced that their moral character coincided with that of the dictator's private motives.

I believe that, historically, there may have been such cases – perhaps Pol Pot's Cambodia was one. But the important point here is that such cases involve the collapse into a single point of typically distinct explanatory vectors – vectors that remain distinct even where most *dictatorships* are concerned. The aforementioned philosopher's thesis would seem to be true of such singularities, but not, in my view, of less singular political situations.

The situation with the gunman is very different. Almost by definition, he has ripped himself out of the sort of social matrix within which a government official, even a dictator, operates. This establishes a relatively direct connection between the moral character of his act and that of the psychological cause of that act – a connection much more direct than the one found in the case of the dictator.

[170] In the philosophy of history, there is a debate as to whether historical developments are to be understood in terms of the operation of impersonal social forces of which individuals are merely the exponents, or in terms of the actions of individuals whose behavior can be accounted for entirely in terms of *sub*-social (e.g. psychological and physical) forces. See Tucker (1981: 27-30). The latter is sometimes known as the "great man theory"; the former is

sometimes known as “historicism.”

According to the one theory, if Hitler hadn’t created a totalitarian state in Germany, then somebody else would have. So the march of history is like the behavior of a corporation. If the CEO resigns, his place will promptly be filled by someone else, who will do more or less the same things. (The limitations of the theory are already becoming obvious: one doesn’t *always* do the same things as one’s predecessor; the new CEO may take the company in a new direction.)

According to the other theory, Hitler was (or at least may have been) a lone actor, but for whom World History might have gone very differently.

The historicist analysis is widely held to be wrong, even *categorically* incapable of being true. See Popper (1957) and Nagel (1962: chapter 15). But seems to me that many of the *a priori* objections to the historicist approach must be reconsidered in light of what Jackson and Pettit say. It may even turn out these two conceptions of history are not opposed to each other.

Social forces *program for* certain behaviors on the part of individuals; they are *causally relevant* to those behaviors, even though the forces that are *causally efficacious* are sub-social.

Consider the following case of causal explanation:

(#) x and y are made of the same kind of metal. x is hotter than y. The reason is that x has been sitting in the hot sun, whereas y has been sitting in a cool basement.

(#) contains no mention of the micro-events that are causally efficacious in x’s having such and such temperature, and y’s having thus and such temperature. Still (#) is, at least in some contexts, a perfectly good causal explanation. There is no need for a stand-off between the micro- and the

macro-physicist, no need to say that the one's explanations are categorically false. The one is giving causally efficacious, the other causally relevant, conditions.

It is worth looking into the idea that, for analogous reasons, there is no need for a stand-off between the advocate of historicism and the advocate of the "great man" theory. This is not to say that specific historicist theories are always right; it is a distinct possibility that Marx's theories of historical development were simply wrong. But, it seems, the distinction between causal relevance and efficacy removes some of the *a priori* objections to theories like Marx's, and may be relevant to the assessment of anti-historicist work like Popper (1957). See Wilkins (1978) for a cogent and original discussion of related issues.

Incidentally, if my extremely limited knowledge of this area is not misleading me, Marx himself made some points rather like those of Jackson and Pettit as regards the relation of causal explanation at the individual level and causal explanation at the social or historical level.

[171] It also seems to me that extensions of recent externalist theories of content could conceivably be relevant here. Drawing on Putnam's (1975) work, Tyler Burge (1979) famously argues that Smith can be thinking about arthritis, while twin-Smith – somebody who is atom-for-atom just like Smith, but who is embedded in a different social environment – can be thinking about a different disorder ("tharthritis"). Presumably the sorts of liaisons discussed by Burge can affect, not just one's beliefs, but also one's *intentions.* Smith intends to find a cure for *arthritis,* while twin-Smith intends to find a cure for *tharthritis.*

In light of this, it is an idea worth investigating whether that same sort of externalism might infect the *moral* character of one's intentions. Brown is the

president of Nation X. Twin-Brown is the president of nation twin-X. Brown and twin-Brown are atom-for-atom duplicates of each other. Brown intends to bring democracy to some part of his country currently governed by warlords, whereas twin-Brown intends to bring *shmemocracy* to the corresponding part of his country. But schmemocracy is actually an evil political order. Therefore, what Brown intends is good whereas what twin-Brown intends is evil, even though subjectively (i.e. considering only the forces lodged within their crania, leaving out the distal causes thereof), Brown and twin-Brown are indistinguishable.

These are obviously not water-tight arguments. My own personal view is that the argument given in the last paragraph must *not* be sound. I believe it to be a datum that if Brown is moral, then so is twin-Brown. And I believe that, if externalism conflicts with this (alleged) datum, then it is externalism that must be rejected. But obviously this viewpoint is far from self-evidently correct.

In any case, it does seem that *if* Burge-style externalism is correct, that fact does, at least in principle, considerably complicate the moral assessments of actions. This, in turn, would suggest that the moral character of an act constitutively depends on the sorts of social liaisons that the dictator typically has and that the gunman typically lacks.

In conclusion, it is not out of the question to suppose that Burge-style externalism could drive yet another wedge between the dictator and gunman -- granting that from an internalist, strictly psychological perspective they are sometimes quite comparable.

[172] The distinction between being a law of nature and being *backed* by a law of nature is discussed in Davidson (1984). According to Davidson, no proposition expressing a psycho-physical concomitance – e.g. *whenever*

Jones' sensory surfaces are disturbed in such and such a manner, Jones has an experience of thus and such a kind – expresses a law of nature, even though such a concomitance may be nomically necessary and thus *backed* by natural law. According to Davidson, the problem is that the vocabulary of mentalistic description is "incommensurable" (to use Kuhn's term) with that of physicalistic description. This has the consequence that psycho-physical statements are in some sense ill-formed and thus incapable of expressing laws.

One of Davidson's reasons for thinking that mentalistic and physicalistic vocabularies are incommensurable is that, in Davidson's view, psychological description is irremediably holistic whereas physical explanation is not. Wilkins (1978) argues persuasively that when we carefully look at physical description and explanation, it turns out to be scarcely less holistic than mentalistic description. Kim (1993) also produces some sharp criticisms of Davidson's view.

But Davidson is surely quite right to distinguish propositions that *express* laws from those that, while not themselves expressing laws, nonetheless hold in virtue of laws. As I hope we are seeing, this has significant consequences for the analysis of law.

[173]

Coleman and Murphy (1990: 14-15) make a similar point.

[174]

In fact, there are some laws, e.g. the law issued by Stalin that we just discussed, such that the moral imperative to obey them is, nearly enough, *systematically* outweighed by contrary imperatives. This creates the illusion that those laws are completely without moral foundation.

[175]

Of course, a government can then use the dependence that people

subsequently have on it as the basis for acts of supreme wickedness towards those very people. (Since a child is dependent on its mother, the mother is in some ways better positioned than anything else to be wicked to the child.) But that sort of wickedness *presupposes* a fundamentally protective relationship. Immoral conduct that is *not* superimposed on an essentially protective relationship is just garden-variety, as opposed to distinctively *governmental,* evil. So *ultimately* a government is an inherently moral entity.

[176]

It is, perhaps, an open question whether there is a one-one correspondence between the presence of government and the presence of law. But it seems clear that the connection between the two is not accidental. Given a government, it is optional whether there is a bicameral legislature or even a multi-party system. But, given a government, it is not optional, at least not in remotely the same sense, whether there are laws.

Perhaps there are *conceivable* societies where there is some kind of authority but no laws. To my limited knowledge, some thinkers and revolutionaries have described, and even aspired to create, societies where there were no laws, but where there is nonetheless *some* kind of administration of public affairs (and therefore, one would think, *some* kind of government). On a possibly related note, some have held that there are tribal societies where there are no laws -- only custom, unstated agreements and understandings: "forms of life" that, in certain respects, take the place of government.

Others have spoken of societies where, paradoxically, the government was *so* authoritarian that there were no laws:

“The thing he was about to do was open a diary. This was not illegal (nothing was illegal, since there were no longer any laws), but if detected it was reasonably certain that it would be punished by death, or at least by twenty-five years in a forced-labor camp.” (Orwell 1949: 6).

Arguably, the presence of law betokens a relatively liberal government. Where there is a *purely* authoritarian or autocratic system, there are no laws (or so it could be argued). Some have said that there aren’t really *laws* in societies like Stalin’s Russian and the present day North Korea.

So given the possible existence of societies whose governments are either *too* authoritarian, or whose governments are not authoritarian at all, a case could be made that there could be government without law and that, consequently, it is a *contingent* fact that there are governments that issue laws.

There are several points to make in response to this. First, supposing that there really have existed tribal societies of the sort just described, it seems fair to say that such societies had *no* government, at least not in any sense relevant to our inquiry.

Second, it would seem that the existence of a hyper-authoritarian government of the sort described by Orwell presupposes the existence of law. Wherever there is any sort of government administration, there is *ipso facto* some kind of law. There cannot be government, in any reasonable sense of the term, where there are no codified procedures -- where all decisions are made in a completely *ad hoc* basis. It may be that the relevant protocols fail to embody the principles of fairness that we now tend to associate with laws. But in that case, the right thing to say is that the laws in question are unfair, not that they don’t exist.

Also, if there were a society as oppressive as the one Orwell describes, it

isn’t entirely clear whether that society would have a *government* in quite the sense in which, say, the U.S. or even North Korea has a government. So the conditions that permit lawlessness also seem to demand the absence of government as we know it.

The hypothetical societies where there are governments but no laws seem to be in the nature of *singularities* -- situations where conditions are so extreme that ordinarily true nomic generalizations are no longer applicable. So the fact that our hypothesis is refuted, if at all, by singularities doesn’t really warrant its rejection. If a hypothesis is counter-exampled only by limiting cases – cases where it is in question whether the antecedent of the conditional expressing the hypothesis is fulfilled – that may *confirm*, rather than disconfirm, the hypothesis.

Consider the society described by Orwell. From a political viewpoint, that society is completely pathological; there has been a total corruption of the organs of government. If there really are no laws in that society, that is *itself* an indication of pathology. (This is clearly Orwell’s position.) Given that there is such a high degree of pathology, there is no reason to expect the usual principles to apply. Facts about an organ, e.g. a lung, that has been completely ravaged by pathology don’t necessarily refute generalizations about lungs; they don’t refute generalizations about what a lung *qua* lung does. If the failure of such an organ to conform to some generalization is *itself* rooted in the pathology in question, then it is less a counter-instance than an instance of that generalization.

So given only that a government as pathological as Orwell’s lacks law, it isn’t clear that we have any genuine counter-example to the thesis that, as a rule, there is law where there is government. We may conclude, then, that there is a *principled* relationship between law and government.

[177] It is also relevant that when a political situation becomes *too* much like the gunman-scenario, it becomes questionable whether the concept of law is applicable. Were there *laws* in Pol Pot's Cambodia? Are there *laws* in North Korea? The more a political system approximates to complete immorality, the more questionable it becomes whether the concept of *law* is applicable. That in itself rather confirms our moralistic analysis of law.

[178] Contrary to what I argued in an earlier version of this work. I reversed myself under the pressure of acute criticisms due to Thomas Holden.

[179]

Contrary to what I alleged in an earlier version of the present work.

[180]

I say "practically" in acknowledgement of trivial exceptions. For example, there might be a law to the effect that one cannot kill everyone on the planet. There are probably no counter-factual worlds (or, at any rate, no counter-factuals worlds that are sufficiently like ours to be, in this context, "accessible" to it) where it is morally permissible to kill everyone.

[181]

There is another, rather more accurate way of articulating the assurance in question. For any law L_i, there are legal consequences to one's *not* conforming to C_i (the condition or mode of conduct associated with L_i). In some cases, the legal consequences of not conforming are undesired. For example, if L_{36} is the law prohibiting drunk driving, then the legal consequences of not conforming to L_{36} are undesirable (these being, say, a heavy fine and loss of one's license). In other cases, one's not conforming to C_i has consequences that are to be understood in a purely privative sense. If one is *not* married, then one does *not* receive such and such tax-exemptions;

but it would probably be going too far to say that one was *penalized* for not being married. If one is *not* handicapped, then one is *not* allowed to park in such and such places; but it would probably be going too far to say that one was *penalized* for not being handicapped. For each i, let K_i be the legal consequences – be they desirable or undesirable – visited upon one for *not* conforming to C_i.

Given this terminology, we can see our legal system as being given by the following assurance:

(*1) There is an assurance that, for anyone x satisfying certain conditions relating to place of residence, biology, age …, and for any i, if x does not conform to C_i, then *ceteris paribus* the consequence for x will be K_i.

Let me explain the *raison d'être* for this. It would be an exaggeration to say that, for any i, if x conforms to C_i, then the consequences for x will definitely be K_i. For example, if I am married, then I am entitled to certain tax-exemptions – *unless*, of course, I am found guilty of income-tax evasion, and consequently fined; *unless* I enter into some kind of agreement with the government whereby I pay *extra* taxes for some unrelated crime; *unless*…

Similarly it would be an exaggeration to say that, if I am found guilty of (say) arson, there will *definitely* be such and such legal consequences. After all, the prosecutor might give me immunity in exchange for testimony against a mafia boss. What would be correct to say is that, *ceteris paribus*, the legal consequences for committing arson are such and such.

(*1) does a better job than (*) in capturing the fact that the legal consequences of a given mode of conduct are typically given by *ceteris paribus* propositions. But it seems to me that (*) is a more wieldy proposition

than (*1) and that, given that our concerns here are highly general, the loss of precision incurred as a result of our operating with the former would be more than off-set by the gain in expository fluidity and over-all understanding of the relevant issues.

[182] To my knowledge, this corollary has not been drawn by Jackson and Pettit themselves. But, I grant, I may well be wrong about this.

[183]

Even though the government could give individualized attention to *any one* problem.

[184]

This might seem to bear against our contention that laws are assurances. For, it might be thought, my assurance to Bill that I will pick him up is an event. I think that this is a mistake. The act of my *giving* Bill this assurance is an event. But the assurance itself is a standing condition. Between the time of my saying “I promise to pick you up tomorrow”, and my actually doing so (or failing to do so), Bill *has an assurance*: the assurance is thus an enduring structure, not an event.

[185]

I say “help” because not all laws are enforced. The fact that there are certain laws in this country *helps* program for the fact that there are certain consequences when (for example) somebody is discovered to have committed an act of racial discrimination. But the fact that such a law exists does not *by itself* entirely program for those consequences. After all, the law could be unenforced. At the same time, it would be absurd to deny that the existence of such a law *partially* programs for certain events. It is virtually impossible to give causally sufficient conditions for anything. (According to some, it is sometimes possible. But, by all appearances, it is possible only in certain

branches of physics, and even then only under rather special conditions. See Russell (1927: 231) for a discussion of the conditions under which causally *sufficient* conditions can be given.) So given only that laws are not *sufficient* program-causes of certain events, given certain initial conditions, it doesn't follow that what we are saying about them here is wrong.

[186]

Up to a point, the inefficiencies in the law correspond to the margin of error in people's judgments. So a certain inefficiency or redundancy in the law is actually necessary.

[187]

There is a related point that is best introduced through a story. Jones lives in the woods. He is entirely self-contained. He grows his own food. He lives in a part of the wilderness that is never visited by gamesmen, hikers, or representatives of timber-corporations. Jones never imposes on anyone. He does not need or want the protection of American law.

One day, Jones goes on a shooting spree in a crowded urban area. After the spree, he makes it clear to everyone concerned that he does not want the protection of law. He tries to convince law-enforcement and judicial personnel to forego deployment of the legal apparatus that would otherwise be set in motion. He makes it clear that he is willing to be treated like a rabid rodent, or any other non-human threat to public safety.

Jones has done awful damage. It is irrelevant that Jones is willing to forego the protection of law. So long as there are people like Jones – people who refuse both the protections and limitations of law – there is no guarantee of a tenable, and a near guarantee of an *un*tenable, civil situation. So the government has a moral duty, and therefore a right, not only to conditionalize the receipt of protections on compliance with certain codes of conduct, but to demand such compliance from *everyone* in the relevant geographical area –

even from those who are willing to forego those protections in exchange for freedoms.

It seems, then, that legal obligation is not a form of *contractual* obligation. Jones hasn't entered into any sort of *contract* with the government. Even if we broaden the concept of a contract – even if allow that people may enter into a contract *without* consciously doing so, even if we allow that one may become party to a contract merely by being embedded in the right "forms of life" and undergoing the relevant acculturation – it would *still* be hard to make a case that Jones had entered into such a contract. Jones has made a deliberate effort *not* to enter into *any* sort of contract with the U.S. government; and he has, it would seem, done a good job of *not* becoming a part of any cultural practices that, given a suitably broadened conception of a contract, could be seen as binding him to some sort of agreement with the government. At the same time, it is, I believe, quite clear that Jones has *legal* obligations under the U.S. government. Jones is *breaking the law* if he starts killing endangered species of animals – animals that it is ordinarily illegal to kill. Jones could not be *breaking* the law if he weren't bound by it.

This last point may need clarification. A diplomat who has immunity is not *breaking* the law when he does something that would *ordinarily* be an act of breaking the law. To have diplomatic immunity is not to be able to *break* the law with impunity; on the contrary, it is to *not* be able to *break* the law for the reason that the law doesn't bind you to begin with.

Jones' position is not that of someone with diplomatic immunity. It is a datum, I believe, that Jones is *breaking the law* when he kills a bald eagle. (It is illegal to kill bald eagles, since they are an endangered species.) He is not breaking the law in a derivative or *ad hoc* sense that would be concocted by legislatures who wanted to curtail Jones' activities, but in the pregnant sense

in which I would be breaking the law if I were to kill a bald eagle.

One cannot exempt oneself from the law at will. I don't mean that one cannot *morally* do so: in some cases, such self-exemption would perhaps be morally *de rigueur* if it could be accomplished. I mean that it isn't *actually* possible to exempt oneself from a legal system just by refusing to enjoy its protections.

Given this fact, a certain way of looking at legal obligation proves untenable. Legal obligation is *not* a contract between citizen and government. Legal obligation is not the moral obligation to live up to one's agreements. Of course, there are no doubt people who *would* enter into such contracts. But here the point is that, even when a person does not do so, she is still (ceteris paribus) bound by the legal system in question.

It should also be pointed out that it is only relative to an extremely permissive conception of contractual obligation – one that borders on the metaphorical -- that it could be held that, as a rule, citizens enter into any sort of *contract* with the government. The supposition that such a contract exists is presumably meant to be an evaluative tool, not an empirical hypothesis (In his *History of Western Philosophy*, Bertrand Russell says that, according to Locke, the existence of the making of such a contract is a historical fact, and not (merely) a heuristic device. Russell himself takes it as obvious that it is not a historical fact.)

[188] In the relevant area.

[189]

I say "practically" because it is possible to hypothesize laws that ought never to be broken (e.g. a law prohibiting the destruction of the whole planet).

[190]

Virtually no legal system is completely fair to anybody or completely unfair

to anybody. Total unfairness amounts to a failure of legal recognition, as we just saw. A legal system would be *totally* fair to x, i.e. would never involve *any* abridgment of x's right to flourish, only if that legal system were entirely at x's control. But this sort of situation occurs seldom, if ever. (And if a social structure were *completely* at some one person's control, it is not entirely clear to me that *laws* would be involved.)

[191]

This is not inconsistent with the fact that L is unfair compared to the legal systems found in Sweden or the United Kingdom. Nor is it inconsistent with the fact that G would be doing much more good if it were to impose a different legal system.

[192] More accurately, even though it may not be necessary that *Jones specifically* comply, it is necessary that *some* person x do so, such that L is unfair to x in the way that it is unfair to Jones. So all of the points that we are making about Jones would be true of x. Since x is not relevantly different from Jones, our argument would go through unchanged.

[193] It was urged by Kant in the *Groundwork for the Metaphysics of Morals.*

[194]

I am not referring to any articles or books, but rather to a talk that I heard at a conference in Hawaii. The speaker was an assistant professor at a university in West Virginia. I don't remember her name or the name of the institution; and (not for lack of trying) I have had no success finding out any of this. I believe that the paper in question is of profound importance, and I am sorry that I cannot identify its author. I wish to stress that I am not to be given credit for that speaker's insights. Had I not heard that talk, I probably would not have questioned the idea that "ought implies can." In fact, I remember

instantly trying to find some fallacy in the argument, in the hopes of holding onto what I then thought was an obvious truth.

[195]

This example was not present in the aforementioned talk.

[196]

Cf. Hintikka (1969: 203): "I give an honest promise (let this act be described by *p*) to bring it about that *q* (say, have a cup of coffee with you). Unknownst [sic] to me, however, my father has fallen ill, which creates an obligation to see him that overrules my earlier promise. It seems to me that moral philosophers have felt somewhat uncomfortable in discussing this kind of situation, and perhaps one can also see why. For the fact that the obligation created by my promise is overruled means that it is false to say *simpliciter* that *Oq*, i.e. that I am obliged to fulfill the promise…Yet it is clear that not everything is morally all right if I have to break my promise, however firmly this course of action may be prescribed to me by the norms I abide by. I have somehow done something wrong. This 'moral failure' is what easily makes one hesitant to say that in such a case there is no absolute duty to keep the promise."

Having made these excellent remarks, Hintikka then goes on to say that, under the circumstances just described, the obligation to keep my promise to you is a "prima facie" obligation, whereas my obligation to tend to my dying father is an "actual" obligation. This choice of terminology suggests that the obligation to keep my promise has been reduced to unreality (a conclusion that Hintikka does not himself quite accept: see below). But the very dilemma that Hintikka is describing arises precisely because the duty to keep my promise is very much alive. So in this context it is hard to see the justification for the term "prima facie", since that term suggests the obligation to keep my

promise is apparent as opposed to actual.

Hintikka (1969: 203) says that my obligation to keep my promise, though overridden by my obligation to tend to my father, remains existent in the sense that:

"*in a deontically perfect world* [Hintikka's emphasis] such a promise cannot be given without keeping it. In such a world, *p* cannot be realized without bringing it about that *q*. Even if the act of promising does not give rise to an actual duty to keep the promise (e.g. because of other duties), it none the less remains true that in this sense giving a promise commits one to keeping it."

A deontically perfect world is a world where everything that *ought* to be done *is* done. So to say that "in a deontically perfect world" I keep my promise to you is just a way of saying that I *ought* to keep my promise to you (even though I may also have a weightier contrary obligation).

So by Hintikka's own lights, my obligation to keep my promise to you – though *overridden* (as Hintikka himself puts it) by a contrary obligation – is still quite real. This calls into question the appropriateness of the term "prima facie" in this connection. We must assume, I believe, that this term is being used in a purely technical sense that is only associatively related to its ordinary meaning.

Hintikka's notion of a deontically – or, more generally, a morally -- perfect world provides us with a way of substantiating, or at least illustrating, some of the moral conceits involved in the present work. In a *morally* perfect world, everyone would have the *freedom* to do anything, but people would freely choose to exercise that freedom only in the best possible way. A world where people freely choose to do good 80% of the time is better than one where people are forced to do good 100% of the time. (In any case, this is

what is reasonably presupposed by the so-called “free-will” theodicy. See Mackie (1982) and also the earlier footnote in the present work on that theodicy.) So *ceteris paribus* a world of absolute freedom is morally better than one where freedom is abridged.

This implies that any abridgment of freedom is a departure from the ideal, and is thus immoral in an absolute sense, though it may be needed to prevent a greater departure from the ideal, and is thus moral in a comparative sense. Thus our position that *by itself* any freedom is a good thing (though perhaps not as good as some other, incompatible good) would seem to be vindicated. So if, as we have argued, all laws protect *somebody’s* freedom, then it straightforwardly follows that laws always protect *some* good and therefore that laws always have some basis in morality. It would then have to be concluded that, when a law is evil, that is not because it protects *no* good, but because the good it destroys outweighs the good it protects

[197]

One of the stock arguments against utilitarianism is that it doesn’t make any moral allowance for the fact that we care about our friends and family more than we do utter strangers. According to utilitarianism, an act is good in so far as it promotes happiness, and bad in so far as it fails to do so. It is irrelevant *whose* happiness is promoted. But we feel that only a person with serious deficits would make no more effort to promote the happiness of his children than that of people he’d never met. It seems to me that comparable deficits must be attributed to a psychotherapist who feels *no* remorse about violating the trust that a patient has placed in him, even if such a violation is necessary for a greater good.

[198]

See Baker (1985) and Fletcher (1990).

[199]

Albeit ones that may wriggle through the net of the models of morality that ethicists create – much as many aspects of language wriggle through the net of the egregious, but sometimes useful, simplifications of Frege's logic.
[200]
The negation of "I ought to do X" is not "I ought to *not* do X", but "it is not the case that I ought to do X." The psychotherapist ought to *not* keep his patient's confidence (since doing so would lead to the deaths of the people whom his patient is planning to kill). But the psychotherapist also ought to keep his patient's confidence. (So it would be *false* to say: "it is not the case that the psychotherapist ought to keep his patient's confidence", even though it would be true to say: "the psychotherapist ought to *not* keep his patient's confidence.") It is no more self-contradictory to affirm the co-existence of these two conflicting oughts than it is to say that Jim wants to punch Bob and also to be kind to Bob. The negation-operator hasn't been given sufficiently wide-scope to generate a contradiction.

Here is a comparison. The negation of "I can do X" is not "I can *not* do X" (i.e. "I am able to *not* do X"), but rather "it is not the case that I can do X". Obviously "I can go to the store" is compatible with "I can *not* go to the store" (but not with "it is not the case that I can go to the store" or, more idiomatically, "I cannot go to the store.") The presence of modal, moral, and epistemic operators demands that we pay special attention to the way we form negations.
[201]
Except perhaps in extreme cases (e.g. "don't force an unconsenting person to have sexual relations with you").
[202]
One could hold onto the idea that "ought implies can" by saying that, even though it is impossible to do all of the things that one ought to do, it is still

possible to do any given one of them. The psychotherapist can uphold his patient's confidence and he can also go to the authorities – but he cannot do both. Thus, "ought implies can" survives the argument we've considered.

A comparison may be in order. It is no contradiction to say that Tom can be in Arizona at time t and Tom can also be in Delaware at time t, but that Tom cannot be in both places at t. Similarly, it is no contradiction to say that Smith ought to do X and he ought to do Y, but that he is under no obligation to do both X and Y, since it is logically impossible for anyone to do so. So "ought" (like "can") doesn't "factor out."

This counter-response is not without its problems. To save your dying child, you must be in place A (the location of a hospital) at time t. To honor an important promise to your best friend, you must be in place B at time t. You decide to save your child. So at time t, you are in A. At that very time, you have an obligation to be in place B. (For reasons we've discussed, this obligation was outweighed, but not cancelled, by your incompatible obligation to be in A.) Of course, at a time prior to t, you *could* have gone to B instead of A. But that doesn't change the fact that *right now*, at time t, you have an obligation to be in B. But it is not now possible for this to be the case. So even though it is now impossible for you to be in B, you are still under an obligation (albeit one that was overridden) to be there. It is not as though your obligation to be in B ceased to exist the moment it *stopped* being possible for you to live up to it. Your friend could call you on your cell phone and say truly: "you are supposed to be here right now."

The counter-response would be to say that "ought implies can" means that one ought to do only those things which *at some point* one was able to do. So you ought to be in B, even though it is now logically impossible for this to occur, because you *could* have gone to B. But then the right slogan is not "ought implies can" but rather "ought implies could have."

But even this weaker thesis is problematic. You lend me a million dollars. For whatever reason, I have no choice but to spend it. I cannot make it back. At no point does it become the case that I *could have* paid it back. I still owe it to you. So I still *ought* to pay it to you. So I ought to do something which I cannot do and which it will never be the case that I *could have* done. So "ought implies can (or could have)" is wrong.

The only way to break this argument is to say that "I owe you money" does not entail "I ought to pay you money." But, it may safely be said, such a view borders on self-contradiction.

It might seem that some support for the view that "ought implies can" is to be found in our moral attitudes. Obviously somebody who cannot repay a debt, but *would* if he could, is (to that extent) morally better than somebody who can repay a debt but won't. We don't have contempt for the former, but we do have contempt for the latter. We regard the latter, but not the former, as *blameworthy*. This might seduce us into thinking that a debt that cannot be repaid ceases to exist: one cannot owe what one cannot pay – so "ought" *does* imply "can", contrary to what I've been arguing.

But that is not the right response to have. It is true that those who would repay their debts, but cannot, are not worthy of contempt. But it doesn't follow that their debts have ceased to exist. We must once again distinguish the property of being *condemnable* from the more general property of being bad. Let me use fiction to illustrate my meaning.

I owe a million dollars to Smith. When I borrowed the money from Smith, I had every reason to believe that I would be able to pay him back, and I have done everything in my power to do so. But circumstances being what they are, I simply cannot pay him back. Here something *bad* is happening: I am not able to pay Smith the money I owe him. But nobody is to be condemned or blamed. Everybody has done everything humanly possible

to ensure the best possible outcome. It just happens that, in this case, our best efforts were not good enough. In this case, the badness associated with my not being able to repay my debt doesn't correspond to anyone's being blameworthy. That is why nobody deserves to be held in contempt. But it is still a fact that I have incurred a debt: even though I am not blameworthy, I still owe Smith money.

Now let us change the story. Smith lent me money. I have the money and I could easily pay him back. But I choose not to do so. The problem isn't simply that I am not paying Smith the money I owe him, but that I am *choosing* to not do this. Here my behavior *is* blameworthy. An attitude of contempt is appropriate.

In both scenarios, I am failing to pay back a debt. The difference is that, in the one case, my failure to do this is deliberate, whereas in the other case it is not.

Thus, the lack of contempt we have for those who cannot pay back debts, but would do so if they could, corresponds to the lack of *blameworthiness* on the part of such people. It does not correspond to the fact that they no longer *ought* to repay the money.

It would seem, then, that the right slogan is not *"ought" implies "can"*, but *rather "blameworthiness" implies "can."* If I have the option of paying the money back, but I choose not to, I am blameworthy. If I don't have that option, I am not blameworthy. But, as we saw, my lack of blameworthiness does not void my debt: I still owe the money, and therefore ought to pay it.

For reasons having to do with the problem of free will – briefly discussed below -- I am not even sure that *"blameworthiness" implies "can"* is true. But it is a more resilient hypothesis than *"ought" implies "can."*

Let us sum up. In some cases, one ought to do things that one cannot. In any case, if we deny this, then we end up having to make some rather

arbitrary distinctions. For example, we must say that since I cannot repay the money that I owe, I am under no obligation to pay it: so I owe money but am under no obligation to pay it. But such a position is absurd.

It should also be pointed out that there are well known problems posed by the truth (or, where the macroscopic world is concerned, approximate truth) of determinism for the idea that we *can* do anything other than what we in fact do. Given the (approximate) truth of determinism, the idea that "ought implies can" seems to show that we are not morally responsible for *any* of the choices we make. Everything that does happen has to happen: nothing *could* have gone differently. It would seem to follow (though some contest this: see below) that nobody *could* have done anything other than what she actually did. So if "ought implies can", then nobody *ought* to do anything other what she in fact does, since nobody *can* do anything other than what she in fact does. See Chisholm (1961) and Taylor (1961). It follows vacuously that everybody does everything that she ought to: everybody is morally perfect. See Hook (1961), especially the section entitled "Determinism and Responsibility in Law and Ethics" for the most complete discussion of this issue known to the present author.

Compatibilists have a response to this. They advocate a counter-factual analysis of freedom: one is free to do X if, *had* one chosen to do X, one would have done so. See Pap (1961: 214) and Van Inwagen (1993). At first, this might seem to eliminate the problem just discussed. Many true counter-factuals involve scenarios where some law is broken. "If the vase had fallen, it would have shattered" is true, even though in actuality it was deterministically impossible for the vase to have fallen. Similarly, "if I had chosen to play tennis, I would have done so" is true, even though, in actuality, it was deterministically impossible for that choice on my part to have occurred. So it seems that the compatibilist conception of freedom

alleviates the problem discussed in the previous paragraph.

Nonetheless, given the (approximate) truth of determinism, this apparently reasonable move seems to have a paradoxical consequence. Suppose that Smith does not actually do X, and that her not doing X was deterministically necessary. In that case, *if* she had done X, then some natural law would have been broken. So if we say that she *could* have done X, we seem to be saying that she *could* have broken a natural law. But surely no one can break a natural law.

Lewis (1986) discusses this issue. His position is that, even though doing X would *involve* a violation of natural law, *Smith* herself would not have broken any law. The violation of natural law would be a *pre*-condition for Smith's choosing to do X – it would have preceded Smith's making that choice and not been a consequence of it. Thus the just proposed analysis of freedom doesn't have the unpalatable consequence that people are capable of breaking the laws of physics.

In any case, despite Lewis' brilliant analysis, which may (or may not) be correct as far as it goes, the seemingly truistic thesis that "ought implies can" is clearly under attack from a number of different quarters. In fact, I would like to end by discussing yet another reason to doubt that thesis – a reason that, to my knowledge, is nowhere in the literature.

A general remark on the content of counterfactual statements is in order. Given the approximate truth of determinism, if we see counterfactuals as making statements about alternative *scenarios*, then we must see practically every counter-factual as involving the occurrence of a *miracle*. See Lewis (1973), Lange (2000), and Kuczynski (2005b). Even if the world is not completely deterministic, it is still deterministic *enough* that, given an account of counter-factuals that tethers their truth-conditions to occurrences in alternative or possible scenarios, we must posit miracles every time we

wish to run a counter-factual. Actually, as I argue elsewhere, such a conception of counter-factuals forces us to posit innumerably many miracles if we are to run even the most pedestrian of counter-factuals. For various reasons, this is unfeasible or, at the very least, extremely unpalatable. See Kuczynski (2005) for a full discussion of this point.

In that same article, I argue that to avoid this unattractive option, we must see counterfactuals as making statements about relations of *probabilification* or confirmation among propositions. So we must revive a version of Goodman's (misleadingly so-called) "meta-linguistic" analysis of counter-factuals.

If this last point is right, then to say that Smith *could* have done such and such is to make a purely *epistemic statement*. It is to say something only about what is not ruled out by our limited knowledge of natural law and contingent circumstance. It is not to say anything about what is possible in some more robust, non-epistemic sense. So given a non-Lewisian analysis of counter-factuals, that statement merely registers some confirmation-relation holding among the propositions embodying our limited knowledge of the structure of our world. (By a "non-Lewisian" account of counter-factuals, I mean one that doesn't see such statements as making statements about occurrences in alternative worlds.) Roughly speaking, when we say that Smith "could have" done X, whereas Jones could not have, we are not saying anything about what the causal structure of the world permits, but only about the gaps in our knowledge of that structure.

But if that is right, then it is hard to see how anything of *moral* consequence could follow from the truth of "could have" statements; it is hard to see how our ignorance of causal structure could be of such deep moral account. Surely the meritoriousness of an action is intrinsic to it and doesn't come and go with changes in what we know. In any case, this is the

traditional, non-nihilistic view. Surely it is simply wrong to kill babies: the wrongness of such an act is not a projection into the external world of cavities in our belief-system. There is little doubt that, if we knew more about the brain-chemistry of serial killers, we would find that their behavior was nomically necessary. Consequently, that when we say of a serial killer that he "could have done otherwise", we are not talking about what is deterministically possible, but only about what is epistemically possible given our extremely superficial knowledge of the causal realities underlying the serial killer's behavior. But it is a datum that serial murder is morally wrong. So the serial killer "could have" done otherwise only in a sense that is of dubious relevance to morality. But it is a datum that the serial killer *should* have acted otherwise. (Of course, it has been said that *tout savoir est tout pardoner…*)

In conclusion, if we cleave to the slogan that "ought" implies "can", then given even the approximate truth of determinism, along with some defensible and (before David Lewis) widely held views concerning counter-factuals, we are committed to some rather dubious views regarding the moral evaluation of our actions. So we have yet one more reason to doubt the truth of the principle that "ought" implies "can."

Let us end with a historical point. Kant was acutely aware of the apparent problem posed by determinism, in conjunction with the thesis that "ought" implies "can", for the possibility of behaving ethically. See Kant (1785). He was a staunch believer both in the truth of that thesis, and also in the truth of determinism. He resolved the contradiction by putting into play his distinction between "empirical" and "transcendental" levels of reality. This is not the place to assess the merits of that solution or of the metaphysical architectonic presupposed thereby.

[203] Of course, we are leaving aside cases where laws unfairly demand the death of those who are subject to them.

[204]

See Searle (1970) for a classic discussion of the nature of speech acts. Most of what is said in this particular discussion is found there. See also Searle (1965).

[205] There are many cases where questions, assertions, imperatives, and so on are not issued by sentences of the corresponding *semantic category*. For example, there are rhetorical questions (sentence-tokens that are *semantically* questions but are pragmatically assertions). But this fact – though important in some contexts – is not of central importance here. It will dramatically simplify exposition if we focus on cases where these semantic and pragmatic classifications coincide. It would be easy in principle – but very tedious and distracting in practice - to reconstruct the discussion so that it didn't involve this pretence.

[206]

With some of the trivial exceptions mentioned in earlier footnotes.

[207]

We might even go further in this direction. Obviously the intentions of the law-maker may be to hurt those to whom the law is subject. Nonetheless, I think a case can be made that, *qua* law-maker, he must *intend* to protect those very people. So far as he intends otherwise, it is not *as* a law-maker. There is thus *a* sense, though not a psychologically pregnant one, in which it is tautologically true that any *law-maker* intends to protect those who are subject to his laws.

In the first paragraph of his most famous work, Rawls (1971: 3) writes: "Justice is the first virtue of social institutions, as truth is of systems of

thought." As we said at the very beginning of the present work, there is no sense in which an unjust military campaign is, on that account, less successful *as a military campaign* than a just one. But there is a sense in which an unjust law is worse *as a law* than a just one. This, in turn, suggests that, from some viewpoint, the creation of a law is *an attempt* to do something just.

Of course, there may be law-makers who have no intention of creating just laws. Their objectives may be entirely political. Relative to those political objectives, the evil laws they design may be great successes. But in such a case, it seems to me, the law-maker is not successful *as a law-maker*, but rather as a political operator whose true objectives (personal gain, political advancement) happen to involve his creating laws.

A comparison may be helpful. Suppose that my main objective at the moment is to curry favor with my wealthy and vain father-in-law. To this end, I consistently let him beat me at tennis. Just as I had hoped, this endears me to my father-in-law, and he rewards me with a lucrative job in his firm.

My performance on the tennis court was a *kind* of success. But it was not *athletic* success. It was athletic failure, in fact. Similarly, a law-maker who designs evil laws for political reasons is a failure *as a law-maker*, though he may be a success as a political operator.

This is not to say that the law-maker's failure to make good laws embodies any sort of *incompetence* as a law-maker. Such a failure may embody supreme prowess. Our tennis-example may once again be helpful. When I lose to my father-in-law, I must do so convincingly. I must place the ball in just such a way that he can return the shot while having to make *some* effort to do so: otherwise his vanity will not be suitably flattered, since he will feel that his victory was cheap and meaningless. When I mis-hit the ball, I must hit it *just* outside the court: otherwise my father-in-law will feel patronized

and manipulated, since he will know that I am deliberately throwing the game. So considerable athletic *competence* is needed for my failure on the court to serve my interests.

Similarly, a lawmaker may sometimes need considerable *ability* as a law-maker to create the right sort of evil laws to serve his objectives – considerable ability may be needed to *fail in the right way* as a law-maker. Nonetheless, it seems a datum that, *qua* law-maker, one is failing in so far as the laws one creates don't serve the interests of those subject to them. We've already discussed (pp. 138-140) why I am not begging any questions in accepting this so-called datum at face-value.

Given that laws and legal systems succeed precisely to the extent that they serve the interests of those subject to them, it is not unreasonable to say that such things are, in some sense (though not a psychologically pregnant one), *attempts* to serve those interests.

A comparison may help. In the game of soccer, success is measured entirely in terms of how many times one puts a certain ball between a certain pair of posts. Given this, it follows that what a soccer player, *qua* soccer player, is *trying* to do is to put the ball between those posts (or, at any rate, is to help his team accomplish that objective). Similarly, given that the success of a legal system is measured entirely in terms of the extent to which it promotes the welfare of those subject to it, it seems to follow that a legal system is an *attempt* to serve those interests.

Of course, not every soccer player has the intention of helping his team win. Team-player Smith has been hired by an opposing team to lose the game for his own team. But it is not *as* a soccer player that Smith has this objective. Rather, it is *as* an agent of some opposing team. A law-maker may have no intention of creating laws that help those who are subject to them. But, I will now argue, it is not *as* a law-maker that he has this intention.

There obviously are, or at least could be, legal systems that are designed with the intention of *harming* those subject to them – of hobbling them, so that they don't pose a threat to the government. Given this, how can it possibly be maintained that, as a rule, legal systems *aspire* to promote the interests of those subject to them? How can it possibly be maintained that laws are *attempts* to serve those interests?

At first, this seems to weigh against our analysis. But, as I will now argue, it supports our analysis, or is at least consistent with it. To see this, we must make some general points about the concept of promise-making.

Suppose you make a promise to Jones – you promise to pick him up at the airport. Your success *as a promise maker* is a function solely of whether you keep that promise. Of course, it may be to your great advantage *in some other respect* to break that promise. Perhaps you must break it to prevent some horrible catastrophe. Perhaps your breaking that promise would bring you great success *relative to* your secret hatred of your "friend" Jones, and your consequent determination to make Jones' life as unpleasant as possible while operating under the cloak of friendship. Given your larger objectives, it may be supremely unimportant to you whether you fail in the narrow arena of promise-making. Nonetheless, there is no doubt that, within that arena, you have failed.

When a government issues a law – i.e. when it issues assurance of protection of a right – its *objectives* may be nefarious. But its success *as a promise-maker* – its success *as* something that is promising to protect the rights of the governed – is a function of whether it keeps that promise. So even though the psychological state of the law-maker may be one of *not* wanting to help those to whom the law applies, there is still a sense in which *as a law-maker* he is aspiring to help them. In so far as he intends *not* to help them, he is not intending that *as* law-maker, but *as* somebody who wishes to

use his position as law-maker as a mere instrument in some larger pursuit – just as the soccer-player who intends to lose the game for his own team does not intend that *as* a soccer-player, but as somebody who wishes to use his place on the team as a mere means to some end. Once again, we are reminded of the need to understand the functioning of institutions in terms that are not wholly internalistic. Similarly, when a law-maker intention uses his position to perpetrate some iniquity, he does not have that intention *as* a law-maker, but rather *as* somebody who is using his position as a law-maker as a means towards some larger end.

[208]

We must distinguish between a non-trivial protection of a trivial good and a trivial protection of a non-trivial good. This distinction may help us defend our conception of law against criticisms that are similar in spirit to that put forth by the objector.

Let L be a law that allows people to wear green shirts indoors. L is trivial on a number of levels. First, the existence of that law isn't necessary to allow people to wear green shirts indoors, at least not in our cultural context. Second, even if such a law were necessary to allow people to wear green shirts indoors (because, let us suppose, recent cultural developments had led to the occurrence of hate-crimes against those who wore green shirts), the freedom protected is so trivial, so hollow morally, that it seems a stretch to say that L is a protection (or an assurance of a protection) of a *moral good.* But there obviously could be laws like L.

We could go even further in this direction. Let L* be a law that allows people to wear green wooden clogs indoors. (So L* punishes anyone who harasses a person x on the grounds that x is wearing green wooden clogs indoors. Of course, L* allows that there might be some other reason to treat x negatively – for example, x might be a burglar who is wearing green-clogs.)

In effect, L* protects nothing, since nobody wears wooden clogs. (More accurately, L* makes no *wide*-scope promise: there is no actual good x such that L* assures protection of x.) Second, even if people *did* wear green wooden clogs indoors, their quality of life surely wouldn't decrease if they were no longer permitted to do so. So even if L* *did* give some actual person some actual protection, the protection would be so trivial that it would be a stretch to see L* as an assurance of a protection of a *moral* right. The moral protections assured by L and L* are thus either entirely, or vanishingly close to, non-existent. This diminishes the credibility of our analysis of law.

There are two points to make here, one of which we've already discussed at considerable length. The right protected by L* is indeed non-existent. But, as we've seen, that has no bearing on whether our analysis is correct, since we hold laws to be narrow-scope assurances of protections of rights.

The situation with L is different. The obvious thing to say in connection with this law is that what it protects is trivial. That is indeed one way of describing the situation. But it would no less correct to say that L does a trivial job of protecting something that is itself of supreme moral value. L protects the individual's freedom to live the way she wants to live.

Like all laws, L must protect that freedom against some *specific* threat. There couldn't be laws, at least not practically feasible ones, that just protect *freedom*. It must always be a specific freedom that is protected – one's right to have an attorney present when being questioned by the police, one's right to see the financial records of the company whose stock one is purchasing, and so on. L protects the individual's freedom to live as she desires against a trivial threat. But this means *not* that L is a protection of a trivial good, but rather that it is trivial protection of a non-trivial good.

An analogy may be useful. Every measure taken to ensure good health is

directed against some specific threat to good health, or is directed towards some *specific* augmentation of health. You avoid fatty foods to avoid developing blockages of your arteries and high-blood pressure due to obesity. You avoid smoking to avoid lung-damage and carcinogenesis.

In some cases, the protection might seem extremely trivial: you take vitamin A to have strong nails. But that measure could also be described as trivial protection of a larger, non-trivial good. In fact, I think that there is some independent corroboration for this viewpoint. I would be very surprised to find someone who smoked three packs of cigarettes a day, but was a stickler about taking care of some minute aspect of his health (e.g. the condition of his nails). When someone fusses over such details, it is in the context of a concern for some over-arching good.

As we've discussed, laws always occur in the context of what is *ultimately* a relationship of protection. If there is a law, like L, that appears to protect some trivial good, it is probably embedded in a legal system that issues some non-trivial protections. So, like the taking of vitamin A for strong nails, L can be seen as miniscule protection of an important good, and needn't be seen as protection of a miniscule good.

[209] Actually, as we'll see in the next section, any law must be a *promise* to protect a moral good. Just as not all promises are kept, so not all laws are enforced and, therefore, the corresponding rights fail to be protected.

[210] Coleman and Murphy (1990: 19-33).

[211]

See David Lewis' "Language and Languages." Lewis argues that the expression "language" is ambiguous, sometimes denoting an abstract, function-theoretic pairing of spatiotemporal tokens (ink-marks, sounds, etc.)

and meanings, other times denoting patterns of activity. For reasons that I give elsewhere ("Function-theoretic versus psychological conceptions of semantic rules", unpublished), I think that Lewis' position is quite erroneous and that, even though linguistic expressions sometimes encode timeless abstracta, languages themselves are *always* spatiotemporal entities (notwithstanding that their boundaries are considerably more amorphous than those of a paradigmatic spatiotemporal entity, like a car or a boat).
[212]
Some content-externalists would disagree with this, saying that if I have a *de re* awareness of Alpha Centauri, then that star is a constituent not only of the state of affairs of which I am aware, but of my very awareness of it. In *Mental Content,* Colin McGinn argues persuasively (and, to my mind, cogently) that we must accept this view, so long as we accept certain now widely accepted externalist views on the individuation of mental content.

McGinn himself says that, in his view, Alpha Centauri *is* a constituent of my awareness. He then tries to take the sting out of this conclusion by saying that, since the mind is not a "substance", the fact that my mind engulfs remote parts of the cosmos is somehow innocuous. I disagree with all of this *except* for McGinn's conditional point that *if* we accept content-externalism, then we must accept the dubious point that my mind has stars and chairs as proper parts. I take it as a datum that our minds are not spread out over the cosmos and that, if indeed content-externalism requires that they be thus spread out, we must reject that doctrine.

In *Conceptual atomism and the computational theory of mind* (forthcoming), and also in *Literal meaning and cognitive content* (unpublished), I argue why those very externalist theories are to be rejected, and I try to show that the intuitions that motivate them can be accommodated within a traditional internalist theory of content-individuation. But these are

obviously not matters than can be adjudicated within the present discussion.

[213]

In any case, this is Tyler Burge's (1982) view. Burge's view strikes me as extremely plausible, and I agree with it.

[214]

See Lewis (1972).

[215] Here my bias towards an Aristotelian ethics is apparent. If you reject my point that a right is a right to flourish, then see this discussion as an attempt to establish the conditional claim that *given* a certain reasonable (though perhaps false) view of rights, it follows that nations have rights.

[216]

It isn't clear whether, in order for x to be capable of flourishing, it is *necessary* that x have aspirations. An adaptation of a question posed by John Stuart Mill can help put this problem into focus. Suppose that x is a primitive life-form that can feel intense pleasure and intense pain, but doesn't have the cognitive wherewithal to have aspirations or, indeed, to have any propositional attitudes. Does the concept of flourishing apply to x? It seems to me that x *would* be capable of a rudimentary sort of flourishing: it would be flourishing in so far as, and only in so far as, it felt good. But, admittedly, this last point is debatable.

In this context, we don't have to settle the question just posed. Given that – as we will see in a moment – nations have aspirations, all we need to know is whether anything that has aspirations has rights. For reasons previously discussed, the answer seems to be "yes."

[217] There is one possible objection that we should address:

"It is a fact of life that hollow assurances can be made. Used-car salesmen make them all the time. What you just said about assurances may be true of

meaningful (non-hollow) assurances, but it is true of assurances *simpliciter*. You are confusing the concept of an assurance *simpliciter* with the concept of a *credible* assurance."

In my view, this point is either purely terminological or simply false. It seems to me that an assurance is the successful *providing* of a certainty (a surety). "Assure" – like "know" – is a success-verb. As before, fiction may help expose the relevant principles. I believe that Smith is a pathological liar. At the same time, I believe that no one other than Smith is capable of refunding me my money. Under these circumstances, Smith says "I assure you that you will receive a refund." It seems to me that, under those circumstances, I have *no* real assurance that I will receive a refund. Therefore, Smith hasn't given me an assurance. A meaningless assurance – one that has no credibility – is no assurance at all.

This is not hard to corroborate. In saying "I assure you that your money will be refunded", what is Smith *trying* to do? He is *trying* to instill in you a feeling of confidence that your money will be refunded. So far as he fails to do that, his speech-act is abortive. This suggests, if it doesn't actually entail, that he has *failed* to tender an actual assurance.

Of course, we often use the expression "assurance" to refer to meaningless assurances, i.e. to refer to failed attempts to give assurances. But, as we said earlier, that is like referring to counterfeit money as "money." If I have *no* faith that Smith will refund my money, then I am either speaking ironically or being willfully misleading if I tell everyone that I have Smith's assurance that my money will be refunded. We do speak of "meaningless assurances." But such talk is parasitic on that of kept assurances, suggesting that "meaningless assurance" is in the same category as "counterfeit money"

and “decoy duck.”

In conclusion, the existence of so-called empty assurances is no real threat to our principle that x can assure y of z only if y has a reasonable (though perhaps defeasible) expectation that x can and will provide z to y.

[218] Even highly unjust legal systems consistently provide certain basic protections to those who fall under them. Before the Civil Rights movement, U.S. law was grossly unfair to African Americans. But they were still consistently provided with certain basic protections. Those who robbed or killed them would, as a rule (though not, of course, on every single occasion), be jailed. It was because U.S. law consistently provided these basic protections that African Americans had legal obligations under the system, wicked and unjust though it was to them.

Obviously governments commit innumerable unspeakable abuses. But there is no denying that, as a rule, people can count on certain protections from their governments, so long as they behave in certain ways. And it is because governments provide individuals with consistent protections that individuals have legal obligations towards governments.

[219] Of course, it is an entirely empirical question whether Guyana does enjoy such protections. I may be quite wrong to deny that it does. But our purpose here is to expose the logical structure of propositions like “there is international law” and “Guyana has various legal obligations towards France.” If the analysis just put forth is correct, we know exactly what must be the case for such statements to be correct; and ascertaining their truth-values becomes a straightforward empirical matter. If nothing else, our analysis gives us a way of cutting through the sludge of rhetorically, or otherwise non-logically, driven discourse about international law (so-called).

[220] More exactly, it is a situation that is so extreme that there is a failure to

satisfy the boundary conditions that those governing principles presuppose. A singularity is not an *exception* or a *counterexample* to such a principle, except in the degenerate sense in which regarding points as limiting-cases of lines would create an exception to the principle that lines have length. It should also be pointed out that even if there are *many* biker-bars of the sort just described, they are still singularities. The concept of a singularity is to be understood in dynamic, and not strictly numerical, terms.

[221] Popper argues powerfully for this view in *The Open Universe* and also in *"Of Clouds and Clocks"* (in his book *Objective Knowledge*).
[222] This is, of course, an empirical proposition. I may be quite wrong to assert it; and I don't have the knowledge-base to assert it with conviction. What I should say is: "Given what my limited empirical knowledge of international conditions suggests…"
[223] It is fact worthy of study that Stalin was himself Georgian, not Russian.

[224] Obviously if one takes nations to be individuated by their political systems, then it follows trivially that Georgia (and Poland and Hungary…) ceased to exist the moment they become satellites of the Soviet Union. But this conception of nation-hood doesn't correspond very well to that involved in day-to-day thinking about nations. Surely *Poland* has pre-existed, and now post-exists, the Soviet era, at least on the most natural conception of what Poland is.
[225] Obviously there are murders, physical assaults, and so forth. But, statistically, these are extremely rare. Do *you* know anyone who has been murdered?
[226] Strictly speaking, flourishing cannot be *identified* with the attainment of objectives. Suppose that Smith's objective is to smoke crack. It seems that,

in attaining that objective, Smith would be undermining himself, and thus not flourishing.

Also, the flourishing of an infant – its growing physical and psychologically – obviously cannot be understood exclusively in terms of its attaining its objectives. It would seem that its growing, and thus flourishing, is a *pre*condition for its having the degree of psychological sophistication needed to have objectives. So it is true that flourishing cannot be categorically identified with attainment of one's objectives.

Flourishing pulls apart from fulfillment of one's objectives only in cases involving pathology (e.g. an addition to crack) or involving a creature that doesn't have the wherewithal to have objectives to begin with. So the right rule seems to be that *ceteris paribus* flourishing and attainment of objectives coalesce. They coalesce except where degenerate cases (a crack-addict, a pre-verbal infant) are concerned. As we discussed earlier, a hypothesis that is counter-exampled only by degenerate or limiting cases is not, on that account, a false hypothesis.

At any rate, the connection between flourishing, on the one hand, and attaining one's objectives, on the other, is sufficiently tight that, in this context, we may safely identify them.

[227] Presumably, anything capable of establishing itself as a protector of nations' rights would *ipso facto* have such coercive power. So, strictly speaking, this last point is probably redundant.

[228] There is a qualification. Insects aren't subject to our laws, but that doesn't nullify those laws. The correct statement of the principle in question is this: where there is one thing *of the relevant type*. Where human laws are concerned, a thing of the relevant type is a human being who is living in a

certain geographical area. To avoid verbosity, I will henceforth leave implicit the qualification “of the relevant type.”

[229] If I am not mistaken, Hobbes made similar points in *Leviathan*, both in connection with national and international law.

[230]

The need to distinguish between universal and non-universal international laws was brought to my attention by Peter J. Rudinskas. (Apart from that, everything said here about that distinction is due to the present author, as is the use of the words “universal” and “non-universal” to mark that distinction.) Agreeing with me that there are no universal international laws, Mr. Rudinskas suggested that there probably non-universal laws. He suggested that the E.U. is governed by such laws.

[231] The following remarks owe much to Hart’s (1977) work.

[232] For obvious reasons, I am setting aside as irrelevant the fact that, given relativity theory, there are no genuinely rigid objects and also that, given that same theory, there are no objects that are straight in the Euclidean sense.

[233] Of course, some of these freedoms will be of a low and wretched character, e.g. the freedom to torture others. But others will not. Consider the various would-be Tolstoys whose dreams must be sacrificed for the preservation of our political institutions.

[234] It intuitively seems that, the moment Jews stopped enjoying the protection of the German government, they stopped having legal obligations towards it. This obviously confirms our analysis.

[235]

Dworkin himself (1986: 2) writes: “the law often becomes what judges say it is.”

[236]

I say “so-called” because they should really be called the legal *anti*-realists. What the judge says is *ipso facto* correct. This means that the judge’s “interpretation” isn’t answerable to any objective reality.

But, of course, the self-described legal “realists” meant to convey that *they* were realists about the nature of the law, not that legal interpretation (so-called) *itself* embodied some kind of realism about the law.

[237]

Of course, given the points made earlier regarding the non-binary nature of morality, each of the occurrences here of the terms “perfidious” and “virtuous” must be taken to be preceded by a phonetically unrealized “relatively.”

[238] In the next section, we will scrutinize this claim, and find that it isn’t quite the right way of putting the matter. But, in this context, the inaccuracy is harmless and is necessary for exposition.

[239]

Strictly speaking, whether someone is bald isn’t a function solely of how many hairs he has. It is also a function of, possibly among other things, the thickness of his hair-follicles and of the size of his scalp. Even if n is smaller than m, someone with m hairs may be bald whereas someone with n hairs is not bald. This could be the case if, for example, the hair-follicles of the person with m hairs were extremely thin, whereas the hair follicles of the person with n hairs were extremely thin. But in this discussion, for the purposes of exposition, let us pretend that whether one is bald is a function *solely* of the numbers of hairs that one has.

[240] Consider the sense in which Albanians do, whereas monolingual English speakers do not, understand Albanian sentences. *That*, it seems to me, is a paradigmatic and non-metaphorical use of the term “understand.” Given that interpretation of the word “understand”, it is very easy to

understand the words “the right to bear arms shall not be infringed.” The debates concerning the Constitutionality of gun-regulation don’t have anything to do with understanding in this sense or even with anything that can naturally be put on a continuum with this sort of understanding.

Consider the case of a doctor who interprets the results of a test of respiratory capacity. A layperson doesn’t know what is meant, in a quite literal sense of “meant”, by the numbers and graphs on the computer printout of the test. The layman doesn’t know what the “87” refers to. (Does it refer to his blood-oxygen level, to his ability to expel air from his lungs, to his ability to take air into his lungs, to his ability to keep air in his lungs, or to none of the above?) The layperson just doesn’t know why there are discontinuities in the graph or why the graph slopes downwards. The doctor isn’t there (merely) to give deep insight. He is there to say what those graphs and figures mean. He is there to say that the “87” refers to your blood-oxygen level, and so on.

Of course, *once* the doctor says what those graphs and figures mean – in the narrowly symbolic sense of “mean” just discussed – he then goes on to give the *meanings* of those meanings, so to speak. In other words, he says what the consequences are for your health that your blood-oxygen level is such and such, what you must do to improve your ability to intake air, and so on. So the doctor tells you what the print-out means in at least two senses of the word “means.”

But, in both cases, the doctor’s act of meaning-giving consists in his *identifying*, not creating, some existing reality. He is *identifying* the fact that your blood-oxygen level is low. He is *identifying* the fact that, if you do not increase it, your health will suffer. He is not to any degree legislating or otherwise creating any of these realities.

In “interpreting” the law, the judge isn’t doing anything comparable to

any of those kinds of meaning-giving. Obviously the judge isn't saying what the words "abortion is legal" mean in the narrow symbolic sense. Nor, as we have seen, is the judge identifying some fact about what the laws *actually*, though obscurely, assure people of. As we've seen, an obscure assurance is, to that extent, *no* assurance. So in so far as the law is obscure there is nothing to know. There is nothing comparable to a hidden tumor or an exotic virus.

Further, in "interpreting" the law, the judge *is* creating a new reality. If he rules one way, that is the law. If he rules the other way, *that* is the law. His interpretation never fails to line up with legal reality, meaning that it creates that reality. But a doctor's interpretation of some test-results can fail dramatically to line up with medical reality.

Given *any* paradigm-case of interpretation – e.g. a physicist's interpretation of meter-readings, a therapist's interpretation of a patient's behavior – everything we just said (mutatis mutandis) shows that it bears no significant resemblance to legal interpretation (so-called).

[241]

Let me explain the presence of the words "at least not necessarily" in this sentence. We could obviously imagine a case where an erroneous judicial interpretation does create a contradiction within the relevant legal system. But, as we've seen, such interpretations do not *categorically* do so. Also nothing of relevance follows from the fact that erroneous judicial decisions sometimes create legal contradictions, since we could equally imagine a case where a correct interpretation of some law has the same result.

Here is a case in point. At time t, in legal system S, there is one law L saying that people cannot smoke cigarettes in restaurants. The letter of L leaves it open whether smoking cigars or pipes is permissible. At the same time, there is some law L* saying that the *owner* of a property is allowed to smoke cigars (not cigarettes, not pipes – but specifically cigars) while he is

on that property. Given only what we've said, there is no reason to believe that, strictly speaking, S is guilty of having inconsistent laws (i.e. of having a pair of laws such that the one permitted what the other forbad). There is a law permitting the smoking of *cigars* (not cigarettes) on one's own property; and there is another law forbidding the smoking of cigarettes in restaurants.

One restaurant owner Smith is eating at his own restaurant, and he lights up a cigar. The case goes to judge Brown. Brown decides, in its spirit though not its letter, L *does* forbid the smoking of cigars in restaurants. Suppose that L was made at a time when cigarettes were the only tobacco-products in existence. So the makers of L had no idea that there would ever be cigars or pipes. In that case, the assurance that they issued did not *itself* guarantee anyone protection from the ravages of cigar-smoke. But it is clear that, under the circumstances, the appropriate way to *extend* that assurance – the right way to modify it, so as to make it applicable to a society where there are not only cigarettes, but also cigars and pipes – is to treat it as forbidding the smoking of cigars in restaurants.

So Brown is surely right to see Smith's behavior as a violation of L – or, rather, Brown is right to extend or delineate L in such a way that Smith's behavior is in violation of that delineation. At the same time, in thus delineating L, Brown has created a contradiction within S. Thanks to Brown's ruling, S comprises a set of laws such that one forbids what the other permits. Brown's correct delineation of L forbids the smoking of cigars in restaurants. At the same time, L* allows people to smoke cigars while on their own property, and thus permits Smith to smoke cigars in his own restaurant. A brief review the of the case makes it clear that, prior to Brown's ruling, there was no contradiction in S and that after Brown's (correct) ruling, there was such a contradiction. So it is not always rulings that embody erroneous delineations of the law that create contradictions within a legal

system. In fact, there is no reason to believe that false rulings are more likely than correct ones to do so.

The reason that Brown's ruling results in a contradiction isn't that the ruling embodies a mis-delineation of the law. The reason lies in the nature of the system prior to that delineation. (Though not formally inconsistent, that system embodied a certain lack of foresight.) In general, if a ruling results in a contradiction, that doesn't necessarily reflect anything about the ruling itself. In particular, it doesn't mean that the ruling was the wrong one.

Of course, there is an obvious objection to the argument just given:

"When a judge makes a ruling, he has to consider the legal system as a whole. He can't just focus on one law. G*iven only L*, Jones is probably right to say that Smith is in violation of the law. But the legal system in question doesn't just include L – it includes various other laws, including L*. A judge's job is to do right by the legal system as a whole. This means that a judge can *never* rule in a way that creates a contradiction. Since Jones' ruling did create a contradiction, that ruling was *ipso facto* erroneous."

But this response involves a completely *ad hoc* characterization of what judges are supposed to do. The traditional picture is this. When the letter of law is silent, the judge is supposed to rule in a way that is maximally consistent with the spirit of the law. The objector's picture is this. When the letter of law is silent, the judge is supposed to rule in a way that is maximally consistent with the spirit of the law – except when doing so creates a contradiction: in that case, the judge's ruling should *not* be consistent with the spirit of the laws. If the objector's picture is right, the concept of legal interpretation is not a very organic one. A good conceptual analysis of a

fundamental concept is never given by a proposition that has an “except” in it. One pre-Fregean analysis of number was (approximately) this: a number is a “measure of multiplicity.” But zero and one are not “measures of multiplicity”, on any delineation of that expression. So if it is to avoid falsely saying that zero and one are not numbers, that analysis must be given by the statement: a number is anything *other than zero and one* that is a measure of multiplicity. (We are setting aside the fact that negatives and irrationals, and other forms of number, are not happily seen as “measures of multiplicity.”) But the italicized expression is a red flag. With some trivial exceptions, no correct conceptual analysis is given by a proposition containing an exception-clause. The exceptions concern cases where the concept in question is an artificial one – e.g. x is shmoo if it is blue *unless* x is on Mt. Olympus, in which case x is shmoo if it is green. It is clear why the existence of exceptions of this kind rather confirms our point that correct conceptual are not given by propositions containing exception-clauses.

The objector’s analysis of the concept of legal interpretation is given by just such a proposition: x is a correct legal interpretation if x is true to the spirit of the laws, *unless* x results in a contradiction, in which case x is a correct legal interpretation if it is untrue to the sprit of the law but preserves the consistency of the relevant system. The concept of legal interpretation is not in the same category as the concept *shmoo.* So it is clear that the objector’s is point false and, consequently, that we *have* shown that a correct legal interpretation can create a contradiction in a hitherto consistent legal system.

[242] Of course, some will reject my belief that laws are categorically assurances, and will prefer to think of law in terms of the concepts of permission and prohibition. But in this context, that makes no difference. Like “assure”, “permit”, and “prohibit” are success-verbs. Suppose that you

don’t speak English. In that case, in saying “I prohibit you from smoking in my house”, I am not succeeding in prohibiting you from doing anything – I am merely *trying* to do so. A cryptic prohibition is no prohibition; and, for similar reasons, a cryptic permission-giving is no permission-giving. So if laws are prohibitions or permissions, an exact analogue of our argument shows that legal interpretation (so-called) doesn’t identify existing law, but rather replaces existing law with new law.

I agree, by the way, that laws *are* prohibitions and permissions. Obviously it is not wrong to say that L is a law only if L is a certain kind of permission or a certain kind of prohibition. But I think that, in all cases, the permissions and prohibitions in question are assurances of rights. Further, we cannot give a disjunctive analysis of law; we cannot say that L is a law iff L is *either* a permission (of a certain kind) *or* a prohibition (of a certain kind). The concept of law is a unified one, and the analysis of that concept is therefore not given any disjunctive proposition. For reasons given in chapter 5, I believe that the right non-disjunctive proposition is one of the form: L is a law iff L is an assurance (of a certain kind).

[243] Of course, we could certainly *imagine* a case where a Constitutional (or, more generally, legal) interpretation creates a contradiction – where it creates a pair of laws such that one forbids what the other permits. Suppose that an insane judge “interpreted” the Constitution, in its present (2006) form, to mean Rutherford was not allowed to vote on the grounds that she is a female. That decision would create a contradiction. Police would have a legal duty to comply with the judge’s decision, and also to ensure that Rutherford be allowed to vote.

But most legal misinterpretations aren’t in this category. While the Supreme Court has made many decisions that are inconsistent with the Spirit

of the Constitution, and are thus “misinterpretations” of it, those decisions seldom (if ever) resulted in situations where some one act was simultaneously legal and illegal. (As we’ve seen, this means that Fuller is wrong to say that such misinterpretations necessarily create contradictions within U.S. law.) For the reasons already given, it follows that the Supreme Court is not *identifying* the assurances actually provided by the Constitution, but rather replaces those assurances with other assurances that are more complete than they are, but are otherwise coincident with them. For this reason, judicial interpretation is not *mere* legislation. Further, some of those completions are more correct than others, in the social-psychological sense of “consistent” discussed earlier. That is why, even though the judge’s completion of the law *always* coincides with subsequent legal reality, some judicial rulings are more correct than others.

[244] These are tendentious examples, of course. And, of course, it is pretty clear (I think) that the U.S.S.R. didn’t have a constitution except in a purely nominal sense, whereas the U.S. really does have a constitution. Nonetheless, it is suggestive that so-called constitutions can diverge so much from legal reality; and we don’t have to look at extreme cases, like the U.S.S.R, to find examples of that sort of divergence. In any case, I am only trying to make a *prima facie* case, not to provide an airtight proof.

[245]

Like me, Timothy Williams (*Vagueness*) and Roy Sorensen (*Blindspots*) believe that there is no objective indeterminacy. Bertrand Russell also held this view. The argument given here for that view is not in the literature, and it therefore hasn’t been criticized in the literature.

[246] In order, for the argument currently being presented to be completely

cogent, there are some subtleties that must be dealt with. But so as not to obscure the basic structure of my argument, I have decided to put these subtleties in a footnote. Also, those subtleties would be of virtually no interest to anyone who is not interested in logical minutiae that have only the most tenuous connection to matters of law or ethics.

Whether a person is bald obviously *supervenes on* facts about the physical make-up of the top of his head – how many hairs he has, how thick they are, and so on. In other words, if a person with n hairs (of such and such thickness, on a head of such and such size…) is bald, then *necessarily* a person of n hairs (same qualifications) is bald. (Henceforth, to avoid verbosity, I will omit the qualification "of such and such thickness…") When they hold, supervenience relations hold necessarily.

This brings us to a crucial joint in our argument. Arguably, some supervenience relations – and, more generally, some necessary relations – are *a posteriori.* The state of affairs described by the proposition *Smith is happy* supervenes on (let us suppose) that described by *such and such neural events are occurring*. But the connection between these two propositions, though necessary, is *a posteriori.* In general, any concomitances between the neural and the mental can known only *a posteriori* and, consequently, there is no analytic or *a priori* inference from *such and such neural events are occurring* to *Smith is happy* (or *vice versa*).

But the relationship between P and Q is not in this category. Suppose you know that Fred has zero hairs. In that case, you *ipso facto* have enough empirical information to determine whether Fred is bald. Of course, you must do empirical work to know that somebody has zero hairs. But you don't have to do empirical research to discover that *if* a person has zero hairs, *then* that person is bald. If you know that Fred has zero hairs, but you don't know that Fred is bald, then you either have a serious lack of deductive horse-power or

you simply don't have a grasp of the concept of baldness. Thus, the connection between *x has zero hairs* and *x is bald* is analytic. The same is true for *x has one hair* and *x has two hairs* and so on…

Here we must deal with a rather delicate point. For any number n, *if* a person with n hairs is bald, then *x is bald* follows both necessarily and analytically from *x has n hairs*. Let me give the raison d'être for this claim. We just say that, for n=0 and n=1…*x has n hairs* entails (i.e. analytically necessitates the truth of) *x is bald.* It would be deeply arbitrary to say that, for n=1 and n=2 and n=3, *x has n hairs* entails *x is bald* but that, for some higher values of n, *x has n hairs* necessitates *but does not entail* the truth of *x is bald.* Given that there is an analytic relationship between *x has n hairs* and *x is bald* for small values of n, it would be arbitrary to say that, for higher values of n such that a person with n hairs is still bald, while saying that the connection between *x has n hairs* and *x is bald* is not analytic. So we may assume that, for any n, in so far as any definite proposition is associated with *x has n hairs*, and if a person with n hairs is bald, then the proposition meant by *x has n hairs* *entails* that x is bald. (In other words, given the truth of the proposition meant by *x has n hairs*, it follows *analytically* or, at least, *a priori* that x is bald.)

So while it may be true that *some* supervenience relations are not analytic, the relation under investigation is not one of them. (As before, let Q be the proposition that is meant by "Harry is bald", on the supposition that a single proposition is indeed associated with that sentence; and let P be a proposition detailing Harry's hair-situation, i.e. identifying the number of hairs on his head, as well as the ratio of hair-covered to hairless skin, and so on.) *If* the truth of Q is necessitated by the truth of P, the necessitation consists in an analytic, or a priori, relation – one whose presence (or absence) is to be decided entirely on the basis of the contents of those two propositions alone,

and not on the basis of empirical information over and above such as is embodied in them. Obviously each of P and Q is a non-analytic, empirical proposition. But, if true, the proposition *if P, then Q* is analytic, the same holding of *if P, then not Q* and also, consequently, of *it is not the case that, if P then Q*. (Given that P is a consistent proposition, i.e. one that doesn't entail a contradiction, the truth of *if P, then not Q*, entails *it is not the case that, if P, then Q*.)

[247] Another way of putting it is this. The *identity* of law L is indeterminate to the extent that it is indeterminate what it is that is assured by L or who precisely is on the receiving end of that assurance.

[248]

There is one other point to make in this context. It seems pretty clear that, for some values of n, English semantics leave it open whether the sentence *a person with n hairs is bald* is true or not. It would be absurd to suppose that one of the semantic rules of English assigned truth to the sentence "a person with 43,243 hairs is bald" and falsity to "a person with 43,244 hairs is bald."

At the same time, we must also explain the fact that, setting aside all issues relating to linguistic competence and to empirical knowledge, we sometimes feel genuine hesitation as to whether to categorize somebody as bald. Two people who both speak English perfectly and who are both maximally well-informed about the state of Smith's coiffure – about the number of hairs on his head, and their distribution and thickness, and so on – can have at least some tendency to disagree about whether it is correct to describe Smith as "bald." *Prima facie*, this seems inconsistent with my contention that, in such cases, English semantics determinately leave it *open* whether Smith is bald. *Prima facie*, the fact that this sort of situation arises favors the view that English semantics doesn't determinately *fail* to yield a verdict, but rather that it is *open* what kind of verdict, *if any*, is yielded. In

other words, my position seems inconsistent with the fact that it is often open whether it is open whether somebody is bald; my position seems to demand that, on each occasion, it is an open and shut case: *x is bald* is definitely true, definitely false, or definitely without truth-value. Put yet another way, my position seems inconsistent with the (apparent) fact that it is open whether it is open whether *x is bald* is true.

I would now like to say why my analysis *is* consistent with the existence of situations of the sort just described, and is not guilty of the errors just described. We've stressed that, strictly speaking, delineations are neither true nor false. But, as we've also stressed, some delineations are better founded than others. As the phenomenon of heat came to be better understood, physicists were in reasonable agreement that the range of phenomena denoted by expressions like "temperature of 78°" (and, of course, their counterparts in other languages) had to be extended. But it was a matter of great theoretical controversy in physics exactly *what* those extensions were to be. (For expository reasons, let us henceforth speak of these past controversies in the present tense.) For some phenomena x, it is known among these physicists that, given its existing meaning, it is neither correct nor incorrect to apply the expression "temperature of 78°" to x. But, as the physicists involved well know, that doesn't mean that it is a matter of caprice whether to apply that expression to x. There is an objective fact as to whether extending the meaning of that expression will, or will not, result in a more accurate and powerful theory. Given the laws of physics, given the logical structure of existing physical theory, and given what is *already* meant by the expression "temperature of 78°", there is an objective fact as to whether choosing to apply that expression to x will or will not make our theories more accurate. That fact has nothing to do with anyone's opinion; it is a function entirely of mind-independent logical and physical facts.

Of course, the expression “temperature of 78°”only means what we let it mean. But once it has been assigned a certain meaning, it is a matter of mind-independent objective fact whether a certain theory *T* is, or is not, enhanced by choosing to extend the meaning of that expression in a given way. When physicists argue about whether “temperature of 78°” ought to apply to x, they obviously aren’t quibbling about semantics; they are in agreement as to the fact that semantic facts are neutral. What they disagree about is whether, *given* those semantic facts and *given* the relevant mind-independent empirical and logical realities, a particular theory would be made more accurate (or otherwise better) by extending the meaning of that expression in a particular way.

At the same time, it would be easy to fall prey to the illusion that the debate concerns whether, given its *existing* meaning, the expression “temperature of 78°” *already* applies to x. In their heated back and forth, the physicists are probably not going to concern themselves with verbal niceties; they are not going to say anything like:

“Although, given its current semantics, it is strictly speaking an open question whether the expression ‘temperature of 78°’ applies to x, it is nonetheless advantageous on theoretical grounds to delineate that expression in such a way that it *does* apply to x.” Rather they will just say “x is 78°” or “x is not 78°.”

For reasons of brevity and efficiency, the structure of their debate will be falsely made to seem as though there is some fact about what is meant by the expression “temperature of 78°” and that the question is merely whether that existing meaning applies to x or not. Consequently, what is, in fact, a debate about how to delineate an expression – how to extend its meaning in a

principled way -- comes to look like a debate as to whether the *existing* meaning of some expression in fact applies to some phenomenon.

So while it remains clear that there is *some* kind of objective fact at issue, the actual nature of that fact is obscured: instead of being seen for what it is – a debate about whether, given the existing (incomplete) semantic realities, and given also the relevant, mind-independent physical and logical realities, theory *T* would be made better by extending the meaning of a certain expression in a certain way – the debate comes to appear to be one about what lies in the depths of the *existing* meaning of that expression.

These points apply to the situation described earlier – the situation where two linguistically competent and empirically well-informed people disagree about whether Smith ought to be described as "bald." First of all, *there is* a real issue. Their debate is not a debate about nothing. Their respective positions are answerable to objective facts. But – despite what the participants themselves may say about it - the debate is not whether, given its *existing* meaning, the term "bald" applies to Smith or not. It would be absurd to suppose that English semantics was so complete; it would be absurd to suppose that, given English semantics, Smith's having 43,243 hairs makes "Smith is bald" true and Smith's having 43,244 hairs makes that same sentence false.

Whether the participants know it or not, the debate is given by a question having at least approximately the following basic form:

"Given what is *currently* meant by the expression "bald"; given the state of Smith's coiffure (the number, thickness, and distribution of his hairs); and given the various desiderata that a system of communication is supposed to satisfy (e.g. it is supposed to facilitate communication and also expedite thought); would those desiderata be *better* met by extending the meaning of

"bald" in such a way that it did apply to Smith?"

Once the desiderata just mentioned are clearly identified, there is an objective answer to that question. In other words, that question cannot be answered through stipulation, but only through examination of the relevant mind-independent logical (and empirical) relations.

Here we must echo a caveat made earlier. Obviously the word "bald" means whatever we let it mean. But *given* that it has that meaning; and *given* the various desiderata that we would like the English language to fulfill; and, finally, *given* the laws of psychology and possibly even of physics: *given* all of these things, it is entirely a matter of mind-independent fact whether or not certain extensions of the meaning of "bald" conduce to the fulfillment of those desiderata. When the two people debate about whether or not Smith is to be described as "bald", they *are* having a substantive debate, and their positions *are* evaluable in terms of objective (albeit extremely delicate and theoretical) facts.

At the same time, their debate might *appear* to be a strictly semantic one; it might *appear* to be a debate as to whether, given its *current* meaning, "bald" applies to Smith or not. Here we need only repeat (*mutatis mutandis*) what we said earlier. In their heated back and forth, the two participants are not going to observe the kinds of niceties that are the focus of theoretical treatises, like the present one. What comes out of their mouths is not going to bear any resemblance to:

"Although, given the existing semantics of the term 'bald', it is neither correct nor incorrect to predicate that expression of Smith, nonetheless, given the various desiderata that the English language is meant to satisfy, it is

necessary to extend the semantics of the expression ‘bald’ so that it *would* be correct to apply it to Smith; for, in light of the relevant empirical and logical interrelations, not making that extension would force the English language to deteriorate in respect of its ability to express such and such information while only expending thus and such effort…”

Rather, what comes out of their mouths will be more like: “Smith is bald!” or “I’m not so sure – he’s kind of bald, but I think that, in the end, he probably isn’t bald.” In general, for reasons of communicative efficiency, the actual structure of the debates between these two fellows will be suppressed, and what is actually a multi-dimensional debate about delicate mind-independent relations among semantic, logical, and empirical facts will come to seem like a fruitless and one-dimensional debate about whether the *actual* meaning of the term “bald” applies to Smith or not.

Given the obvious fact that the debate in question is *not* to be settled by looking into the depths of English semantics, and given that (by hypothesis) both participants know exactly how many hairs Smith has, many are quick to assume that the debate concerns propositions themselves: many are quick to believe that the *proposition* meant by “Smith is bald” is such that there is no fact as to whether it is true or not.

But such a position is absurd, given that the *identity* of a proposition is indistinguishable from its entailment-relations – given that entailment-relations are content-internal. Further, such a position warrants a massive and *ad hoc* complication of logic. Those who believe that *propositions* are objectively indeterminate, i.e. that *objectively* there is no fact as to whether they are true or false, are thereby rejecting the law of excluded middle, along with everything that it entails. There are a couple of reasons why, in my judgment, this would not be a wise course of action.

First of all, the law of excluded middle is equivalent with the law of non-contradiction, at least within classical logic. This suggests that, at the very least, those principles are not likely to be independent of each other (though the degree of interdependence depends on the system in question). Given that the law of non-contradiction is *uncontroversially* true, this in turn suggests that, if we discard bivalence (the principle that any proposition is determinately true or false), we are likely to end up with a system that *admits* true contradictions.

There is more to be said about why *not* to reject bivalence. When that principle is rejected, it is (often, though not always) because we feel hesitation or doubt about whether "bald" applies to Smith or not – because, in general, we aren't sure about whether some *expression* is applicable to some situation. But, it seems to me, rejecting bivalence for that reason is like a bookkeeper inventing a new system of arithmetic because he can't convert dollars into yen. Feelings of hesitation concerning words are not enough to warrant a reconstruction of logic or of our conception of meaning.

There is also the fact that, once you pick a logical system, you are stuck with it. You can't say that classical logic is wrong when we are talking about baldness, but right when we are talking about number theory. So if you accept non-classical logic to deal with the fact that we aren't comfortable describing Smith as "bald" or as "not bald", then you must stick with your non-classical logic when you are dealing with situations that obviously *are* binary; and your new system will inevitably fail to do justice to the intuitions that you have in connection with *those* situations. For example, depending on your specific choice of non-classical logic, you might end up saying that *no* statements (other than tautologies or tautological falsehoods) are either entirely true or entirely false.

(Indeed, I have long suspected that if – with many philosophers and

physicists – we take quantum uncertainties to warrant a rejection of two-valued logic, then given that all spatiotemporal phenomena are characterized, to some degree, by such relations, we end up in the absurd position of stripping *all* empirical statements of complete truth or complete falsity. In *Experience and Prediction,* Reichenbach made a similar point. Since he believed that quantum-theory did warrant a rejection of bivalence, Reichenbach was forced to say that *no* propositions besides tautologies were completely true or completely false. Reichenbach accepted this position, and developed a whole philosophical system to accommodate it. But it seems to me a datum that some empirical propositions are simply true and that others are simply false.)

This is not to mention the point stressed earlier, namely, that non-classical logic is flatly incompatible with the fact that entailment relations are obviously relations of content. Given their circumstance-invariant character, there is no other plausible analysis of such relations.

We have seen, I hope, that by carefully distinguishing between delineation, on the one hand, and meaning-identification, on the other, we can accommodate all of the intuitions that motivate the belief in indeterminacy, without sacrificing the benefits with which an acceptance of bivalence (*any proposition is determinately true or determinately false*) provides us.